

Assignment_4_Code_906213559

1 ECE-6524 / CS-6524 Deep Learning

2 Assignment 4 [100 pts]

In this assignment, we will explore Recurrent Neural Networks (RNN) to deal with the data with temporal sequence. Specifically, we will generate captions for images. We will design an encoder-decoder architecture to achieve this. This homework is inspired by Stanford CS231n and UCSD CSE253.

2.1 Submission guideline for the coding part (Jupyter Notebook)

1. Click the Save button at the top of the Jupyter Notebook
2. Please make sure to have entered your Virginia Tech PID below
3. Once you've completed everything (make sure output for all cells are visible), select File -> Download as -> PDF via LaTeX
4. Look at the PDF file and make sure all your solutions are displayed correctly there
5. Zip all the files along with this notebook (Please don't include the data). Name it as Assignment_3_Code_[YOUR PID NUMBER].zip
6. Name your PDF file as Assignment_4_NB_[YOUR PID NUMBER].pdf
7. **Submit your zipped file and the PDF SEPARATELY**

Note: if facing issues with step 3 refer: <https://pypi.org/project/notebook-as-pdf/>

2.2 Submission guideline for the coding part (Google Colab)

1. Click the Save button at the top of the Notebook
2. Please make sure to have entered your Virginia Tech PID below
3. Follow last two cells in this notebook for guidelines to download pdf file of this notebook
4. Look at the PDF file and make sure all your solutions are displayed correctly there
5. Zip all the files along with this notebook (Please don't include the data). Name it as Assignment_2_Code_[YOUR PID NUMBER].zip
6. Name your PDF file as Assignment_4_NB_[YOUR PID NUMBER].pdf
7. **Submit your zipped file and the PDF SEPARATELY**

While you are encouraged to discuss with your peers, all work submitted is expected to be your own. If you use any information from other resources (e.g. online materials), you are required to cite it below you VT PID. Any violation will result in a 0 mark for the assignment.

2.2.1 Please Write Your VT PID Here: 906213559

2.2.2 Reference (if any):

<https://towardsdatascience.com/automatic-image-captioning-with-cnn-rnn-aae3cd442d83>

<https://medium.com/@stepanulyanin/captioning-images-with-pytorch-bc592e5fd1a3>

<https://towardsdatascience.com/pre-trained-word-embeddings-or-embedding-layer-a-dilemma-8406959fd76c>

In this homework, you would need to use **Python 3.6+** along with the following packages (**need to update**):

1. pytorch 1.2
2. torchvision
3. numpy
4. matplotlib
5. nltk

To install pytorch, please follow the instructions on the [Official website](#). In addition, the [official document](#) could be very helpful when you want to find certain functionalities.

```
[1]: # import necessary packages and modules
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
from torchvision import transforms
from torch.nn.utils.rnn import pack_padded_sequence
from torch.autograd import Variable
import numpy as np
import matplotlib.pyplot as plt
import os
```

3 Image Captioning Using Encoder-Decoder Architecture

Simply, the encoder will take the image as input and encode it into a vector of feature values. The decoder will take this output from encoder as hidden state and starts to predict next words at each step. The following figure illustrates this:

Figure 1. An overview of the encoder-decoder architecture (image credit: Deep Neural Network Based Image Captioning)

You will use a pre-trained CNN as the encoder and Vanilla RNN/LSTM as decoder to predict the captions.

4 Section 1.1 Data

Download dataset following the instructions. We are gonna use Flickr30k dataset, which consists of 31783 images and 158,915 captions. However, instead of using the whole dataset and

“results.csv”, we will follow Karpathy’s Flickr30k annotations and use “dataset_flickr30k.json”. Please move “dataset_flickr30k.json” to “flickr30k_images” folder.

4.1 How to download the data (ARC)

Since the original Flickr30K dataset requires application, we will use the dataset from Kaggle.

Step 1: Register a Kaggle account. <https://www.kaggle.com/>

Step 2: Log into ARC server in the terminal, e.g. Huckleberry.

Step 3: Install required packages.

```
- if you're gonna use powerai on Huckleberry
  # step 1: request for GPU nodes
  salloc --partition=normal_q --nodes=1 --tasks-per-node=10 --gres=gpu:1 bash
  # or if you don't want a GPU, it will be faster to get a job
  salloc --partition=normal_q --nodes=1 --tasks-per-node=10 bash
  # step 2: load all necessary modules
  module load gcc cuda Anaconda3 jdk
  # step 3: activate the virtual environment
  source activate powerai16_ibm
  # step 4: for new packages(take tqdm for example)
  pip install --user kaggle nltk # on hulogin1/hulogin2

- if you're gonna use your own conda environment, simply type
  pip install kaggle nltk
```

Step 4: Make sure you have kaggle. Type kaggle.

Step 5: Download your kaggle.json file from https://www.kaggle.com/Your_Username/account. In API section, click Create New API Token. Then move you kaggle.json file to the path /home/your_name_space/.kaggle/kaggle.json.

Step 6: Download the dataset. Type kaggle datasets download hsankesara/flickr-image-dataset.

Step 7: Unzip the dataset. Type unzip flickr-image-dataset.zip
-x "flickr30k_images/flickr30k_images/flickr30k_images/*.jpg" -d
"/path-to-Assignment_4/Assignment_4/"

Step 8: You should have your dataset in /path-to-Assignment_4/Assignment_4/flickr30k_images/

Step 9: Move “dataset_flickr30k.json” to “flickr30k_images” folder.

Note that you might want to use `nltk.download('punkt')` and `torchvision.models.resnet50(pretrained=True)` before you enter the GPU node.

4.2 How to download the data (Google Colab)

Step 1: Register a Kaggle account. <https://www.kaggle.com/>

Step 2: Download your kaggle.json file from https://www.kaggle.com/Your_Username/account. In API section, click Create New API Token.

Step 3: As we did before, upload all files on Google Drive and open Google Colab.

Step 4: Install required packages.

```
! pip install -q kaggle nltk
```

Step 5: Insert a cell.

```
from google.colab import files
files.upload()
```

Upload `kaggle.json` you just downloaded.

Step 6: Move kaggle.json to the right place,

```
! mkdir ~/.kaggle
! cp kaggle.json ~/.kaggle/
```

Step 7: Change the permission.

```
! chmod 600 ~/.kaggle/kaggle.json
```

Step 8: Download.

```
!kaggle datasets download hsankeasara/flickr-image-dataset
```

Step 9: Move it to your drive and unzip it.

```
unzip flickr-image-dataset.zip -x "flickr30k_images/flickr30k_images/flickr30k_images/*.jpg" -d
```

Step 10: Move "dataset_flickr30k.json" to "flickr30k_images" folder.

```
[ ]: ! pip install -q kaggle nltk
from google.colab import files
files.upload()

! mkdir ~/.kaggle
! cp kaggle.json ~/.kaggle/
! chmod 600 ~/.kaggle/kaggle.json
!kaggle datasets download hsankeasara/flickr-image-dataset
```

<IPython.core.display.HTML object>

Saving kaggle.json to kaggle.json

Downloading flickr-image-dataset.zip to /content

100% 8.16G/8.16G [03:08<00:00, 33.7MB/s]

100% 8.16G/8.16G [03:08<00:00, 46.5MB/s]

4.2.1 Colab Setup:

- Below are some basic steps for colab setup.
- Make changes based on requirements.
- Comment out in case of ARC or your local device with powerful GPU.

Note: For Google Colab give proper paths in this notebook and in dataloader.py if required.

```
[3]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: ! unzip flickr-image-dataset.zip -x "flickr30k_images/flickr30k_images/
      ↳flickr30k_images/*.jpg" -d "/content/drive/My Drive/DL_Fall_2020/Assignment_4/"
```

```
[3]: import sys
      # modify "customized_path_to_homework", path of folder in drive, where you
      ↳uploaded your homework
      path_to_homework = "/content/drive/My Drive/DL_Fall_2020/Assignment_4/"
      sys.path.append(path_to_homework)
```

```
[4]: from dataloader import Flickr30k, get_loader
import numpy as np
import os
import torch.nn as nn
import torch.nn.functional as F
from torchvision import transforms
import torch
import torchvision
import matplotlib.pyplot as plt
```

5 Section 1.2 Take a look at the data

```
[5]: # visualize images and captions
flickr = Flickr30k(split='val', root=path_to_homework+'flickr30k_images/') #
      ↳load validation set as an example
flickr()
```

```
-----flickr30k-----
image root: /content/drive/My
Drive/DL_Fall_2020/Assignment_4/flickr30k_images/flickr30k_images
dataset split: val
the length of the dataset: 1014
```

```
[ ]: # show a random image and its captions
img_id = np.random.randint(len(flickr))
img = flickr.get_img(img_id)
captions = flickr.get_captions(img_id)

plt.figure()
plt.imshow(img)
plt.axis('off')
```

```
print(captions)
```

```
['The boy in the gray shirt is holding the boy in the red shirt on his back.',  
'Two boys are back to back as one holds up the other.', 'One child lifts another  
on his back, inside a room.', 'One boy hoists another boy up on his back.', 'Boy  
pulls other boy over back.']
```



```
[6]: del flickr
```

6 Section 1.3 Build vocabulary

We need to build a vocabulary for our dataset. The vocabulary stores all the words and their indices. We will use it to embed and recover the words.

```
[7]: import nltk  
import pickle  
import json  
from tqdm import tqdm  
from collections import Counter  
nltk.download('punkt') # You can comment this line once you've downloaded 'punkt'  
  
class Vocabulary(object):  
    """Simple vocabulary wrapper."""  
    def __init__(self):
```

```

        self.word2idx = {'<pad>': 0, '<unk>': 1, '<start>': 2, '<end>': 3} #
→follow Pytorch padding rules: pad sentence with zero.
        self.idx = 4
        self.idx2word = {v: k for k, v in self.word2idx.items()}

    def __call__(self, key):
        if key not in self.word2idx:
            return self.word2idx['<unk>']
        return self.word2idx[key]

    def __len__(self):
        return len(self.word2idx)

    def add_word(self, word):
        """
        Add new words
        :param word: word
        """
        if word not in self.word2idx:
            self.word2idx[word] = self.idx # add a new word
            self.idx2word[self.idx] = word
            self.idx += 1

    def reverse(self, value):
        """
        From idx to words.
        :param value: index
        :return:
        """
        if value not in self.idx2word:
            return self.idx2word[1] # return '<unk>' if the word is unseen
→before.
        return self.idx2word[value]

def build_vocab(json_file=path_to_homework+ '/flickr30k_images/dataset_flickr30k.
→json', threshold=3):
    with open(json_file) as f:
        data = json.load(f)
    f.close()
    counter = Counter()
    for img_idx in tqdm(range(len(data['images']))):
        img_annos = data['images'][img_idx]
        for sent_idx in range(len(img_annos['sentids'])):
            #         tokens = img_annos['sentences'][sent_idx]['tokens'] # directly
→load tokens

            caption = img_annos['sentences'][sent_idx]['raw']

```

```

tokens = nltk.tokenize.word_tokenize(caption.lower())

counter.update(tokens)

# If the number of words is less than threshold we don't count it.
words = [word for word, cnt in counter.items() if cnt >= threshold]

# create a Vocabulary class
vocab = Vocabulary()

# add words to Vocab
for i, word in enumerate(words):
    vocab.add_word(word)

return vocab

```

[nltk_data] Downloading package punkt to /root/nltk_data...

[nltk_data] Unzipping tokenizers/punkt.zip.

```

[8]: # let's create a vocabulary for future usage
vocab_path = path_to_homework + '/flickr30k_images/vocab.pkl'
if not os.path.isfile(vocab_path): # if we don't have vocab, create one
    vocab = build_vocab(json_file=path_to_homework + '/flickr30k_images/
→dataset_flickr30k.json', threshold=3)
    with open(vocab_path, 'wb') as f:
        pickle.dump(vocab, f)
    print("Total vocabulary size: {}".format(len(vocab)))
    print("Saved the vocabulary wrapper to '{}'".format(vocab_path))
else: # if we have, load the existing vocab
    with open(vocab_path, 'rb') as f:
        vocab = pickle.load(f)
    print('vocab loaded!')
    print('the size of vocab:', len(vocab))
f.close()

```

vocab loaded!

the size of vocab: 9991

```

[9]: vocab_path = path_to_homework + '/flickr30k_images/vocab.pkl'
with open(vocab_path, 'rb') as f:
    vocab = pickle.load(f)
print('vocab loaded!')
print('the size of vocab:', len(vocab))
# print(vocab.word2idx.keys())
# print(vocab.idx2word)

```



```
# check some random words
for i in range(3):
    random_idx = np.random.randint(len(vocab))
    print('word: {}, index: {}'.format(list(vocab.word2idx.keys())[random_idx],
    ↪vocab(list(vocab.word2idx.keys())[random_idx])))
```

```
vocab loaded!
the size of vocab: 9991
word: enrolling, index: 8889
word: fighters, index: 3907
word: altima, index: 9866
```

7 Section 2 Vanilla RNN [45 pts]

8 Section 2.1 Design the Network: Encoder [5 pts]

Implement the baseline model by using pre-trained ResNet-50 as the encoder and Vanilla RNN as the decoder. Note that we will remove the last layer (fc layer) of ResNet-50 and add a trainable linear layer to finetune it for our task. During the training, we will **freeze** the layer before the fc layer. The encoder should output a feature vector of a fixed size for each image.

```
[10]: class Encoder(nn.Module):
    def __init__(self, emb_dim):
        """
        Use ResNet-50 as encoder.
        :param emb_dim: output size of ResNet-50.
        """
        super(Encoder, self).__init__()
        self.resnet = torchvision.models.resnet50(pretrained=True)
        #####Your code#####
        # freeze the parameters
        for param in self.resnet.parameters():
            param.requires_grad = False
        # replace the last layer (fc layer) with a trainable layer for finetuning
        self.resnet.fc = nn.Linear(resnet.fc.in_features, emb_dim)

    def forward(self, x):
        x = self.resnet(x) # output shape: [N, emb_dim]
        return x
```

9 Section 2.2 Design the Network: Decoder [10 pts]

During decoding, we will train a RNN (<https://pytorch.org/docs/stable/generated/torch.nn.RNN.html#torch.nn.RNN>) to learn the structure of the caption text through “**Teacher Forcing**”. Teacher forcing works by using the teaching signal from the training dataset at the current time step, $target(t)$, as input in the next time step $x(t+1) = target(t)$, rather than the output $y(t)$ generated by the network.

As shown in Figure 1 above, RNN will take three inputs: the *current feature*, hidden state (h_0) and cell state (c_0). The *current feature* for the first step should be the output of encoder to predict '<start>' word. Hidden states for this step should be set to None. Then in the second step '<start>' will be passed into RNN as the input, and so on.

To use '<start>' or any subsequent word as current feature, get its index from the vocabulary you created, convert it to one-hot vector and pass it through a linear layer to embed into a feature (or you can take advantage of Pytorch's `nn.Embedding` which does one-hot encoding + linear layer for you).

For convenience, you might want to 'pad' the captions in a mini-batch to convert them into fixed length. You can use 'pack_padded_sequence' function.

```
[11]: class Decoder(nn.Module):
    def __init__(self, vocab_size, emb_dim, hidden_dim, num_layers=1, dropout=0):
        """
        Use RNN as decoder for captions.
        :param emb_dim: Embedding dimensions.
        :param hidden_dim: Hidden states dimensions.
        :param num_layers: Number of RNN layers.
        :param vocab_size: The size of Vocabulary.
        :param dropout: the probability for dropout.
        """
        super(Decoder, self).__init__()
        self.max_length = 30 # the maximum length of a sentence, in case it's
→trapped
        self.hidden_dim = hidden_dim
        self.vocab_size = vocab_size

        self.temp = 1
        #####Your code#####
        # you need to implement a Vanilla RNN for the decoder. Take a look at
→the official documentation.
        # https://pytorch.org/docs/stable/generated/torch.nn.RNN.html#torch.nn.
→RNN

        # one-hot encoding + linear layer
        self.embed = nn.Embedding(self.vocab_size, emb_dim)
        # vanilla rnn network
        self.rnn = nn.RNN(emb_dim, hidden_dim, num_layers)
        # output layer
        self.out = nn.Linear(hidden_dim, vocab_size)

    def forward(self, encode_features, captions, lengths):
        """
        Feed forward to generate captions. Note that you need to pad the input
→so they have the same length
        :param encode_features: output of encoder, size [N, emb_dim]
```

```

        :param captions: captions, size [N, max(lengths)]
        :param lengths: a list indicating valid length for each caption. size is
        → (batch_size).
        """
        #####Your Code#####
        # compute the embedding using one-hot technique and linear function
        caption_embedded = self.embed(captions)

        # concatenate the encoded features from encoder and embeddings
        encode_features = torch.cat((encode_features.unsqueeze(1),
        →caption_embedded), dim = 1)

        # feed into RNN.
        encode_features_pack = pack_padded_sequence(encode_features, lengths,
        →batch_first=True)
        output, hidden = self.rnn(encode_features_pack)

        # output layer
        outputs = self.out(output[0])

        return outputs

```

10 Encoder-decoder [10 pts]

Now we need to put our encoder and decoder together.

In the sample_generate stage, the idea is to “let the network run on its own”, predicting the next word, and then use the network’s prediction to obtain the next input word. There are at least two ways to obtain the next word.

- **Deterministic:** Take the maximum output at each step.
- **Stochastic:** Sample from the probability distribution. To get the distribution, we need to compute the weighted softmax of the outputs: $y^i = \exp(o^i/\tau) / \sum_n \exp(o^n/\tau)$, where o^i is the output from the last layer, n is the size of the vocabulary, and τ is the so-called “temperature”. By doing this, you should get a different caption each time.

```

[12]: class Vanilla_rnn(nn.Module):
        def __init__(self, vocab_size, emb_dim, hidden_dim, num_layers=1, dropout=0):
            """
            Encoder-decoder vanilla RNN.
            :param vocab_size: the size of Vocabulary.
            :param emb_dim: the dimensions of word embedding.
            :param hidden_dim: the dimensions of hidden units.
            :param num_layers: the number of RNN layers.
            :param dropout: dropout probability
            """
            super(Vanilla_rnn, self).__init__()

```

```

#####Your Code#####
# Encoder: ResNet-50
self.Vanilla_encoder = Encoder(emb_dim)
# Decoder: RNN
self.Vanilla_decoder = Decoder(vocab_size, emb_dim, hidden_dim,
→num_layers, dropout)
self.max_length = self.Vanilla_decoder.max_length
self.temp = 1
self.softmax = nn.Softmax(dim=1)

def forward(self, x, captions, lengths):
    """
    Feed forward.
    :param x: Images, [N, 3, H, W]
    :param captions: encoded captions, [N, max(lengths)]
    :param lengths: a list indicating valid length for each caption. length_
→is (batch_size).
    :return: output logits, usually followed by a softmax layer.
    """
    #####Your code#####
    # forward passing
    encoder_features = self.Vanilla_encoder(x)
    x = self.Vanilla_decoder(encoder_features, captions, lengths)

    return x

def sample_generate(self, x, states=None, mode='Deterministic',
→temperature=5.0):
    """
    Generate samples during the evaluation.

    :param x: input image
    :param states: rnn states
    :param mode: which mode we use.
        - 'Deterministic': Take the maximum output at each step.
        - 'Stochastic': Sample from the probability distribution from the
→output layer.
    :param temperature: will be used in the stochastic mode
    :return: sample_idxes. Word indices. We can use vocab to recover the
→sentence later.
    """
    sample_idxes = [] # record the index of your generated words

    #####Your Code#####
    # compute the encoded features
    in_feature = self.Vanilla_encoder.forward(x)

```

```

in_feature = in_feature.unsqueeze(1)
batch_size = in_feature.shape[0]

for i in range(self.max_length):
    outputs, states = self.Vanilla_decoder.rnn(in_feature, states)
    outputs = outputs.squeeze(1)
    outputs = self.Vanilla_decoder.out(outputs)
    outputs = self.softmax(outputs/temperature)

    # decide which mode we use
    if mode == 'Deterministic':
        # take the maximum index after the softmax
        max_val, predicted = outputs.max(1)
        sample_idxes.append(predicted)

    elif mode == 'Stochastic':
        # sample from the probability distribution after the softmax
        # Hint: use torch.multinomial() to sample from a distribution.
        predicted = torch.multinomial(outputs, 1)
        sample_idxes.append(predicted)

    x = self.Vanilla_decoder.embed(predicted)
    x = x.unsqueeze(1)
    sample_idxes = torch.stack(sample_idxes, 1)
    sample_idxes = sample_idxes.squeeze()

return sample_idxes

```

11 Section 2.3 Training [10 pts]

Train your encoder-decoder. You might also want to check the output sentence every epoch.

```

[13]: # some hyperparameters, you can change them
      ## training parameters
      batch_size = 256
      lr = 1e-2
      num_epochs = 50
      weight_decay = 0.0
      log_step = 50

      ## network architecture
      emb_dim = 1024
      hidden_dim = 256
      num_layers = 1 # number of RNN layers
      dropout = 0.0

```

```

## image transformation
transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    # transforms.RandomCrop(224, pad_if_needed=True),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.485, 0.456, 0.406),
                          std=(0.229, 0.224, 0.225))])

## Output directory
output_dir = path_to_homework + '/checkpoints/rnn/'
os.makedirs(output_dir, exist_ok=True)

## device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

```

```

[14]: # Validation code here. We are gonna use this during the training.
def val(model, data_loader, vocab):
    """
    Inputs:
    :param model: the encoder-decoder network.
    :param data_loader: validation data loader
    :param vocab: pre-built vocabulary
    Output:
    the mean value of validation losses
    """
    print('Validating...')
    val_loss = []
    total_step = len(data_loader)
    criterion = nn.CrossEntropyLoss()
    for itr, (images, captions, lengths) in enumerate(data_loader):
        #####Your Code#####
        # forward inputs and compute the validation loss
        captions = captions.to(device)
        outputs = model(images.to(device), captions, lengths)
        outputs_pack = pack_padded_sequence(captions, lengths,
        ↪batch_first=True)[0]
        outputs = outputs.view(-1, len(vocab))
        loss = criterion(outputs, outputs_pack)

        # record the validation loss
        val_loss.append(float(loss))

    # Print current loss
    if itr % log_step == 0:
        print('Step [{}/{}], Loss: {:.4f}, Perplexity: {:.54f}'

```

```

        .format(itr, total_step, loss.item(), np.exp(loss.item()))
    val_loss = np.array(val_loss)

    # (optional) you might also want to print out the sentence to see the
    → qualitative performance of your model.
    # You can use deterministic mode to generate sentences

    return np.mean(val_loss)

```

```

[ ]: # Training code here

train_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    → split='train', vocab=vocab,
                                transform=transform, batch_size=batch_size,
    → shuffle=True, num_workers=4)
val_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    → split='val', vocab=vocab,
                                transform=transform, batch_size=8, shuffle=True,
    → num_workers=4)

model = Vanilla_rnn(vocab_size=len(vocab), emb_dim=emb_dim,
    → hidden_dim=hidden_dim, num_layers=1, dropout=dropout).to(device) # build a
    → model

vocab_size=len(vocab)

# loss and optimizer
criterion = nn.CrossEntropyLoss() # CE loss
optimizer = torch.optim.Adam(model.parameters(), lr=lr,
    → weight_decay=weight_decay) # optimizer
scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                step_size=5,
                                gamma=0.5) # decay LR by a factor of 0.5
    → every 10 epochs. You can change this

# logs
Train_Losses = [] # record average training loss each epoch
Val_Losses = [] # record average validation loss each epoch
total_step = len(train_data_loader) # number of iterations each epoch
best_val_loss = np.inf

# start training
print('Start training...')
import time
tic = time.time()

```

```

for epoch in range(num_epochs):
# for epoch in range(3):
    print('Switch to training...')
    model.train()
    Train_loss_iter = [] # record the the training loss each iteration
    for itr, (images, captions, lengths) in enumerate(train_data_loader):
        #####Your Code#####
        # train your model
        model.zero_grad()
        images = Variable(images).to(device)
        captions = Variable(captions).to(device)
        predicted_cap_pack = pack_padded_sequence(captions, lengths,
→batch_first=True)[0]
        predicted_cap = model(images, captions, lengths)
        predicted_cap = predicted_cap.view(-1, vocab_size)
        loss = criterion(input = predicted_cap, target=predicted_cap_pack)
        loss.backward()
        optimizer.step()

        # record the training loss
        Train_Losses.append(float(loss))

        # print log info
        if itr % log_step == 0:
            # print current loss and perplexity
            print('Epoch [{} / {}], Step [{} / {}], Loss: {:.4f}, Perplexity: {:.5.
→4f}')
            .format(epoch, num_epochs, itr, total_step, loss.item(),
→np.exp(loss.item()))
            scheduler.step()
            Train_Losses.append(np.mean(Train_loss_iter))
            np.save(os.path.join(output_dir, 'TrainingLoss_rnn.npy'), Train_Losses) #
→save the training loss

        model.eval()
        # (optional) generate a sample during the training, you can use
→deterministic mode
        # Your code

        # validation
        Val_Losses.append(val(model, val_data_loader, vocab))
        np.save(os.path.join(output_dir, 'ValLoss_rnn.npy'), Val_Losses) # save the
→val loss

        # save model

```



```

    if Val_Losses[-1] < best_val_loss:
        best_val_loss = Val_Losses[-1]
        print('updated best val loss:', best_val_loss)
        print('Save model weights to...', output_dir)
        torch.save(model.state_dict(),
                    os.path.join(output_dir, 'vanilla_rnn-best.pth'.format(epoch_
→+ 1, itr + 1)))

print('It took: {} s'.format(time.time() - tic))

```

Downloading: "<https://download.pytorch.org/models/resnet50-19c8e357.pth>" to
/root/.cache/torch/hub/checkpoints/resnet50-19c8e357.pth

HBox(children=(FloatProgress(value=0.0, max=102502400.0), HTML(value='')))

Start training...

Switch to training...

Epoch [0/50], Step [0/114], Loss: 9.2990, Perplexity: 10927.0114

Epoch [0/50], Step [50/114], Loss: 3.8386, Perplexity: 46.4584

Epoch [0/50], Step [100/114], Loss: 3.6042, Perplexity: 36.7507

Validating...

/usr/local/lib/python3.6/dist-packages/numpy/core/fromnumeric.py:3335:

RuntimeWarning: Mean of empty slice.

out=out, **kwargs)

/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:161:

RuntimeWarning: invalid value encountered in double_scalars

ret = ret.dtype.type(ret / rcount)

Step [0/127], Loss: 4.1656, Perplexity: 64.4300

Step [50/127], Loss: 4.0322, Perplexity: 56.3838

Step [100/127], Loss: 3.2433, Perplexity: 25.6184

updated best val loss: 3.671107530593872

Save model weights to... /content/drive/My

Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/

Switch to training...

Epoch [1/50], Step [0/114], Loss: 3.5300, Perplexity: 34.1254

Epoch [1/50], Step [50/114], Loss: 3.5267, Perplexity: 34.0123

Epoch [1/50], Step [100/114], Loss: 3.5725, Perplexity: 35.6045

Validating...

Step [0/127], Loss: 3.8533, Perplexity: 47.1464

Step [50/127], Loss: 4.3051, Perplexity: 74.0762

Step [100/127], Loss: 3.4891, Perplexity: 32.7554

updated best val loss: 3.601075799446406

Save model weights to... /content/drive/My

Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/

Switch to training...

Epoch [2/50], Step [0/114], Loss: 3.4513, Perplexity: 31.5418
 Epoch [2/50], Step [50/114], Loss: 3.5018, Perplexity: 33.1766
 Epoch [2/50], Step [100/114], Loss: 3.5181, Perplexity: 33.7187
 Validating...
 Step [0/127], Loss: 3.4302, Perplexity: 30.8815
 Step [50/127], Loss: 4.0189, Perplexity: 55.6414
 Step [100/127], Loss: 3.8191, Perplexity: 45.5614
 updated best val loss: 3.5668688045711967
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [3/50], Step [0/114], Loss: 3.4842, Perplexity: 32.5955
 Epoch [3/50], Step [50/114], Loss: 3.5505, Perplexity: 34.8302
 Epoch [3/50], Step [100/114], Loss: 3.4859, Perplexity: 32.6529
 Validating...
 Step [0/127], Loss: 3.2147, Perplexity: 24.8964
 Step [50/127], Loss: 4.0613, Perplexity: 58.0522
 Step [100/127], Loss: 3.8623, Perplexity: 47.5727
 Switch to training...
 Epoch [4/50], Step [0/114], Loss: 3.5081, Perplexity: 33.3854
 Epoch [4/50], Step [50/114], Loss: 3.4329, Perplexity: 30.9649
 Epoch [4/50], Step [100/114], Loss: 3.4119, Perplexity: 30.3242
 Validating...
 Step [0/127], Loss: 3.7977, Perplexity: 44.5978
 Step [50/127], Loss: 2.9505, Perplexity: 19.1151
 Step [100/127], Loss: 3.3779, Perplexity: 29.3080
 updated best val loss: 3.512489335743461
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [5/50], Step [0/114], Loss: 3.4232, Perplexity: 30.6674
 Epoch [5/50], Step [50/114], Loss: 3.2039, Perplexity: 24.6274
 Epoch [5/50], Step [100/114], Loss: 3.2004, Perplexity: 24.5415
 Validating...
 Step [0/127], Loss: 3.5653, Perplexity: 35.3500
 Step [50/127], Loss: 3.3220, Perplexity: 27.7157
 Step [100/127], Loss: 3.7276, Perplexity: 41.5801
 updated best val loss: 3.337581520005474
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [6/50], Step [0/114], Loss: 3.2339, Perplexity: 25.3792
 Epoch [6/50], Step [50/114], Loss: 3.2582, Perplexity: 26.0024
 Epoch [6/50], Step [100/114], Loss: 3.1173, Perplexity: 22.5864
 Validating...
 Step [0/127], Loss: 3.4112, Perplexity: 30.3019
 Step [50/127], Loss: 2.9932, Perplexity: 19.9488
 Step [100/127], Loss: 3.0383, Perplexity: 20.8707

Switch to training...

Epoch [7/50], Step [0/114], Loss: 3.1146, Perplexity: 22.5242

Epoch [7/50], Step [50/114], Loss: 3.0713, Perplexity: 21.5708

Epoch [7/50], Step [100/114], Loss: 3.0943, Perplexity: 22.0722

Validating...

Step [0/127], Loss: 3.2245, Perplexity: 25.1406

Step [50/127], Loss: 3.2410, Perplexity: 25.5605

Step [100/127], Loss: 3.1710, Perplexity: 23.8301

updated best val loss: 3.288263664470883

Save model weights to... /content/drive/My

Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/

Switch to training...

Epoch [8/50], Step [0/114], Loss: 3.1889, Perplexity: 24.2626

Epoch [8/50], Step [50/114], Loss: 3.0698, Perplexity: 21.5368

Epoch [8/50], Step [100/114], Loss: 3.1861, Perplexity: 24.1932

Validating...

Step [0/127], Loss: 3.3716, Perplexity: 29.1255

Step [50/127], Loss: 4.2271, Perplexity: 68.5157

Step [100/127], Loss: 3.4140, Perplexity: 30.3867

Switch to training...

Epoch [9/50], Step [0/114], Loss: 3.0014, Perplexity: 20.1132

Epoch [9/50], Step [50/114], Loss: 3.0874, Perplexity: 21.9203

Epoch [9/50], Step [100/114], Loss: 3.1656, Perplexity: 23.7035

Validating...

Step [0/127], Loss: 3.2300, Perplexity: 25.2794

Step [50/127], Loss: 3.2347, Perplexity: 25.3999

Step [100/127], Loss: 3.4253, Perplexity: 30.7326

Switch to training...

Epoch [10/50], Step [0/114], Loss: 3.1566, Perplexity: 23.4901

Epoch [10/50], Step [50/114], Loss: 3.0670, Perplexity: 21.4782

Epoch [10/50], Step [100/114], Loss: 2.9891, Perplexity: 19.8669

Validating...

Step [0/127], Loss: 3.3540, Perplexity: 28.6171

Step [50/127], Loss: 3.4342, Perplexity: 31.0058

Step [100/127], Loss: 3.4286, Perplexity: 30.8341

updated best val loss: 3.1965027155838612

Save model weights to... /content/drive/My

Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/

Switch to training...

Epoch [11/50], Step [0/114], Loss: 2.9665, Perplexity: 19.4236

Epoch [11/50], Step [50/114], Loss: 3.0302, Perplexity: 20.7013

Epoch [11/50], Step [100/114], Loss: 2.9597, Perplexity: 19.2922

Validating...

Step [0/127], Loss: 2.5568, Perplexity: 12.8939

Step [50/127], Loss: 3.1069, Perplexity: 22.3521

Step [100/127], Loss: 3.0165, Perplexity: 20.4200

Switch to training...

Epoch [12/50], Step [0/114], Loss: 2.9106, Perplexity: 18.3684

Epoch [12/50], Step [50/114], Loss: 2.9672, Perplexity: 19.4369
 Epoch [12/50], Step [100/114], Loss: 2.9683, Perplexity: 19.4594
 Validating...
 Step [0/127], Loss: 2.8816, Perplexity: 17.8419
 Step [50/127], Loss: 2.8213, Perplexity: 16.7983
 Step [100/127], Loss: 3.4687, Perplexity: 32.0965
 updated best val loss: 3.1893812953017826
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [13/50], Step [0/114], Loss: 2.9594, Perplexity: 19.2862
 Epoch [13/50], Step [50/114], Loss: 2.9048, Perplexity: 18.2609
 Epoch [13/50], Step [100/114], Loss: 3.0572, Perplexity: 21.2681
 Validating...
 Step [0/127], Loss: 2.9277, Perplexity: 18.6846
 Step [50/127], Loss: 3.1788, Perplexity: 24.0182
 Step [100/127], Loss: 3.4187, Perplexity: 30.5299
 Switch to training...
 Epoch [14/50], Step [0/114], Loss: 2.8939, Perplexity: 18.0634
 Epoch [14/50], Step [50/114], Loss: 2.9014, Perplexity: 18.1993
 Epoch [14/50], Step [100/114], Loss: 2.9725, Perplexity: 19.5414
 Validating...
 Step [0/127], Loss: 2.9496, Perplexity: 19.0987
 Step [50/127], Loss: 2.9935, Perplexity: 19.9547
 Step [100/127], Loss: 2.5781, Perplexity: 13.1719
 updated best val loss: 3.1814778770987444
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [15/50], Step [0/114], Loss: 2.9933, Perplexity: 19.9514
 Epoch [15/50], Step [50/114], Loss: 2.8711, Perplexity: 17.6560
 Epoch [15/50], Step [100/114], Loss: 2.8944, Perplexity: 18.0730
 Validating...
 Step [0/127], Loss: 3.4155, Perplexity: 30.4334
 Step [50/127], Loss: 3.0547, Perplexity: 21.2146
 Step [100/127], Loss: 3.4499, Perplexity: 31.4958
 updated best val loss: 3.1385314502115325
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [16/50], Step [0/114], Loss: 2.8434, Perplexity: 17.1738
 Epoch [16/50], Step [50/114], Loss: 2.8848, Perplexity: 17.8997
 Epoch [16/50], Step [100/114], Loss: 2.9165, Perplexity: 18.4764
 Validating...
 Step [0/127], Loss: 2.9856, Perplexity: 19.7976
 Step [50/127], Loss: 3.4809, Perplexity: 32.4885
 Step [100/127], Loss: 3.3999, Perplexity: 29.9615
 Switch to training...

Epoch [17/50], Step [0/114], Loss: 2.7847, Perplexity: 16.1955
 Epoch [17/50], Step [50/114], Loss: 2.8731, Perplexity: 17.6912
 Epoch [17/50], Step [100/114], Loss: 2.8951, Perplexity: 18.0853
 Validating...
 Step [0/127], Loss: 3.1474, Perplexity: 23.2749
 Step [50/127], Loss: 3.1053, Perplexity: 22.3153
 Step [100/127], Loss: 3.0071, Perplexity: 20.2291
 updated best val loss: 3.1380890542127955
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [18/50], Step [0/114], Loss: 2.9423, Perplexity: 18.9585
 Epoch [18/50], Step [50/114], Loss: 2.9052, Perplexity: 18.2689
 Epoch [18/50], Step [100/114], Loss: 2.8824, Perplexity: 17.8568
 Validating...
 Step [0/127], Loss: 3.1679, Perplexity: 23.7564
 Step [50/127], Loss: 3.0608, Perplexity: 21.3446
 Step [100/127], Loss: 2.9254, Perplexity: 18.6414
 Switch to training...
 Epoch [19/50], Step [0/114], Loss: 2.7966, Perplexity: 16.3892
 Epoch [19/50], Step [50/114], Loss: 2.8129, Perplexity: 16.6585
 Epoch [19/50], Step [100/114], Loss: 2.8923, Perplexity: 18.0342
 Validating...
 Step [0/127], Loss: 3.5365, Perplexity: 34.3450
 Step [50/127], Loss: 3.5281, Perplexity: 34.0578
 Step [100/127], Loss: 2.8435, Perplexity: 17.1759
 Switch to training...
 Epoch [20/50], Step [0/114], Loss: 2.8256, Perplexity: 16.8713
 Epoch [20/50], Step [50/114], Loss: 2.7618, Perplexity: 15.8289
 Epoch [20/50], Step [100/114], Loss: 2.7990, Perplexity: 16.4283
 Validating...
 Step [0/127], Loss: 3.1563, Perplexity: 23.4844
 Step [50/127], Loss: 3.4632, Perplexity: 31.9180
 Step [100/127], Loss: 2.7674, Perplexity: 15.9170
 updated best val loss: 3.117585886181809
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
 Switch to training...
 Epoch [21/50], Step [0/114], Loss: 2.8209, Perplexity: 16.7913
 Epoch [21/50], Step [50/114], Loss: 2.7216, Perplexity: 15.2044
 Epoch [21/50], Step [100/114], Loss: 2.8026, Perplexity: 16.4871
 Validating...
 Step [0/127], Loss: 3.0410, Perplexity: 20.9268
 Step [50/127], Loss: 3.0794, Perplexity: 21.7462
 Step [100/127], Loss: 3.1394, Perplexity: 23.0901
 updated best val loss: 3.100470419005146
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/

Switch to training...

Epoch [22/50], Step [0/114], Loss: 2.8679, Perplexity: 17.6006
Epoch [22/50], Step [50/114], Loss: 2.7297, Perplexity: 15.3283
Epoch [22/50], Step [100/114], Loss: 2.7226, Perplexity: 15.2199
Validating...

Step [0/127], Loss: 2.7374, Perplexity: 15.4467
Step [50/127], Loss: 3.0840, Perplexity: 21.8460
Step [100/127], Loss: 3.3896, Perplexity: 29.6538
Switch to training...

Epoch [23/50], Step [0/114], Loss: 2.7286, Perplexity: 15.3122
Epoch [23/50], Step [50/114], Loss: 2.7789, Perplexity: 16.1018
Epoch [23/50], Step [100/114], Loss: 2.7830, Perplexity: 16.1681
Validating...

Step [0/127], Loss: 4.0628, Perplexity: 58.1362
Step [50/127], Loss: 3.2350, Perplexity: 25.4059
Step [100/127], Loss: 3.2519, Perplexity: 25.8391
Switch to training...

Epoch [24/50], Step [0/114], Loss: 2.8066, Perplexity: 16.5541
Epoch [24/50], Step [50/114], Loss: 2.6985, Perplexity: 14.8574
Epoch [24/50], Step [100/114], Loss: 2.8212, Perplexity: 16.7977
Validating...

Step [0/127], Loss: 3.9105, Perplexity: 49.9225
Step [50/127], Loss: 3.0598, Perplexity: 21.3238
Step [100/127], Loss: 2.3015, Perplexity: 9.9892
updated best val loss: 3.0731245964530887
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
Switch to training...

Epoch [25/50], Step [0/114], Loss: 2.7803, Perplexity: 16.1242
Epoch [25/50], Step [50/114], Loss: 2.7777, Perplexity: 16.0813
Epoch [25/50], Step [100/114], Loss: 2.7227, Perplexity: 15.2221
Validating...

Step [0/127], Loss: 3.1107, Perplexity: 22.4365
Step [50/127], Loss: 3.2879, Perplexity: 26.7857
Step [100/127], Loss: 3.4780, Perplexity: 32.3938
updated best val loss: 3.0715413844491555
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
Switch to training...

Epoch [26/50], Step [0/114], Loss: 2.8767, Perplexity: 17.7555
Epoch [26/50], Step [50/114], Loss: 2.7009, Perplexity: 14.8925
Epoch [26/50], Step [100/114], Loss: 2.8266, Perplexity: 16.8886
Validating...

Step [0/127], Loss: 3.0918, Perplexity: 22.0173
Step [50/127], Loss: 3.4706, Perplexity: 32.1560
Step [100/127], Loss: 3.0168, Perplexity: 20.4254
Switch to training...

Epoch [27/50], Step [0/114], Loss: 2.7927, Perplexity: 16.3245

Epoch [27/50], Step [50/114], Loss: 2.7265, Perplexity: 15.2792
 Epoch [27/50], Step [100/114], Loss: 2.7699, Perplexity: 15.9568
 Validating...
 Step [0/127], Loss: 3.3283, Perplexity: 27.8922
 Step [50/127], Loss: 3.5669, Perplexity: 35.4055
 Step [100/127], Loss: 2.8196, Perplexity: 16.7697
 Switch to training...
 Epoch [28/50], Step [0/114], Loss: 2.7356, Perplexity: 15.4186
 Epoch [28/50], Step [50/114], Loss: 2.7565, Perplexity: 15.7445
 Epoch [28/50], Step [100/114], Loss: 2.7027, Perplexity: 14.9205
 Validating...
 Step [0/127], Loss: 3.7608, Perplexity: 42.9845
 Step [50/127], Loss: 2.9381, Perplexity: 18.8804
 Step [100/127], Loss: 3.2651, Perplexity: 26.1834
 Switch to training...
 Epoch [29/50], Step [0/114], Loss: 2.7303, Perplexity: 15.3375
 Epoch [29/50], Step [50/114], Loss: 2.7352, Perplexity: 15.4129
 Epoch [29/50], Step [100/114], Loss: 2.6747, Perplexity: 14.5075
 Validating...
 Step [0/127], Loss: 3.0812, Perplexity: 21.7847
 Step [50/127], Loss: 2.6892, Perplexity: 14.7202
 Step [100/127], Loss: 3.1091, Perplexity: 22.4000
 Switch to training...
 Epoch [30/50], Step [0/114], Loss: 2.7688, Perplexity: 15.9401
 Epoch [30/50], Step [50/114], Loss: 2.7326, Perplexity: 15.3724
 Epoch [30/50], Step [100/114], Loss: 2.7484, Perplexity: 15.6174
 Validating...
 Step [0/127], Loss: 2.6037, Perplexity: 13.5142
 Step [50/127], Loss: 2.9568, Perplexity: 19.2357
 Step [100/127], Loss: 2.7856, Perplexity: 16.2100
 Switch to training...
 Epoch [31/50], Step [0/114], Loss: 2.7249, Perplexity: 15.2542
 Epoch [31/50], Step [50/114], Loss: 2.7234, Perplexity: 15.2315
 Epoch [31/50], Step [100/114], Loss: 2.7434, Perplexity: 15.5403
 Validating...
 Step [0/127], Loss: 3.3860, Perplexity: 29.5490
 Step [50/127], Loss: 2.7989, Perplexity: 16.4267
 Step [100/127], Loss: 3.2641, Perplexity: 26.1556
 Switch to training...
 Epoch [32/50], Step [0/114], Loss: 2.7916, Perplexity: 16.3079
 Epoch [32/50], Step [50/114], Loss: 2.7383, Perplexity: 15.4608
 Epoch [32/50], Step [100/114], Loss: 2.7144, Perplexity: 15.0953
 Validating...
 Step [0/127], Loss: 3.4751, Perplexity: 32.2995
 Step [50/127], Loss: 3.3821, Perplexity: 29.4328
 Step [100/127], Loss: 3.2108, Perplexity: 24.7988
 Switch to training...
 Epoch [33/50], Step [0/114], Loss: 2.7368, Perplexity: 15.4375

Epoch [33/50], Step [50/114], Loss: 2.6621, Perplexity: 14.3256
Epoch [33/50], Step [100/114], Loss: 2.7613, Perplexity: 15.8206
Validating...
Step [0/127], Loss: 2.8184, Perplexity: 16.7500
Step [50/127], Loss: 3.2714, Perplexity: 26.3470
Step [100/127], Loss: 3.7577, Perplexity: 42.8496
Switch to training...
Epoch [34/50], Step [0/114], Loss: 2.7287, Perplexity: 15.3122
Epoch [34/50], Step [50/114], Loss: 2.7516, Perplexity: 15.6681
Epoch [34/50], Step [100/114], Loss: 2.6906, Perplexity: 14.7412
Validating...
Step [0/127], Loss: 3.4691, Perplexity: 32.1068
Step [50/127], Loss: 2.9756, Perplexity: 19.6022
Step [100/127], Loss: 3.2287, Perplexity: 25.2458
Switch to training...
Epoch [35/50], Step [0/114], Loss: 2.7312, Perplexity: 15.3511
Epoch [35/50], Step [50/114], Loss: 2.7032, Perplexity: 14.9278
Epoch [35/50], Step [100/114], Loss: 2.7533, Perplexity: 15.6936
Validating...
Step [0/127], Loss: 3.0698, Perplexity: 21.5379
Step [50/127], Loss: 3.3699, Perplexity: 29.0770
Step [100/127], Loss: 3.4739, Perplexity: 32.2633
Switch to training...
Epoch [36/50], Step [0/114], Loss: 2.6774, Perplexity: 14.5477
Epoch [36/50], Step [50/114], Loss: 2.6717, Perplexity: 14.4649
Epoch [36/50], Step [100/114], Loss: 2.6921, Perplexity: 14.7624
Validating...
Step [0/127], Loss: 2.6614, Perplexity: 14.3159
Step [50/127], Loss: 3.5899, Perplexity: 36.2298
Step [100/127], Loss: 2.7627, Perplexity: 15.8429
Switch to training...
Epoch [37/50], Step [0/114], Loss: 2.7478, Perplexity: 15.6083
Epoch [37/50], Step [50/114], Loss: 2.7510, Perplexity: 15.6587
Epoch [37/50], Step [100/114], Loss: 2.6760, Perplexity: 14.5275
Validating...
Step [0/127], Loss: 3.4586, Perplexity: 31.7735
Step [50/127], Loss: 3.2749, Perplexity: 26.4411
Step [100/127], Loss: 3.8873, Perplexity: 48.7769
Switch to training...
Epoch [38/50], Step [0/114], Loss: 2.7379, Perplexity: 15.4548
Epoch [38/50], Step [50/114], Loss: 2.7183, Perplexity: 15.1547
Epoch [38/50], Step [100/114], Loss: 2.7442, Perplexity: 15.5518
Validating...
Step [0/127], Loss: 2.4631, Perplexity: 11.7407
Step [50/127], Loss: 3.1684, Perplexity: 23.7691
Step [100/127], Loss: 3.8304, Perplexity: 46.0823
updated best val loss: 3.0556812286376953
Save model weights to... /content/drive/My


```

Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
Switch to training...
Epoch [39/50], Step [0/114], Loss: 2.7851, Perplexity: 16.2021
Epoch [39/50], Step [50/114], Loss: 2.7309, Perplexity: 15.3461
Epoch [39/50], Step [100/114], Loss: 2.7198, Perplexity: 15.1769
Validating...
Step [0/127], Loss: 3.6130, Perplexity: 37.0757
Step [50/127], Loss: 2.7318, Perplexity: 15.3606
Step [100/127], Loss: 2.8113, Perplexity: 16.6311
updated best val loss: 3.0484642794751746
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/rnn/
Switch to training...
Epoch [40/50], Step [0/114], Loss: 2.7373, Perplexity: 15.4445
Epoch [40/50], Step [50/114], Loss: 2.7213, Perplexity: 15.1999
Epoch [40/50], Step [100/114], Loss: 2.7665, Perplexity: 15.9027
Validating...
Step [0/127], Loss: 3.4611, Perplexity: 31.8508
Step [50/127], Loss: 3.0259, Perplexity: 20.6133
Step [100/127], Loss: 3.0787, Perplexity: 21.7303
Switch to training...
Epoch [41/50], Step [0/114], Loss: 2.7094, Perplexity: 15.0206
Epoch [41/50], Step [50/114], Loss: 2.7782, Perplexity: 16.0903
Epoch [41/50], Step [100/114], Loss: 2.6633, Perplexity: 14.3431
Validating...
Step [0/127], Loss: 3.6181, Perplexity: 37.2682
Step [50/127], Loss: 3.0113, Perplexity: 20.3140
Step [100/127], Loss: 3.0394, Perplexity: 20.8922
Switch to training...
Epoch [42/50], Step [0/114], Loss: 2.7026, Perplexity: 14.9185
Epoch [42/50], Step [50/114], Loss: 2.7593, Perplexity: 15.7886
Epoch [42/50], Step [100/114], Loss: 2.7019, Perplexity: 14.9076
Validating...
Step [0/127], Loss: 2.7947, Perplexity: 16.3580
Step [50/127], Loss: 4.3382, Perplexity: 76.5699
Step [100/127], Loss: 3.0872, Perplexity: 21.9147
Switch to training...
Epoch [43/50], Step [0/114], Loss: 2.7254, Perplexity: 15.2621
Epoch [43/50], Step [50/114], Loss: 2.6250, Perplexity: 13.8040
Epoch [43/50], Step [100/114], Loss: 2.7765, Perplexity: 16.0625
Validating...
Step [0/127], Loss: 3.5173, Perplexity: 33.6938
Step [50/127], Loss: 3.0043, Perplexity: 20.1721
Step [100/127], Loss: 3.4019, Perplexity: 30.0217
Switch to training...
Epoch [44/50], Step [0/114], Loss: 2.6477, Perplexity: 14.1215
Epoch [44/50], Step [50/114], Loss: 2.6653, Perplexity: 14.3728
Epoch [44/50], Step [100/114], Loss: 2.7234, Perplexity: 15.2314

```

```

Validating...
Step [0/127], Loss: 3.2623, Perplexity: 26.1089
Step [50/127], Loss: 3.2486, Perplexity: 25.7533
Step [100/127], Loss: 2.5923, Perplexity: 13.3606
Switch to training...
Epoch [45/50], Step [0/114], Loss: 2.7444, Perplexity: 15.5547
Epoch [45/50], Step [50/114], Loss: 2.7405, Perplexity: 15.4950
Epoch [45/50], Step [100/114], Loss: 2.7032, Perplexity: 14.9269
Validating...
Step [0/127], Loss: 3.0232, Perplexity: 20.5563
Step [50/127], Loss: 3.2143, Perplexity: 24.8846
Step [100/127], Loss: 2.7735, Perplexity: 16.0152
Switch to training...
Epoch [46/50], Step [0/114], Loss: 2.7590, Perplexity: 15.7844
Epoch [46/50], Step [50/114], Loss: 2.6928, Perplexity: 14.7725
Epoch [46/50], Step [100/114], Loss: 2.6443, Perplexity: 14.0729
Validating...
Step [0/127], Loss: 3.7961, Perplexity: 44.5283
Step [50/127], Loss: 3.0172, Perplexity: 20.4331
Step [100/127], Loss: 2.9283, Perplexity: 18.6951
Switch to training...
Epoch [47/50], Step [0/114], Loss: 2.7161, Perplexity: 15.1213
Epoch [47/50], Step [50/114], Loss: 2.6394, Perplexity: 14.0049
Epoch [47/50], Step [100/114], Loss: 2.6928, Perplexity: 14.7724
Validating...
Step [0/127], Loss: 3.5637, Perplexity: 35.2938
Step [50/127], Loss: 3.0865, Perplexity: 21.8998
Step [100/127], Loss: 3.6792, Perplexity: 39.6135
Switch to training...
Epoch [48/50], Step [0/114], Loss: 2.7707, Perplexity: 15.9698
Epoch [48/50], Step [50/114], Loss: 2.6973, Perplexity: 14.8396
Epoch [48/50], Step [100/114], Loss: 2.7517, Perplexity: 15.6699
Validating...
Step [0/127], Loss: 3.6207, Perplexity: 37.3621
Step [50/127], Loss: 3.2497, Perplexity: 25.7830
Step [100/127], Loss: 2.8976, Perplexity: 18.1298
Switch to training...
Epoch [49/50], Step [0/114], Loss: 2.6939, Perplexity: 14.7899
Epoch [49/50], Step [50/114], Loss: 2.7550, Perplexity: 15.7213
Epoch [49/50], Step [100/114], Loss: 2.7109, Perplexity: 15.0423
Validating...
Step [0/127], Loss: 3.3984, Perplexity: 29.9159
Step [50/127], Loss: 3.0992, Perplexity: 22.1798
Step [100/127], Loss: 3.0750, Perplexity: 21.6491
It took: 13557.006758928299 s

```

12 Section 2.4 Evaluation [10 pts]

```
[15]: ## evaluation code
from tqdm import tqdm, tqdm_notebook
from nltk.translate.bleu_score import sentence_bleu
from nltk.translate.bleu_score import SmoothingFunction
smoother = SmoothingFunction()

def caption_generator(model, images, vocab, img_ids, captions,
    mode='Deterministic', temperature=1.0):
    """
    Generate captions.
    :param mode:
    :return:
    """
    sample_idxes = model.sample_generate(images, mode=mode,
                                         temperature=temperature).data.cpu().
    numpy() # [N, max_length]
    for i, sentence in enumerate(sample_idxes): # every sentence in this batch
        # for sentence in sample_idxes:
        sentence_caption = ''
        for word_idx in sentence:
            word = vocab.idx2word[word_idx]
            if word != '<start>' and word != '<end>':
                if word == '.':
                    sentence_caption += '.'
                else:
                    sentence_caption += word + ' '
            if word == '<end>':
                break
        captions.append({'caption': sentence_caption})
        # captions.append(sentence_caption)

    return captions

def run_test(model, data_loader, vocab, mode='Deterministic', temperature=1.0):
    """
    Run your model on the test set.
    Inputs:
    :param model: the model you use
    :param data_loader: the data_loader
    :param mode: use 'deterministic' or 'stochastic'
    Outputs:
    :param predictions
    """
    predictions = []
    for itr, (images, captions, lengths) in enumerate(tqdm(data_loader)):
```

```

        images = Variable(images).to(device)
        captions = Variable(captions).to(device)
        outputs = model(images, captions, lengths)

        img_ids = list(range(itr * data_loader.batch_size, (itr + 1) *
→data_loader.batch_size))
        predictions = caption_generator(model, images, vocab, img_ids,
                                       predictions, mode=mode,
→temperature=temperature)

    return predictions

def evaluation(model, vocab, data_path=path_to_homework + '/flickr30k_images/',
→mode='Deterministic', temperature=1.0,
        split='test'):
    """
    Evaluate the performance of your model on the test set using BLEU scores.
    Inputs:
    :param model: the model you use
    :param weight_path: the directory to the weights of your model
    :param vocab: vocabulary
    :param data_path: the directory to the dataset
    :param mode: use 'deterministic' or 'stochastic'
    Outputs:
    :param predictions
    """
    # data loader
    test_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
→split=split, vocab=vocab,
                                transform=transform, batch_size=8,
→shuffle=False, num_workers=4)

    # run your model on the test set
    print('Run on the test set...')
    preds = run_test(model, test_data_loader, vocab, mode, temperature)

    # load the groundtruth
    gt = test_data_loader.dataset.annos

    # evaluate the performance using BLEU score
    score1 = 0
    score2 = 0
    score3 = 0
    score4 = 0

    print('Computing BLEU')
    for itr in tqdm(range(len(gt))):

```

```

        candidate = preds[itr]['caption']
        reference = [sent['raw'] for sent in gt[itr]['sentences']]
        score1 += sentence_bleu(reference, candidate, weights=(1, 0, 0, 0),
→smoothing_function=smoother.method1)
        score2 += sentence_bleu(reference, candidate, weights=(0, 1, 0, 0),
→smoothing_function=smoother.method1)
        score3 += sentence_bleu(reference, candidate, weights=(0, 0, 1, 0),
→smoothing_function=smoother.method1)
        score4 += sentence_bleu(reference, candidate, weights=(0, 0, 0, 1),
→smoothing_function=smoother.method1)

    bleu1 = 100 * score1/len(gt)
    bleu2 = 100 * score2/len(gt)
    bleu3 = 100 * score3/len(gt)
    bleu4 = 100 * score4/len(gt)

    return bleu1, bleu2, bleu3, bleu4

```

- Test your outputs in the **Deterministic** way by using BLEU scores. You should at achieve a BLEU 4 of 25.

```

[17]: ## Evaluate your model using BLEU score. Use Deterministic mode.

## Image transformation
transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    # transforms.RandomCrop(224, pad_if_needed=True),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.485, 0.456, 0.406),
                          std=(0.229, 0.224, 0.225))])

## Evaluate your model using BLEU score. Use Deterministic mode
model = Vanilla_rnn(vocab_size=len(vocab), emb_dim=emb_dim,
→hidden_dim=hidden_dim,
                    num_layers=1, dropout=dropout).to(device) # build a model
model.load_state_dict(torch.load(path_to_homework + '/checkpoints/rnn/
→vanilla_rnn-best.pth', map_location=torch.device('cpu'))))
model.eval()
bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic')
print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}".format(bleu1, bleu2, bleu3,
→bleu4))

```

```
0%|          | 0/125 [00:00<?, ?it/s]
```

Run on the test set...

1%	1/125 [02:37<5:25:21, 157.43s/it]
2%	2/125 [02:37<3:45:58, 110.23s/it]
3%	4/125 [02:37<2:35:39, 77.18s/it]
6%	7/125 [02:37<1:46:17, 54.04s/it]
7%	9/125 [02:39<1:13:35, 38.06s/it]
9%	11/125 [02:39<50:41, 26.68s/it]
10%	12/125 [02:40<35:51, 19.04s/it]
10%	13/125 [02:41<25:04, 13.43s/it]
11%	14/125 [02:41<17:28, 9.44s/it]
12%	15/125 [02:41<12:11, 6.65s/it]
13%	16/125 [02:42<09:08, 5.03s/it]
14%	17/125 [02:42<06:25, 3.57s/it]
14%	18/125 [02:43<04:46, 2.68s/it]
16%	20/125 [02:44<03:36, 2.07s/it]
17%	21/125 [02:44<02:34, 1.48s/it]
18%	22/125 [02:45<01:57, 1.15s/it]
18%	23/125 [02:45<01:32, 1.10it/s]
19%	24/125 [02:46<01:37, 1.04it/s]
21%	26/125 [02:46<01:09, 1.42it/s]
22%	27/125 [02:47<01:17, 1.27it/s]
22%	28/125 [02:48<01:07, 1.43it/s]
24%	30/125 [02:48<00:50, 1.90it/s]
25%	31/125 [02:49<00:58, 1.61it/s]
26%	32/125 [02:50<01:02, 1.48it/s]
28%	35/125 [02:51<00:50, 1.77it/s]
29%	36/125 [02:51<00:56, 1.58it/s]
30%	38/125 [02:52<00:45, 1.90it/s]
31%	39/125 [02:52<00:41, 2.06it/s]
32%	40/125 [02:54<01:00, 1.40it/s]
34%	43/125 [02:54<00:45, 1.82it/s]
35%	44/125 [02:56<01:06, 1.21it/s]
38%	47/125 [02:56<00:47, 1.66it/s]
38%	48/125 [02:57<01:06, 1.15it/s]
40%	50/125 [02:58<00:48, 1.56it/s]
42%	52/125 [02:59<00:48, 1.51it/s]
43%	54/125 [02:59<00:37, 1.89it/s]
44%	55/125 [03:00<00:30, 2.26it/s]
45%	56/125 [03:01<00:53, 1.29it/s]
46%	58/125 [03:01<00:38, 1.73it/s]
47%	59/125 [03:02<00:40, 1.62it/s]
48%	60/125 [03:03<00:46, 1.39it/s]
50%	62/125 [03:04<00:35, 1.76it/s]
50%	63/125 [03:04<00:29, 2.08it/s]
51%	64/125 [03:05<00:37, 1.62it/s]
53%	66/125 [03:05<00:29, 1.99it/s]
54%	67/125 [03:06<00:28, 2.00it/s]
54%	68/125 [03:06<00:29, 1.94it/s]

56%		70/125	[03:07<00:25,	2.18it/s]
57%		71/125	[03:07<00:25,	2.13it/s]
58%		72/125	[03:08<00:24,	2.12it/s]
59%		74/125	[03:09<00:24,	2.09it/s]
60%		75/125	[03:09<00:19,	2.56it/s]
61%		76/125	[03:10<00:21,	2.32it/s]
62%		78/125	[03:11<00:22,	2.10it/s]
63%		79/125	[03:11<00:20,	2.24it/s]
64%		80/125	[03:11<00:17,	2.57it/s]
66%		82/125	[03:12<00:18,	2.29it/s]
66%		83/125	[03:13<00:16,	2.53it/s]
67%		84/125	[03:14<00:21,	1.93it/s]
69%		86/125	[03:14<00:17,	2.19it/s]
70%		87/125	[03:14<00:14,	2.55it/s]
70%		88/125	[03:15<00:21,	1.75it/s]
72%		90/125	[03:16<00:16,	2.09it/s]
73%		91/125	[03:16<00:14,	2.42it/s]
74%		92/125	[03:17<00:20,	1.59it/s]
75%		94/125	[03:17<00:14,	2.18it/s]
76%		95/125	[03:18<00:14,	2.01it/s]
77%		96/125	[03:19<00:21,	1.35it/s]
79%		99/125	[03:20<00:15,	1.72it/s]
80%		100/125	[03:21<00:19,	1.31it/s]
82%		102/125	[03:21<00:12,	1.82it/s]
82%		103/125	[03:22<00:12,	1.74it/s]
83%		104/125	[03:23<00:14,	1.45it/s]
85%		106/125	[03:23<00:09,	2.01it/s]
86%		107/125	[03:24<00:09,	1.96it/s]
86%		108/125	[03:24<00:10,	1.56it/s]
88%		110/125	[03:25<00:06,	2.15it/s]
89%		111/125	[03:25<00:07,	1.79it/s]
90%		112/125	[03:26<00:08,	1.59it/s]
91%		114/125	[03:26<00:04,	2.20it/s]
92%		115/125	[03:27<00:05,	1.78it/s]
93%		116/125	[03:38<00:32,	3.59s/it]
94%		118/125	[03:38<00:17,	2.53s/it]
96%		120/125	[03:39<00:09,	2.00s/it]
98%		123/125	[03:39<00:02,	1.41s/it]
100%		125/125	[03:41<00:00,	1.77s/it]
0%		0/1000	[00:00<?,	?it/s]
2%		23/1000	[00:00<00:04,	228.74it/s]

Computing BLEU

5%		49/1000	[00:00<00:04,	234.46it/s]
7%		74/1000	[00:00<00:03,	237.87it/s]
10%		100/1000	[00:00<00:03,	242.78it/s]

```

13%|      | 126/1000 [00:00<00:03, 247.23it/s]
15%|      | 152/1000 [00:00<00:03, 249.10it/s]
18%|      | 179/1000 [00:00<00:03, 253.85it/s]
20%|      | 205/1000 [00:00<00:03, 255.30it/s]
23%|      | 231/1000 [00:00<00:02, 256.40it/s]
26%|      | 257/1000 [00:01<00:02, 257.29it/s]
28%|      | 283/1000 [00:01<00:02, 251.11it/s]
31%|      | 310/1000 [00:01<00:02, 255.81it/s]
34%|      | 336/1000 [00:01<00:02, 256.14it/s]
36%|      | 362/1000 [00:01<00:02, 253.53it/s]
39%|      | 389/1000 [00:01<00:02, 257.44it/s]
42%|      | 415/1000 [00:01<00:02, 250.40it/s]
44%|      | 441/1000 [00:01<00:02, 252.25it/s]
47%|      | 467/1000 [00:01<00:02, 254.04it/s]
49%|      | 493/1000 [00:01<00:01, 255.67it/s]
52%|      | 519/1000 [00:02<00:01, 252.39it/s]
55%|      | 545/1000 [00:02<00:01, 247.25it/s]
57%|      | 572/1000 [00:02<00:01, 251.92it/s]
60%|      | 598/1000 [00:02<00:01, 252.22it/s]
62%|      | 624/1000 [00:02<00:01, 250.13it/s]
65%|      | 650/1000 [00:02<00:01, 246.84it/s]
68%|      | 675/1000 [00:02<00:01, 245.52it/s]
70%|      | 700/1000 [00:02<00:01, 243.98it/s]
72%|      | 725/1000 [00:02<00:01, 245.54it/s]
75%|      | 750/1000 [00:02<00:01, 245.24it/s]
78%|      | 775/1000 [00:03<00:00, 246.46it/s]
80%|      | 800/1000 [00:03<00:00, 240.17it/s]
82%|      | 825/1000 [00:03<00:00, 242.48it/s]
85%|      | 850/1000 [00:03<00:00, 239.53it/s]
87%|      | 874/1000 [00:03<00:00, 239.36it/s]
90%|      | 898/1000 [00:03<00:00, 237.11it/s]
92%|      | 922/1000 [00:03<00:00, 235.86it/s]
95%|      | 946/1000 [00:03<00:00, 225.74it/s]
97%|      | 969/1000 [00:03<00:00, 222.37it/s]
100%|     | 1000/1000 [00:04<00:00, 245.38it/s]

```

```

BLEU 1:69.2213376024829, BLEU 2:44.871153595427984, BLEU 3:25.158089893489194,
BLEU 4:16.818247623859282

```

- Try different temperatures (e.g. 0.1, 0.2, 0.5, 1.0, 1.5, 2, etc.) during the generation. Report BLEU scores for at least 3 different temperatures.

```

[ ]: ## Use at least 3 different temperatures to generate captions on the test set.
    →Report the BLEU scores.
    # Your code here
    bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic',
    →temperature=0.5)

```



```

print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}"
      .format(bleu1, bleu2, bleu3, bleu4))

bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic',
      temperature=1.0)
print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}"
      .format(bleu1, bleu2, bleu3, bleu4))

bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic',
      temperature=2.0)
print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}"
      .format(bleu1, bleu2, bleu3, bleu4))

# End of code

```

```

0%|          | 0/125 [00:00<?, ?it/s]

Run on the test set...

100%|| 125/125 [00:12<00:00, 10.05it/s]
 2%|          | 24/1000 [00:00<00:04, 233.05it/s]

Computing BLEU

100%|| 1000/1000 [00:03<00:00, 260.82it/s]

BLEU 1:69.2213376024829, BLEU 2:44.871153595427984, BLEU 3:25.158089893489194,
BLEU 4:16.818247623859282

0%|          | 0/125 [00:00<?, ?it/s]

Run on the test set...

100%|| 125/125 [00:12<00:00, 10.06it/s]
 3%|          | 26/1000 [00:00<00:03, 259.91it/s]

Computing BLEU

100%|| 1000/1000 [00:03<00:00, 263.60it/s]

BLEU 1:69.2213376024829, BLEU 2:44.871153595427984, BLEU 3:25.158089893489194,
BLEU 4:16.818247623859282

0%|          | 0/125 [00:00<?, ?it/s]

Run on the test set...

100%|| 125/125 [00:12<00:00, 10.14it/s]
 3%|          | 27/1000 [00:00<00:03, 269.62it/s]

Computing BLEU

100%|| 1000/1000 [00:03<00:00, 263.76it/s]

```

BLEU 1:69.2213376024829, BLEU 2:44.871153595427984, BLEU 3:25.158089893489194,
BLEU 4:16.818247623859282

13 Section 3 Variations [55 pts]

13.1 Section 3.1 LSTM [35 pts]

13.2 Section 3.1.1 Decoder: LSTM [5 pts]

This time, replace the RNN module with an LSTM module.

```
[18]: class Decoder(nn.Module):
    def __init__(self, vocab_size, emb_dim, hidden_dim, num_layers=1, dropout=0):
        """
        Use LSTM as decoder for captions.
        :param emb_dim: Embedding dimensions.
        :param hidden_dim: Hidden states dimensions.
        :param num_layers: Number of LSTM layers.
        :param vocab_size: The size of Vocabulary.
        :param dropout: dropout probability
        """
        super(Decoder, self).__init__()
        #####Your code#####
        # you need to implement a Vanilla RNN for the decoder. Take a look at
        → the official documentation.
        # https://pytorch.org/docs/stable/generated/torch.nn.RNN.html#torch.nn.
        → RNN
        self.max_length = 30 # the maximum length of a sentence, in case it's
        → trapped
        self.hidden_dim = hidden_dim
        self.vocab_size = vocab_size
        self.temp = 1

        # one-hot encoding + linear layer
        self.embed = nn.Embedding(self.vocab_size, emb_dim)
        # LSTM network
        self.LSTM = nn.LSTM(emb_dim, hidden_dim, num_layers)
        # output layer
        self.out = nn.Linear(hidden_dim, vocab_size)

    def forward(self, encode_features, captions, lengths):
        """
        Feed forward to generate captions.
        :param encode_features: output of encoder, size [N, emb_dim]
        :param captions: captions, size [N, max(lengths)]
```

```

        :param lengths: a list indicating valid length for each caption. length_
→is (batch_size).
        """
        #####Your Code#####
        # compute the embedding using one-hot technique and linear function
        caption_embedded = self.embed(captions)

        # concatenate the encoded features from encoder and embeddings
        encode_features = torch.cat((encode_features.unsqueeze(1),
→caption_embedded), dim = 1)

        # feed into RNN.
        encode_features_pack = pack_padded_sequence(encode_features, lengths,
→batch_first=True)

        # output layer
        output, hidden = self.LSTM(encode_features_pack)
        outputs = self.out(output[0])

        return outputs

```

13.3 Encoder-Decoder [5 pts]

```

[19]: class LSTM(nn.Module):
    def __init__(self, vocab_size, emb_dim, hidden_dim, num_layers=1, dropout=0):
        """
        Encoder-decoder vanilla RNN.
        :param vocab_size: the size of Vocabulary.
        :param emb_dim: the dimensions of word embedding.
        :param hidden_dim: the dimensions of hidden units.
        :param num_layers: the number of RNN layers.
        """
        super(LSTM, self).__init__()

        #####Your Code#####
        # Encoder: ResNet-50
        self.lstm_encoder = Encoder(emb_dim)
        # Decoder: LSTM
        self.lstm_decoder = Decoder(vocab_size, emb_dim, hidden_dim, num_layers,
→dropout)
        self.max_length = self.lstm_decoder.max_length
        self.temp = 1
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x, captions, lengths):
        """

```

```

    Feed forward.
    :param x: Images, [N, 3, H, W]
    :param captions: encoded captions, [N, max(lengths)]
    :param lengths: a list indicating valid length for each caption. length_
    → is (batch_size).

    :return: output logits, usually followed by a softmax layer.
    """
    #####Your code#####
    # forward passing
    encoder_features = self.lstm_encoder(x)
    x = self.lstm_decoder(encoder_features, captions, lengths)

    return x

def sample_generate(self, x, states=None, mode='Deterministic',
    → temperature=5.0):
    """
    Generate samples during the evaluation.

    :param x: input image
    :param states: rnn states
    :param mode: which mode we use.
        - 'Deterministic': Take the maximum output at each step.
        - 'Stochastic': Sample from the probability distribution from the_
    → output layer.

    :param temperature: will be used in the stochastic mode
    :return: sample_idx. Word indices. We can use vocab to recover the_
    → sentence.
    """
    sample_idx = []
    #####Your Code#####
    # compute the encoded features
    in_feature = self.lstm_encoder.forward(x)
    in_feature = in_feature.squeeze(1)
    batch_size = in_feature.shape[0]

    for i in range(self.max_length):
        outputs, states = self.lstm_decoder.rnn(in_feature, states)
        outputs = outputs.squeeze(1)
        outputs = self.lstm_decoder.out(outputs)
        outputs = self.softmax(outputs/temperature)

    # decide which mode we use
    if mode == 'Deterministic':
        # take the maximum index after the softmax
        max_val, predicted = outputs.max(1)
        sample_idx.append(predicted)

```

```

elif mode == 'Stochastic':
    # sample from the probability distribution after the softmax
    # Hint: use torch.multinomial() to sample from a distribution.
    predicted = torch.multinomial(outputs, 1)
    sample_idxes.append(predicted)

    x = self.lstm_decoder.embed(predicted)
    x = x.unsqueeze(1)
    sample_idxes = torch.stack(sample_idxes, 1)
    sample_idxes = sample_idxes.squeeze()

return sample_idxes

```

13.4 Section 3.1.2 Training [10 pts]

Use the same set of hyper-parameters (hidden units, optimizer, learning rate etc.) for both models.

```

[20]: # some hyperparameters, you can change them
      ## training parameters
      batch_size = 256
      lr = 1e-2
      num_epochs = 50
      weight_decay = 0.0
      log_step = 50

      ## network architecture
      emb_dim = 1024
      hidden_dim = 256
      num_layers = 1 # number of RNN layers
      dropout = 0.0

      ## image transformation
      transform = transforms.Compose([
          transforms.Resize(256),
          transforms.CenterCrop(224),
          # transforms.RandomCrop(224, pad_if_needed=True),
          transforms.RandomHorizontalFlip(),
          transforms.ToTensor(),
          transforms.Normalize(mean=(0.485, 0.456, 0.406),
                               std=(0.229, 0.224, 0.225))]

      ## Output directory
      output_dir = path_to_homework + '/checkpoints/lstm/'
      os.makedirs(output_dir, exist_ok=True)

```

```

[ ]: # Training code here

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

train_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    ↪split='train', vocab=vocab,
                                transform=transform, batch_size=batch_size,
    ↪shuffle=True, num_workers=12)
val_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    ↪split='val', vocab=vocab,
                                transform=transform, batch_size=8, shuffle=True,
    ↪num_workers=4)

model = LSTM(vocab_size=len(vocab), emb_dim=emb_dim, hidden_dim=hidden_dim,
              num_layers=1, dropout=dropout).to(device) # build a model

vocab_size=len(vocab)

# loss and optimizer
criterion = nn.CrossEntropyLoss().to(device) # CE loss
optimizer = torch.optim.Adam(model.parameters(), lr=lr,
    ↪weight_decay=weight_decay) # optimizer
scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                              step_size=5,
                                              gamma=0.5) # decay LR by a factor of 0.5
    ↪every 10 epochs. You can change this

# logs
Train_Losses = [] # record average training loss each epoch
Val_Losses = [] # record average validation loss each epoch
total_step = len(train_data_loader) # number of iterations each epoch
best_val_loss = np.inf

# start training
print('Start training...')
import time
tic = time.time()
for epoch in range(num_epochs):
    print('Switch to training...')
    model.train()
    Train_loss_iter = [] # record the the training loss each iteration
    for itr, (images, captions, lengths) in enumerate(train_data_loader):
        #####Your Code#####
        model.zero_grad()
        images = Variable(images).to(device)
        captions = Variable(captions).to(device)

```

```

        predicted_cap_pack = pack_padded_sequence(captions, lengths,
→batch_first=True)[0]
        predicted_cap = model(images, captions, lengths)
        predicted_cap = predicted_cap.view(-1, vocab_size)
        loss = criterion(input = predicted_cap, target=predicted_cap_pack)
        loss.backward()
        optimizer.step()

        # record the training loss
        Train_Losses.append(float(loss))

        # print log info
        if itr % log_step == 0:
            # print current loss and perplexity
            print('Epoch [{} / {}], Step [{} / {}], Loss: {:.4f}, Perplexity: {:.
→4f}'
                    .format(epoch, num_epochs, itr, total_step, loss.item(),
→np.exp(loss.item()))
            scheduler.step()
            Train_Losses.append(np.mean(Train_loss_iter))
            np.save(os.path.join(output_dir, 'TrainingLoss_lstm.npy'), Train_Losses) #
→save the training loss

        model.eval()
        # (optional) generate a sample during the training, you can use
→deterministic mode
        # Your code

        # validation
        Val_Losses.append(val(model, val_data_loader, vocab))
        np.save(os.path.join(output_dir, 'ValLoss_lstm.npy'), Val_Losses) # save the
→val loss

        # save model
        if Val_Losses[-1] < best_val_loss:
            best_val_loss = Val_Losses[-1]
            print('updated best val loss:', best_val_loss)
            print('Save model weights to...', output_dir)
            torch.save(model.state_dict(),
                        os.path.join(output_dir, 'lstm-best.pth'.format(epoch + 1,
→itr + 1)))

    print('It took: {} s'.format(time.time() - tic))

```

Start training...

Switch to training...

```

Epoch [0/50], Step [0/114], Loss: 9.2092, Perplexity: 9988.7733
Epoch [0/50], Step [50/114], Loss: 3.6335, Perplexity: 37.8462
Epoch [0/50], Step [100/114], Loss: 3.4302, Perplexity: 30.8838
Validating...

/usr/local/lib/python3.6/dist-packages/numpy/core/fromnumeric.py:3335:
RuntimeWarning: Mean of empty slice.
  out=out, **kwargs)
/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:161:
RuntimeWarning: invalid value encountered in double_scalars
  ret = ret.dtype.type(ret / rcount)

Step [0/127], Loss: 3.2063, Perplexity: 24.6869
Step [50/127], Loss: 3.3955, Perplexity: 29.8297
Step [100/127], Loss: 3.0332, Perplexity: 20.7628
updated best val loss: 3.3872082646437516
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [1/50], Step [0/114], Loss: 3.3987, Perplexity: 29.9261
Epoch [1/50], Step [50/114], Loss: 3.2961, Perplexity: 27.0077
Epoch [1/50], Step [100/114], Loss: 3.2551, Perplexity: 25.9223
Validating...
Step [0/127], Loss: 3.6364, Perplexity: 37.9546
Step [50/127], Loss: 3.6200, Perplexity: 37.3364
Step [100/127], Loss: 3.3697, Perplexity: 29.0694
updated best val loss: 3.2531456909780427
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [2/50], Step [0/114], Loss: 3.1212, Perplexity: 22.6726
Epoch [2/50], Step [50/114], Loss: 3.2186, Perplexity: 24.9925
Epoch [2/50], Step [100/114], Loss: 3.2312, Perplexity: 25.3089
Validating...
Step [0/127], Loss: 3.4916, Perplexity: 32.8397
Step [50/127], Loss: 2.8498, Perplexity: 17.2848
Step [100/127], Loss: 3.2790, Perplexity: 26.5484
updated best val loss: 3.2111755844176284
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [3/50], Step [0/114], Loss: 3.1079, Perplexity: 22.3742
Epoch [3/50], Step [50/114], Loss: 3.0588, Perplexity: 21.3021
Epoch [3/50], Step [100/114], Loss: 3.1670, Perplexity: 23.7362
Validating...
Step [0/127], Loss: 3.6116, Perplexity: 37.0271
Step [50/127], Loss: 3.1199, Perplexity: 22.6446
Step [100/127], Loss: 3.7624, Perplexity: 43.0537
updated best val loss: 3.182656032832589

```



```

Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [4/50], Step [0/114], Loss: 3.0856, Perplexity: 21.8806
Epoch [4/50], Step [50/114], Loss: 3.0058, Perplexity: 20.2029
Epoch [4/50], Step [100/114], Loss: 3.0509, Perplexity: 21.1336
Validating...
Step [0/127], Loss: 3.3157, Perplexity: 27.5411
Step [50/127], Loss: 3.0505, Perplexity: 21.1266
Step [100/127], Loss: 4.0848, Perplexity: 59.4312
Switch to training...
Epoch [5/50], Step [0/114], Loss: 2.9088, Perplexity: 18.3355
Epoch [5/50], Step [50/114], Loss: 2.9563, Perplexity: 19.2261
Epoch [5/50], Step [100/114], Loss: 2.9170, Perplexity: 18.4861
Validating...
Step [0/127], Loss: 2.9846, Perplexity: 19.7778
Step [50/127], Loss: 3.3383, Perplexity: 28.1708
Step [100/127], Loss: 3.2997, Perplexity: 27.1052
updated best val loss: 3.1018434697248805
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [6/50], Step [0/114], Loss: 2.9192, Perplexity: 18.5269
Epoch [6/50], Step [50/114], Loss: 2.8708, Perplexity: 17.6518
Epoch [6/50], Step [100/114], Loss: 2.9233, Perplexity: 18.6022
Validating...
Step [0/127], Loss: 3.3216, Perplexity: 27.7057
Step [50/127], Loss: 2.9402, Perplexity: 18.9193
Step [100/127], Loss: 3.2715, Perplexity: 26.3507
updated best val loss: 3.0702739343868464
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [7/50], Step [0/114], Loss: 2.8481, Perplexity: 17.2543
Epoch [7/50], Step [50/114], Loss: 2.9175, Perplexity: 18.4947
Epoch [7/50], Step [100/114], Loss: 2.8678, Perplexity: 17.5988
Validating...
Step [0/127], Loss: 3.7062, Perplexity: 40.6993
Step [50/127], Loss: 3.4933, Perplexity: 32.8936
Step [100/127], Loss: 3.1970, Perplexity: 24.4583
Switch to training...
Epoch [8/50], Step [0/114], Loss: 2.9043, Perplexity: 18.2528
Epoch [8/50], Step [50/114], Loss: 2.7768, Perplexity: 16.0680
Epoch [8/50], Step [100/114], Loss: 2.8306, Perplexity: 16.9558
Validating...
Step [0/127], Loss: 2.9622, Perplexity: 19.3402
Step [50/127], Loss: 3.7252, Perplexity: 41.4795
Step [100/127], Loss: 3.0451, Perplexity: 21.0113

```

Switch to training...

Epoch [9/50], Step [0/114], Loss: 2.8373, Perplexity: 17.0701
Epoch [9/50], Step [50/114], Loss: 2.7605, Perplexity: 15.8081
Epoch [9/50], Step [100/114], Loss: 2.8063, Perplexity: 16.5490
Validating...

Step [0/127], Loss: 2.8764, Perplexity: 17.7495
Step [50/127], Loss: 3.4441, Perplexity: 31.3136
Step [100/127], Loss: 3.5390, Perplexity: 34.4314
Switch to training...

Epoch [10/50], Step [0/114], Loss: 2.7690, Perplexity: 15.9426
Epoch [10/50], Step [50/114], Loss: 2.7048, Perplexity: 14.9510
Epoch [10/50], Step [100/114], Loss: 2.7803, Perplexity: 16.1232
Validating...

Step [0/127], Loss: 2.8100, Perplexity: 16.6098
Step [50/127], Loss: 3.0468, Perplexity: 21.0481
Step [100/127], Loss: 2.6617, Perplexity: 14.3207
updated best val loss: 3.0513561748144196
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...

Epoch [11/50], Step [0/114], Loss: 2.7291, Perplexity: 15.3188
Epoch [11/50], Step [50/114], Loss: 2.7631, Perplexity: 15.8495
Epoch [11/50], Step [100/114], Loss: 2.7919, Perplexity: 16.3122
Validating...

Step [0/127], Loss: 2.5162, Perplexity: 12.3811
Step [50/127], Loss: 2.6523, Perplexity: 14.1873
Step [100/127], Loss: 3.1476, Perplexity: 23.2792
updated best val loss: 3.0223752795241947
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...

Epoch [12/50], Step [0/114], Loss: 2.7140, Perplexity: 15.0899
Epoch [12/50], Step [50/114], Loss: 2.6831, Perplexity: 14.6300
Epoch [12/50], Step [100/114], Loss: 2.7021, Perplexity: 14.9104
Validating...

Step [0/127], Loss: 4.0060, Perplexity: 54.9252
Step [50/127], Loss: 2.9865, Perplexity: 19.8162
Step [100/127], Loss: 2.7126, Perplexity: 15.0684
Switch to training...

Epoch [13/50], Step [0/114], Loss: 2.6712, Perplexity: 14.4579
Epoch [13/50], Step [50/114], Loss: 2.6996, Perplexity: 14.8741
Epoch [13/50], Step [100/114], Loss: 2.7723, Perplexity: 15.9951
Validating...

Step [0/127], Loss: 3.2097, Perplexity: 24.7706
Step [50/127], Loss: 2.7496, Perplexity: 15.6361
Step [100/127], Loss: 3.0757, Perplexity: 21.6656
Switch to training...

Epoch [14/50], Step [0/114], Loss: 2.6783, Perplexity: 14.5608

Epoch [14/50], Step [50/114], Loss: 2.6717, Perplexity: 14.4641
 Epoch [14/50], Step [100/114], Loss: 2.6644, Perplexity: 14.3599
 Validating...
 Step [0/127], Loss: 3.3568, Perplexity: 28.6965
 Step [50/127], Loss: 3.2253, Perplexity: 25.1610
 Step [100/127], Loss: 3.7404, Perplexity: 42.1160
 Switch to training...
 Epoch [15/50], Step [0/114], Loss: 2.6977, Perplexity: 14.8462
 Epoch [15/50], Step [50/114], Loss: 2.6689, Perplexity: 14.4247
 Epoch [15/50], Step [100/114], Loss: 2.6214, Perplexity: 13.7552
 Validating...
 Step [0/127], Loss: 3.5084, Perplexity: 33.3938
 Step [50/127], Loss: 2.4390, Perplexity: 11.4616
 Step [100/127], Loss: 3.0438, Perplexity: 20.9849
 Switch to training...
 Epoch [16/50], Step [0/114], Loss: 2.6392, Perplexity: 14.0018
 Epoch [16/50], Step [50/114], Loss: 2.5708, Perplexity: 13.0759
 Epoch [16/50], Step [100/114], Loss: 2.6436, Perplexity: 14.0642
 Validating...
 Step [0/127], Loss: 3.0589, Perplexity: 21.3043
 Step [50/127], Loss: 3.1339, Perplexity: 22.9625
 Step [100/127], Loss: 2.7640, Perplexity: 15.8631
 updated best val loss: 3.0058976946853275
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
 Switch to training...
 Epoch [17/50], Step [0/114], Loss: 2.6318, Perplexity: 13.8986
 Epoch [17/50], Step [50/114], Loss: 2.6380, Perplexity: 13.9859
 Epoch [17/50], Step [100/114], Loss: 2.6412, Perplexity: 14.0297
 Validating...
 Step [0/127], Loss: 3.0537, Perplexity: 21.1934
 Step [50/127], Loss: 3.0606, Perplexity: 21.3414
 Step [100/127], Loss: 3.4126, Perplexity: 30.3429
 Switch to training...
 Epoch [18/50], Step [0/114], Loss: 2.5540, Perplexity: 12.8584
 Epoch [18/50], Step [50/114], Loss: 2.6107, Perplexity: 13.6085
 Epoch [18/50], Step [100/114], Loss: 2.6359, Perplexity: 13.9558
 Validating...
 Step [0/127], Loss: 3.1912, Perplexity: 24.3181
 Step [50/127], Loss: 2.2678, Perplexity: 9.6578
 Step [100/127], Loss: 2.7417, Perplexity: 15.5133
 updated best val loss: 2.983503542547151
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
 Switch to training...
 Epoch [19/50], Step [0/114], Loss: 2.5924, Perplexity: 13.3615
 Epoch [19/50], Step [50/114], Loss: 2.6096, Perplexity: 13.5934
 Epoch [19/50], Step [100/114], Loss: 2.6161, Perplexity: 13.6825

```

Validating...
Step [0/127], Loss: 3.2313, Perplexity: 25.3130
Step [50/127], Loss: 2.5452, Perplexity: 12.7460
Step [100/127], Loss: 2.6936, Perplexity: 14.7842
Switch to training...
Epoch [20/50], Step [0/114], Loss: 2.5843, Perplexity: 13.2536
Epoch [20/50], Step [50/114], Loss: 2.5379, Perplexity: 12.6534
Epoch [20/50], Step [100/114], Loss: 2.5176, Perplexity: 12.3989
Validating...
Step [0/127], Loss: 2.9751, Perplexity: 19.5912
Step [50/127], Loss: 3.1422, Perplexity: 23.1538
Step [100/127], Loss: 3.4476, Perplexity: 31.4255
Switch to training...
Epoch [21/50], Step [0/114], Loss: 2.6178, Perplexity: 13.7061
Epoch [21/50], Step [50/114], Loss: 2.5632, Perplexity: 12.9774
Epoch [21/50], Step [100/114], Loss: 2.5819, Perplexity: 13.2225
Validating...
Step [0/127], Loss: 3.0761, Perplexity: 21.6736
Step [50/127], Loss: 3.1204, Perplexity: 22.6560
Step [100/127], Loss: 3.2029, Perplexity: 24.6048
Switch to training...
Epoch [22/50], Step [0/114], Loss: 2.5458, Perplexity: 12.7540
Epoch [22/50], Step [50/114], Loss: 2.5651, Perplexity: 13.0015
Epoch [22/50], Step [100/114], Loss: 2.5611, Perplexity: 12.9506
Validating...
Step [0/127], Loss: 3.7708, Perplexity: 43.4148
Step [50/127], Loss: 3.3769, Perplexity: 29.2794
Step [100/127], Loss: 3.0838, Perplexity: 21.8421
Switch to training...
Epoch [23/50], Step [0/114], Loss: 2.5454, Perplexity: 12.7489
Epoch [23/50], Step [50/114], Loss: 2.5369, Perplexity: 12.6400
Epoch [23/50], Step [100/114], Loss: 2.6144, Perplexity: 13.6585
Validating...
Step [0/127], Loss: 2.9872, Perplexity: 19.8300
Step [50/127], Loss: 3.1717, Perplexity: 23.8490
Step [100/127], Loss: 2.8599, Perplexity: 17.4605
updated best val loss: 2.983343030524066
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/lstm/
Switch to training...
Epoch [24/50], Step [0/114], Loss: 2.5215, Perplexity: 12.4475
Epoch [24/50], Step [50/114], Loss: 2.5822, Perplexity: 13.2268
Epoch [24/50], Step [100/114], Loss: 2.4814, Perplexity: 11.9575
Validating...
Step [0/127], Loss: 3.0142, Perplexity: 20.3719
Step [50/127], Loss: 2.6591, Perplexity: 14.2830
Step [100/127], Loss: 2.8667, Perplexity: 17.5794
Switch to training...

```

Epoch [25/50], Step [0/114], Loss: 2.5205, Perplexity: 12.4354
 Epoch [25/50], Step [50/114], Loss: 2.5571, Perplexity: 12.8979
 Epoch [25/50], Step [100/114], Loss: 2.5156, Perplexity: 12.3736
 Validating...
 Step [0/127], Loss: 2.7126, Perplexity: 15.0691
 Step [50/127], Loss: 3.1860, Perplexity: 24.1922
 Step [100/127], Loss: 2.6225, Perplexity: 13.7698
 Switch to training...
 Epoch [26/50], Step [0/114], Loss: 2.5530, Perplexity: 12.8459
 Epoch [26/50], Step [50/114], Loss: 2.5031, Perplexity: 12.2197
 Epoch [26/50], Step [100/114], Loss: 2.4986, Perplexity: 12.1653
 Validating...
 Step [0/127], Loss: 3.2609, Perplexity: 26.0735
 Step [50/127], Loss: 3.2875, Perplexity: 26.7750
 Step [100/127], Loss: 3.5123, Perplexity: 33.5251
 Switch to training...
 Epoch [27/50], Step [0/114], Loss: 2.5535, Perplexity: 12.8517
 Epoch [27/50], Step [50/114], Loss: 2.5071, Perplexity: 12.2692
 Epoch [27/50], Step [100/114], Loss: 2.5041, Perplexity: 12.2327
 Validating...
 Step [0/127], Loss: 3.0362, Perplexity: 20.8269
 Step [50/127], Loss: 3.5058, Perplexity: 33.3072
 Step [100/127], Loss: 2.7043, Perplexity: 14.9443
 Switch to training...
 Epoch [28/50], Step [0/114], Loss: 2.5539, Perplexity: 12.8576
 Epoch [28/50], Step [50/114], Loss: 2.4982, Perplexity: 12.1606
 Epoch [28/50], Step [100/114], Loss: 2.5779, Perplexity: 13.1696
 Validating...
 Step [0/127], Loss: 2.6160, Perplexity: 13.6811
 Step [50/127], Loss: 3.1596, Perplexity: 23.5613
 Step [100/127], Loss: 3.1996, Perplexity: 24.5224
 Switch to training...
 Epoch [29/50], Step [0/114], Loss: 2.5220, Perplexity: 12.4540
 Epoch [29/50], Step [50/114], Loss: 2.4859, Perplexity: 12.0114
 Epoch [29/50], Step [100/114], Loss: 2.5433, Perplexity: 12.7219
 Validating...
 Step [0/127], Loss: 3.0019, Perplexity: 20.1231
 Step [50/127], Loss: 2.7442, Perplexity: 15.5528
 Step [100/127], Loss: 3.3640, Perplexity: 28.9038
 Switch to training...
 Epoch [30/50], Step [0/114], Loss: 2.5744, Perplexity: 13.1239
 Epoch [30/50], Step [50/114], Loss: 2.5333, Perplexity: 12.5948
 Epoch [30/50], Step [100/114], Loss: 2.5483, Perplexity: 12.7856
 Validating...
 Step [0/127], Loss: 3.3189, Perplexity: 27.6290
 Step [50/127], Loss: 3.2925, Perplexity: 26.9112
 Step [100/127], Loss: 2.9456, Perplexity: 19.0226
 Switch to training...

Epoch [31/50], Step [0/114], Loss: 2.6138, Perplexity: 13.6510
 Epoch [31/50], Step [50/114], Loss: 2.5185, Perplexity: 12.4096
 Epoch [31/50], Step [100/114], Loss: 2.5460, Perplexity: 12.7563
 Validating...
 Step [0/127], Loss: 2.4137, Perplexity: 11.1747
 Step [50/127], Loss: 3.0566, Perplexity: 21.2553
 Step [100/127], Loss: 3.1902, Perplexity: 24.2939
 Switch to training...
 Epoch [32/50], Step [0/114], Loss: 2.5405, Perplexity: 12.6857
 Epoch [32/50], Step [50/114], Loss: 2.5056, Perplexity: 12.2505
 Epoch [32/50], Step [100/114], Loss: 2.5094, Perplexity: 12.2976
 Validating...
 Step [0/127], Loss: 3.2712, Perplexity: 26.3431
 Step [50/127], Loss: 3.2343, Perplexity: 25.3874
 Step [100/127], Loss: 3.3417, Perplexity: 28.2672
 Switch to training...
 Epoch [33/50], Step [0/114], Loss: 2.4663, Perplexity: 11.7788
 Epoch [33/50], Step [50/114], Loss: 2.5156, Perplexity: 12.3735
 Epoch [33/50], Step [100/114], Loss: 2.5652, Perplexity: 13.0030
 Validating...
 Step [0/127], Loss: 2.8283, Perplexity: 16.9172
 Step [50/127], Loss: 2.6631, Perplexity: 14.3411
 Step [100/127], Loss: 3.5133, Perplexity: 33.5590
 Switch to training...
 Epoch [34/50], Step [0/114], Loss: 2.4987, Perplexity: 12.1671
 Epoch [34/50], Step [50/114], Loss: 2.5707, Perplexity: 13.0747
 Epoch [34/50], Step [100/114], Loss: 2.5569, Perplexity: 12.8959
 Validating...
 Step [0/127], Loss: 3.0982, Perplexity: 22.1587
 Step [50/127], Loss: 2.3166, Perplexity: 10.1416
 Step [100/127], Loss: 2.9464, Perplexity: 19.0381
 Switch to training...
 Epoch [35/50], Step [0/114], Loss: 2.5227, Perplexity: 12.4619
 Epoch [35/50], Step [50/114], Loss: 2.5738, Perplexity: 13.1158
 Epoch [35/50], Step [100/114], Loss: 2.5375, Perplexity: 12.6480
 Validating...
 Step [0/127], Loss: 3.4064, Perplexity: 30.1555
 Step [50/127], Loss: 2.8421, Perplexity: 17.1519
 Step [100/127], Loss: 3.3082, Perplexity: 27.3356
 Switch to training...
 Epoch [36/50], Step [0/114], Loss: 2.4501, Perplexity: 11.5897
 Epoch [36/50], Step [50/114], Loss: 2.4917, Perplexity: 12.0812
 Epoch [36/50], Step [100/114], Loss: 2.5452, Perplexity: 12.7453
 Validating...
 Step [0/127], Loss: 2.9687, Perplexity: 19.4657
 Step [50/127], Loss: 2.8037, Perplexity: 16.5049
 Step [100/127], Loss: 2.8883, Perplexity: 17.9627
 Switch to training...

Epoch [37/50], Step [0/114], Loss: 2.5329, Perplexity: 12.5901
 Epoch [37/50], Step [50/114], Loss: 2.5315, Perplexity: 12.5717
 Epoch [37/50], Step [100/114], Loss: 2.4917, Perplexity: 12.0815
 Validating...
 Step [0/127], Loss: 2.7629, Perplexity: 15.8464
 Step [50/127], Loss: 2.9824, Perplexity: 19.7358
 Step [100/127], Loss: 3.2267, Perplexity: 25.1963
 Switch to training...
 Epoch [38/50], Step [0/114], Loss: 2.5105, Perplexity: 12.3113
 Epoch [38/50], Step [50/114], Loss: 2.5200, Perplexity: 12.4287
 Epoch [38/50], Step [100/114], Loss: 2.5904, Perplexity: 13.3353
 Validating...
 Step [0/127], Loss: 3.1892, Perplexity: 24.2698
 Step [50/127], Loss: 3.1871, Perplexity: 24.2178
 Step [100/127], Loss: 3.2751, Perplexity: 26.4462
 Switch to training...
 Epoch [39/50], Step [0/114], Loss: 2.5328, Perplexity: 12.5883
 Epoch [39/50], Step [50/114], Loss: 2.5508, Perplexity: 12.8173
 Epoch [39/50], Step [100/114], Loss: 2.5119, Perplexity: 12.3284
 Validating...
 Step [0/127], Loss: 3.4229, Perplexity: 30.6579
 Step [50/127], Loss: 3.0678, Perplexity: 21.4945
 Step [100/127], Loss: 3.3362, Perplexity: 28.1129
 Switch to training...
 Epoch [40/50], Step [0/114], Loss: 2.5090, Perplexity: 12.2931
 Epoch [40/50], Step [50/114], Loss: 2.4719, Perplexity: 11.8447
 Epoch [40/50], Step [100/114], Loss: 2.4979, Perplexity: 12.1570
 Validating...
 Step [0/127], Loss: 2.9359, Perplexity: 18.8392
 Step [50/127], Loss: 3.7445, Perplexity: 42.2876
 Step [100/127], Loss: 3.6316, Perplexity: 37.7740
 Switch to training...
 Epoch [41/50], Step [0/114], Loss: 2.5570, Perplexity: 12.8977
 Epoch [41/50], Step [50/114], Loss: 2.5487, Perplexity: 12.7902
 Epoch [41/50], Step [100/114], Loss: 2.4495, Perplexity: 11.5822
 Validating...
 Step [0/127], Loss: 2.3570, Perplexity: 10.5588
 Step [50/127], Loss: 2.8187, Perplexity: 16.7556
 Step [100/127], Loss: 3.6568, Perplexity: 38.7383
 Switch to training...
 Epoch [42/50], Step [0/114], Loss: 2.6121, Perplexity: 13.6282
 Epoch [42/50], Step [50/114], Loss: 2.5222, Perplexity: 12.4561
 Epoch [42/50], Step [100/114], Loss: 2.4684, Perplexity: 11.8038
 Validating...
 Step [0/127], Loss: 3.2650, Perplexity: 26.1799
 Step [50/127], Loss: 2.7989, Perplexity: 16.4267
 Step [100/127], Loss: 3.4870, Perplexity: 32.6885
 Switch to training...

Epoch [43/50], Step [0/114], Loss: 2.5799, Perplexity: 13.1961
 Epoch [43/50], Step [50/114], Loss: 2.4939, Perplexity: 12.1084
 Epoch [43/50], Step [100/114], Loss: 2.5316, Perplexity: 12.5734
 Validating...
 Step [0/127], Loss: 3.1979, Perplexity: 24.4809
 Step [50/127], Loss: 3.1743, Perplexity: 23.9109
 Step [100/127], Loss: 3.1573, Perplexity: 23.5081
 Switch to training...
 Epoch [44/50], Step [0/114], Loss: 2.4935, Perplexity: 12.1037
 Epoch [44/50], Step [50/114], Loss: 2.4493, Perplexity: 11.5798
 Epoch [44/50], Step [100/114], Loss: 2.5211, Perplexity: 12.4422
 Validating...
 Step [0/127], Loss: 2.9847, Perplexity: 19.7800
 Step [50/127], Loss: 3.0636, Perplexity: 21.4040
 Step [100/127], Loss: 3.5115, Perplexity: 33.4978
 Switch to training...
 Epoch [45/50], Step [0/114], Loss: 2.5255, Perplexity: 12.4976
 Epoch [45/50], Step [50/114], Loss: 2.5100, Perplexity: 12.3049
 Epoch [45/50], Step [100/114], Loss: 2.5052, Perplexity: 12.2461
 Validating...
 Step [0/127], Loss: 3.1976, Perplexity: 24.4740
 Step [50/127], Loss: 3.2894, Perplexity: 26.8267
 Step [100/127], Loss: 2.5607, Perplexity: 12.9447
 Switch to training...
 Epoch [46/50], Step [0/114], Loss: 2.5380, Perplexity: 12.6541
 Epoch [46/50], Step [50/114], Loss: 2.5248, Perplexity: 12.4882
 Epoch [46/50], Step [100/114], Loss: 2.5729, Perplexity: 13.1038
 Validating...
 Step [0/127], Loss: 3.6066, Perplexity: 36.8405
 Step [50/127], Loss: 2.6696, Perplexity: 14.4342
 Step [100/127], Loss: 3.0579, Perplexity: 21.2820
 Switch to training...
 Epoch [47/50], Step [0/114], Loss: 2.4714, Perplexity: 11.8393
 Epoch [47/50], Step [50/114], Loss: 2.5356, Perplexity: 12.6240
 Epoch [47/50], Step [100/114], Loss: 2.5347, Perplexity: 12.6121
 Validating...
 Step [0/127], Loss: 2.8675, Perplexity: 17.5925
 Step [50/127], Loss: 3.0249, Perplexity: 20.5913
 Step [100/127], Loss: 2.8748, Perplexity: 17.7220
 Switch to training...
 Epoch [48/50], Step [0/114], Loss: 2.4894, Perplexity: 12.0543
 Epoch [48/50], Step [50/114], Loss: 2.5115, Perplexity: 12.3238
 Epoch [48/50], Step [100/114], Loss: 2.5197, Perplexity: 12.4246
 Validating...
 Step [0/127], Loss: 3.1014, Perplexity: 22.2296
 Step [50/127], Loss: 3.2681, Perplexity: 26.2616
 Step [100/127], Loss: 2.8346, Perplexity: 17.0233
 Switch to training...

Epoch [49/50], Step [0/114], Loss: 2.4606, Perplexity: 11.7118
 Epoch [49/50], Step [50/114], Loss: 2.4959, Perplexity: 12.1329
 Epoch [49/50], Step [100/114], Loss: 2.5045, Perplexity: 12.2377
 Validating...
 Step [0/127], Loss: 3.1538, Perplexity: 23.4241
 Step [50/127], Loss: 3.4654, Perplexity: 31.9896
 Step [100/127], Loss: 2.8784, Perplexity: 17.7861
 It took: 13621.590710639954 s

13.5 Section 3.1.3 Evaluation [10 pts]

Evaluate your model on the test set by perplexity score or BLEU score

```
[21]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

train_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    ↪split='train', vocab=vocab,
                                transform=transform, batch_size=batch_size,
    ↪shuffle=True, num_workers=12)
val_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    ↪split='val', vocab=vocab,
                                transform=transform, batch_size=8, shuffle=True,
    ↪num_workers=4)

model = LSTM(vocab_size=len(vocab), emb_dim=emb_dim, hidden_dim=hidden_dim,
              num_layers=1, dropout=dropout).to(device) # build a model

vocab_size=len(vocab)

# loss and optimizer
criterion = nn.CrossEntropyLoss().to(device) # CE loss
optimizer = torch.optim.Adam(model.parameters(), lr=lr,
    ↪weight_decay=weight_decay) # optimizer
scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                              step_size=5,
                                              gamma=0.5) # decay LR by a factor of 0.5
    ↪every 10 epochs. You can change this
```

```
[22]: ## Evaluate your model using BLEU score. Use Deterministic mode.
# Your code here
## Image transformation
transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    # transforms.RandomCrop(224, pad_if_needed=True),
    transforms.RandomHorizontalFlip(),
```

```

        transforms.ToTensor(),
        transforms.Normalize(mean=(0.485, 0.456, 0.406),
                               std=(0.229, 0.224, 0.225)))

## Evaluate your model using BLEU score. Use Deterministic mode
model = LSTM(vocab_size=len(vocab), emb_dim=emb_dim, hidden_dim=hidden_dim,
              num_layers=1, dropout=dropout).to(device) # build a model
model.load_state_dict(torch.load(path_to_homework + '/checkpoints/lstm/lstm-best.
→pth', map_location=torch.device('cpu'))))
model.eval()
bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic')
print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}" .format(bleu1, bleu2, bleu3,
→bleu4))
# End of code

```

```
0%|          | 0/125 [00:00<?, ?it/s]
```

Run on the test set...

```

1%|          | 1/125 [00:00<01:09, 1.80it/s]
2%|          | 2/125 [00:00<00:56, 2.18it/s]
3%|          | 4/125 [00:00<00:42, 2.86it/s]
4%|          | 5/125 [00:01<00:34, 3.52it/s]
5%|          | 6/125 [00:01<00:29, 4.06it/s]
6%|          | 7/125 [00:01<00:24, 4.78it/s]
7%|          | 9/125 [00:01<00:20, 5.59it/s]
8%|          | 10/125 [00:01<00:17, 6.41it/s]
9%|          | 11/125 [00:01<00:16, 6.93it/s]
10%|         | 13/125 [00:01<00:13, 8.02it/s]
12%|         | 15/125 [00:02<00:12, 8.81it/s]
14%|         | 17/125 [00:02<00:12, 8.45it/s]
15%|         | 19/125 [00:02<00:12, 8.81it/s]
17%|         | 21/125 [00:02<00:11, 9.28it/s]
18%|         | 23/125 [00:03<00:10, 9.35it/s]
20%|         | 25/125 [00:03<00:10, 9.67it/s]
21%|         | 26/125 [00:03<00:11, 8.73it/s]
22%|         | 28/125 [00:03<00:09, 9.89it/s]
24%|         | 30/125 [00:03<00:09, 9.88it/s]
26%|         | 32/125 [00:03<00:09, 9.94it/s]
27%|         | 34/125 [00:04<00:09, 9.81it/s]
29%|         | 36/125 [00:04<00:08, 10.37it/s]
30%|         | 38/125 [00:04<00:09, 9.52it/s]
32%|         | 40/125 [00:04<00:08, 9.70it/s]
33%|         | 41/125 [00:04<00:09, 9.16it/s]
34%|         | 43/125 [00:04<00:07, 10.30it/s]
36%|         | 45/125 [00:05<00:07, 10.37it/s]

```

38%	47/125 [00:05<00:07, 10.30it/s]
39%	49/125 [00:05<00:07, 10.32it/s]
41%	51/125 [00:05<00:07, 9.79it/s]
42%	53/125 [00:05<00:06, 10.82it/s]
44%	55/125 [00:06<00:06, 10.41it/s]
46%	57/125 [00:06<00:06, 11.09it/s]
47%	59/125 [00:06<00:06, 10.28it/s]
49%	61/125 [00:06<00:05, 11.03it/s]
50%	63/125 [00:06<00:06, 9.92it/s]
52%	65/125 [00:07<00:05, 10.15it/s]
54%	67/125 [00:07<00:05, 9.80it/s]
55%	69/125 [00:07<00:05, 10.37it/s]
57%	71/125 [00:07<00:05, 10.03it/s]
58%	73/125 [00:07<00:05, 10.34it/s]
60%	75/125 [00:08<00:05, 9.86it/s]
62%	77/125 [00:08<00:04, 10.37it/s]
63%	79/125 [00:08<00:04, 9.97it/s]
65%	81/125 [00:08<00:04, 9.69it/s]
66%	82/125 [00:08<00:04, 9.61it/s]
67%	84/125 [00:09<00:04, 9.98it/s]
69%	86/125 [00:09<00:03, 10.36it/s]
70%	88/125 [00:09<00:03, 10.20it/s]
72%	90/125 [00:09<00:03, 10.36it/s]
74%	92/125 [00:09<00:03, 9.96it/s]
75%	94/125 [00:09<00:02, 10.95it/s]
77%	96/125 [00:10<00:03, 9.24it/s]
78%	97/125 [00:10<00:03, 8.93it/s]
79%	99/125 [00:10<00:02, 10.08it/s]
81%	101/125 [00:10<00:02, 10.16it/s]
82%	103/125 [00:10<00:02, 9.64it/s]
84%	105/125 [00:11<00:02, 9.39it/s]
86%	107/125 [00:11<00:01, 9.83it/s]
87%	109/125 [00:11<00:01, 9.99it/s]
89%	111/125 [00:11<00:01, 10.10it/s]
90%	113/125 [00:11<00:01, 10.66it/s]
92%	115/125 [00:12<00:00, 11.11it/s]
94%	117/125 [00:12<00:00, 9.89it/s]
95%	119/125 [00:12<00:00, 11.22it/s]
97%	121/125 [00:12<00:00, 12.83it/s]
100%	125/125 [00:12<00:00, 9.76it/s]
0%	0/1000 [00:00<?, ?it/s]
3%	29/1000 [00:00<00:03, 289.61it/s]

Computing BLEU

6%	57/1000 [00:00<00:03, 285.33it/s]
8%	84/1000 [00:00<00:03, 278.51it/s]

```

11%|      | 113/1000 [00:00<00:03, 279.99it/s]
14%|      | 142/1000 [00:00<00:03, 282.40it/s]
17%|      | 171/1000 [00:00<00:02, 284.04it/s]
20%|      | 201/1000 [00:00<00:02, 285.61it/s]
23%|      | 230/1000 [00:00<00:02, 286.31it/s]
26%|      | 260/1000 [00:00<00:02, 289.26it/s]
29%|      | 288/1000 [00:01<00:02, 281.92it/s]
32%|      | 316/1000 [00:01<00:02, 280.52it/s]
34%|      | 344/1000 [00:01<00:02, 280.35it/s]
37%|      | 372/1000 [00:01<00:02, 280.07it/s]
40%|      | 402/1000 [00:01<00:02, 284.48it/s]
43%|      | 432/1000 [00:01<00:01, 288.61it/s]
46%|      | 461/1000 [00:01<00:01, 287.35it/s]
49%|      | 491/1000 [00:01<00:01, 289.24it/s]
52%|      | 520/1000 [00:01<00:01, 271.71it/s]
55%|      | 548/1000 [00:01<00:01, 272.44it/s]
58%|      | 576/1000 [00:02<00:01, 274.08it/s]
60%|      | 604/1000 [00:02<00:01, 272.48it/s]
63%|      | 632/1000 [00:02<00:01, 271.35it/s]
66%|      | 660/1000 [00:02<00:01, 269.76it/s]
69%|      | 688/1000 [00:02<00:01, 269.72it/s]
72%|      | 715/1000 [00:02<00:01, 267.82it/s]
74%|      | 743/1000 [00:02<00:00, 270.93it/s]
77%|      | 771/1000 [00:02<00:00, 273.12it/s]
80%|      | 799/1000 [00:02<00:00, 270.34it/s]
83%|      | 828/1000 [00:02<00:00, 272.32it/s]
86%|      | 856/1000 [00:03<00:00, 268.39it/s]
88%|      | 883/1000 [00:03<00:00, 262.02it/s]
91%|      | 910/1000 [00:03<00:00, 260.18it/s]
94%|      | 937/1000 [00:03<00:00, 259.80it/s]
96%|      | 965/1000 [00:03<00:00, 264.35it/s]
100%|     | 1000/1000 [00:03<00:00, 274.07it/s]

```

```

BLEU 1:43.43058246704518, BLEU 2:32.799554560543, BLEU 3:19.94015934517099, BLEU
4:13.527134409104775

```

```

[ ]: ## Use at least 3 different temperatures to generate captions on the test set.
    →Report the BLEU scores.
    # Your code here
    bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic',
    →temperature=0.5)
    print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}"
    →bleu4))

    bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic',
    →temperature=1.0)

```

```

print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}".format(bleu1, bleu2, bleu3,
    bleu4))

bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic',
    temperature=2.0)
print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}".format(bleu1, bleu2, bleu3,
    bleu4))

# End of code

```

```

0%|          | 0/125 [00:00<?, ?it/s]

Run on the test set...

100%|| 125/125 [00:12<00:00, 9.86it/s]
 2%|          | 25/1000 [00:00<00:03, 244.64it/s]

Computing BLEU

100%|| 1000/1000 [00:03<00:00, 285.80it/s]

BLEU 1:43.43058246704518, BLEU 2:32.799554560543, BLEU 3:19.94015934517099, BLEU
4:13.527134409104775

0%|          | 0/125 [00:00<?, ?it/s]

Run on the test set...

100%|| 125/125 [00:12<00:00, 9.86it/s]
 3%|          | 29/1000 [00:00<00:03, 289.80it/s]

Computing BLEU

100%|| 1000/1000 [00:03<00:00, 289.54it/s]

BLEU 1:43.43058246704518, BLEU 2:32.799554560543, BLEU 3:19.94015934517099, BLEU
4:13.527134409104775

0%|          | 0/125 [00:00<?, ?it/s]

Run on the test set...

100%|| 125/125 [00:12<00:00, 9.93it/s]
 3%|          | 29/1000 [00:00<00:03, 285.22it/s]

Computing BLEU

100%|| 1000/1000 [00:03<00:00, 287.14it/s]

BLEU 1:43.43058246704518, BLEU 2:32.799554560543, BLEU 3:19.94015934517099, BLEU
4:13.527134409104775

```

13.6 Section 3.1.4 Discussion [5 pts]

What's the difference between Vanilla RNN and LSTM (training loss, evaluation results, etc)?

Your comments:

Training loss: The training loss of Vanilla RNN are slightly larger than LSTM's.

Evaluation results: The evaluation result of Vanilla RNN is better than LSTM's.

Running time: The running time of Vanilla RNN is slightly less than LSTM's.

13.7 Section 3.2 Using pre-trained word embeddings [20 pts]

For now, the decoder uses a word as input by converting it into a fixed size embedding, and our networks learn these word embeddings by training. In this experiment, you will use pre-trained word embeddings like Word2Vec or GloVe in LSTM. If you use Pytorch's `nn.Embedding` layer, you can initialize its weights with a matrix containing pre-trained word embeddings for all words in your vocabulary, and freeze the weights (i.e. don't train this layer). You can find these embeddings online.

Some resources: - GloVe: <https://nlp.stanford.edu/projects/glove/> - Word2Vec: <http://jalamar.github.io/illustrated-word2vec/>

In case you don't know how to get one, we've already provided a light GloVe embedding: `wm_06.npy`, which can produce 300-d word embeddings.

13.8 Section 3.2.1 Encoder-decoder [10 pts]

```
[23]: class Decoder(nn.Module):
    def __init__(self, vocab_size, emb_dim, hidden_dim, pretrained_emb,
        num_layers=1, dropout=0):
        """
        Use LSTM as decoder for captions.
        :param emb_dim: Embedding dimensions.
        :param hidden_dim: Hidden states dimensions.
        :param pretrained_emb: the path to the pretrained embedding
        :param num_layers: Number of LSTM layers.
        :param vocab_size: The size of Vocabulary.
        :param dropout: dropout probability
        """
        super(Decoder, self).__init__()
        #####Your code#####
        # you need to implement a Vanilla RNN for the decoder. Take a look at
        the official documentation.
        # https://pytorch.org/docs/stable/generated/torch.nn.RNN.html#torch.nn.
        RNN
        self.max_length = 30 # the maximum length of a sentence, in case it's
        trapped
        self.hidden_dim = hidden_dim
        self.vocab_size = vocab_size
```

```

self.temp = 1

# one-hot encoding + linear layer
self.embed = nn.Embedding(self.vocab_size, emb_dim)
self.embed.weight.requires_grad = False
# LSTM network
self.lstm = nn.LSTM(emb_dim, hidden_dim, num_layers)
# output layer
self.out = nn.Linear(hidden_dim, vocab_size)

def forward(self, encode_features, captions, lengths):
    """
    Feed forward to generate captions.
    :param encode_features: output of encoder, size [N, emb_dim]
    :param captions: captions, size [N, max(lengths)]
    :param lengths: a list indicating valid length for each caption. length_
→is (batch_size).
    """
    #####Your Code#####
    # compute the embedding using one-hot technique and linear function
    caption_embedded = self.embed(captions)
    # concatenate the encoded features from encoder and embeddings
    encode_features = torch.cat((encode_features.unsqueeze(1),
→caption_embedded), dim=1)
    packed = pack_padded_sequence(encode_features, lengths, batch_first=True)
    # feed into RNN
    output, hidden = self.lstm(packed)
    # output layer
    outputs = self.fc(output[0])

    return outputs

```

```

[24]: class Word_embeddings(nn.Module):
    def __init__(self, vocab_size, emb_dim, hidden_dim, pretrained_emb,
→num_layers=1, dropout=0):
        """
        Encoder-decoder baseline.
        :param vocab_size: the size of Vocabulary.
        :param emb_dim: the dimensions of word embedding.
        :param hidden_dim: the dimensions of hidden units.
        :param pretrained_emb: the path to the pretrained embedding
        :param num_layers: the number of LSTM layers.
        :param dropout: dropout probability.
        """
        super(Word_embeddings, self).__init__()
        # self.max_length = self.decoder.max_length
        #####Your Code#####

```

```

# Encoder: ResNet-50
self.lstm_encoder = Encoder(emb_dim)
# Decoder: LSTM
self.lstm_decoder = Decoder(vocab_size, emb_dim, hidden_dim, num_layers,
→dropout)

self.max_length = self.lstm_decoder.max_length
self.temp = 1
self.softmax = nn.Softmax(dim=1)

def forward(self, x, captions, lengths):
    """
    Feed forward.
    :param x: Images, [N, 3, H, W]
    :param captions: encoded captions, [N, max(lengths)]
    :param lengths: a list indicating valid length for each caption. length_
→is (batch_size).
    :return: output logits, usually followed by a softmax layer.
    """
    #####Your code#####
    # forward passing
    encoder_features = self.lstm_encoder(x)
    x = self.lstm_decoder(encoder_features, captions, lengths)

    return x

def sample_generate(self, x, states=None, mode='Deterministic',
→temperature=5.0):
    """
    Generate samples.
    :param x:
    :return:
    """
    sample_idx = []
    #####Your Code#####
    # compute the encoded features
    in_feature = self.lstm_encoder.forward(x)
    in_feature = in_feature.squeeze(1)
    batch_size = in_feature.shape[0]

    for i in range(self.max_length):
        outputs, states = self.lstm_decoder.rnn(in_feature, states)
        outputs = outputs.squeeze(1)
        outputs = self.lstm_decoder.out(outputs)
        outputs = self.softmax(outputs/temperature)

    # decide which mode we use
    if mode == 'Deterministic':

```



```

        # take the maximum index after the softmax
        max_val, predicted = outputs.max(1)
        sample_idxes.append(predicted)

    elif mode == 'Stochastic':
        # sample from the probability distribution after the softmax
        # Hint: use torch.multinomial() to sample from a distribution.
        predicted = torch.multinomial(outputs, 1)
        sample_idxes.append(predicted)

    x = self.lstm_decoder.embed(predicted)
    x = x.unsqueeze(1)
    sample_idxes = torch.stack(sample_idxes, 1)
    sample_idxes = sample_idxes.squeeze()

    # return sample_idxes

```

13.9 Section 3.2.2 Training [5 pts]

```

[25]: # some hyperparameters, you can change them
      ## training parameters
      batch_size = 256
      lr = 1e-2
      num_epochs = 50
      weight_decay = 0.0
      log_step = 50

      ## network architecture
      emb_dim = 300
      hidden_dim = 256
      num_layers = 1 # number of RNN layers
      dropout = 0.0

      ## image transformation
      transform = transforms.Compose([
          transforms.Resize(256),
          transforms.CenterCrop(224),
          # transforms.RandomCrop(224, pad_if_needed=True),
          transforms.RandomHorizontalFlip(),
          transforms.ToTensor(),
          transforms.Normalize(mean=(0.485, 0.456, 0.406),
                               std=(0.229, 0.224, 0.225))]

      ## Output directory
      output_dir = path_to_homework + '/checkpoints/pretrained_emb/'
      os.makedirs(output_dir, exist_ok=True)

```

```

[26]: # Training code here
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

train_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    →split='train', vocab=vocab,
                                transform=transform, batch_size=batch_size,
    →shuffle=True, num_workers=12)
val_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    →split='val', vocab=vocab,
                                transform=transform, batch_size=8, shuffle=True,
    →num_workers=4)

# pretrained embedding weights
pre_emb_path = '/content/drive/My Drive/DL_Fall_2020/Assignment_4/wm_06.npy' #
    →type the path to the pretrained embedding you find

model = Word_embeddings(vocab_size=len(vocab), emb_dim=emb_dim,
    →hidden_dim=hidden_dim, pretrained_emb=pre_emb_path,
                                num_layers=1, dropout=dropout).to(device) # build a model

# loss and optimizer
criterion = nn.CrossEntropyLoss().to(device) # CE loss
optimizer = torch.optim.Adam(model.parameters(), lr=lr,
    →weight_decay=weight_decay) # optimizer
scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                step_size=5,
                                gamma=0.5) # decay LR by a factor of 0.5
    →every 10 epochs. You can change this

# logs
Train_Losses = [] # record average training loss each epoch
Val_Losses = [] # record average validation loss each epoch
total_step = len(train_data_loader) # number of iterations each epoch
best_val_loss = np.inf

# start training
print('Start training...')
import time
tic = time.time()
for epoch in range(num_epochs):
    # for epoch in range(2):
        print('Switch to training...')
        model.train()
        Train_loss_iter = [] # record the the training loss each iteration
        for itr, (images, captions, lengths) in enumerate(train_data_loader):

```

```

#####Your Code#####
model.zero_grad()
images = Variable(images).to(device)
captions = Variable(captions).to(device)
predicted_cap_pack = pack_padded_sequence(captions, lengths,
→batch_first=True)[0]
predicted_cap = model(images, captions, lengths)
predicted_cap = predicted_cap.view(-1, vocab_size)
loss = criterion(input = predicted_cap, target=predicted_cap_pack)
loss.backward()
optimizer.step()

# record the training loss
Train_Losses.append(float(loss))

# print log info
if itr % log_step == 0:
    # print current loss and perplexity
    print('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}, Perplexity: {:.
→4f}')
    .format(epoch, num_epochs, itr, total_step, loss.item(),
→np.exp(loss.item()))
    scheduler.step()
    Train_Losses.append(np.mean(Train_loss_iter))
    np.save(os.path.join(output_dir, 'TrainingLoss_lstm.npy'), Train_Losses) #
→save the training loss

model.eval()
# (optional) generate a sample during the training, you can use
→deterministic mode
# Your code

# validation
Val_Losses.append(val(model, val_data_loader, vocab))
np.save(os.path.join(output_dir, 'ValLoss_lstm.npy'), Val_Losses) # save the
→val loss

# save model
if Val_Losses[-1] < best_val_loss:
    best_val_loss = Val_Losses[-1]
    print('updated best val loss:', best_val_loss)
    print('Save model weights to...', output_dir)
    torch.save(model.state_dict(),
    os.path.join(output_dir, 'pretrain-best.pth'.format(epoch +
→1, itr + 1)))

```

```
print('It took: {} s'.format(time.time() - tic))
```

```
/content/drive/My Drive/DL_Fall_2020/Assignment_4/wm_06.npy
Start training...
Switch to training...
Epoch [0/50], Step [0/114], Loss: 9.1967, Perplexity: 9864.4292
Epoch [0/50], Step [50/114], Loss: 3.5588, Perplexity: 35.1226
Epoch [0/50], Step [100/114], Loss: 3.2991, Perplexity: 27.0884

/usr/local/lib/python3.6/dist-packages/numpy/core/fromnumeric.py:3335:
RuntimeWarning: Mean of empty slice.
  out=out, **kwargs)
/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:161:
RuntimeWarning: invalid value encountered in double_scalars
  ret = ret.dtype.type(ret / rcount)

Validating...
Step [0/127], Loss: 3.3755, Perplexity: 29.2382
Step [50/127], Loss: 3.2646, Perplexity: 26.1707
Step [100/127], Loss: 3.6123, Perplexity: 37.0526
updated best val loss: 3.3000397306727614
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
Switch to training...
Epoch [1/50], Step [0/114], Loss: 3.1994, Perplexity: 24.5174
Epoch [1/50], Step [50/114], Loss: 3.1185, Perplexity: 22.6130
Epoch [1/50], Step [100/114], Loss: 3.1364, Perplexity: 23.0215
Validating...
Step [0/127], Loss: 2.6516, Perplexity: 14.1773
Step [50/127], Loss: 3.0158, Perplexity: 20.4056
Step [100/127], Loss: 2.5862, Perplexity: 13.2787
updated best val loss: 3.1186115685410387
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
Switch to training...
Epoch [2/50], Step [0/114], Loss: 3.1360, Perplexity: 23.0118
Epoch [2/50], Step [50/114], Loss: 3.1355, Perplexity: 22.9998
Epoch [2/50], Step [100/114], Loss: 3.1685, Perplexity: 23.7714
Validating...
Step [0/127], Loss: 2.8438, Perplexity: 17.1813
Step [50/127], Loss: 2.9001, Perplexity: 18.1763
Step [100/127], Loss: 3.2208, Perplexity: 25.0483
updated best val loss: 3.0789894832400826
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
Switch to training...
Epoch [3/50], Step [0/114], Loss: 3.0870, Perplexity: 21.9111
```

Epoch [3/50], Step [50/114], Loss: 3.0810, Perplexity: 21.7810
 Epoch [3/50], Step [100/114], Loss: 2.9672, Perplexity: 19.4374
 Validating...
 Step [0/127], Loss: 3.0122, Perplexity: 20.3331
 Step [50/127], Loss: 2.8196, Perplexity: 16.7701
 Step [100/127], Loss: 2.8023, Perplexity: 16.4828
 updated best val loss: 3.0533739037401095
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [4/50], Step [0/114], Loss: 2.9306, Perplexity: 18.7392
 Epoch [4/50], Step [50/114], Loss: 2.9195, Perplexity: 18.5326
 Epoch [4/50], Step [100/114], Loss: 2.9324, Perplexity: 18.7724
 Validating...
 Step [0/127], Loss: 2.8991, Perplexity: 18.1572
 Step [50/127], Loss: 3.1722, Perplexity: 23.8595
 Step [100/127], Loss: 2.9768, Perplexity: 19.6242
 updated best val loss: 3.0393130779266357
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [5/50], Step [0/114], Loss: 2.9816, Perplexity: 19.7193
 Epoch [5/50], Step [50/114], Loss: 2.7993, Perplexity: 16.4329
 Epoch [5/50], Step [100/114], Loss: 2.8397, Perplexity: 17.1100
 Validating...
 Step [0/127], Loss: 2.4462, Perplexity: 11.5444
 Step [50/127], Loss: 3.1410, Perplexity: 23.1261
 Step [100/127], Loss: 2.8435, Perplexity: 17.1750
 updated best val loss: 2.9582198060403657
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [6/50], Step [0/114], Loss: 2.8194, Perplexity: 16.7661
 Epoch [6/50], Step [50/114], Loss: 2.8207, Perplexity: 16.7882
 Epoch [6/50], Step [100/114], Loss: 2.8846, Perplexity: 17.8965
 Validating...
 Step [0/127], Loss: 3.6719, Perplexity: 39.3281
 Step [50/127], Loss: 2.5758, Perplexity: 13.1415
 Step [100/127], Loss: 3.0980, Perplexity: 22.1525
 updated best val loss: 2.9457708137241876
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [7/50], Step [0/114], Loss: 2.7077, Perplexity: 14.9950
 Epoch [7/50], Step [50/114], Loss: 2.7242, Perplexity: 15.2441
 Epoch [7/50], Step [100/114], Loss: 2.8001, Perplexity: 16.4469
 Validating...
 Step [0/127], Loss: 2.4566, Perplexity: 11.6654

Step [50/127], Loss: 3.0743, Perplexity: 21.6348
 Step [100/127], Loss: 2.8221, Perplexity: 16.8114
 updated best val loss: 2.9399302306137685
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [8/50], Step [0/114], Loss: 2.7430, Perplexity: 15.5332
 Epoch [8/50], Step [50/114], Loss: 2.7839, Perplexity: 16.1814
 Epoch [8/50], Step [100/114], Loss: 2.8609, Perplexity: 17.4771
 Validating...
 Step [0/127], Loss: 2.8387, Perplexity: 17.0928
 Step [50/127], Loss: 3.3202, Perplexity: 27.6664
 Step [100/127], Loss: 3.3835, Perplexity: 29.4732
 Switch to training...
 Epoch [9/50], Step [0/114], Loss: 2.7565, Perplexity: 15.7444
 Epoch [9/50], Step [50/114], Loss: 2.7470, Perplexity: 15.5954
 Epoch [9/50], Step [100/114], Loss: 2.7525, Perplexity: 15.6821
 Validating...
 Step [0/127], Loss: 2.7507, Perplexity: 15.6536
 Step [50/127], Loss: 3.3305, Perplexity: 27.9525
 Step [100/127], Loss: 2.7179, Perplexity: 15.1490
 updated best val loss: 2.905920849071713
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [10/50], Step [0/114], Loss: 2.6438, Perplexity: 14.0668
 Epoch [10/50], Step [50/114], Loss: 2.6761, Perplexity: 14.5287
 Epoch [10/50], Step [100/114], Loss: 2.7600, Perplexity: 15.8005
 Validating...
 Step [0/127], Loss: 2.8030, Perplexity: 16.4943
 Step [50/127], Loss: 2.6401, Perplexity: 14.0145
 Step [100/127], Loss: 3.0468, Perplexity: 21.0482
 Switch to training...
 Epoch [11/50], Step [0/114], Loss: 2.6419, Perplexity: 14.0392
 Epoch [11/50], Step [50/114], Loss: 2.6751, Perplexity: 14.5141
 Epoch [11/50], Step [100/114], Loss: 2.7255, Perplexity: 15.2643
 Validating...
 Step [0/127], Loss: 2.9722, Perplexity: 19.5350
 Step [50/127], Loss: 3.0590, Perplexity: 21.3057
 Step [100/127], Loss: 3.1910, Perplexity: 24.3134
 updated best val loss: 2.8773781577433186
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [12/50], Step [0/114], Loss: 2.6225, Perplexity: 13.7697
 Epoch [12/50], Step [50/114], Loss: 2.6798, Perplexity: 14.5820
 Epoch [12/50], Step [100/114], Loss: 2.6560, Perplexity: 14.2389
 Validating...

Step [0/127], Loss: 3.3561, Perplexity: 28.6782
 Step [50/127], Loss: 2.8389, Perplexity: 17.0964
 Step [100/127], Loss: 2.8091, Perplexity: 16.5949
 Switch to training...
 Epoch [13/50], Step [0/114], Loss: 2.6248, Perplexity: 13.8012
 Epoch [13/50], Step [50/114], Loss: 2.7257, Perplexity: 15.2675
 Epoch [13/50], Step [100/114], Loss: 2.6792, Perplexity: 14.5741
 Validating...
 Step [0/127], Loss: 3.5100, Perplexity: 33.4474
 Step [50/127], Loss: 2.8835, Perplexity: 17.8767
 Step [100/127], Loss: 2.8325, Perplexity: 16.9880
 Switch to training...
 Epoch [14/50], Step [0/114], Loss: 2.6107, Perplexity: 13.6083
 Epoch [14/50], Step [50/114], Loss: 2.6608, Perplexity: 14.3076
 Epoch [14/50], Step [100/114], Loss: 2.6752, Perplexity: 14.5146
 Validating...
 Step [0/127], Loss: 2.8089, Perplexity: 16.5909
 Step [50/127], Loss: 2.9213, Perplexity: 18.5658
 Step [100/127], Loss: 3.2196, Perplexity: 25.0183
 Switch to training...
 Epoch [15/50], Step [0/114], Loss: 2.6033, Perplexity: 13.5081
 Epoch [15/50], Step [50/114], Loss: 2.6393, Perplexity: 14.0029
 Epoch [15/50], Step [100/114], Loss: 2.6312, Perplexity: 13.8898
 Validating...
 Step [0/127], Loss: 3.0927, Perplexity: 22.0364
 Step [50/127], Loss: 2.7312, Perplexity: 15.3508
 Step [100/127], Loss: 2.7424, Perplexity: 15.5239
 updated best val loss: 2.8615870062760482
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [16/50], Step [0/114], Loss: 2.6917, Perplexity: 14.7565
 Epoch [16/50], Step [50/114], Loss: 2.5993, Perplexity: 13.4548
 Epoch [16/50], Step [100/114], Loss: 2.6174, Perplexity: 13.6996
 Validating...
 Step [0/127], Loss: 2.6297, Perplexity: 13.8692
 Step [50/127], Loss: 2.5046, Perplexity: 12.2384
 Step [100/127], Loss: 2.9687, Perplexity: 19.4671
 Switch to training...
 Epoch [17/50], Step [0/114], Loss: 2.5535, Perplexity: 12.8521
 Epoch [17/50], Step [50/114], Loss: 2.5234, Perplexity: 12.4712
 Epoch [17/50], Step [100/114], Loss: 2.6090, Perplexity: 13.5860
 Validating...
 Step [0/127], Loss: 3.1692, Perplexity: 23.7894
 Step [50/127], Loss: 2.6280, Perplexity: 13.8458
 Step [100/127], Loss: 2.5929, Perplexity: 13.3689
 Switch to training...
 Epoch [18/50], Step [0/114], Loss: 2.6862, Perplexity: 14.6758

Epoch [18/50], Step [50/114], Loss: 2.5974, Perplexity: 13.4294
 Epoch [18/50], Step [100/114], Loss: 2.5548, Perplexity: 12.8693
 Validating...
 Step [0/127], Loss: 2.9099, Perplexity: 18.3545
 Step [50/127], Loss: 3.0089, Perplexity: 20.2643
 Step [100/127], Loss: 3.1570, Perplexity: 23.4994
 Switch to training...
 Epoch [19/50], Step [0/114], Loss: 2.5982, Perplexity: 13.4394
 Epoch [19/50], Step [50/114], Loss: 2.6393, Perplexity: 14.0036
 Epoch [19/50], Step [100/114], Loss: 2.5786, Perplexity: 13.1787
 Validating...
 Step [0/127], Loss: 2.1423, Perplexity: 8.5190
 Step [50/127], Loss: 2.8382, Perplexity: 17.0849
 Step [100/127], Loss: 2.8037, Perplexity: 16.5063
 Switch to training...
 Epoch [20/50], Step [0/114], Loss: 2.5285, Perplexity: 12.5352
 Epoch [20/50], Step [50/114], Loss: 2.6854, Perplexity: 14.6647
 Epoch [20/50], Step [100/114], Loss: 2.5220, Perplexity: 12.4537
 Validating...
 Step [0/127], Loss: 2.6534, Perplexity: 14.2019
 Step [50/127], Loss: 2.8397, Perplexity: 17.1105
 Step [100/127], Loss: 3.1112, Perplexity: 22.4485
 updated best val loss: 2.8494993570282703
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [21/50], Step [0/114], Loss: 2.5930, Perplexity: 13.3705
 Epoch [21/50], Step [50/114], Loss: 2.5899, Perplexity: 13.3279
 Epoch [21/50], Step [100/114], Loss: 2.6301, Perplexity: 13.8748
 Validating...
 Step [0/127], Loss: 3.2296, Perplexity: 25.2692
 Step [50/127], Loss: 2.7239, Perplexity: 15.2389
 Step [100/127], Loss: 3.6257, Perplexity: 37.5510
 Switch to training...
 Epoch [22/50], Step [0/114], Loss: 2.4671, Perplexity: 11.7888
 Epoch [22/50], Step [50/114], Loss: 2.6088, Perplexity: 13.5831
 Epoch [22/50], Step [100/114], Loss: 2.5855, Perplexity: 13.2698
 Validating...
 Step [0/127], Loss: 3.0787, Perplexity: 21.7299
 Step [50/127], Loss: 2.5891, Perplexity: 13.3179
 Step [100/127], Loss: 3.0095, Perplexity: 20.2779
 Switch to training...
 Epoch [23/50], Step [0/114], Loss: 2.5248, Perplexity: 12.4884
 Epoch [23/50], Step [50/114], Loss: 2.6101, Perplexity: 13.5997
 Epoch [23/50], Step [100/114], Loss: 2.5515, Perplexity: 12.8265
 Validating...
 Step [0/127], Loss: 2.6921, Perplexity: 14.7632
 Step [50/127], Loss: 2.8838, Perplexity: 17.8828

Step [100/127], Loss: 3.6494, Perplexity: 38.4513
 Switch to training...
 Epoch [24/50], Step [0/114], Loss: 2.6055, Perplexity: 13.5383
 Epoch [24/50], Step [50/114], Loss: 2.5914, Perplexity: 13.3489
 Epoch [24/50], Step [100/114], Loss: 2.5800, Perplexity: 13.1974
 Validating...
 Step [0/127], Loss: 2.8379, Perplexity: 17.0796
 Step [50/127], Loss: 2.8680, Perplexity: 17.6016
 Step [100/127], Loss: 3.2286, Perplexity: 25.2442
 Switch to training...
 Epoch [25/50], Step [0/114], Loss: 2.5142, Perplexity: 12.3572
 Epoch [25/50], Step [50/114], Loss: 2.5643, Perplexity: 12.9916
 Epoch [25/50], Step [100/114], Loss: 2.5188, Perplexity: 12.4142
 Validating...
 Step [0/127], Loss: 3.7119, Perplexity: 40.9306
 Step [50/127], Loss: 2.5537, Perplexity: 12.8551
 Step [100/127], Loss: 2.8592, Perplexity: 17.4476
 Switch to training...
 Epoch [26/50], Step [0/114], Loss: 2.5700, Perplexity: 13.0661
 Epoch [26/50], Step [50/114], Loss: 2.5030, Perplexity: 12.2194
 Epoch [26/50], Step [100/114], Loss: 2.5886, Perplexity: 13.3112
 Validating...
 Step [0/127], Loss: 2.7194, Perplexity: 15.1710
 Step [50/127], Loss: 2.6673, Perplexity: 14.4004
 Step [100/127], Loss: 2.6868, Perplexity: 14.6846
 Switch to training...
 Epoch [27/50], Step [0/114], Loss: 2.5567, Perplexity: 12.8930
 Epoch [27/50], Step [50/114], Loss: 2.4836, Perplexity: 11.9847
 Epoch [27/50], Step [100/114], Loss: 2.5198, Perplexity: 12.4255
 Validating...
 Step [0/127], Loss: 2.5743, Perplexity: 13.1221
 Step [50/127], Loss: 3.0380, Perplexity: 20.8637
 Step [100/127], Loss: 2.9892, Perplexity: 19.8697
 Switch to training...
 Epoch [28/50], Step [0/114], Loss: 2.5173, Perplexity: 12.3957
 Epoch [28/50], Step [50/114], Loss: 2.6102, Perplexity: 13.6023
 Epoch [28/50], Step [100/114], Loss: 2.6124, Perplexity: 13.6314
 Validating...
 Step [0/127], Loss: 2.5525, Perplexity: 12.8398
 Step [50/127], Loss: 2.5027, Perplexity: 12.2155
 Step [100/127], Loss: 3.1519, Perplexity: 23.3793
 Switch to training...
 Epoch [29/50], Step [0/114], Loss: 2.5095, Perplexity: 12.2983
 Epoch [29/50], Step [50/114], Loss: 2.5590, Perplexity: 12.9225
 Epoch [29/50], Step [100/114], Loss: 2.5575, Perplexity: 12.9029
 Validating...
 Step [0/127], Loss: 3.0368, Perplexity: 20.8380
 Step [50/127], Loss: 2.7998, Perplexity: 16.4412

Step [100/127], Loss: 3.1684, Perplexity: 23.7702
 Switch to training...
 Epoch [30/50], Step [0/114], Loss: 2.5088, Perplexity: 12.2903
 Epoch [30/50], Step [50/114], Loss: 2.5926, Perplexity: 13.3642
 Epoch [30/50], Step [100/114], Loss: 2.5754, Perplexity: 13.1360
 Validating...
 Step [0/127], Loss: 2.4757, Perplexity: 11.8902
 Step [50/127], Loss: 2.7512, Perplexity: 15.6612
 Step [100/127], Loss: 3.1640, Perplexity: 23.6650
 Switch to training...
 Epoch [31/50], Step [0/114], Loss: 2.5388, Perplexity: 12.6647
 Epoch [31/50], Step [50/114], Loss: 2.5552, Perplexity: 12.8742
 Epoch [31/50], Step [100/114], Loss: 2.5697, Perplexity: 13.0614
 Validating...
 Step [0/127], Loss: 2.8812, Perplexity: 17.8355
 Step [50/127], Loss: 2.6630, Perplexity: 14.3393
 Step [100/127], Loss: 2.9449, Perplexity: 19.0091
 Switch to training...
 Epoch [32/50], Step [0/114], Loss: 2.5897, Perplexity: 13.3252
 Epoch [32/50], Step [50/114], Loss: 2.5348, Perplexity: 12.6142
 Epoch [32/50], Step [100/114], Loss: 2.5883, Perplexity: 13.3073
 Validating...
 Step [0/127], Loss: 3.3178, Perplexity: 27.6003
 Step [50/127], Loss: 2.8602, Perplexity: 17.4645
 Step [100/127], Loss: 2.7902, Perplexity: 16.2846
 Switch to training...
 Epoch [33/50], Step [0/114], Loss: 2.5153, Perplexity: 12.3705
 Epoch [33/50], Step [50/114], Loss: 2.5228, Perplexity: 12.4637
 Epoch [33/50], Step [100/114], Loss: 2.5394, Perplexity: 12.6720
 Validating...
 Step [0/127], Loss: 2.7785, Perplexity: 16.0941
 Step [50/127], Loss: 2.5469, Perplexity: 12.7669
 Step [100/127], Loss: 2.5493, Perplexity: 12.7982
 updated best val loss: 2.8425030783405454
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [34/50], Step [0/114], Loss: 2.4937, Perplexity: 12.1055
 Epoch [34/50], Step [50/114], Loss: 2.6083, Perplexity: 13.5758
 Epoch [34/50], Step [100/114], Loss: 2.5213, Perplexity: 12.4451
 Validating...
 Step [0/127], Loss: 3.2536, Perplexity: 25.8826
 Step [50/127], Loss: 2.7840, Perplexity: 16.1841
 Step [100/127], Loss: 2.1760, Perplexity: 8.8109
 Switch to training...
 Epoch [35/50], Step [0/114], Loss: 2.5229, Perplexity: 12.4652
 Epoch [35/50], Step [50/114], Loss: 2.5738, Perplexity: 13.1150
 Epoch [35/50], Step [100/114], Loss: 2.5409, Perplexity: 12.6906

Validating...
 Step [0/127], Loss: 2.7848, Perplexity: 16.1970
 Step [50/127], Loss: 2.7138, Perplexity: 15.0864
 Step [100/127], Loss: 3.0755, Perplexity: 21.6617
 Switch to training...
 Epoch [36/50], Step [0/114], Loss: 2.5519, Perplexity: 12.8314
 Epoch [36/50], Step [50/114], Loss: 2.5538, Perplexity: 12.8553
 Epoch [36/50], Step [100/114], Loss: 2.4590, Perplexity: 11.6934
 Validating...
 Step [0/127], Loss: 3.1431, Perplexity: 23.1749
 Step [50/127], Loss: 2.9564, Perplexity: 19.2295
 Step [100/127], Loss: 3.1953, Perplexity: 24.4171
 Switch to training...
 Epoch [37/50], Step [0/114], Loss: 2.4676, Perplexity: 11.7940
 Epoch [37/50], Step [50/114], Loss: 2.5541, Perplexity: 12.8596
 Epoch [37/50], Step [100/114], Loss: 2.4754, Perplexity: 11.8864
 Validating...
 Step [0/127], Loss: 3.4164, Perplexity: 30.4602
 Step [50/127], Loss: 2.7267, Perplexity: 15.2822
 Step [100/127], Loss: 2.5885, Perplexity: 13.3098
 Switch to training...
 Epoch [38/50], Step [0/114], Loss: 2.5341, Perplexity: 12.6045
 Epoch [38/50], Step [50/114], Loss: 2.6179, Perplexity: 13.7069
 Epoch [38/50], Step [100/114], Loss: 2.5301, Perplexity: 12.5549
 Validating...
 Step [0/127], Loss: 2.5736, Perplexity: 13.1123
 Step [50/127], Loss: 2.1873, Perplexity: 8.9109
 Step [100/127], Loss: 3.4548, Perplexity: 31.6509
 Switch to training...
 Epoch [39/50], Step [0/114], Loss: 2.5128, Perplexity: 12.3394
 Epoch [39/50], Step [50/114], Loss: 2.5903, Perplexity: 13.3342
 Epoch [39/50], Step [100/114], Loss: 2.5226, Perplexity: 12.4608
 Validating...
 Step [0/127], Loss: 2.6632, Perplexity: 14.3418
 Step [50/127], Loss: 3.2896, Perplexity: 26.8325
 Step [100/127], Loss: 2.6255, Perplexity: 13.8117
 updated best val loss: 2.841626351273905
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [40/50], Step [0/114], Loss: 2.5148, Perplexity: 12.3647
 Epoch [40/50], Step [50/114], Loss: 2.5257, Perplexity: 12.4990
 Epoch [40/50], Step [100/114], Loss: 2.5529, Perplexity: 12.8449
 Validating...
 Step [0/127], Loss: 2.6444, Perplexity: 14.0757
 Step [50/127], Loss: 3.0954, Perplexity: 22.0967
 Step [100/127], Loss: 3.5971, Perplexity: 36.4925
 Switch to training...

Epoch [41/50], Step [0/114], Loss: 2.5291, Perplexity: 12.5422
 Epoch [41/50], Step [50/114], Loss: 2.5528, Perplexity: 12.8430
 Epoch [41/50], Step [100/114], Loss: 2.5882, Perplexity: 13.3058
 Validating...
 Step [0/127], Loss: 2.9757, Perplexity: 19.6034
 Step [50/127], Loss: 3.1688, Perplexity: 23.7792
 Step [100/127], Loss: 2.7754, Perplexity: 16.0444
 Switch to training...
 Epoch [42/50], Step [0/114], Loss: 2.5562, Perplexity: 12.8869
 Epoch [42/50], Step [50/114], Loss: 2.5088, Perplexity: 12.2906
 Epoch [42/50], Step [100/114], Loss: 2.5362, Perplexity: 12.6322
 Validating...
 Step [0/127], Loss: 2.8588, Perplexity: 17.4413
 Step [50/127], Loss: 2.8259, Perplexity: 16.8769
 Step [100/127], Loss: 2.4660, Perplexity: 11.7750
 Switch to training...
 Epoch [43/50], Step [0/114], Loss: 2.5374, Perplexity: 12.6468
 Epoch [43/50], Step [50/114], Loss: 2.4976, Perplexity: 12.1531
 Epoch [43/50], Step [100/114], Loss: 2.5057, Perplexity: 12.2525
 Validating...
 Step [0/127], Loss: 2.5595, Perplexity: 12.9295
 Step [50/127], Loss: 2.8840, Perplexity: 17.8862
 Step [100/127], Loss: 3.1887, Perplexity: 24.2559
 Switch to training...
 Epoch [44/50], Step [0/114], Loss: 2.5095, Perplexity: 12.2982
 Epoch [44/50], Step [50/114], Loss: 2.5137, Perplexity: 12.3509
 Epoch [44/50], Step [100/114], Loss: 2.5183, Perplexity: 12.4080
 Validating...
 Step [0/127], Loss: 2.6764, Perplexity: 14.5325
 Step [50/127], Loss: 2.5301, Perplexity: 12.5546
 Step [100/127], Loss: 3.7817, Perplexity: 43.8896
 updated best val loss: 2.8387038182085895
 Save model weights to... /content/drive/My
 Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
 Switch to training...
 Epoch [45/50], Step [0/114], Loss: 2.5692, Perplexity: 13.0556
 Epoch [45/50], Step [50/114], Loss: 2.5030, Perplexity: 12.2196
 Epoch [45/50], Step [100/114], Loss: 2.5596, Perplexity: 12.9306
 Validating...
 Step [0/127], Loss: 2.7298, Perplexity: 15.3296
 Step [50/127], Loss: 3.4381, Perplexity: 31.1268
 Step [100/127], Loss: 3.2502, Perplexity: 25.7960
 Switch to training...
 Epoch [46/50], Step [0/114], Loss: 2.5979, Perplexity: 13.4351
 Epoch [46/50], Step [50/114], Loss: 2.5236, Perplexity: 12.4735
 Epoch [46/50], Step [100/114], Loss: 2.4594, Perplexity: 11.6974
 Validating...
 Step [0/127], Loss: 2.8805, Perplexity: 17.8226

```

Step [50/127], Loss: 3.0389, Perplexity: 20.8820
Step [100/127], Loss: 3.4368, Perplexity: 31.0889
Switch to training...
Epoch [47/50], Step [0/114], Loss: 2.6026, Perplexity: 13.4992
Epoch [47/50], Step [50/114], Loss: 2.5164, Perplexity: 12.3841
Epoch [47/50], Step [100/114], Loss: 2.5802, Perplexity: 13.1992
Validating...
Step [0/127], Loss: 2.9822, Perplexity: 19.7320
Step [50/127], Loss: 2.6580, Perplexity: 14.2675
Step [100/127], Loss: 2.8360, Perplexity: 17.0474
Switch to training...
Epoch [48/50], Step [0/114], Loss: 2.5064, Perplexity: 12.2603
Epoch [48/50], Step [50/114], Loss: 2.5741, Perplexity: 13.1199
Epoch [48/50], Step [100/114], Loss: 2.5717, Perplexity: 13.0877
Validating...
Step [0/127], Loss: 2.8604, Perplexity: 17.4687
Step [50/127], Loss: 2.5709, Perplexity: 13.0776
Step [100/127], Loss: 3.0426, Perplexity: 20.9605
Switch to training...
Epoch [49/50], Step [0/114], Loss: 2.5228, Perplexity: 12.4639
Epoch [49/50], Step [50/114], Loss: 2.5407, Perplexity: 12.6885
Epoch [49/50], Step [100/114], Loss: 2.5666, Perplexity: 13.0216
Validating...
Step [0/127], Loss: 2.7909, Perplexity: 16.2962
Step [50/127], Loss: 2.9607, Perplexity: 19.3119
Step [100/127], Loss: 3.2011, Perplexity: 24.5585
updated best val loss: 2.8088028768854816
Save model weights to... /content/drive/My
Drive/DL_Fall_2020/Assignment_4//checkpoints/pretrained_emb/
It took: 14552.241356372833 s

```

13.10 Section 3.2.3 Evaluation [3 pts]

```

[ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

train_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    ↪split='train', vocab=vocab,
                                transform=transform, batch_size=batch_size,
    ↪shuffle=True, num_workers=12)
val_data_loader = get_loader(root=path_to_homework + '/flickr30k_images/',
    ↪split='val', vocab=vocab,
                                transform=transform, batch_size=8, shuffle=True,
    ↪num_workers=4)

model = Word_embeddings(vocab_size=len(vocab), emb_dim=emb_dim,
    ↪hidden_dim=hidden_dim,

```

```

num_layers=1, dropout=dropout).to(device) # build a model

vocab_size=len(vocab)

# loss and optimizer
criterion = nn.CrossEntropyLoss().to(device) # CE loss
optimizer = torch.optim.Adam(model.parameters(), lr=lr,
    ↳weight_decay=weight_decay) # optimizer
scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
    step_size=5,
    gamma=0.5) # decay LR by a factor of 0.5
    ↳every 10 epochs. You can change this

```

```

[29]: # ## Evaluate your model using BLEU score. Use Deterministic mode
model.eval()
bleu1, bleu2, bleu3, bleu4 = evaluation(model, vocab, mode='Deterministic')
print("BLEU 1:{}, BLEU 2:{}, BLEU 3:{}, BLEU 4:{}"
    ↳.format(bleu1, bleu2, bleu3,
    ↳bleu4))

```

```

0%|          | 0/125 [00:00<?, ?it/s]

```

Run on the test set...

```

1%|          | 1/125 [00:00<01:23, 1.48it/s]
2%|          | 2/125 [00:00<01:03, 1.95it/s]
2%|          | 3/125 [00:00<00:48, 2.53it/s]
3%|          | 4/125 [00:01<00:39, 3.10it/s]
4%|          | 5/125 [00:01<00:32, 3.70it/s]
6%|          | 7/125 [00:01<00:24, 4.73it/s]
6%|          | 8/125 [00:01<00:21, 5.45it/s]
7%|          | 9/125 [00:01<00:19, 5.90it/s]
9%|          | 11/125 [00:01<00:15, 7.26it/s]
10%|         | 13/125 [00:01<00:14, 7.61it/s]
12%|         | 15/125 [00:02<00:12, 8.78it/s]
14%|         | 17/125 [00:02<00:12, 8.42it/s]
15%|         | 19/125 [00:02<00:11, 9.54it/s]
17%|         | 21/125 [00:02<00:10, 9.65it/s]
18%|         | 23/125 [00:02<00:10, 9.68it/s]
20%|         | 25/125 [00:03<00:09, 10.25it/s]
22%|         | 27/125 [00:03<00:10, 9.25it/s]
23%|         | 29/125 [00:03<00:09, 10.48it/s]
25%|         | 31/125 [00:03<00:09, 10.17it/s]
26%|         | 33/125 [00:03<00:08, 10.54it/s]
28%|         | 35/125 [00:04<00:08, 10.90it/s]
30%|         | 37/125 [00:04<00:09, 9.39it/s]
31%|         | 39/125 [00:04<00:08, 10.15it/s]

```

33%	41/125	[00:04<00:08, 9.65it/s]
34%	43/125	[00:04<00:07, 10.35it/s]
36%	45/125	[00:05<00:08, 9.63it/s]
38%	47/125	[00:05<00:07, 10.12it/s]
39%	49/125	[00:05<00:07, 10.74it/s]
41%	51/125	[00:05<00:07, 10.54it/s]
42%	53/125	[00:05<00:06, 10.82it/s]
44%	55/125	[00:05<00:05, 11.90it/s]
46%	57/125	[00:06<00:07, 9.68it/s]
47%	59/125	[00:06<00:06, 10.21it/s]
49%	61/125	[00:06<00:06, 10.27it/s]
50%	63/125	[00:06<00:06, 10.30it/s]
52%	65/125	[00:07<00:06, 9.78it/s]
54%	67/125	[00:07<00:05, 10.89it/s]
55%	69/125	[00:07<00:05, 10.39it/s]
57%	71/125	[00:07<00:05, 9.76it/s]
58%	73/125	[00:07<00:04, 10.98it/s]
60%	75/125	[00:08<00:04, 10.01it/s]
62%	77/125	[00:08<00:04, 9.96it/s]
63%	79/125	[00:08<00:04, 10.19it/s]
65%	81/125	[00:08<00:04, 10.58it/s]
66%	83/125	[00:08<00:03, 10.56it/s]
68%	85/125	[00:09<00:04, 9.80it/s]
70%	87/125	[00:09<00:03, 10.16it/s]
71%	89/125	[00:09<00:03, 9.86it/s]
73%	91/125	[00:09<00:03, 10.40it/s]
74%	93/125	[00:09<00:03, 10.31it/s]
76%	95/125	[00:09<00:02, 10.62it/s]
78%	97/125	[00:10<00:02, 10.30it/s]
79%	99/125	[00:10<00:02, 9.66it/s]
80%	100/125	[00:10<00:02, 9.60it/s]
82%	102/125	[00:10<00:02, 10.13it/s]
83%	104/125	[00:10<00:02, 10.04it/s]
85%	106/125	[00:11<00:01, 10.33it/s]
86%	108/125	[00:11<00:01, 10.41it/s]
88%	110/125	[00:11<00:01, 10.49it/s]
90%	112/125	[00:11<00:01, 10.67it/s]
91%	114/125	[00:11<00:01, 10.32it/s]
93%	116/125	[00:12<00:00, 10.26it/s]
94%	118/125	[00:12<00:00, 11.33it/s]
96%	120/125	[00:12<00:00, 12.23it/s]
98%	122/125	[00:12<00:00, 13.78it/s]
100%	125/125	[00:12<00:00, 9.91it/s]
0%	0/1000	[00:00<?, ?it/s]
2%	22/1000	[00:00<00:04, 217.05it/s]

Computing BLEU

5%	47/1000	[00:00<00:04, 224.16it/s]
7%	71/1000	[00:00<00:04, 226.82it/s]
10%	98/1000	[00:00<00:03, 236.23it/s]
12%	123/1000	[00:00<00:03, 240.08it/s]
15%	150/1000	[00:00<00:03, 247.49it/s]
17%	174/1000	[00:00<00:03, 243.82it/s]
20%	197/1000	[00:00<00:03, 236.04it/s]
22%	224/1000	[00:00<00:03, 243.40it/s]
25%	248/1000	[00:01<00:03, 236.40it/s]
27%	272/1000	[00:01<00:03, 229.24it/s]
30%	296/1000	[00:01<00:03, 231.52it/s]
32%	321/1000	[00:01<00:02, 236.00it/s]
35%	346/1000	[00:01<00:02, 237.69it/s]
37%	370/1000	[00:01<00:02, 235.87it/s]
39%	394/1000	[00:01<00:02, 236.87it/s]
42%	419/1000	[00:01<00:02, 239.26it/s]
44%	444/1000	[00:01<00:02, 241.70it/s]
47%	469/1000	[00:01<00:02, 228.49it/s]
50%	496/1000	[00:02<00:02, 238.14it/s]
52%	521/1000	[00:02<00:02, 236.08it/s]
55%	545/1000	[00:02<00:01, 232.18it/s]
57%	569/1000	[00:02<00:01, 230.30it/s]
59%	593/1000	[00:02<00:01, 226.53it/s]
62%	616/1000	[00:02<00:01, 226.29it/s]
64%	640/1000	[00:02<00:01, 228.02it/s]
67%	666/1000	[00:02<00:01, 233.79it/s]
69%	690/1000	[00:02<00:01, 223.42it/s]
72%	715/1000	[00:03<00:01, 229.16it/s]
74%	739/1000	[00:03<00:01, 223.34it/s]
76%	762/1000	[00:03<00:01, 220.82it/s]
79%	786/1000	[00:03<00:00, 225.45it/s]
81%	809/1000	[00:03<00:00, 226.11it/s]
83%	832/1000	[00:03<00:00, 214.77it/s]
86%	855/1000	[00:03<00:00, 217.88it/s]
88%	880/1000	[00:03<00:00, 224.95it/s]
90%	904/1000	[00:03<00:00, 227.34it/s]
93%	927/1000	[00:04<00:00, 223.37it/s]
95%	952/1000	[00:04<00:00, 229.07it/s]
98%	976/1000	[00:04<00:00, 226.87it/s]
100%	1000/1000	[00:04<00:00, 230.83it/s]

BLEU 1:69.3856560848057, BLEU 2:45.79195915761376, BLEU 3:26.821433226826393,
BLEU 4:17.211552439590815

13.11 Section 3.2.4 Discussion [2 pts]

Compared to index embeddings, do pretrained embeddings improve the performance? Try to explain it.

Your Comments: It seems that pretrained embeddings lead to a faster training and a better test results. I think this is because the model can obtain more semantic signals from pretrained embeddings than training data through the index embeddings.

13.11.1 Guidelines for Downloading PDF in Google Colab

- Run below cells only in Google Colab, Comment out in case of Jupyter notebook

```
[ ]: #Run below two lines (in google colab), installation steps to get .pdf of the ↵  
      ↪notebook
```

```
!apt-get install texlive texlive-xetex texlive-latex-extra pandoc  
!pip install py pandoc
```

```
# After installation, comment above two lines and run again to remove ↵  
      ↪installation comments from the notebook.
```

```
[4]: # Find path to your notebook file in drive and enter in below line
```

```
!jupyter nbconvert --to PDF "/content/drive/My Drive/DL_Fall_2020/Assignment_4/  
      ↪Assignment_4.ipynb"
```

```
#Example: "/content/drive/My Drive/DL_Fall_2020/Assignment_4/DL_Assignment_4.  
      ↪ipynb"
```