

Assignment_5_Code_906213559

November 21, 2020

1 ECE-6524 / CS-6524 Deep Learning

2 Assignment 5 (110 pts)

In this assignment, **you will implement the following:** 1. GAN / LSGAN loss functions and training code for MNIST dataset. 2. Train GAN / LSGAN on CelebA dataset with DCGAN architecture.

Again, it may take hours to train on CelebA. So start early. (We don't want you to work during holidays)

2.1 Submission guideline for the coding part (Jupyter Notebook)

1. Click the Save button at the top of the Jupyter Notebook
2. Please make sure to have entered your Virginia Tech PID below
3. Once you've completed everything (make sure output for all cells are visible), select File -> Download as -> PDF via LaTeX
4. Look at the PDF file and make sure all your solutions are displayed correctly there
5. Zip this notebook (Please don't include the data). Name it as Assignment_3_Code_[YOUR PID NUMBER].zip
6. Name your PDF file as Assignment_4_NB_[YOUR PID NUMBER].pdf
7. **Submit your zipped file and the PDF SEPARATELY**

Note: if facing issues with step 3 refer: <https://pypi.org/project/notebook-as-pdf/>

2.2 Submission guideline for the coding part (Google Colab)

1. Click the Save button at the top of the Notebook
2. Please make sure to have entered your Virginia Tech PID below
3. Follow last two cells in this notebook for guidelines to download pdf file of this notebook
4. Look at the PDF file and make sure all your solutions are displayed correctly there
5. Zip this notebook (Please don't include the data). Name it as Assignment_2_Code_[YOUR PID NUMBER].zip
6. Name your PDF file as Assignment_4_NB_[YOUR PID NUMBER].pdf
7. **Submit your zipped file and the PDF SEPARATELY**

While you are encouraged to discuss with your peers, all work submitted is expected to be your own. If you use any information from other resources (e.g. online materials), you are required to cite it below you VT PID. Any violation will result in a 0 mark for the assignment.

2.2.1 Please Write Your VT PID Here: 906213559

2.2.2 Reference (if any):

<https://github.com/Zeleni9/pytorch-wgan>

<https://machinelearningmastery.com/how-to-evaluate-generative-adversarial-networks/>

[https://arxiv.org/abs/1802.03446\(Pros and Cons of GAN Evaluation Measures\)](https://arxiv.org/abs/1802.03446)

2.2.3 Colab Setup:

- Below are some basic steps for colab setup.
- Make changes based on requirements.
- Comment out in case of ARC or your local device with powerful GPU.

```
[ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: import sys  
# modify "path_to_homework", path of folder in drive, where you uploaded your  
→homework files  
path_to_homework = "/content/drive/My Drive/DL_Fall_2020/Assignment_5/"  
sys.path.append(path_to_homework)
```

3 Section 1. Generative Adversarial Networks on MNIST Dataset [54 pts]

In this section, you will need to: 1. Implement two different types of loss functions (GAN / LSGAN) for generative adversarial networks. 2. Build the Discriminator and Generator. 3. Implement training codes for your GAN models. 3. Train your model on MNIST dataset and visualize the generated images.

Now, let's get started!

```
[ ]: import torch  
import torch.nn as nn  
from torchvision import datasets  
from torchvision import transforms  
from torch.utils.data import DataLoader  
from torchvision.datasets import ImageFolder  
  
import numpy as np  
import os  
  
import matplotlib.pyplot as plt  
import matplotlib.gridspec as gridspec
```

```
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

[ ]: def show_images(images, color=False):
    if color:
        sqrtimg = int(np.ceil(np.sqrt(images.shape[2]*images.shape[3])))
    else:
        images = np.reshape(images, [images.shape[0], -1]) # images reshape to (batch_size, D)
        sqrtimg = int(np.ceil(np.sqrt(images.shape[1])))
    sqrtm = int(np.ceil(np.sqrt(images.shape[0])))

    fig = plt.figure(figsize=(sqrtm, sqrtm))
    gs = gridspec.GridSpec(sqrtm, sqrtm)
    gs.update(wspace=0.05, hspace=0.05)

    for i, img in enumerate(images):
        ax = plt.subplot(gs[i])
        plt.axis('off')
        ax.set_xticklabels([])
        ax.set_yticklabels([])
        ax.set_aspect('equal')
        if color:
            plt.imshow(np.swapaxes(np.swapaxes(img, 0, 1), 1, 2))
        else:
            plt.imshow(img.reshape([sqrtimg,sqrtimg]))
    return

def preprocess_img(x):
    return 2 * x - 1.0

def deprocess_img(x):
    return (x + 1.0) / 2.0

def rel_error(x,y):
    return np.max(np.abs(x - y)) / (np.maximum(1e-8, np.abs(x) + np.abs(y)))

def sample_noise(batch_size, dim):
    """
    Generate a PyTorch Tensor of uniform random noise.

```

```

Input:
- batch_size: Integer giving the batch size of noise to generate.
- dim: Integer giving the dimension of noise to generate.

Output:
- A PyTorch Tensor of shape (batch_size, dim) containing uniform
random noise in the range (-1, 1).
"""

to_return = torch.randn((batch_size, dim))
return to_return/torch.max(to_return)

class Flatten(nn.Module):
    def forward(self, x):
        N, C, H, W = x.size() # read in N, C, H, W
        return x.view(N, -1)

```

```
[ ]: # set your device
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

3.1 Section 1.1. Vanilla GAN loss functions

3.1.1 Section 1.1.1 GAN Loss [10 pts]

We start from implementing the vanilla GAN loss from the [original GAN paper](#). Specifically, you need to complete the generator_loss and discriminator_loss in the cell below.

Recalled from the class, the generator loss is written as:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$$

Note that these equations could be slightly different from what we have seen before. This is because that in Pytorch, our optimize will be *minimizing* the loss functions. As a result, we negate the formulas to match pytorch's behavior.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log (1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
[ ]: def discriminator_loss(logits_real, logits_fake):
    """
    Computes the discriminator loss.

    You should use the stable torch.nn.functional.
    →binary_cross_entropy_with_logits
    loss rather than using a separate softmax function followed by the binary
    →cross
    entropy loss.

    Inputs:
    - logits_real: PyTorch Tensor of shape (N,) giving scores for the real data.
    - logits_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.

    Returns:
    - loss: PyTorch Tensor containing (scalar) the loss for the discriminator.
    """
    loss = None

    #####
    #           YOUR CODE HERE           #
    #####
    real = torch.ones_like(logits_real)
    fake = torch.zeros_like(logits_fake)
    loss = torch.nn.functional.binary_cross_entropy_with_logits(logits_real, real,
    →reduction='mean') + torch.nn.functional.
    →binary_cross_entropy_with_logits(logits_fake, fake, reduction='mean')
    #####
    END      #####
    return loss

def generator_loss(logits_fake):
    """
    Computes the generator loss.

    You should use the stable torch.nn.functional.
    →binary_cross_entropy_with_logits
    loss rather than using a separate softmax function followed by the binary
    →cross
    entropy loss.

    Inputs:
    - logits_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.

    Returns:
    - loss: PyTorch Tensor containing the (scalar) loss for the generator.
    """

```

```

loss = None

#####
#           YOUR CODE HERE          #
#####
fake = torch.ones_like(logits_fake)
loss = torch.nn.functional.binary_cross_entropy_with_logits(logits_fake, fake,
→ reduction='mean')
#####
#           END          #####
return loss

```

3.1.2 Section 1.1.2 Build simple model [10 pts]

Build your simple model using below layers:

Discriminator:

- Flatten input (flatten the C x H x W into a single vector per image)
- linear (784, 256)
- Leaky ReLU ()
- Linear(256, 256)
- LeakyReLU()
- Linear(256, 1)

Generator: - Linear(noise_dim, 1024) - ReLU() - Linear(1024, 1024) - ReLU() - linear(1024, 784) - Tanh()

```

[ ]: NOISE_DIM = 100

def discriminator():
    """
    Initialize and return a simple discriminator model.
    """
    # Your code here:
    model = torch.nn.Sequential(Flatten(), torch.nn.Linear(784, 256), torch.nn.
→LeakyReLU(), torch.nn.Linear(256, 256), torch.nn.LeakyReLU(), torch.nn.
→Linear(256, 1))
    return model

def generator(noise_dim=NOISE_DIM):
    """
    Initialize and return a simple generator model.
    """
    # Your code here:
    model = nn.Sequential(torch.nn.Linear(noise_dim, 1024), torch.nn.ReLU(), torch.
→nn.Linear(1024, 1024), torch.nn.ReLU(), torch.nn.Linear(1024, 784), torch.nn.
→Tanh())

```

```
    return model
```

3.1.3 Section 1.1.3 Training code [10 pts]

You can't train a model without a proper training code. Implement the GAN training procedure here following the [original GAN paper](#) and the course slides in the cell below. Note that this code would be reused in the subsequent section, so make sure that it is correctly implemented.

```
[ ]: import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

def train(D, G, D_solver, G_solver, discriminator_loss, generator_loss, show_every=250,
          ckpt_dir = path_to_homework + '/ckpts/Vanilla_gan_MINST/',
          cont_train=False,
          batch_size=128, noise_size=100, num_epochs=10, train_loader=None,
          device=None):
    """
    Train loop for GAN.

    The loop will consist of two steps: a discriminator step and a generator step.
    """

    Train loop for GAN.
```

The loop will consist of two steps: a discriminator step and a generator step.

(1) In the discriminator step, you should zero gradients in the discriminator and sample noise to generate a fake data batch using the generator. Flatten real images

to 784 (28 * 28). Calculate the discriminator output for real and fake data, and use the output to compute discriminator loss. Call backward() on the loss output and take an optimizer step for the discriminator.

(2) For the generator step, you should once again zero gradients in the generator

and sample noise to generate a fake data batch. Get the discriminator output for the fake data batch and use this to compute the generator loss. Once again call backward() on the loss and take an optimizer step.

You will need to reshape the fake image tensor outputted by the generator to be dimensions (batch_size x input_channels x img_size x img_size).

Use the sample_noise function to sample random noise, and the discriminator_loss

and generator_loss functions for their respective loss computations

Inputs:

```

- D, G: PyTorch models for the discriminator and generator
- D_solver, G_solver: torch.optim Optimizers to use for training the
  discriminator and generator.
- discriminator_loss, generator_loss: Functions to use for computing the
  generator and
  discriminator loss, respectively.
- show_every: Show samples after every show_every iterations.
- batch_size: Batch size to use for training.
- noise_size: Dimension of the noise to use as input to the generator.
- num_epochs: Number of epochs over the training dataset to use for training.
- train_loader: image dataloader
- device: PyTorch device
"""

if os.path.exists(os.path.join(ckpt_dir, 'checkpoint.pth')) and cont_train:
    ckpt = torch.load(os.path.join(ckpt_dir, 'checkpoint.pth'))
    start_epoch = ckpt['epoch'] + 1
    iter_count = ckpt['iter_count'] + 1
    G.load_state_dict(ckpt['G'])
    D.load_state_dict(ckpt['D'])
    print('Start from a checkpoint: {}, epoch:{}, iter:{}'.format(os.path.
→join(ckpt_dir, 'checkpoint.pth'),
                                                               str(start_epoch), □
                                                               →str(iter_count)))
else:
    start_epoch = 0
    iter_count = 0

for epoch in range(start_epoch, num_epochs):
    print('EPOCH: ', (epoch+1))
    for x, _ in train_loader:
        _, input_channels, img_size, _ = x.shape

        real_images = preprocess_img(x).to(device) # normalize

        # Store discriminator loss output, generator loss output, and fake
        →image output
        # in these variables for logging and visualization below
        d_error = None
        g_error = None
        fake_images = None

#####
#      Discriminator step      #
#      YOUR CODE HERE          #
#####

D_solver.zero_grad()

```

```

Discriminator = sample_noise(batch_size, noise_size).to(device)
real_logits = D.forward(real_images)
fake_logits = D.forward(G.forward(Discriminator) .
→reshape_as(real_images))
Discriminator_loss = discriminator_loss(real_logits, fake_logits)
Discriminator_loss.backward()
D_solver.step()
d_error = Discriminator_loss.view(-1)
#####
# Generator step #
# YOUR CODE HERE #
#####

G_solver.zero_grad()
Generator = sample_noise(batch_size, noise_size).to(device)
fake_logits = D.forward(G.forward(Generator).reshape_as(real_images))
Generator_loss = generator_loss(fake_logits)
Generator_loss.backward()
G_solver.step()
g_error = Generator_loss.view(-1)
fake_images = G.forward(Generator).reshape_as(real_images)
#####
# END #
#####

if (iter_count % show_every == 0):
    print('Iter: {}, D: {:.4}, G:{:.4}'.format(iter_count,d_error.
→item(),g_error.item()))
    disp_fake_images = deprocess_img(fake_images.data) # denormalize
    imgs_numpy = (disp_fake_images).cpu().numpy()
    show_images(imgs_numpy[0:16], color=input_channels!=1)
    plt.show()
    print()
iter_count += 1

# save checkpoints
os.makedirs(ckpt_dir, exist_ok=True)
print('Saving the model as a checkpoint... ')
torch.save({'epoch': epoch,
            'iter_count': iter_count,
            'G': G.state_dict(),
            'D': D.state_dict()},
            os.path.join(ckpt_dir, 'checkpoint.pth'))

```

MNIST is a simple dataset that contains one hand-written digit in each image. It is usually used for sanity check. So, let's test our loss functions and training code on it!

[]: batch_size = 128

```

mnist = datasets.MNIST('./MNIST_data', train=True, download=True,
                      transform=transforms.ToTensor())
loader_train = DataLoader(mnist, batch_size=batch_size, drop_last=True)

imgs = loader_train.__iter__().next()[0].view(batch_size, 784).numpy().squeeze()
show_images(imgs)

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
./MNIST_data/MNIST/raw/train-images-idx3-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST_data/MNIST/raw/train-images-idx3-ubyte.gz to
./MNIST_data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
./MNIST_data/MNIST/raw/train-labels-idx1-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST_data/MNIST/raw/train-labels-idx1-ubyte.gz to
./MNIST_data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to
./MNIST_data/MNIST/raw/t10k-images-idx3-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

Extracting ./MNIST_data/MNIST/raw/t10k-images-idx3-ubyte.gz to
./MNIST_data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to
./MNIST_data/MNIST/raw/t10k-labels-idx1-ubyte.gz

HBox(children=(FloatProgress(value=1.0, bar_style='info', max=1.0), HTML(value='')))

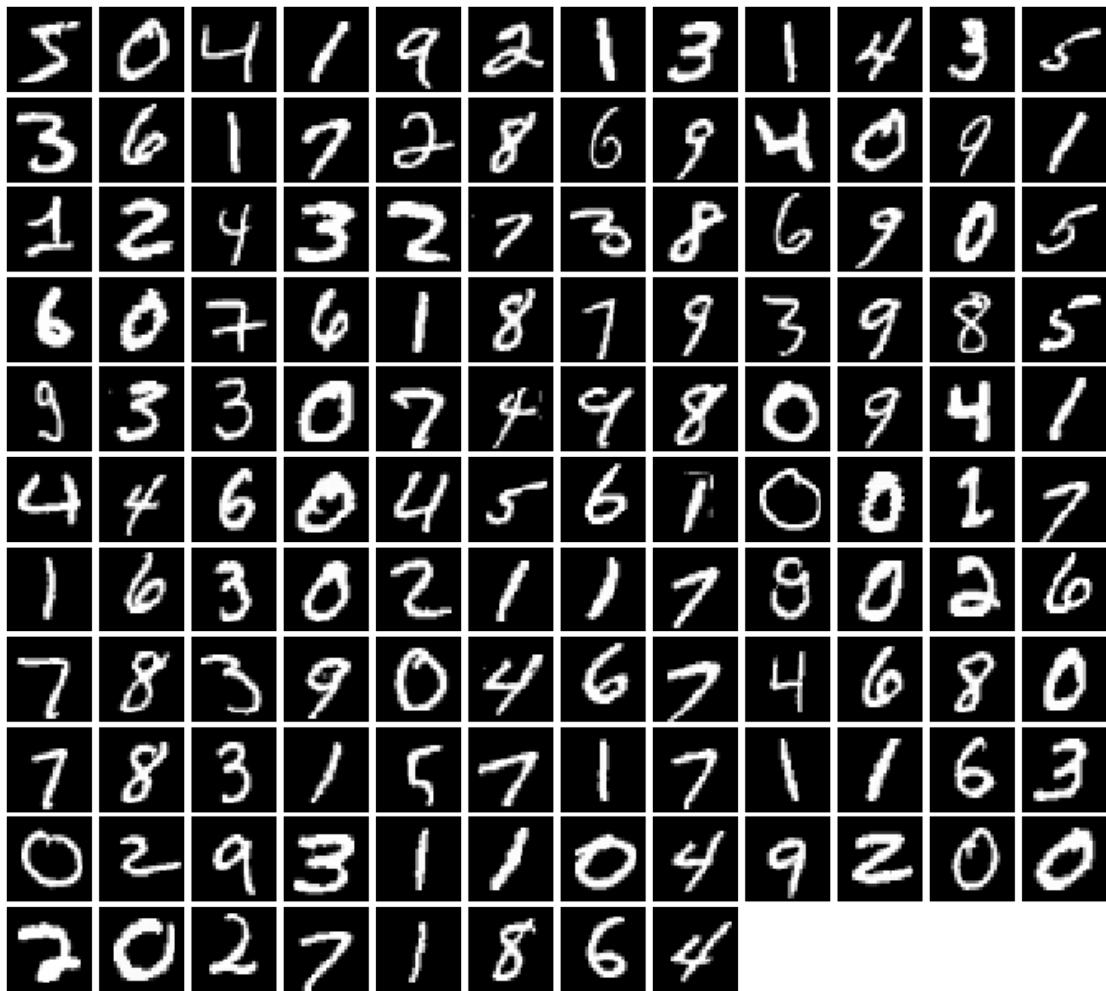
Extracting ./MNIST_data/MNIST/raw/t10k-labels-idx1-ubyte.gz to
./MNIST_data/MNIST/raw

Processing...

Done!

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:480:
UserWarning: The given NumPy array is not writeable, and PyTorch does not
support non-writeable tensors. This means you can write to the underlying
(supposedly non-writeable) NumPy array using the tensor. You may want to copy
the array to protect its data or make it writeable before converting it to a
tensor. This type of warning will be suppressed for the rest of this program.

```
(Triggered internally at  /pytorch/torch/csrc/utils/tensor_numpy.cpp:141.)
return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```



3.1.4 Section 1.1.4 Train your model [7 pts]

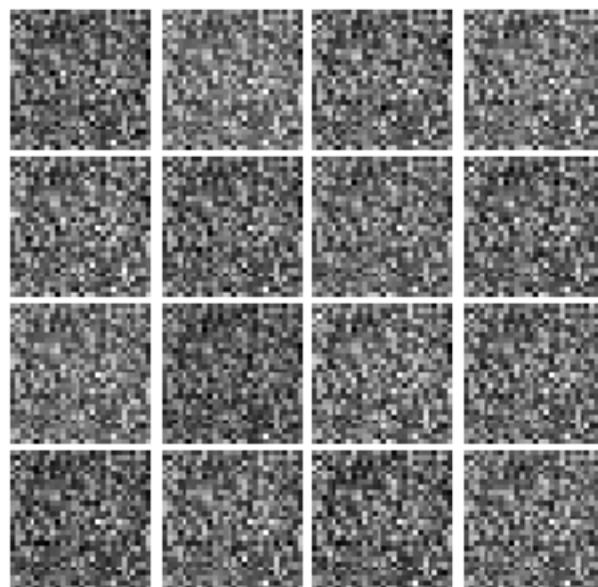
- Call Discriminator and Generator for training.
- Call optimizers for both discriminator and generator for training. (Use Adam with lr=1e-3, betas = (0.5, 0.999))
- Call train function to train.
- Train for 10 epochs.

After training your GAN model, you should expect results that resemble the following if your loss function and training loop implementations are correct:

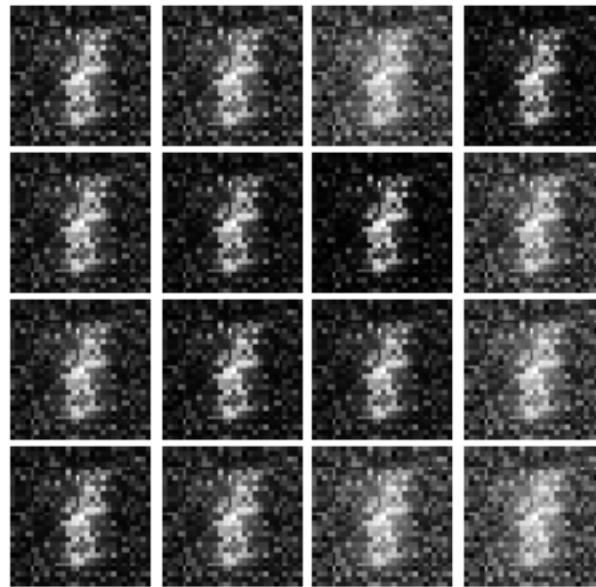
Refer mnist.jpg from gan_samples folder.

```
[ ]: # Add training code here:  
# original GAN  
Discriminator = discriminator().to(device)  
Generator = generator().to(device)  
Discriminator_optimizer = torch.optim.Adam(Discriminator.parameters(), lr=1e-3, betas = (0.5, 0.999))  
Generator_optimizer = torch.optim.Adam(Generator.parameters(), lr=1e-3, betas = (0.5, 0.999))  
train(Discriminator, Generator, Discriminator_optimizer, Generator_optimizer, discriminator_loss, generator_loss, train_loader=loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 1.365, G:0.6973



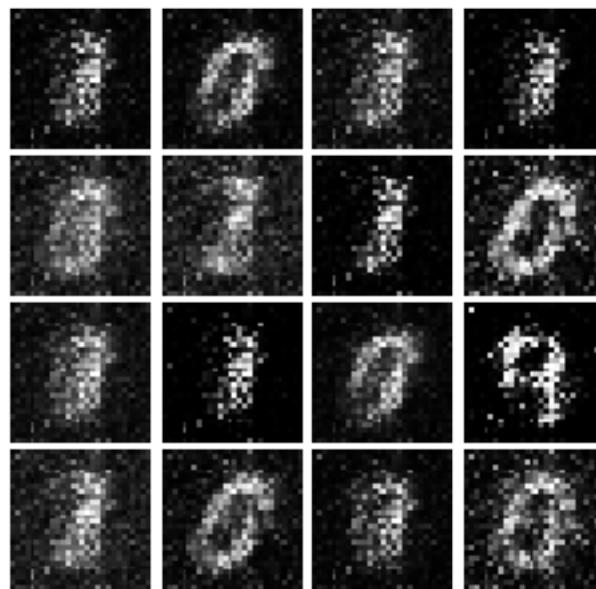
Iter: 250, D: 0.8937, G:0.9792



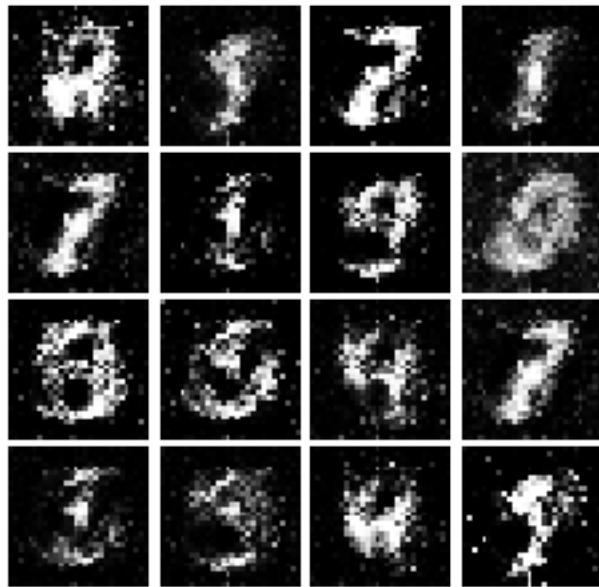
Saving the model as a checkpoint...

EPOCH: 2

Iter: 500, D: 1.257, G:1.247



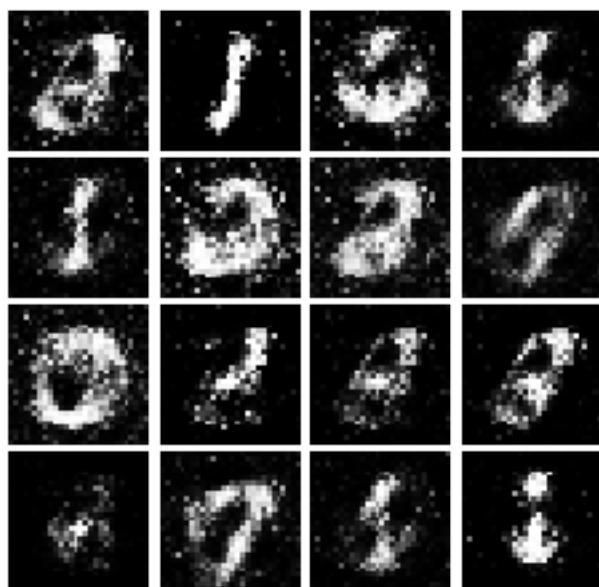
Iter: 750, D: 0.8385, G:1.785



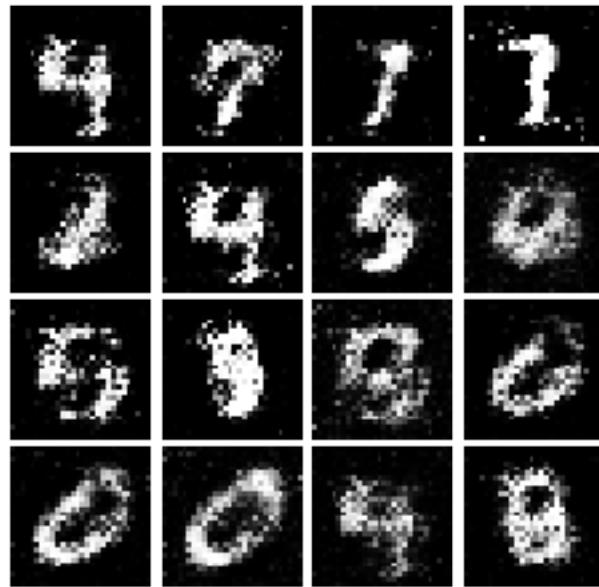
Saving the model as a checkpoint...

EPOCH: 3

Iter: 1000, D: 1.146, G: 1.066



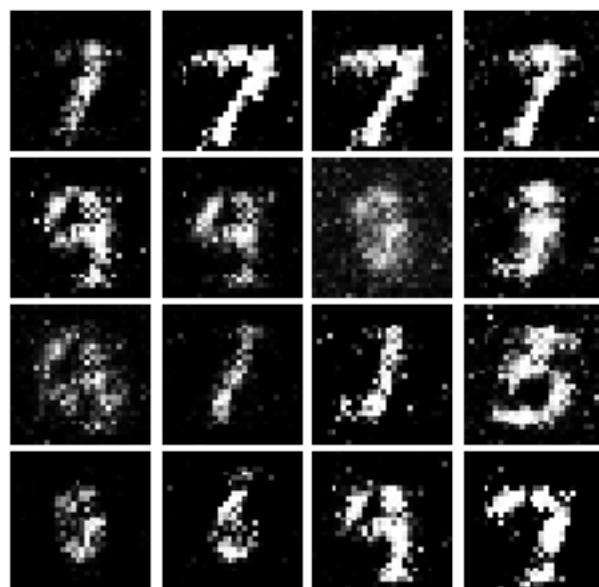
Iter: 1250, D: 1.138, G: 0.9237



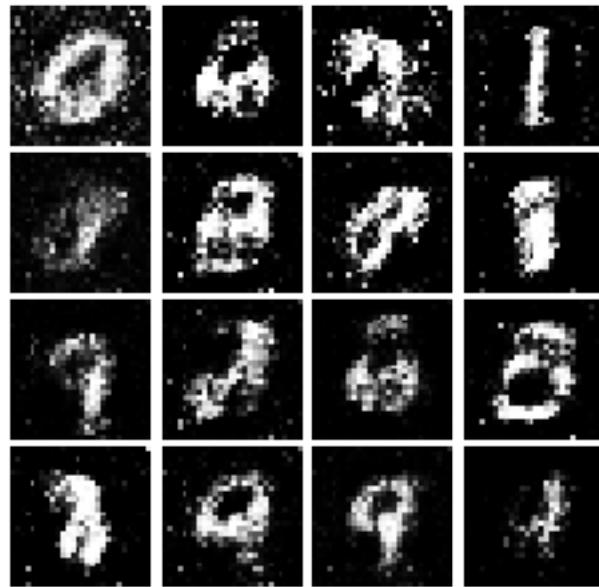
Saving the model as a checkpoint...

EPOCH: 4

Iter: 1500, D: 1.166, G:1.18



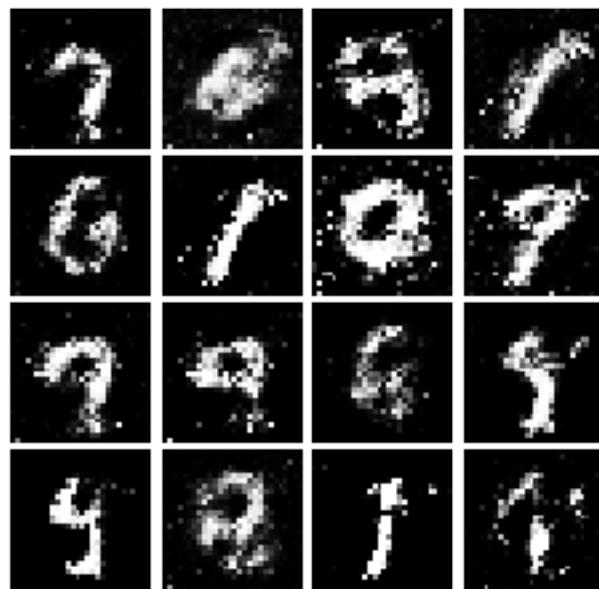
Iter: 1750, D: 1.073, G:0.981



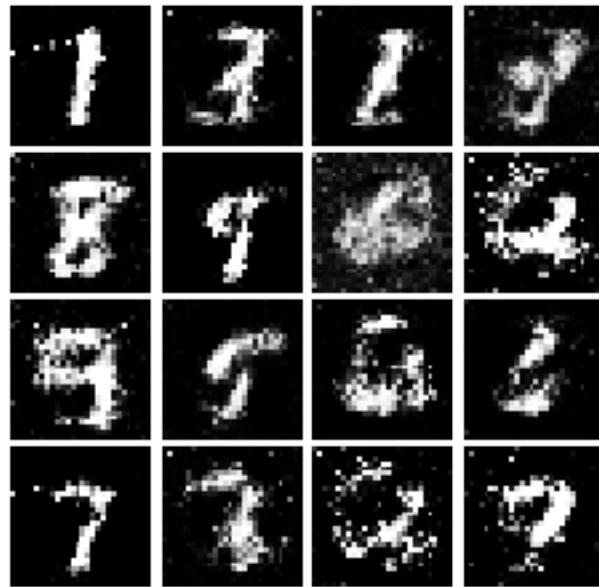
Saving the model as a checkpoint...

EPOCH: 5

Iter: 2000, D: 1.236, G: 0.8874



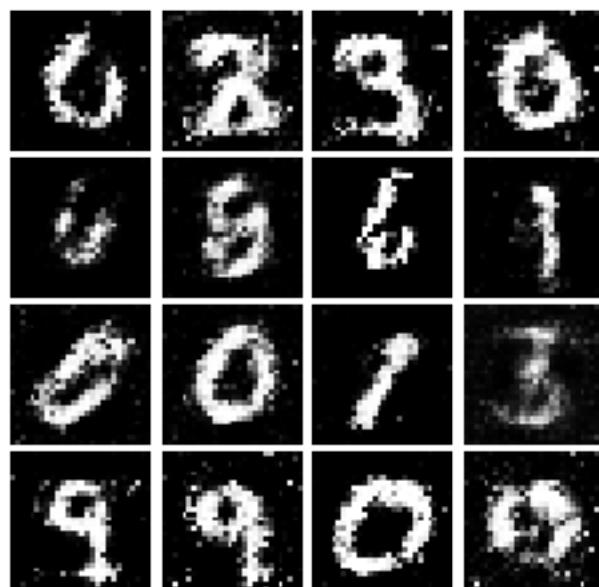
Iter: 2250, D: 1.232, G: 0.8178



Saving the model as a checkpoint...

EPOCH: 6

Iter: 2500, D: 1.258, G: 0.7426



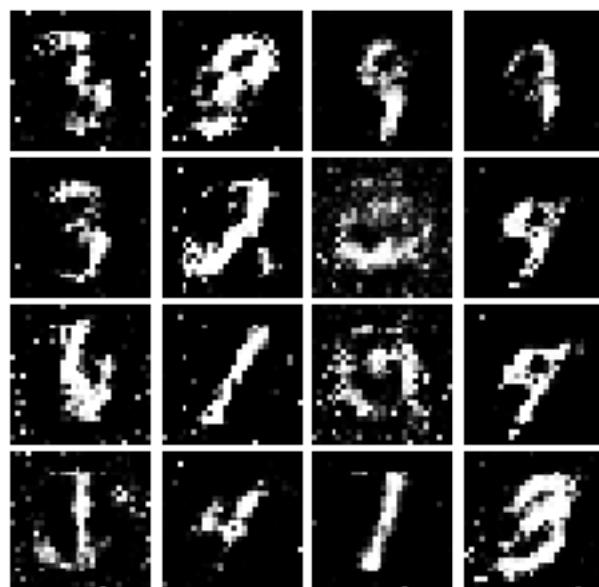
Iter: 2750, D: 1.325, G: 0.9565



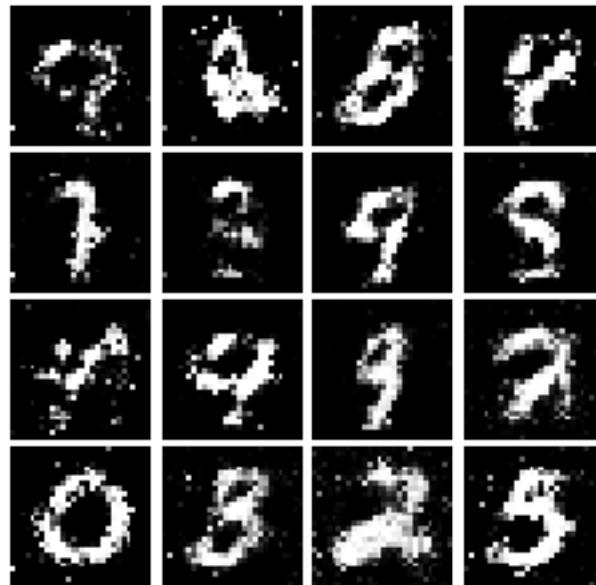
Saving the model as a checkpoint...

EPOCH: 7

Iter: 3000, D: 1.328, G: 0.7802



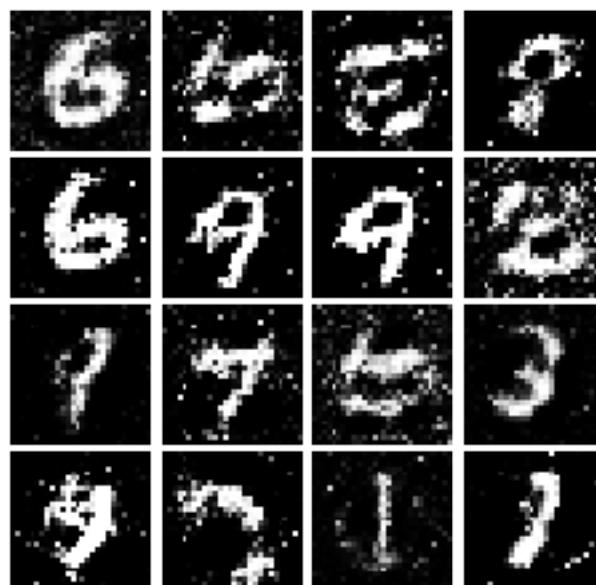
Iter: 3250, D: 1.295, G: 1.034



Saving the model as a checkpoint...

EPOCH: 8

Iter: 3500, D: 1.412, G: 0.7897

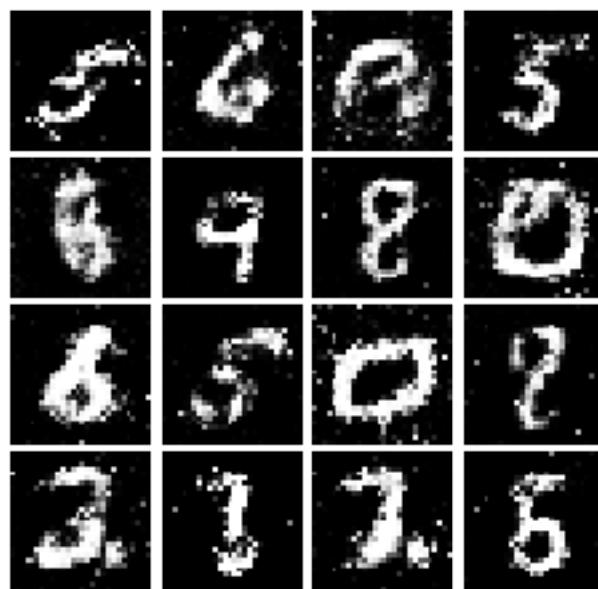


Saving the model as a checkpoint...

EPOCH: 9
Iter: 3750, D: 1.301, G:0.7817

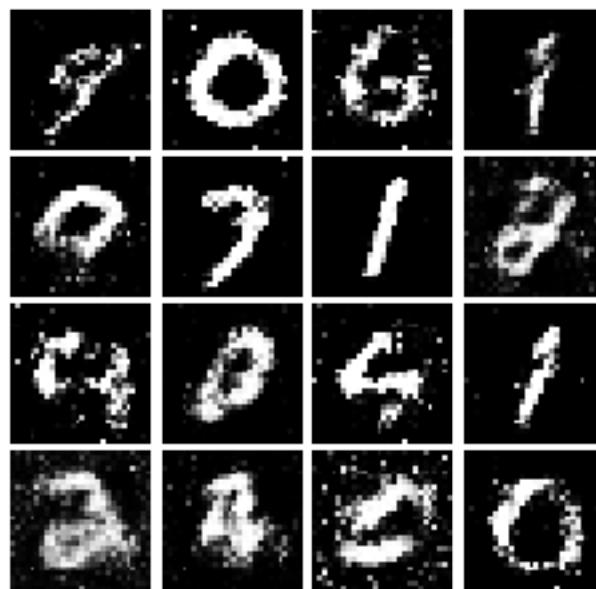


Iter: 4000, D: 1.335, G:0.7885

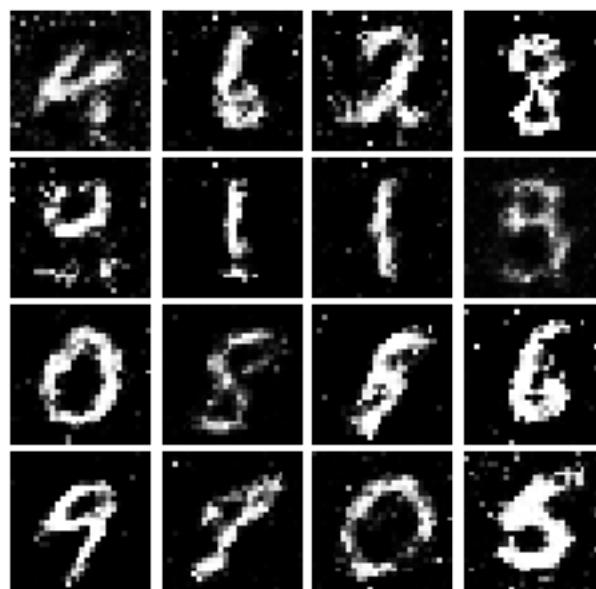


Saving the model as a checkpoint...

EPOCH: 10
Iter: 4250, D: 1.31, G:0.7612



Iter: 4500, D: 1.29, G:0.8414



Saving the model as a checkpoint...

3.2 Section 1.2 Least-square GAN Loss

3.2.1 Section 1.2.1 LSGAN Loss [10 pts]

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

Fill in the `ls_discriminator_loss` and `ls_generator_loss` in the cell below.

HINTS: Instead of computing the expectation, we will be averaging over elements of the mini-batch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
[ ]: def ls_discriminator_loss(logits_real, logits_fake):
    """
    Compute the Least-Squares GAN loss for the discriminator.

    Inputs:
    - scores_real: PyTorch Tensor of shape (N,) giving scores for the real data.
    - scores_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.

    Outputs:
    - loss: A PyTorch Tensor containing the loss.
    """
    loss = None

    #####
    # YOUR CODE HERE
    #####
    loss = torch.mean(0.5*(logits_real-1)**2+0.5*logits_fake**2,dim=0)
    #####
    # END
    #####
    return loss

def ls_generator_loss(logits_fake):
    """
    Computes the Least-Squares GAN loss for the generator.

    Inputs:
    - scores_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.

    Outputs:
    
```

```

- loss: A PyTorch Tensor containing the loss.
"""

loss = None

#####
#           YOUR CODE HERE           #
#####
loss = torch.mean(0.5*(logits_fake-1)**2, dim=0)
#####          END          #####
return loss

```

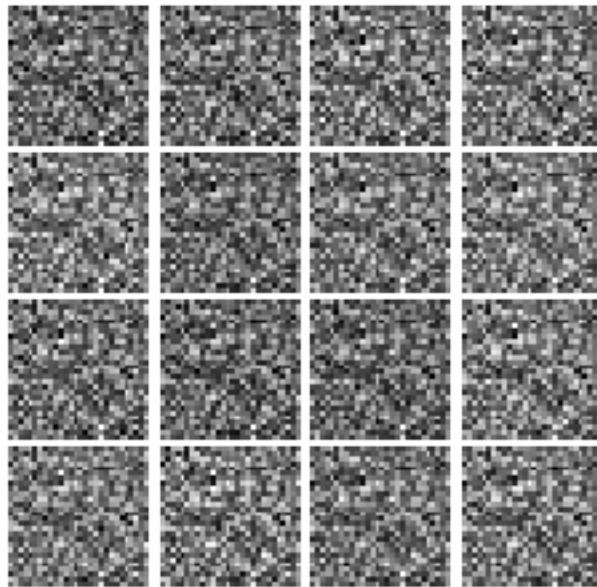
3.2.2 Section 1.2.2 Train model with LSGAN loss[7 pts]

- Call Discriminator and Generator for training.
- Call optimizers for both discriminator and generator for training. (Use Adam with lr=1e-3, betas = (0.5, 0.999))
- Call train function to train.
- Train for 10 epochs.

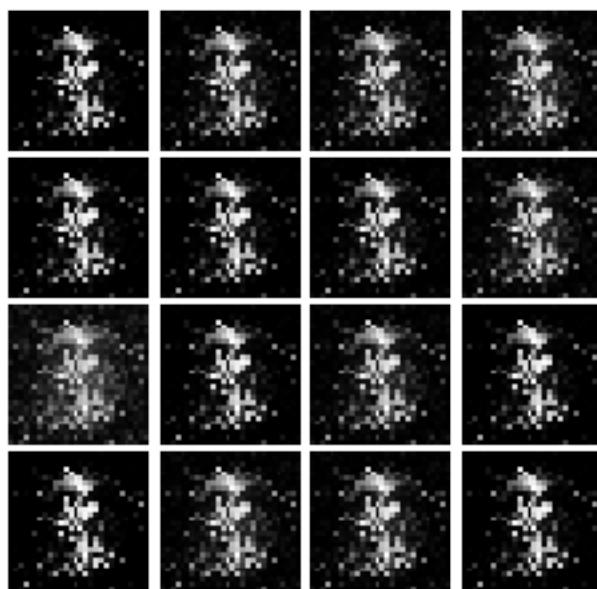
Similarly, train your LSGAN on MNIST dataset. You should expect results that resemble the following if your loss function and training loop implementations are correct:

```
[ ]: # YOUR TRAINING CODE HERE
Discriminator = discriminator().to(device)
Generator = generator().to(device)
Discriminator_optimizer = torch.optim.Adam(Discriminator.parameters(), lr=1e-3, 
                                           betas = (0.5, 0.999))
Generator_optimizer = torch.optim.Adam(Generator.parameters(), lr=1e-3, betas = 
                                         (0.5, 0.999))
train(Discriminator, Generator, Discriminator_optimizer, Generator_optimizer, 
      ls_discriminator_loss, ls_generator_loss, train_loader=loader_train, 
      device=device)
```

EPOCH: 1
Iter: 0, D: 0.5492, G:0.4561



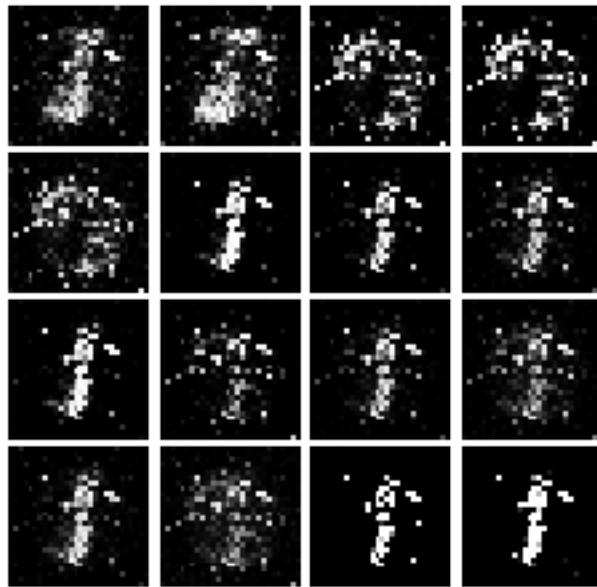
Iter: 250, D: 0.1203, G: 0.6983



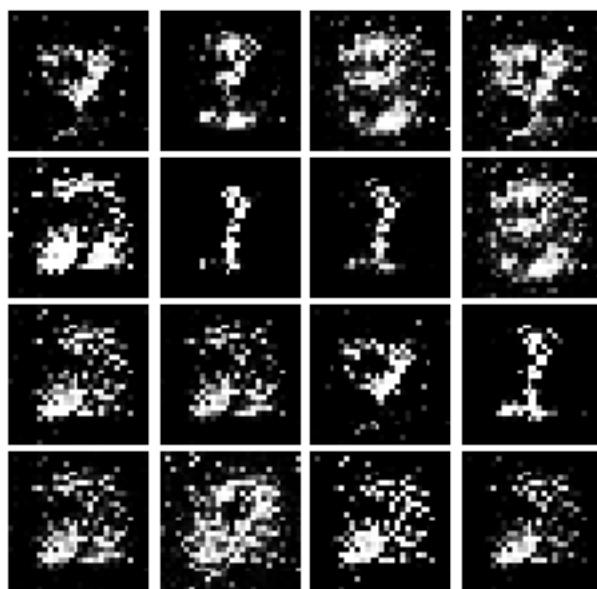
Saving the model as a checkpoint...

EPOCH: 2

Iter: 500, D: 0.2079, G: 0.3382



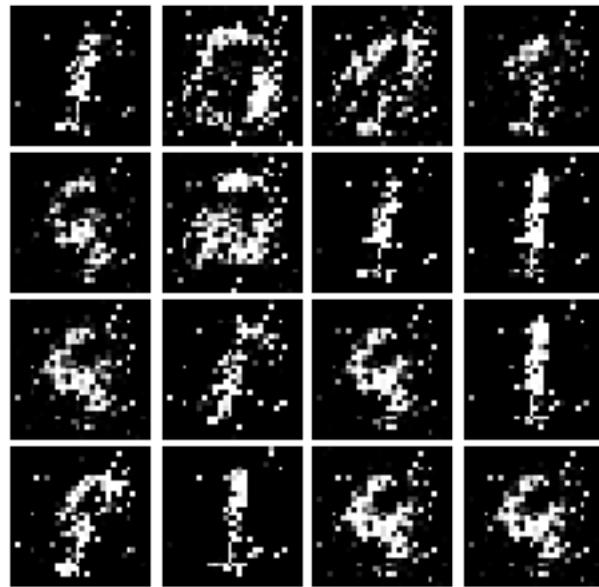
Iter: 750, D: 0.191, G:0.5072



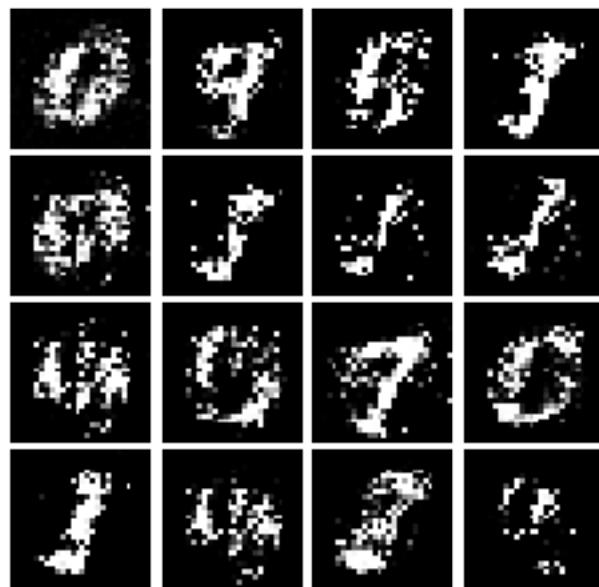
Saving the model as a checkpoint...

EPOCH: 3

Iter: 1000, D: 0.1374, G:0.4243



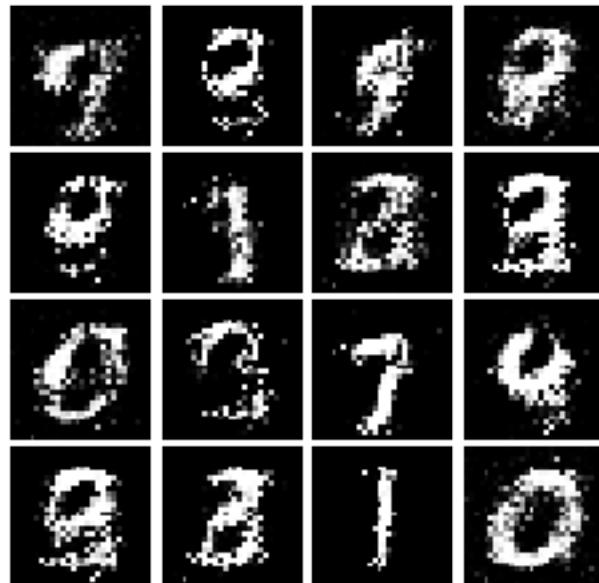
Iter: 1250, D: 0.1803, G: 0.2617



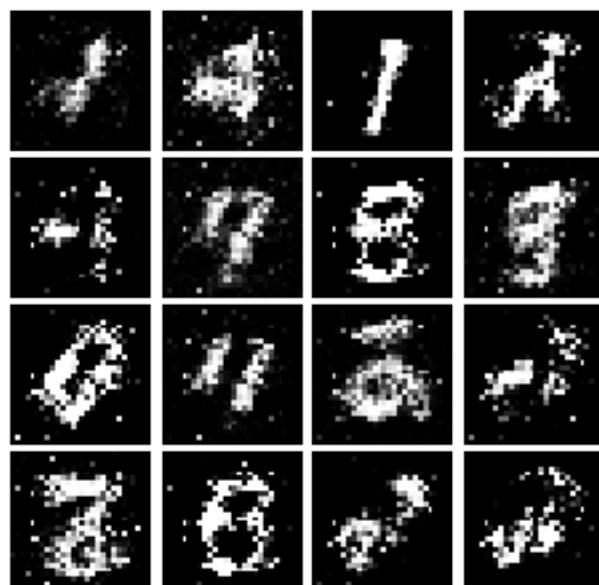
Saving the model as a checkpoint...

EPOCH: 4

Iter: 1500, D: 0.1575, G: 0.2341



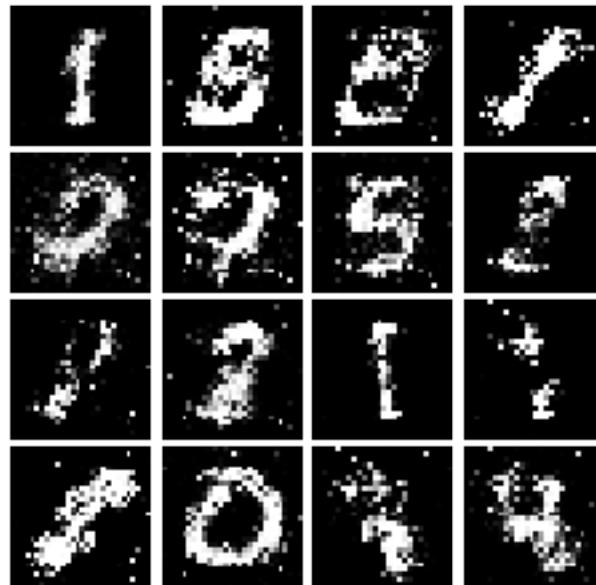
Iter: 1750, D: 0.2054, G: 0.1936



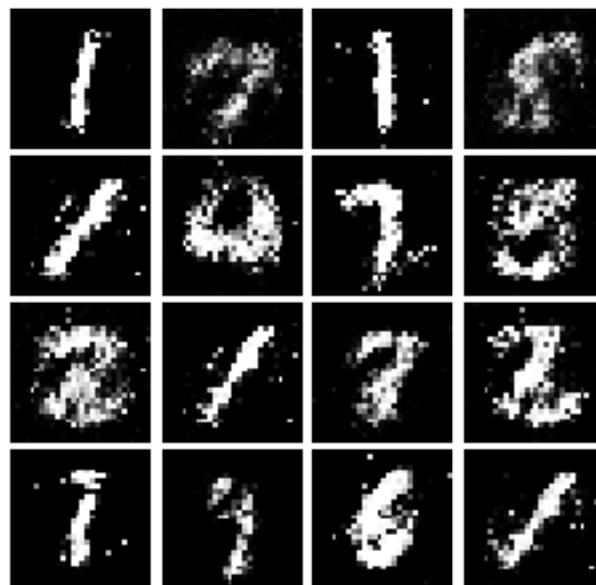
Saving the model as a checkpoint...

EPOCH: 5

Iter: 2000, D: 0.2016, G: 0.2079



Iter: 2250, D: 0.206, G: 0.1949



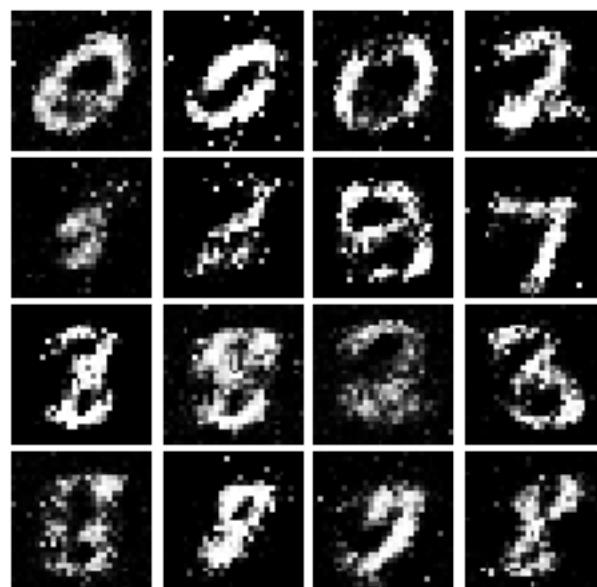
Saving the model as a checkpoint...

EPOCH: 6

Iter: 2500, D: 0.2268, G: 0.1631



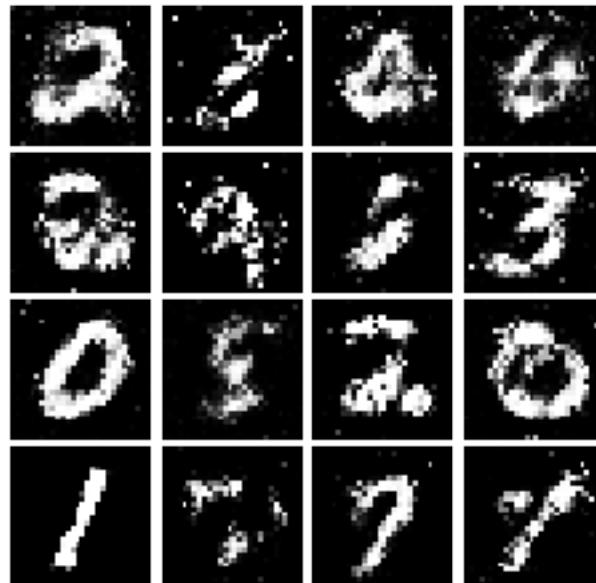
Iter: 2750, D: 0.2355, G: 0.1782



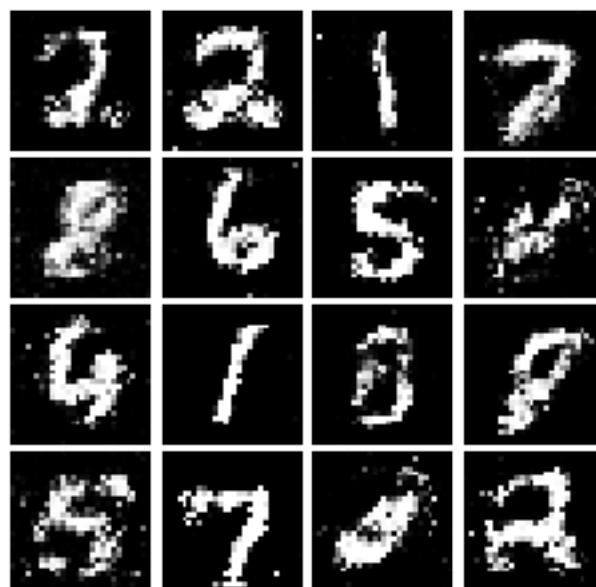
Saving the model as a checkpoint...

EPOCH: 7

Iter: 3000, D: 0.2451, G: 0.2237



Iter: 3250, D: 0.2123, G: 0.1828



Saving the model as a checkpoint...

EPOCH: 8

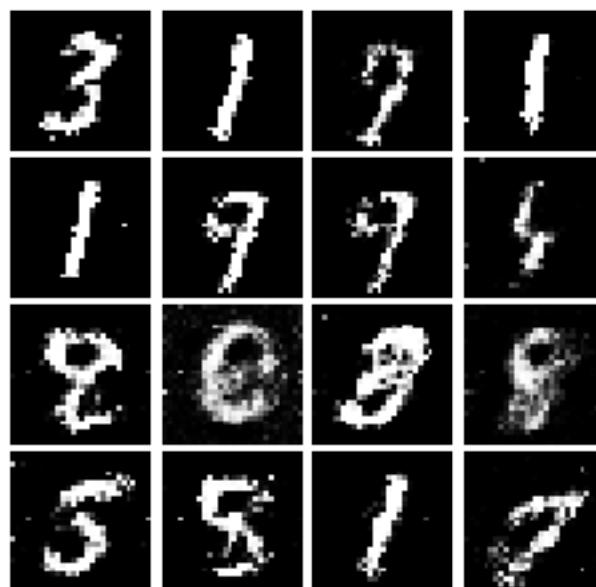
Iter: 3500, D: 0.2365, G: 0.1676



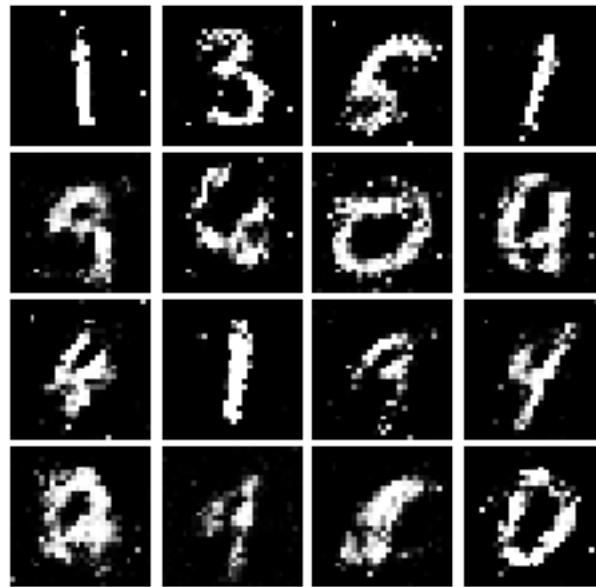
Saving the model as a checkpoint...

EPOCH: 9

Iter: 3750, D: 0.2311, G:0.1876



Iter: 4000, D: 0.2295, G:0.1646



Saving the model as a checkpoint...

EPOCH: 10

Iter: 4250, D: 0.2027, G: 0.1629



Iter: 4500, D: 0.2325, G: 0.1671



Saving the model as a checkpoint...

4 Section 2. Generative Adversarial Networks on CelebA Dataset [56 pts]

In this section, you will need to: 1. Implement DCGAN architecture 2. Train it on CelebA dataset.

We are done with the simple, not-so-challenging MNIST dataset. Now, you need to implement a specific model architecture called [DCGAN](#), and train your model to generate human faces!

4.0.1 Section 2.1. GAN model architecture [20 pts]

Implement your generator and discriminator for generating faces. We recommend the following architectures which are inspired by [DCGAN: Discriminator](#):

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=256, out_channels=512, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=512, out_channels=1024, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=1024, out_channels=1, kernel=4, stride=1`

Instead of Relu we use LeakyReLu throughout the discriminator (we use a negative slope value of 0.2).

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with `in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1`
- batch norm
- transpose convolution with `in_channels=1024, out_channels=512, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=512, out_channels=256, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=256, out_channels=128, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=128, out_channels=3, kernel=4, stride=2`

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a $3 \times 64 \times 64$ tensor for each sample (equal dimensions to the images from the dataset).

```
[ ]: # YOUR GENERATOR/DISCRIMINATOR HERE
class Reshape(nn.Module):
    def __init__(self, *args):
        super(Reshape, self).__init__()
        self.shape = args
    def forward(self, x):
        batch_size = x.shape[0]
        noise_dim = x.shape[1]
        return x.view((batch_size,noise_dim,1,1))

def Discriminator():
    model = nn.Sequential(
        torch.nn.
    →Conv2d(in_channels=3,out_channels=128,kernel_size=4,stride=2,padding=1),
        torch.nn.LeakyReLU(negative_slope=0.2),
        torch.nn.
    →Conv2d(in_channels=128,out_channels=256,kernel_size=4,stride=2,padding=1),
        torch.nn.LeakyReLU(negative_slope=0.2),
        torch.nn.BatchNorm2d(256),
        torch.nn.
    →Conv2d(in_channels=256,out_channels=512,kernel_size=4,stride=2,padding=1),
        torch.nn.LeakyReLU(negative_slope=0.2),
        torch.nn.BatchNorm2d(512),
        torch.nn.
    →Conv2d(in_channels=512,out_channels=1024,kernel_size=4,stride=2,padding=1),
```

```

    torch.nn.LeakyReLU(negative_slope=0.2),
    torch.nn.BatchNorm2d(1024),
    torch.nn.
    →Conv2d(in_channels=1024,out_channels=1,kernel_size=4,stride=1,padding=0),
        torch.nn.LeakyReLU(negative_slope=0.2))
return model

def Generator(noise_dim=NOISE_DIM):
    model = nn.Sequential(Reshape(),
        torch.nn.
    →ConvTranspose2d(in_channels=NOISE_DIM,out_channels=1024,kernel_size=4,stride=1,padding=0),
        torch.nn.ReLU(),
        torch.nn.BatchNorm2d(1024),
        torch.nn.
    →ConvTranspose2d(in_channels=1024,out_channels=512,kernel_size=4,stride=2,padding=1),
        torch.nn.ReLU(),
        torch.nn.BatchNorm2d(512),
        torch.nn.
    →ConvTranspose2d(in_channels=512,out_channels=256,kernel_size=4,stride=2,padding=1),
        torch.nn.ReLU(),
        torch.nn.BatchNorm2d(256),
        torch.nn.
    →ConvTranspose2d(in_channels=256,out_channels=128,kernel_size=4,stride=2,padding=1),
        torch.nn.ReLU(),
        torch.nn.BatchNorm2d(128),
        torch.nn.
    →ConvTranspose2d(in_channels=128,out_channels=3,kernel_size=4,stride=2,padding=1),
        torch.nn.Tanh())
return model

```

4.0.2 Section 2.2 Data loading: Celeb A Dataset

The CelebA images we provide have been filtered to obtain only images with clear faces and have been cropped and downsampled to 128x128 resolution.

Run download_celeba.sh to get dataset.

```
[ ]: !wget https://uofi.box.com/shared/static/q4pf89jtkvjndi4f8ip7wofuulhhphjj.zip
!mkdir celeba_data
!unzip q4pf89jtkvjndi4f8ip7wofuulhhphjj.zip -d celeba_data
!rm q4pf89jtkvjndi4f8ip7wofuulhhphjj.zip
!cp -r /content/celeba_data /content/drive/My\ Drive/DL_Fall_2020/Assignment_5/
```

```
[ ]: import sys
```

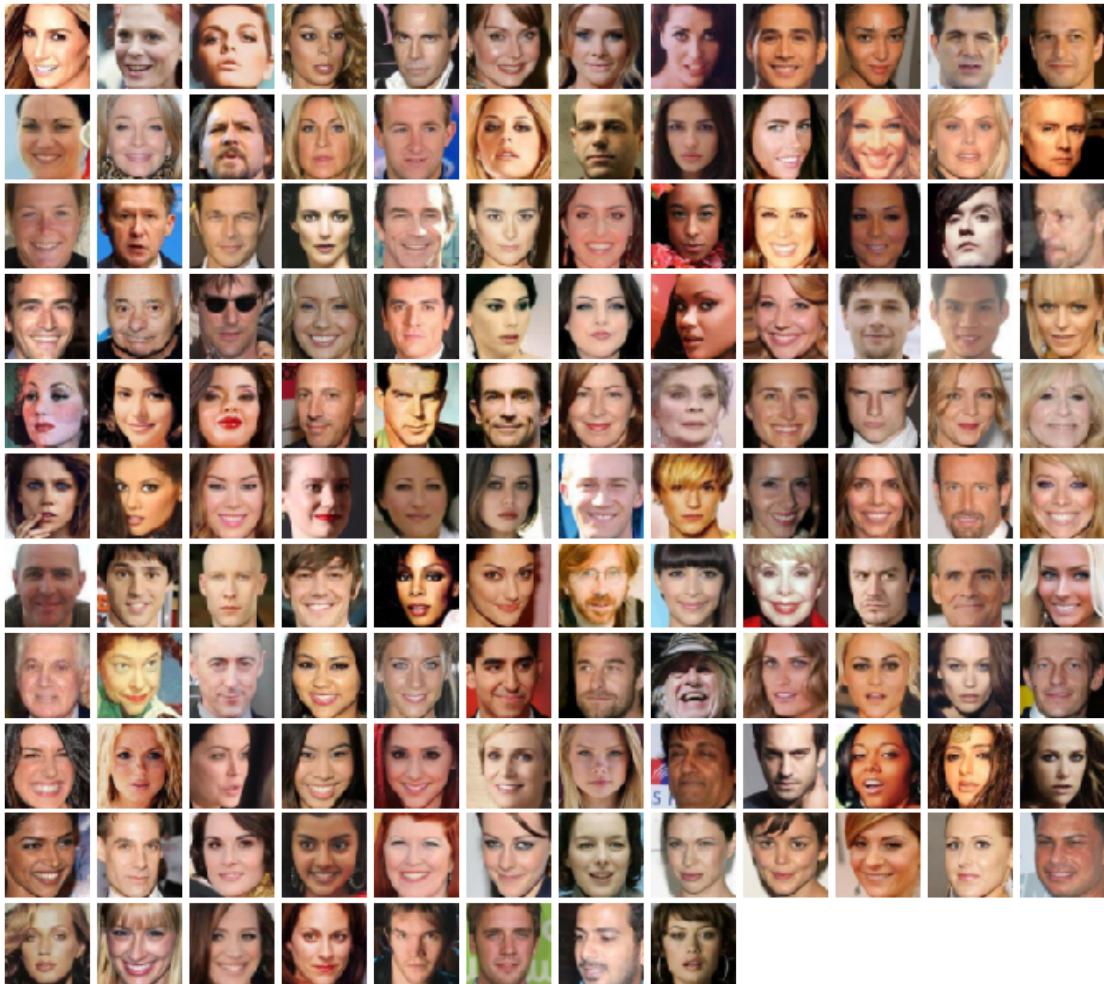
```
sys.path, dirs, files = next(os.walk("/content/drive/My Drive/DL_Fall_2020/  
→Assignment_5/celeba_data/celeba_train_128res"))  
file_count = len(files)  
print(file_count)
```

```
[ ]: batch_size = 128  
scale_size = 64 # We resize the images to 64x64 for training  
celeba_root = 'celeba_data'
```

```
[ ]: celeba_train = ImageFolder(root=celeba_root, transform=transforms.Compose([  
    transforms.Resize(scale_size),  
    transforms.ToTensor(),  
]))  
  
# You can change the num_workers to speed up loading  
celeba_loader_train = DataLoader(celeba_train, batch_size=batch_size, □  
→drop_last=True)
```

Visualize dataset

```
[ ]: imgs = celeba_loader_train.__iter__().next()[0].numpy().squeeze()  
show_images(imgs, color=True)
```



4.0.3 Section 2.3 Train a Vanilla GAN on CelebA [13 pts]

- Call discriminator and generator for training.
- Call optimizers for both discriminator and generator for training. (Use Adam with betas = (0.5, 0.999))
- Call train function to train.
- Train for 30 epochs.

Now, train your GAN model with vanilla GAN loss. If your models are implemented correctly, you should see something like this:

Now, train your model. **Observe the visualized result of your model, and describe what you see.**

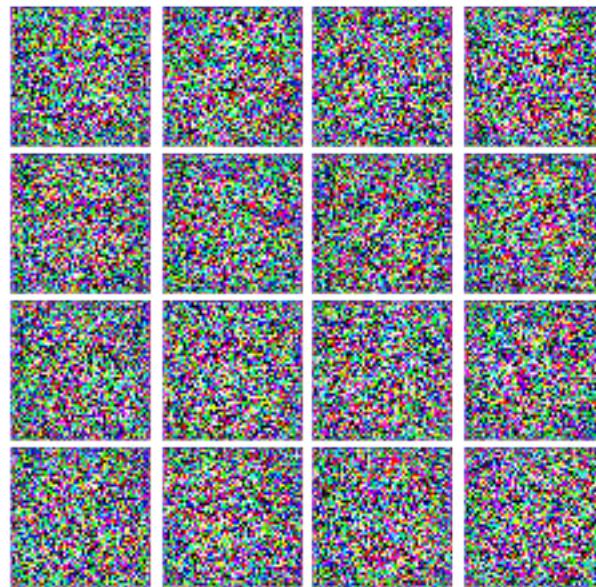
```
[ ]: NOISE_DIM = 100
      NUM_EPOCHS = 30
      learning_rate = 0.0002
```

```
[ ]: # original GAN
# Add code here:
Discriminator = Discriminator().to(device)
Generator = Generator(noise_dim=NOISE_DIM).to(device)

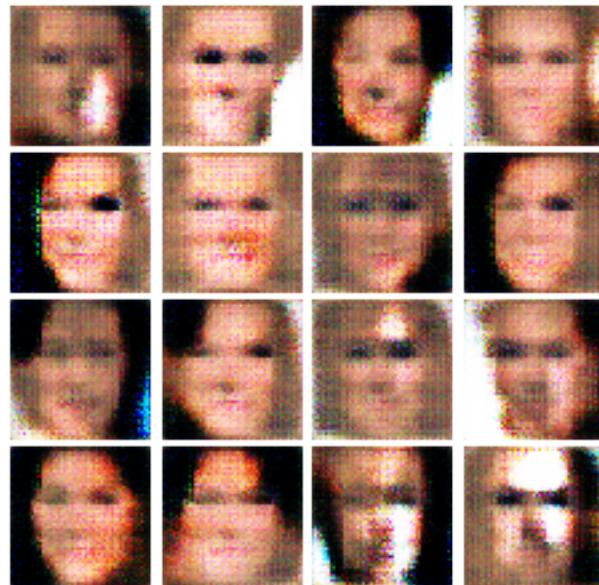
Discriminator_optimizer = torch.optim.Adam(Discriminator.parameters(), lr=learning_rate, betas = (0.5, 0.999))
Generator_optimizer = torch.optim.Adam(Generator.parameters(), lr=learning_rate, betas = (0.5, 0.999))

train(Discriminator, Generator, Discriminator_optimizer, Generator_optimizer, discriminator_loss, generator_loss, num_epochs=NUM_EPOCHS, train_loader=celeba_loader_train, device=device)
```

EPOCH: 1
 Iter: 0, D: 1.438, G:1.058



Iter: 250, D: 0.7416, G:0.933



Iter: 500, D: 1.197, G:2.466



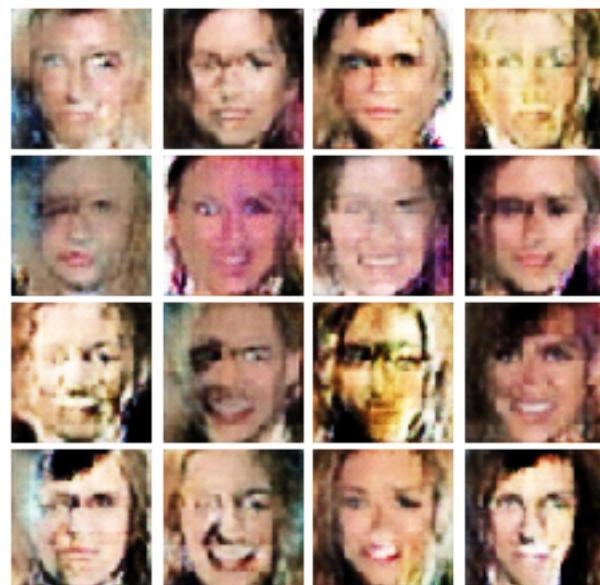
Iter: 750, D: 1.218, G:2.011



Saving the model as a checkpoint...

EPOCH: 2

Iter: 1000, D: 1.706, G:1.731



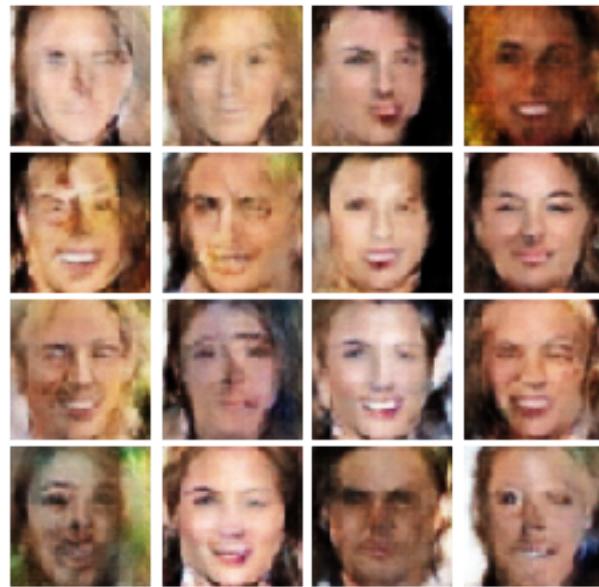
Iter: 1250, D: 1.266, G:1.305



Iter: 1500, D: 1.156, G: 0.8481



Iter: 1750, D: 1.149, G: 0.9253



Saving the model as a checkpoint...

EPOCH: 3

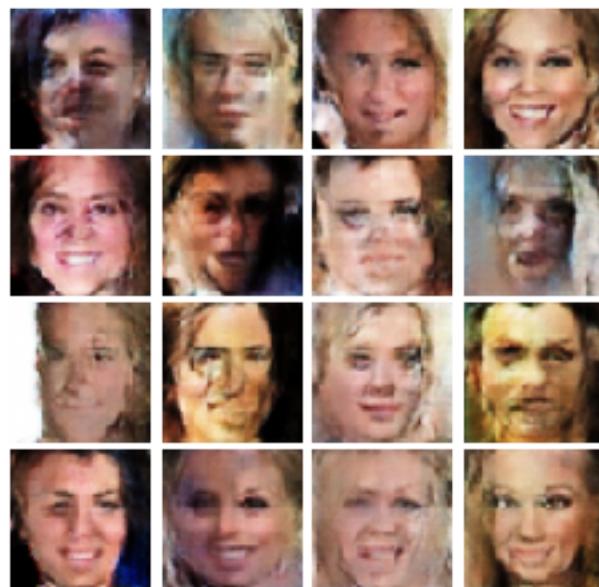
Iter: 2000, D: 1.414, G: 0.9775



Iter: 2250, D: 0.9734, G: 0.8645



Iter: 2500, D: 1.247, G:1.016



Iter: 2750, D: 1.058, G:1.077



Saving the model as a checkpoint...

EPOCH: 4

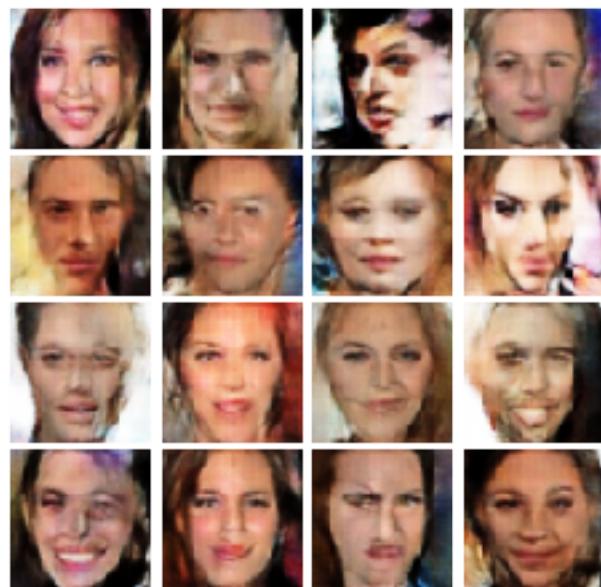
Iter: 3000, D: 1.517, G: 0.9261



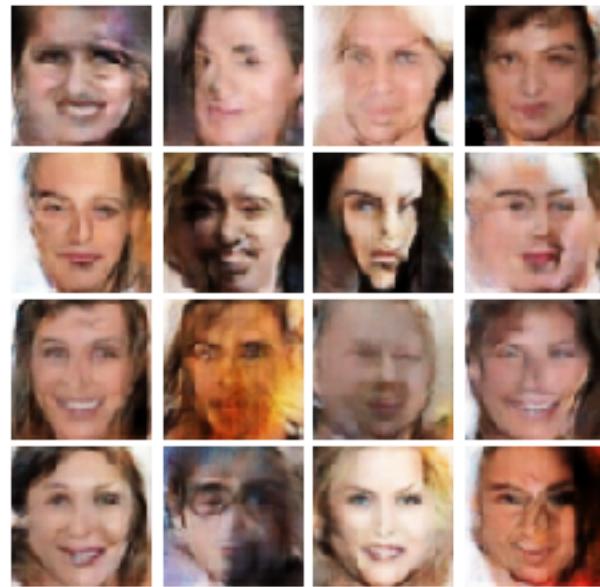
Iter: 3250, D: 1.386, G: 0.8513



Iter: 3500, D: 1.495, G: 1.468



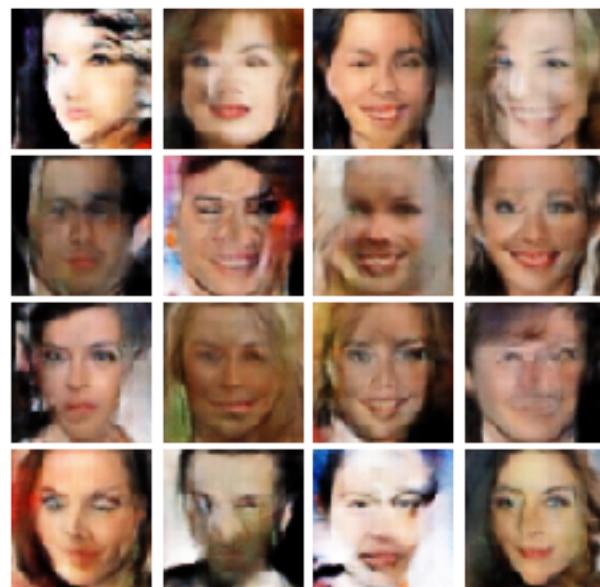
Iter: 3750, D: 1.251, G: 1.161



Saving the model as a checkpoint...

EPOCH: 5

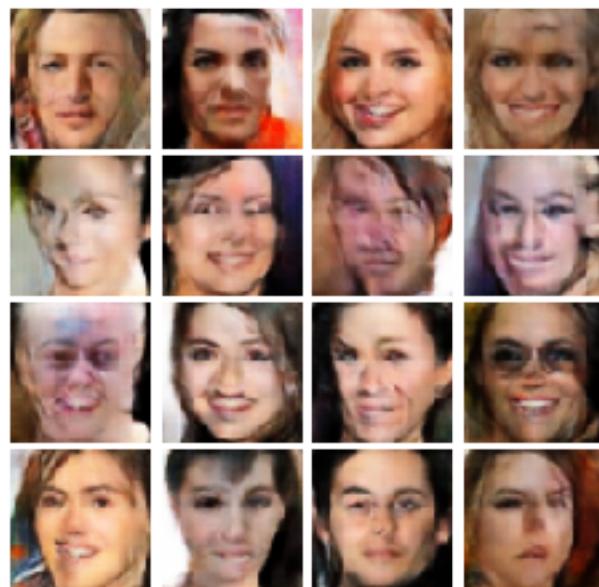
Iter: 4000, D: 1.39, G: 0.9363



Iter: 4250, D: 1.445, G: 0.82



Iter: 4500, D: 1.227, G: 1.336



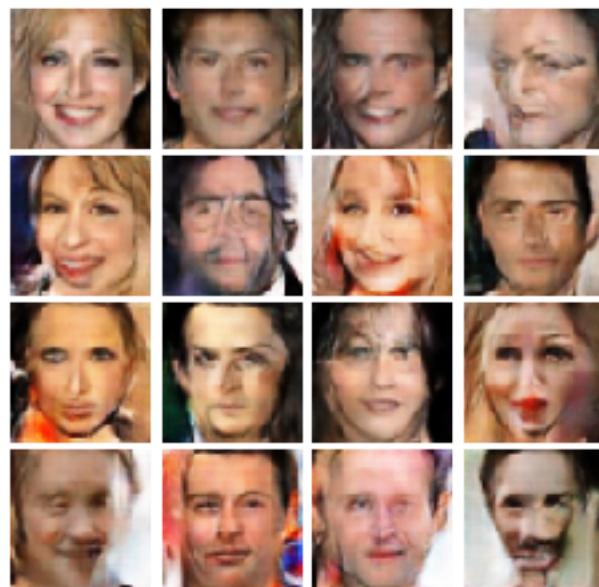
Iter: 4750, D: 1.289, G: 1.154



Saving the model as a checkpoint...

EPOCH: 6

Iter: 5000, D: 1.446, G: 1.214



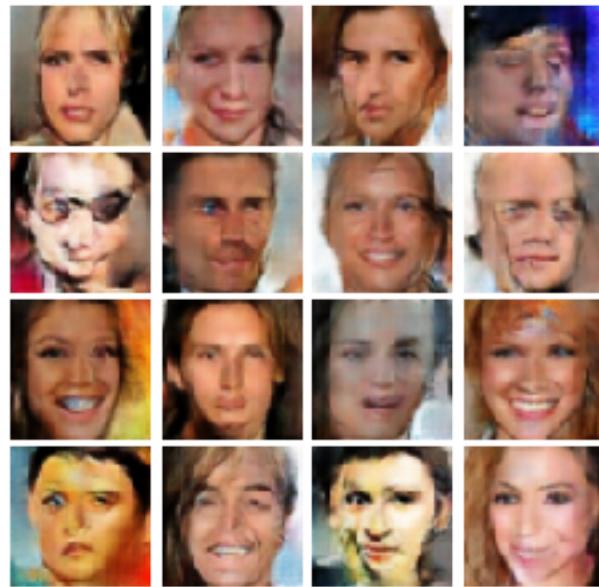
Iter: 5250, D: 1.46, G: 0.9585



Iter: 5500, D: 1.572, G:1.309



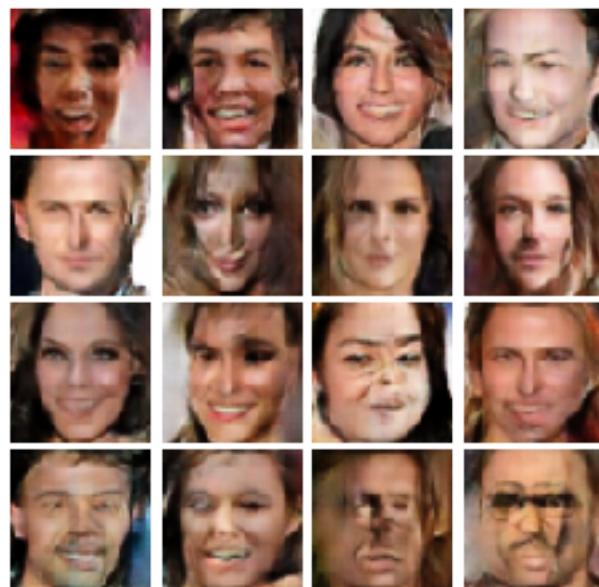
Iter: 5750, D: 1.429, G:1.159



Saving the model as a checkpoint...

EPOCH: 7

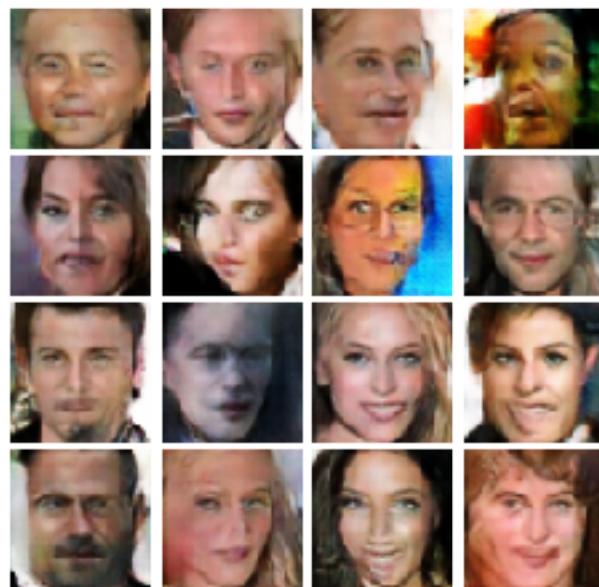
Iter: 6000, D: 1.712, G: 2.163



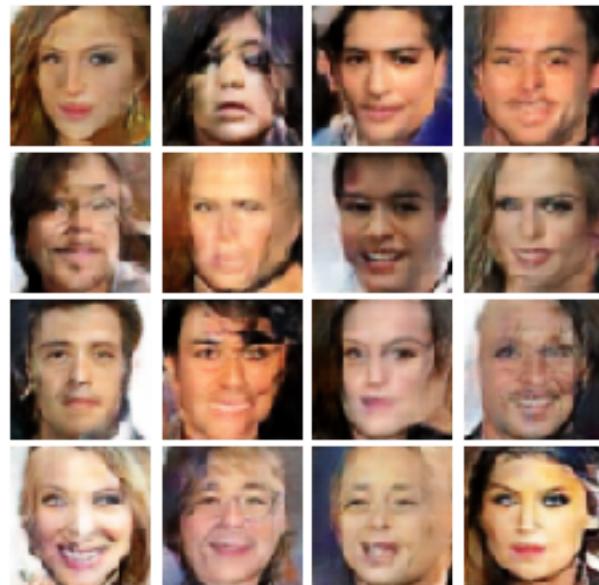
Iter: 6250, D: 0.9849, G: 0.7809



Iter: 6500, D: 1.062, G:0.8662



Iter: 6750, D: 1.388, G:0.8857



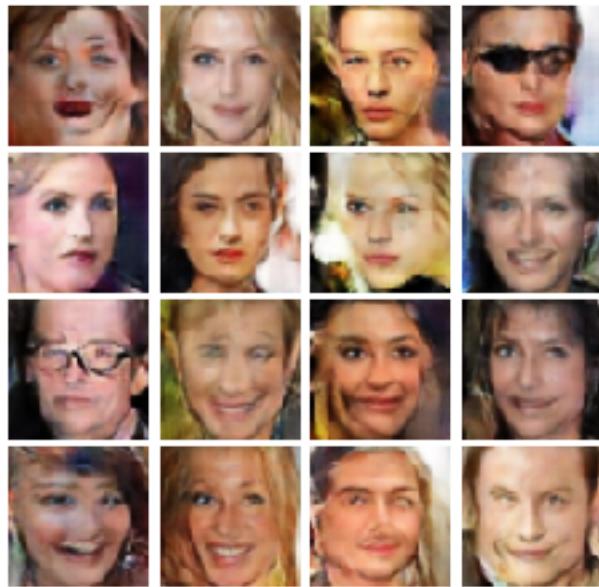
Saving the model as a checkpoint...

EPOCH: 8

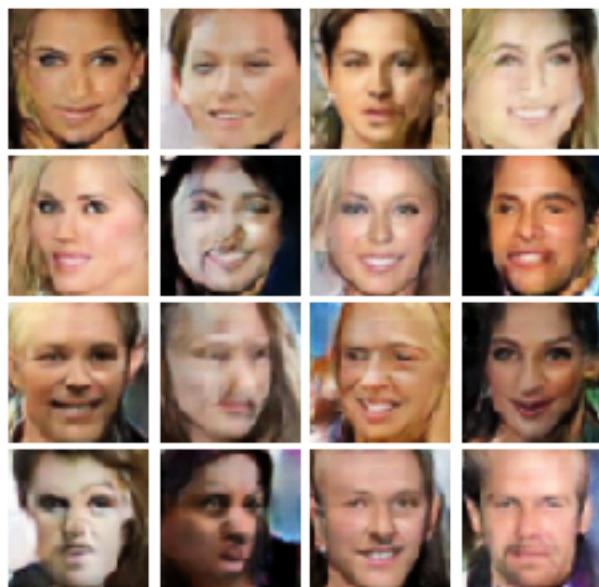
Iter: 7000, D: 1.056, G: 0.7804



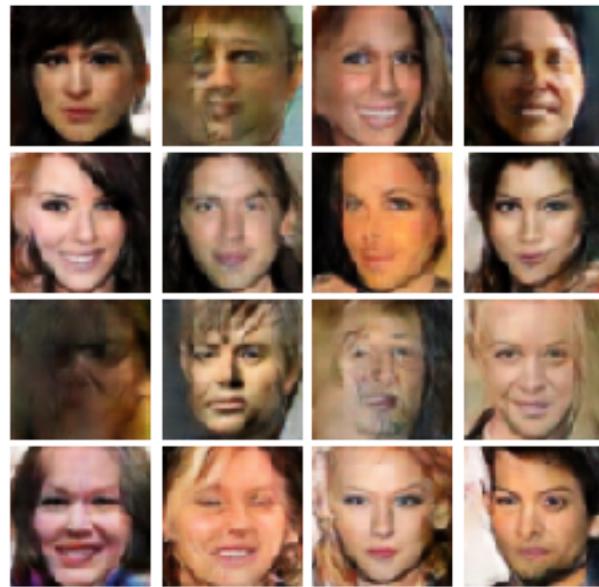
Iter: 7250, D: 0.9631, G: 1.183



Iter: 7500, D: 1.031, G: 0.7642



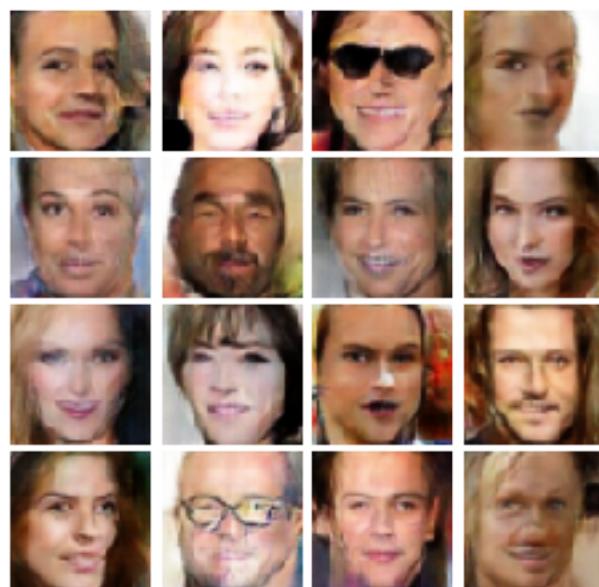
Iter: 7750, D: 1.138, G: 0.6318



Saving the model as a checkpoint...

EPOCH: 9

Iter: 8000, D: 1.162, G: 1.012



Iter: 8250, D: 1.024, G: 0.8713



Iter: 8500, D: 1.417, G: 1.381



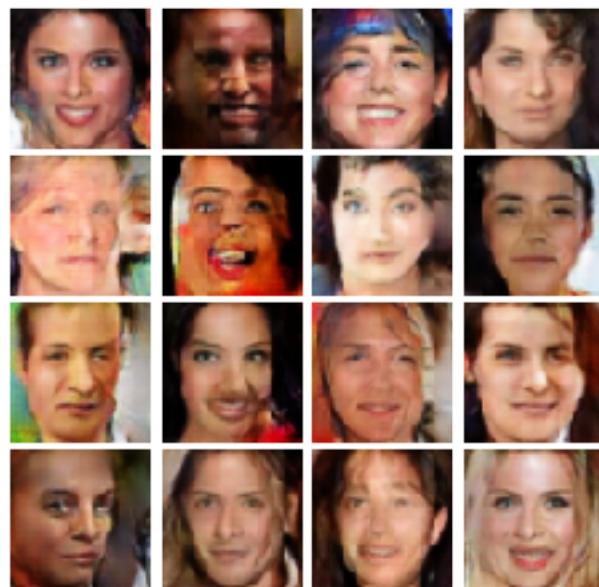
Iter: 8750, D: 0.9345, G: 0.9742



Saving the model as a checkpoint...

EPOCH: 10

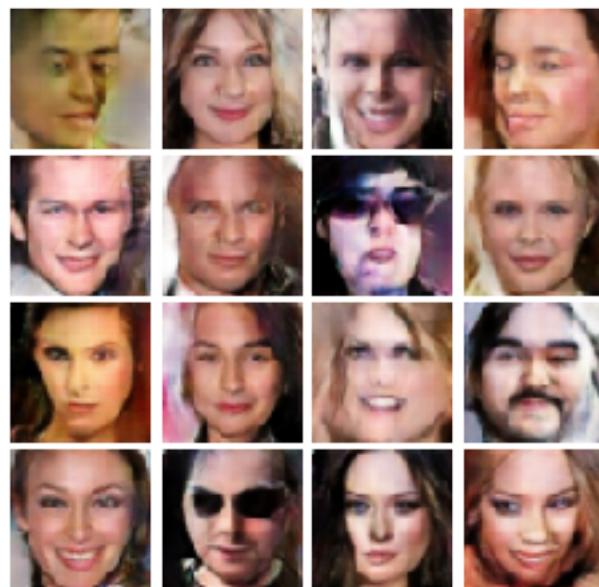
Iter: 9000, D: 1.233, G:1.648



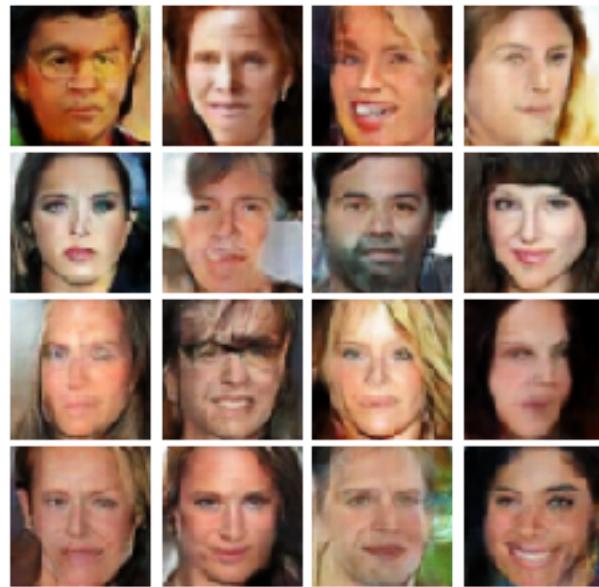
Iter: 9250, D: 0.801, G:1.101



Iter: 9500, D: 1.258, G:1.537



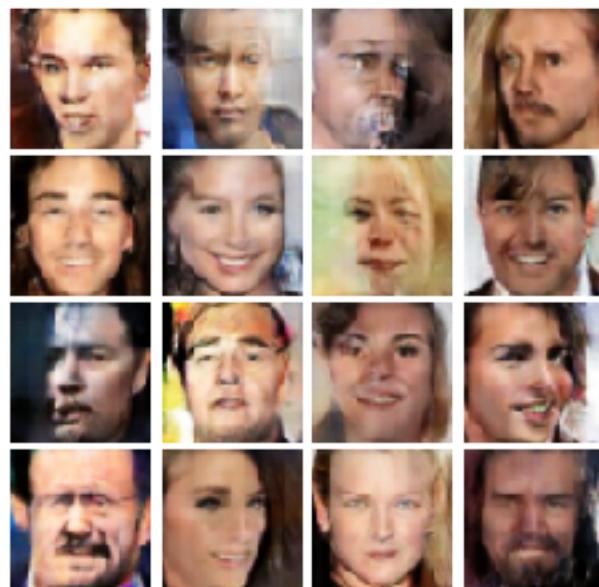
Iter: 9750, D: 0.9569, G:0.6229



Saving the model as a checkpoint...

EPOCH: 11

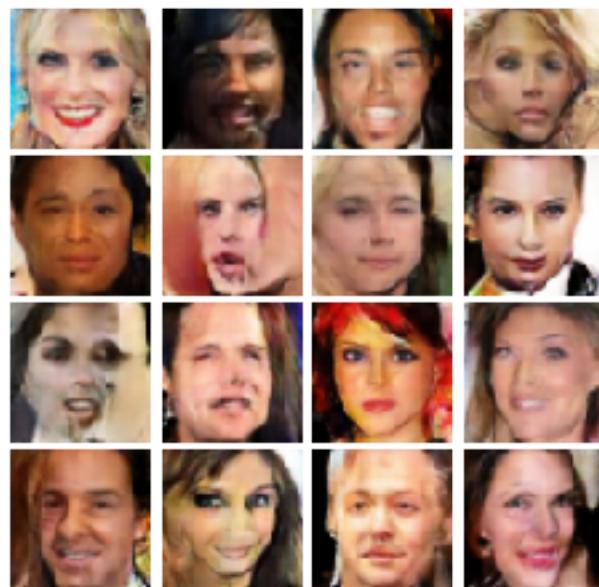
Iter: 10000, D: 0.9802, G:1.189



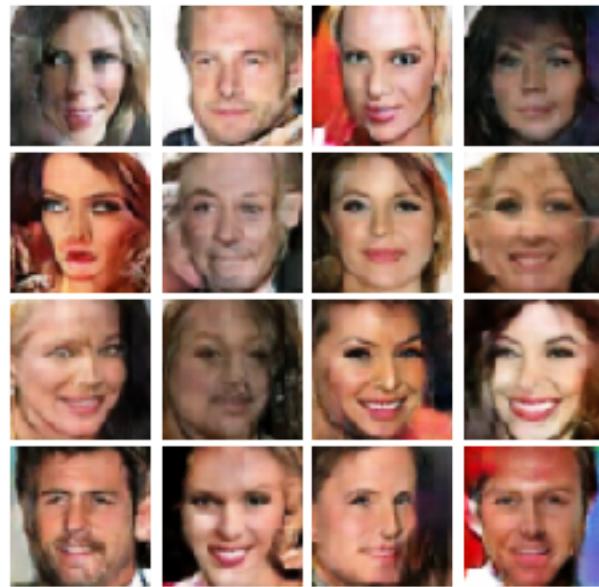
Iter: 10250, D: 1.071, G:1.895



Iter: 10500, D: 0.7206, G:1.142



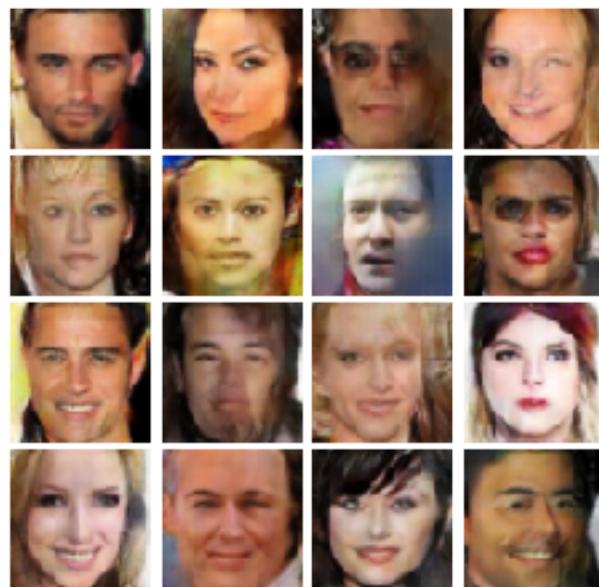
Iter: 10750, D: 0.6457, G:0.9265



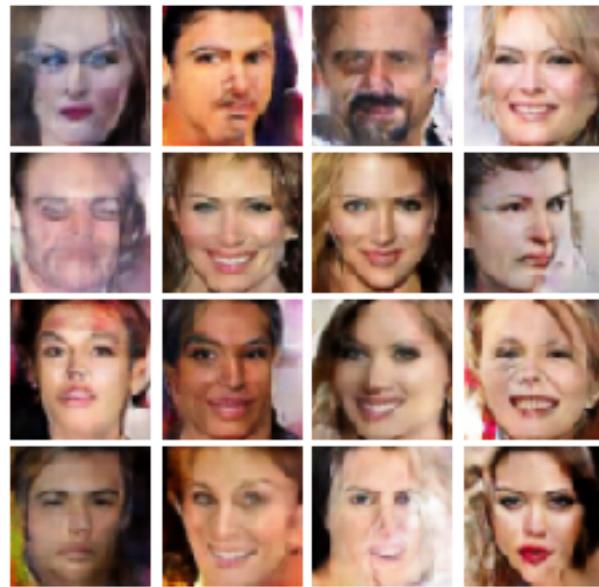
Saving the model as a checkpoint...

EPOCH: 12

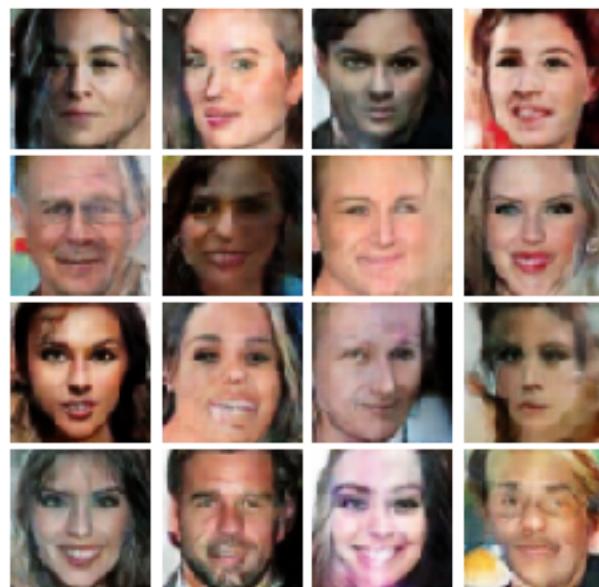
Iter: 11000, D: 0.9569, G: 0.9214



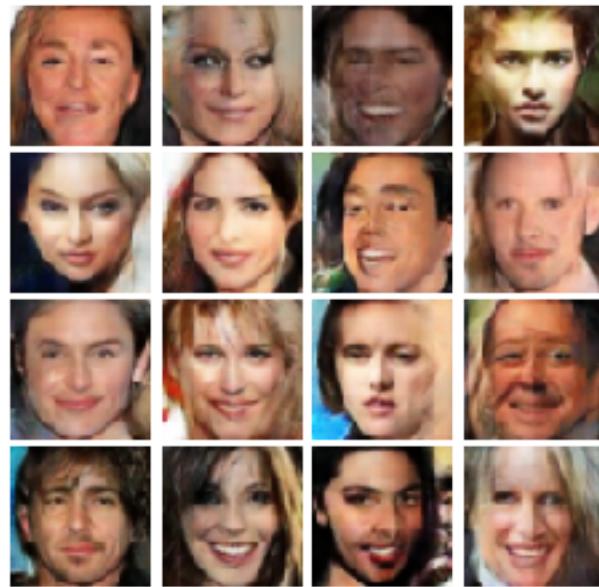
Iter: 11250, D: 1.012, G: 0.92



Iter: 11500, D: 0.5924, G:1.234



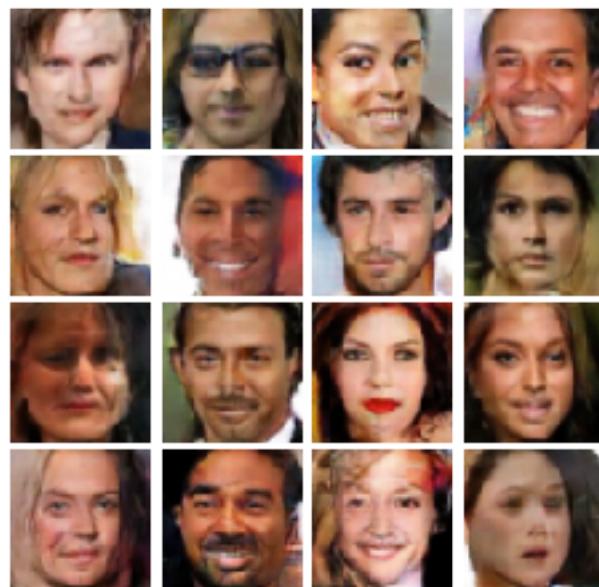
Iter: 11750, D: 0.406, G:1.574



Saving the model as a checkpoint...

EPOCH: 13

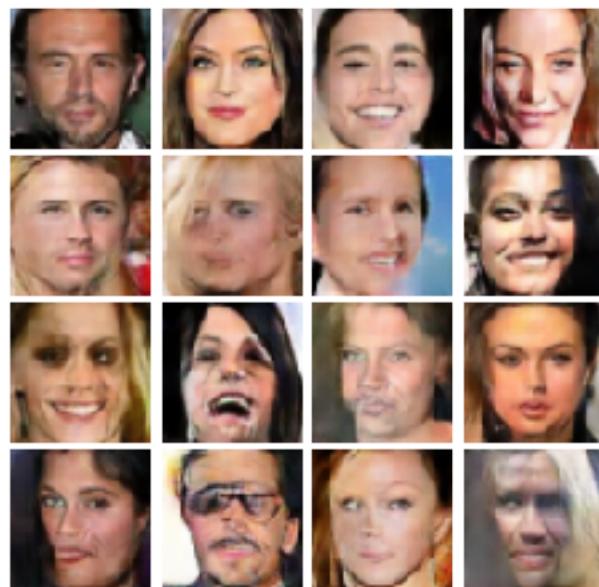
Iter: 12000, D: 0.5151, G:1.596



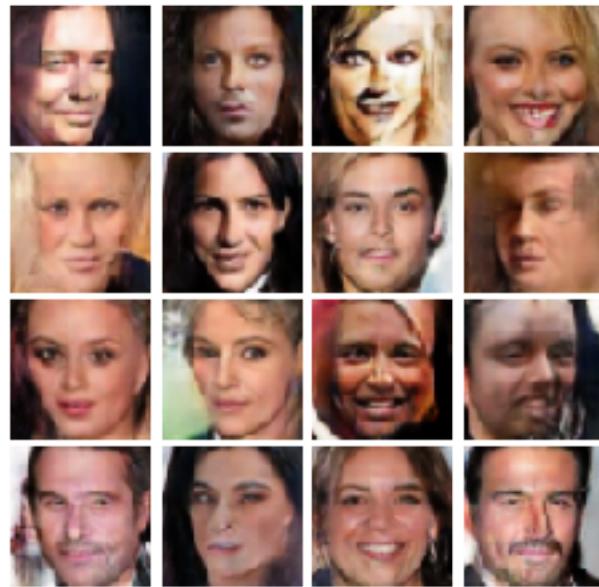
Iter: 12250, D: 1.186, G:3.003



Iter: 12500, D: 0.4939, G:1.223



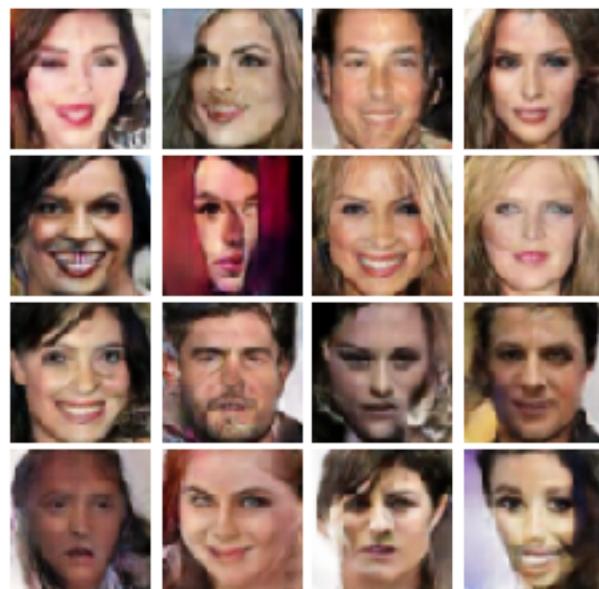
Iter: 12750, D: 0.5383, G:1.524



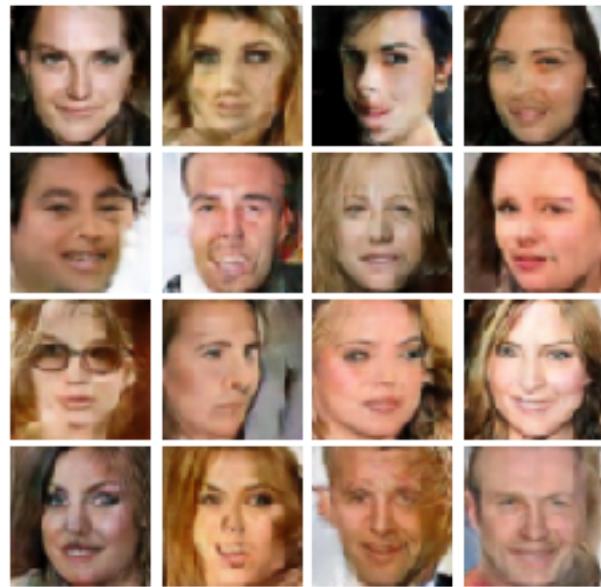
Saving the model as a checkpoint...

EPOCH: 14

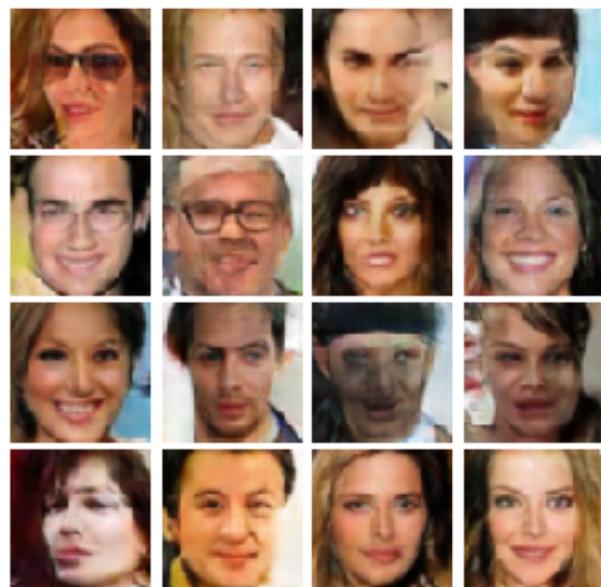
Iter: 13000, D: 0.3773, G:1.71



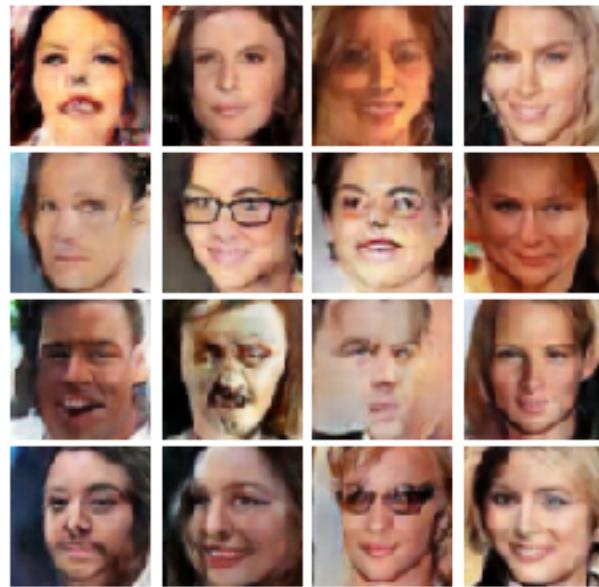
Iter: 13250, D: 0.5989, G:1.105



Iter: 13500, D: 0.5765, G:0.9919



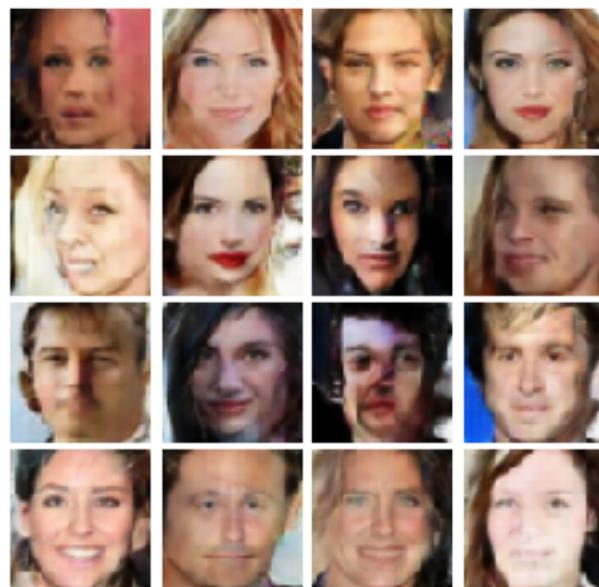
Iter: 13750, D: 0.2428, G:1.466



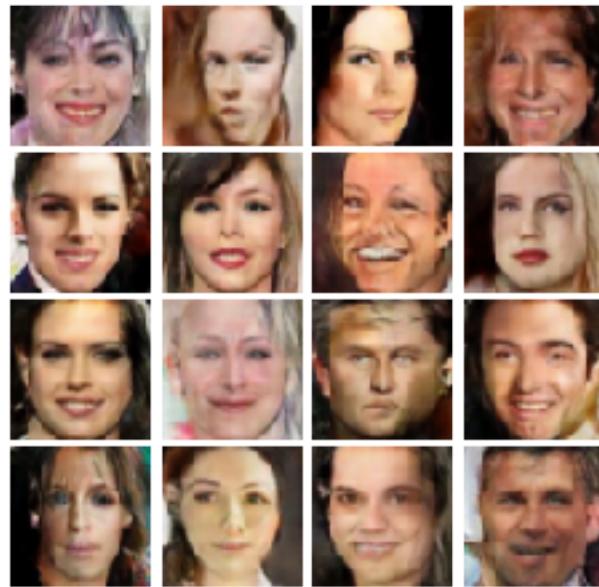
Saving the model as a checkpoint...

EPOCH: 15

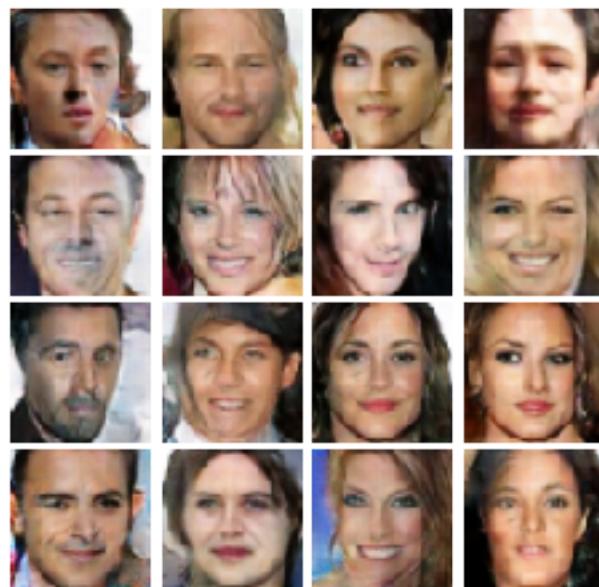
Iter: 14000, D: 1.359, G:2.212



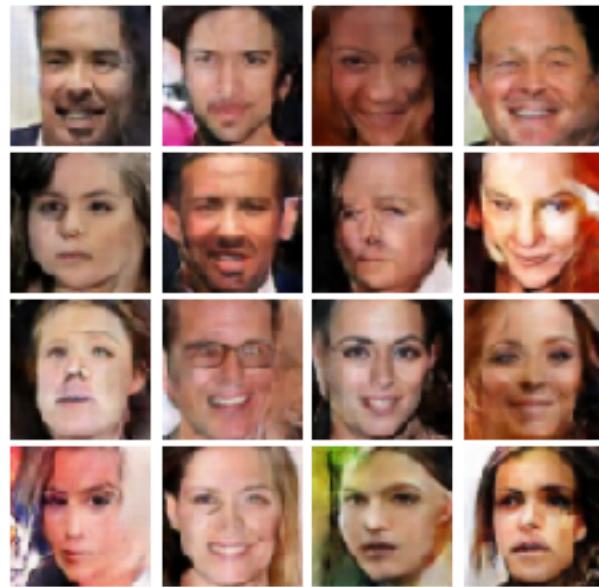
Iter: 14250, D: 0.259, G:1.784



Iter: 14500, D: 0.6733, G:1.474



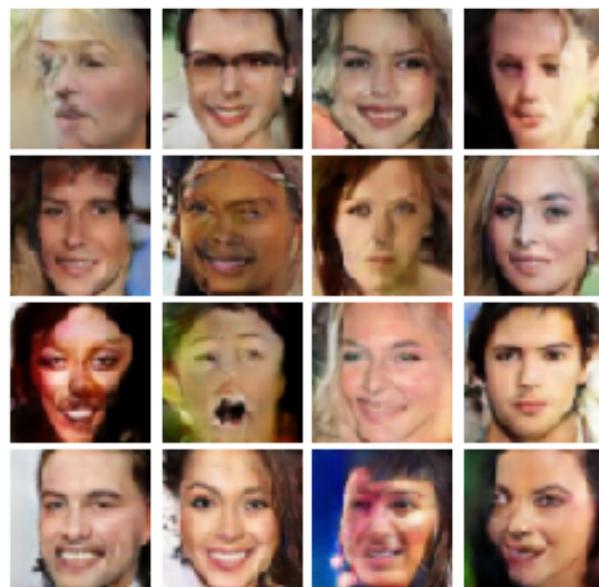
Iter: 14750, D: 0.5327, G:1.539



Saving the model as a checkpoint...

EPOCH: 16

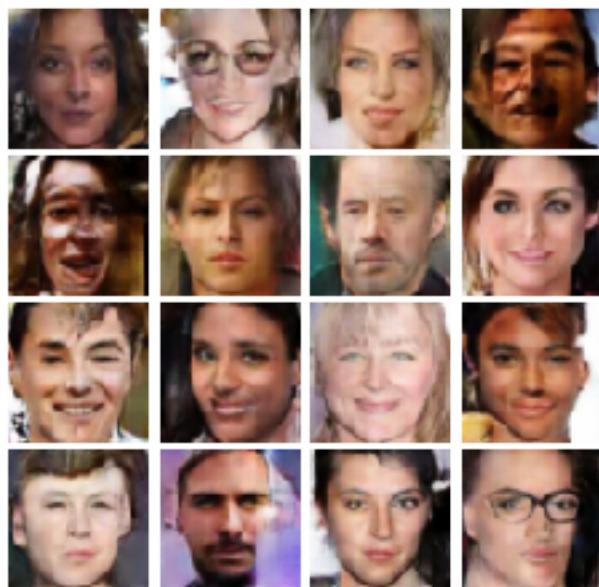
Iter: 15000, D: 0.263, G: 2.219



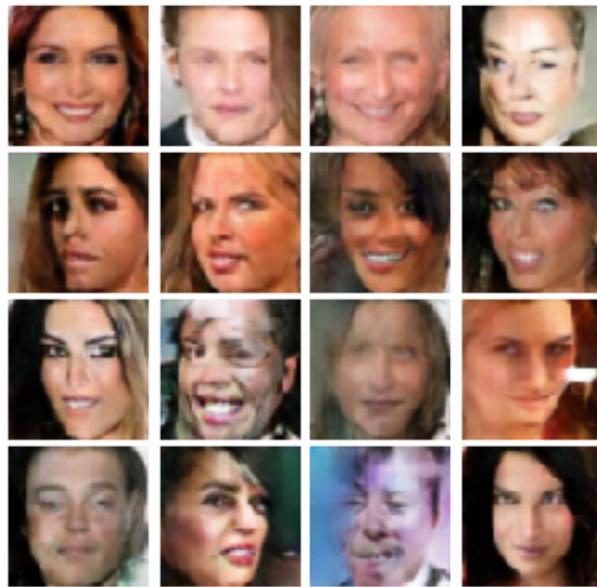
Iter: 15250, D: 0.3132, G: 1.481



Iter: 15500, D: 0.2866, G:1.917



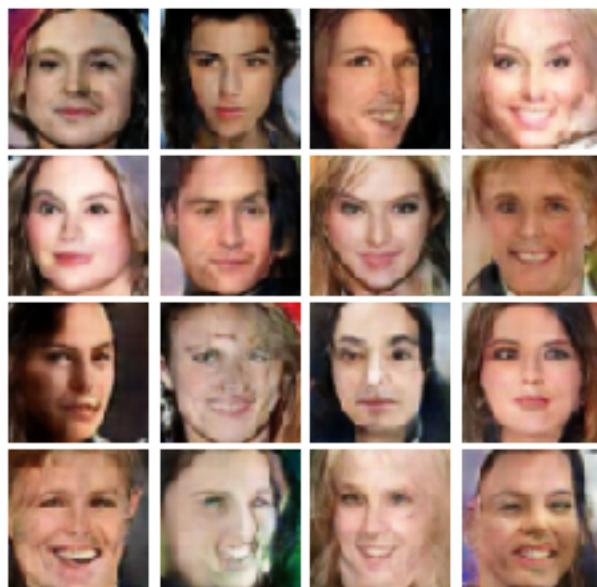
Iter: 15750, D: 0.9266, G:0.8389



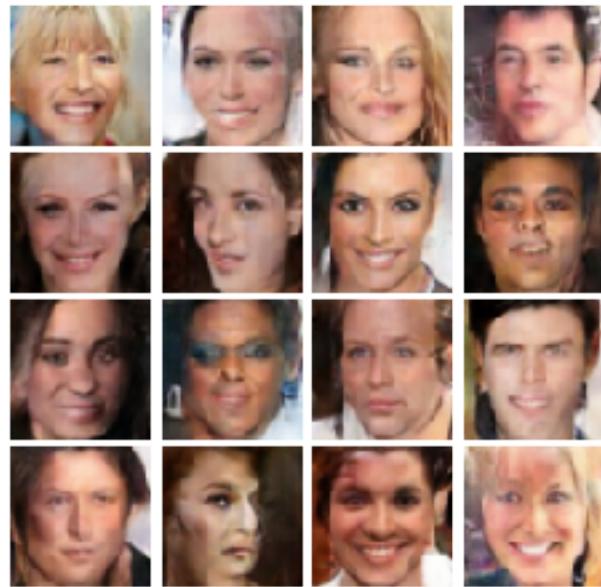
Saving the model as a checkpoint...

EPOCH: 17

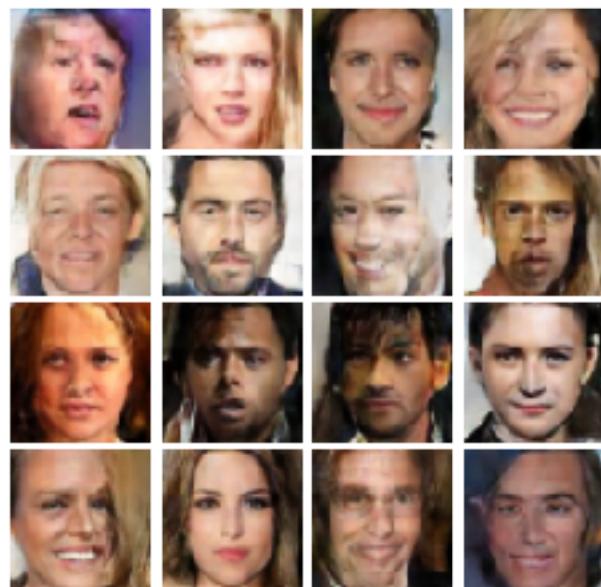
Iter: 16000, D: 0.2755, G: 1.463



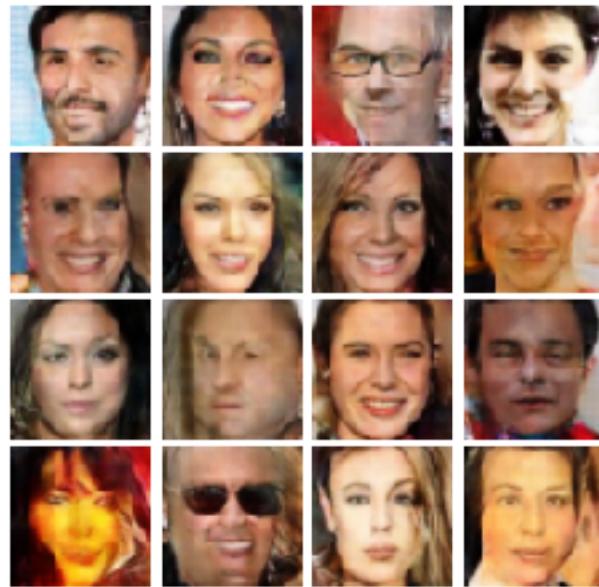
Iter: 16250, D: 2.11, G: 3.539



Iter: 16500, D: 0.1331, G:2.292



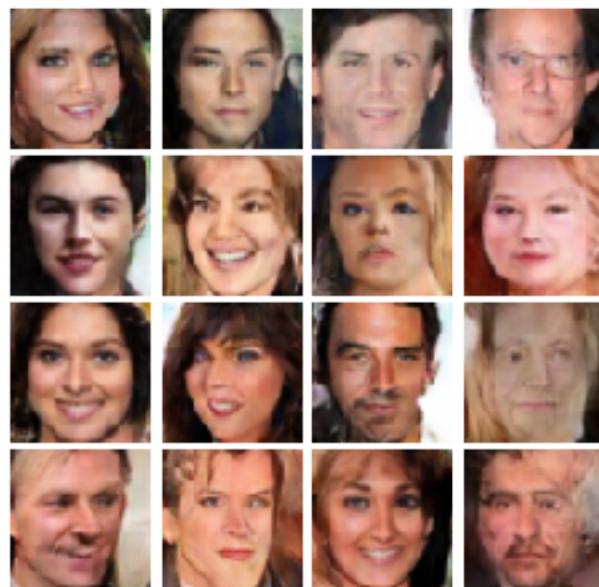
Iter: 16750, D: 0.4863, G:2.086



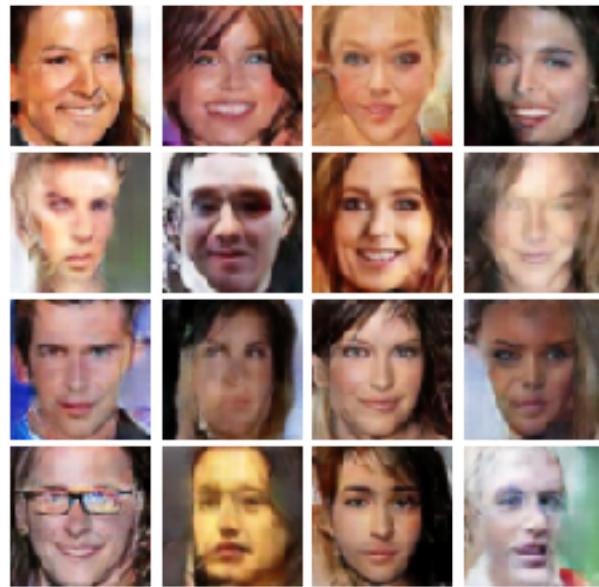
Saving the model as a checkpoint...

EPOCH: 18

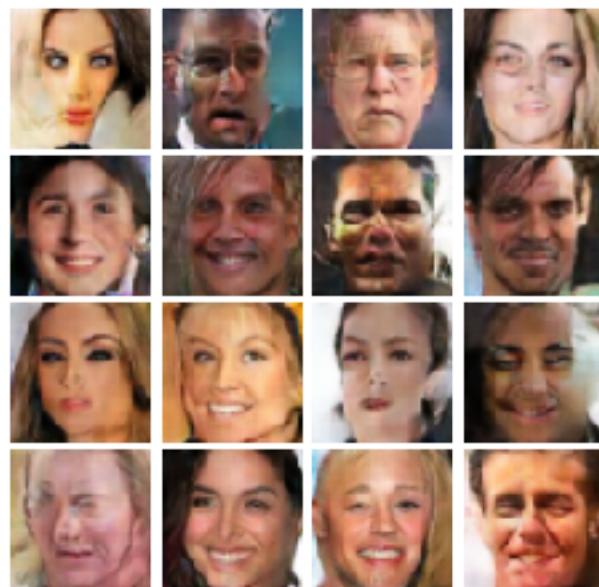
Iter: 17000, D: 0.1523, G:2.005



Iter: 17250, D: 0.3748, G:1.875



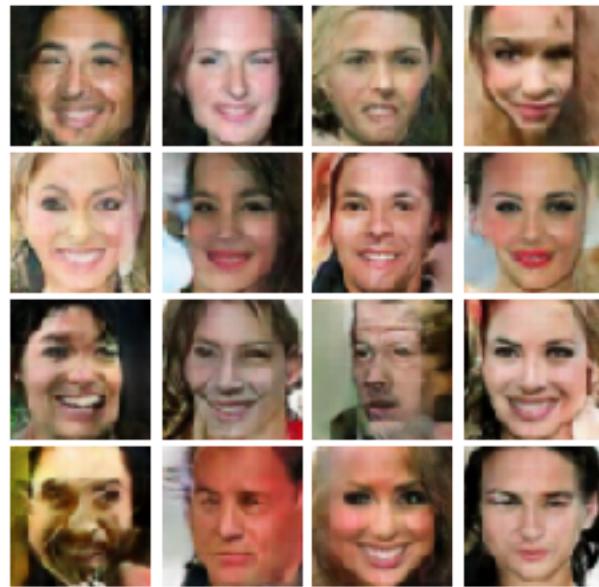
Iter: 17500, D: 0.7205, G: 0.6492



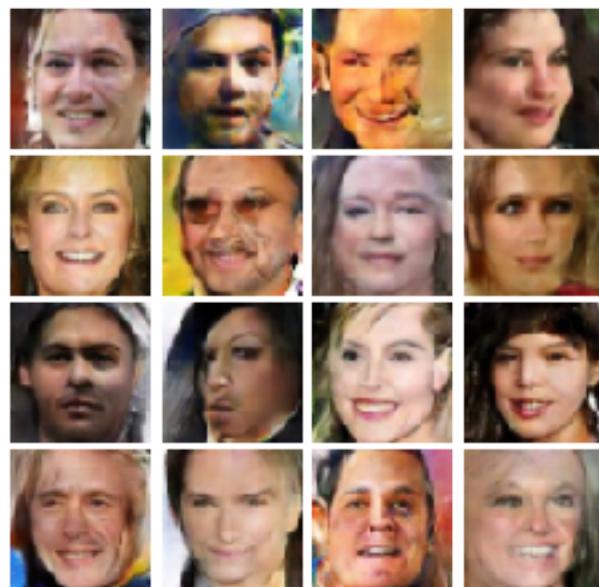
Saving the model as a checkpoint...

EPOCH: 19

Iter: 17750, D: 0.207, G: 2.199



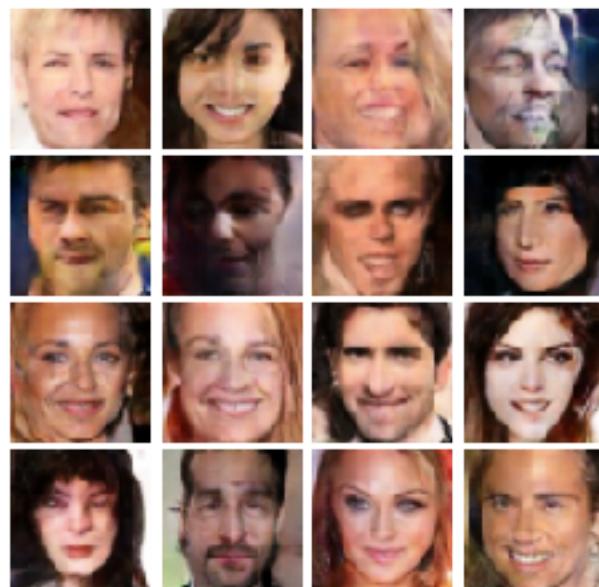
Iter: 18000, D: 0.1269, G:2.616



Iter: 18250, D: 0.3044, G:2.129



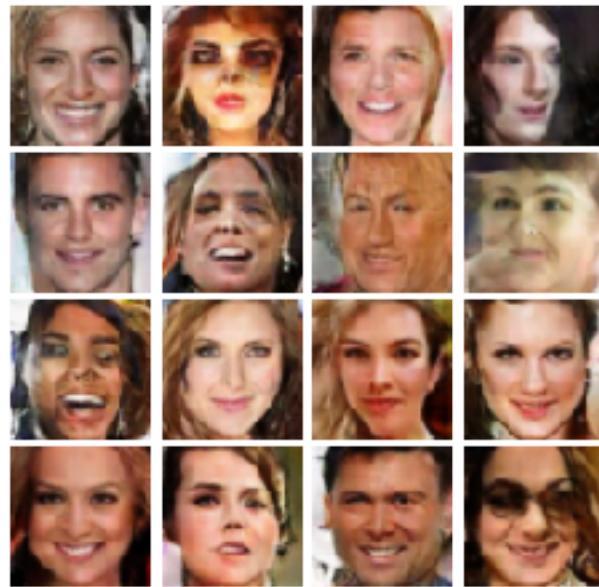
Iter: 18500, D: 0.7016, G: 2.016



Saving the model as a checkpoint...

EPOCH: 20

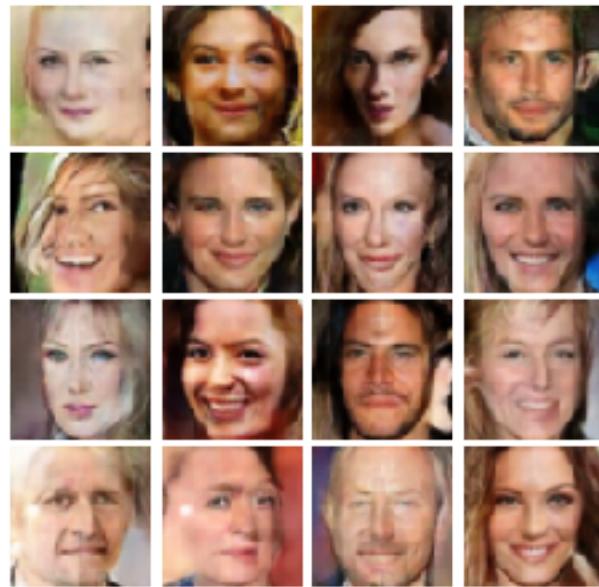
Iter: 18750, D: 0.2275, G: 2.011



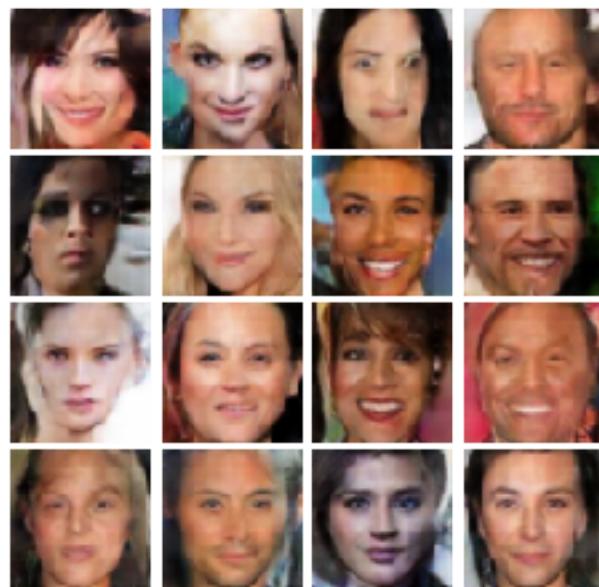
Iter: 19000, D: 0.167, G:2.663



Iter: 19250, D: 0.2132, G:1.849



Iter: 19500, D: 0.8389, G:0.7802



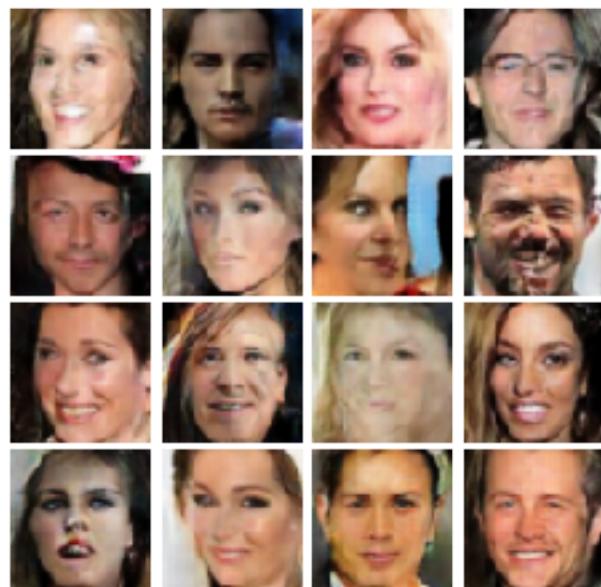
Saving the model as a checkpoint...

EPOCH: 21

Iter: 19750, D: 0.6588, G:1.82



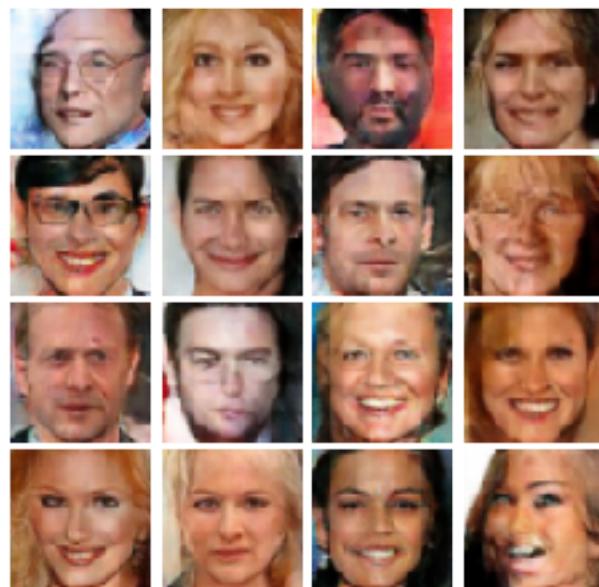
Iter: 20000, D: 0.08408, G:3.15



Iter: 20250, D: 0.2082, G:1.716



Iter: 20500, D: 0.1889, G:2.296



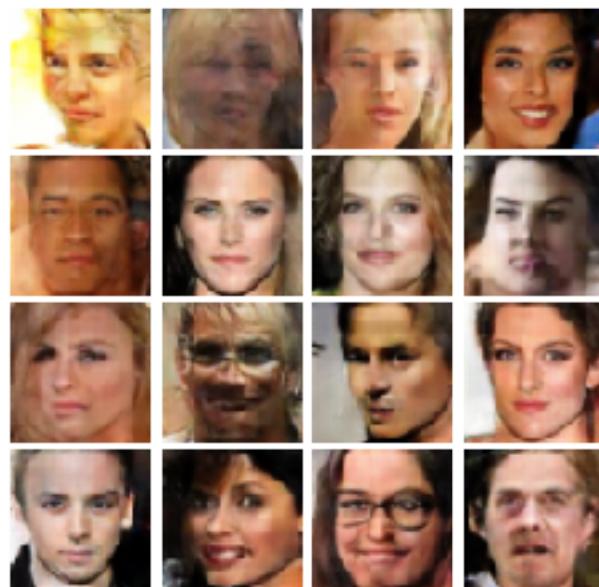
Saving the model as a checkpoint...

EPOCH: 22

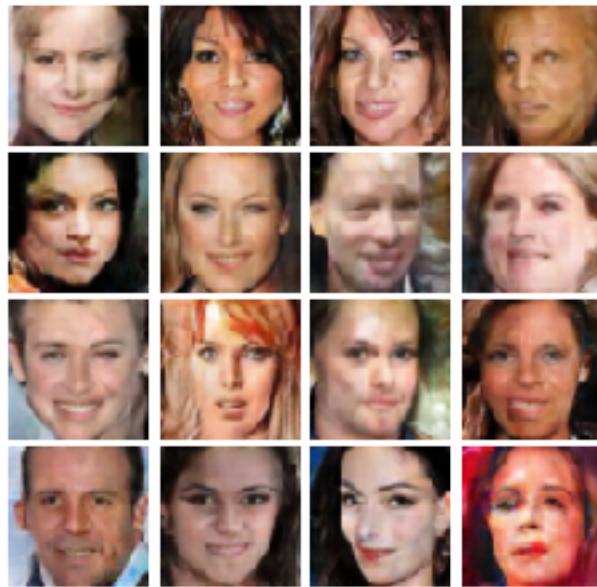
Iter: 20750, D: 0.2209, G:2.381



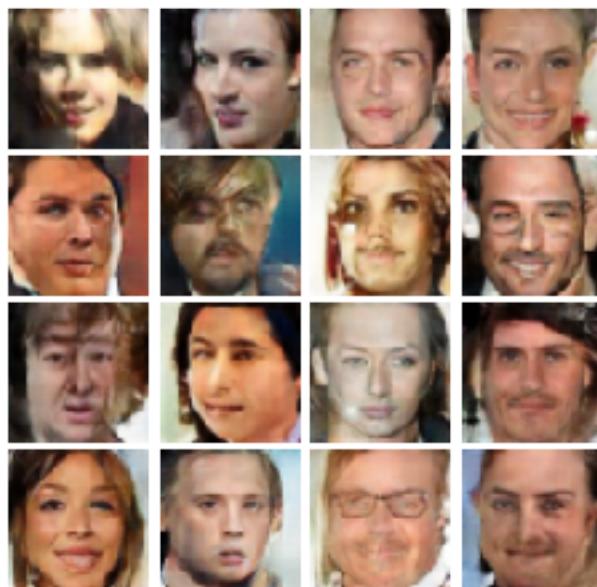
Iter: 21000, D: 0.5584, G:1.91



Iter: 21250, D: 0.2062, G:2.404



Iter: 21500, D: 0.2445, G: 2.333



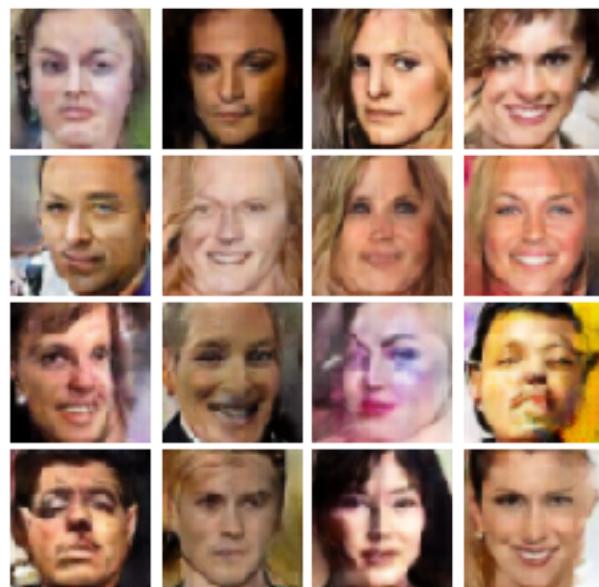
Saving the model as a checkpoint...

EPOCH: 23

Iter: 21750, D: 0.15, G: 2.216



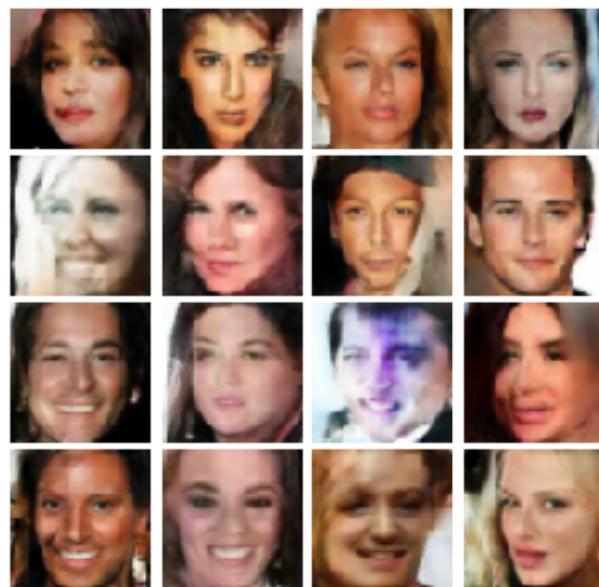
Iter: 22000, D: 0.296, G:2.576



Iter: 22250, D: 0.4117, G:2.505



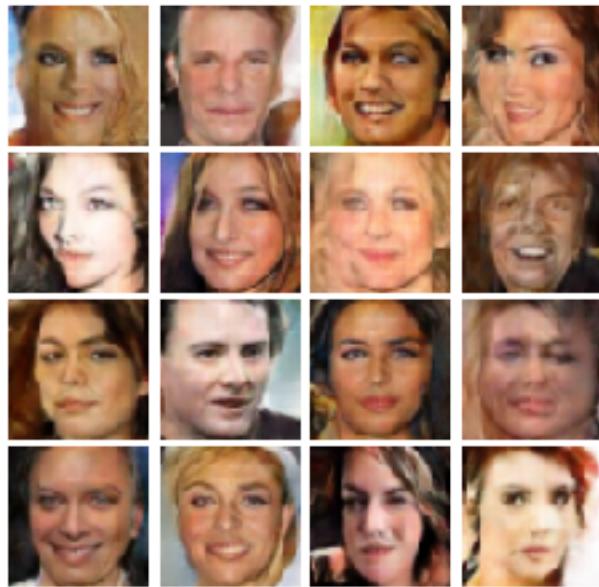
Iter: 22500, D: 0.6837, G:0.5808



Saving the model as a checkpoint...

EPOCH: 24

Iter: 22750, D: 0.4137, G:2.356



Iter: 23000, D: 0.1193, G:2.731



Iter: 23250, D: 1.676, G:2.122



Iter: 23500, D: 0.1137, G:3.327



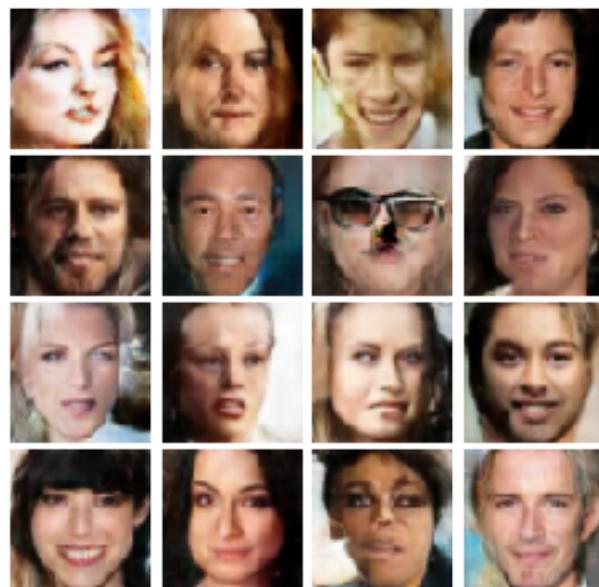
Saving the model as a checkpoint...

EPOCH: 25

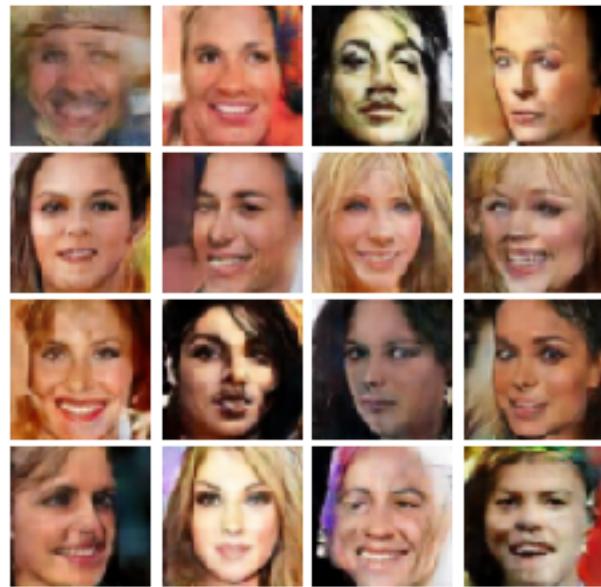
Iter: 23750, D: 0.2707, G:2.544



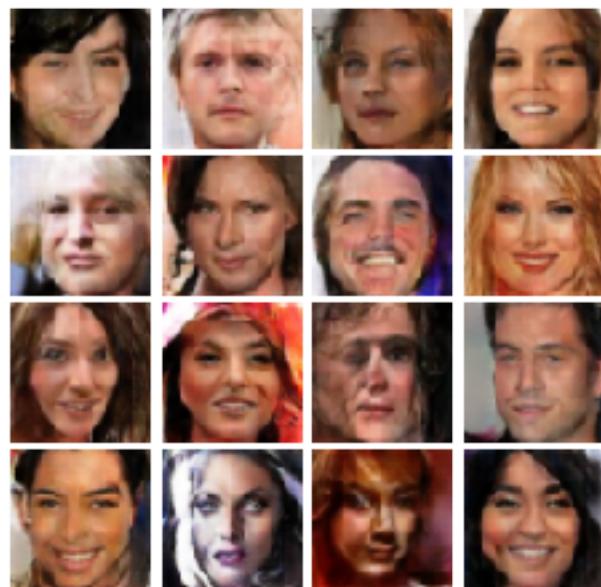
Iter: 24000, D: 1.013, G:1.25



Iter: 24250, D: 0.1149, G:2.67



Iter: 24500, D: 0.245, G:2.074



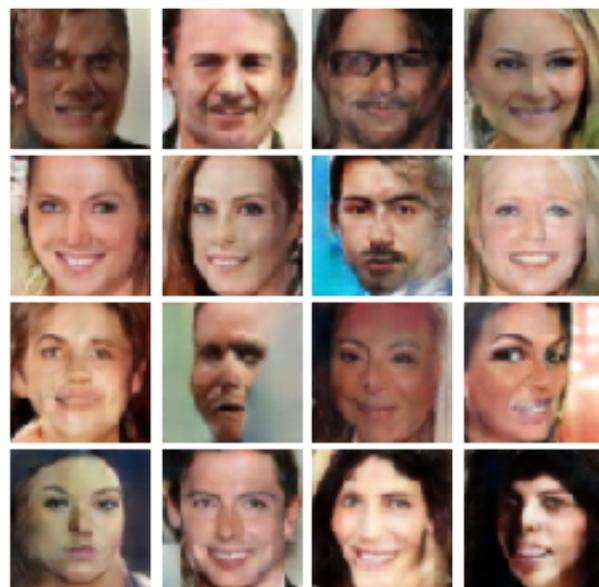
Saving the model as a checkpoint...

EPOCH: 26

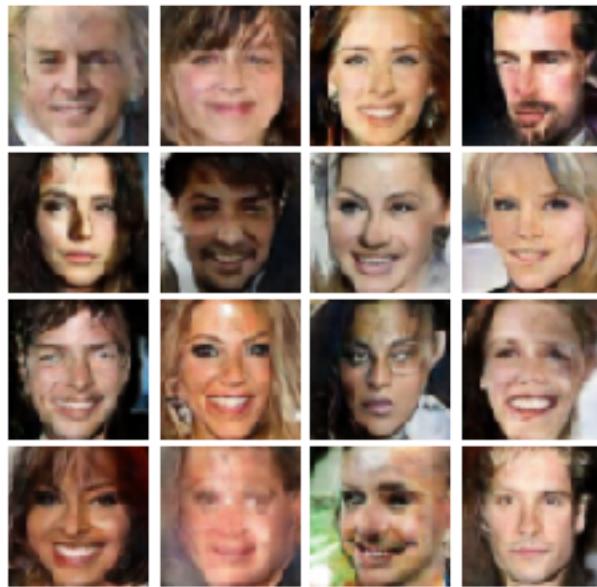
Iter: 24750, D: 0.1378, G:2.561



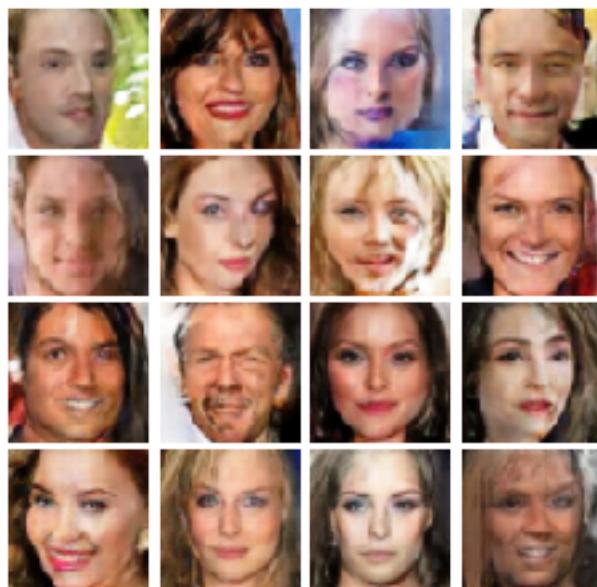
Iter: 25000, D: 0.08143, G:3.013



Iter: 25250, D: 0.1813, G:2.756



Iter: 25500, D: 0.2371, G:1.915



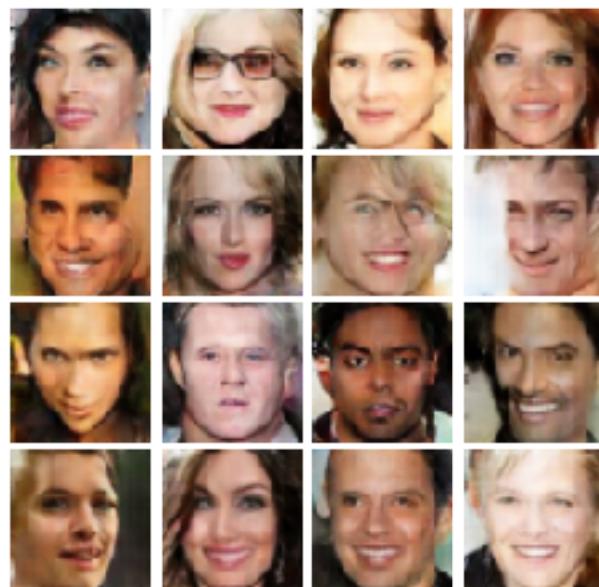
Saving the model as a checkpoint...

EPOCH: 27

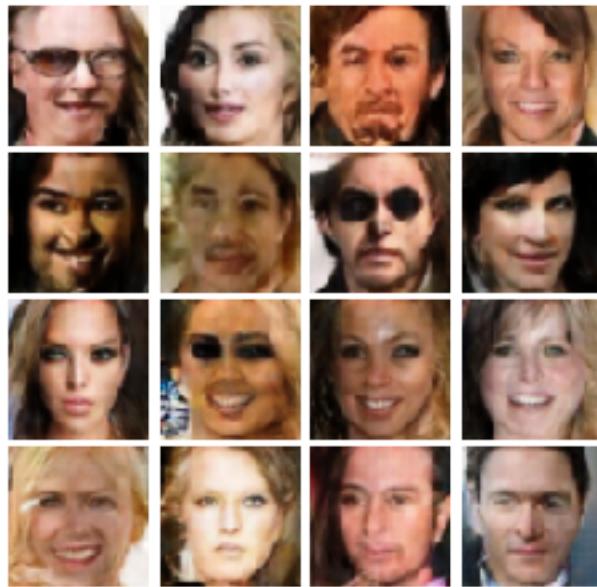
Iter: 25750, D: 0.2815, G:2.083



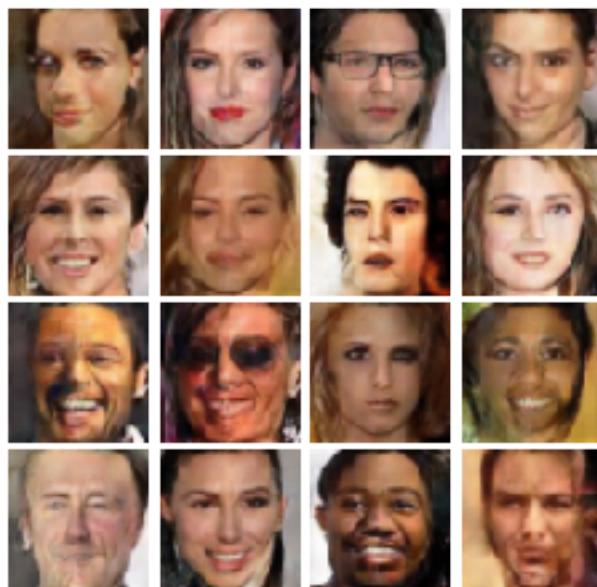
Iter: 26000, D: 0.9508, G:1.187



Iter: 26250, D: 0.6689, G:2.835



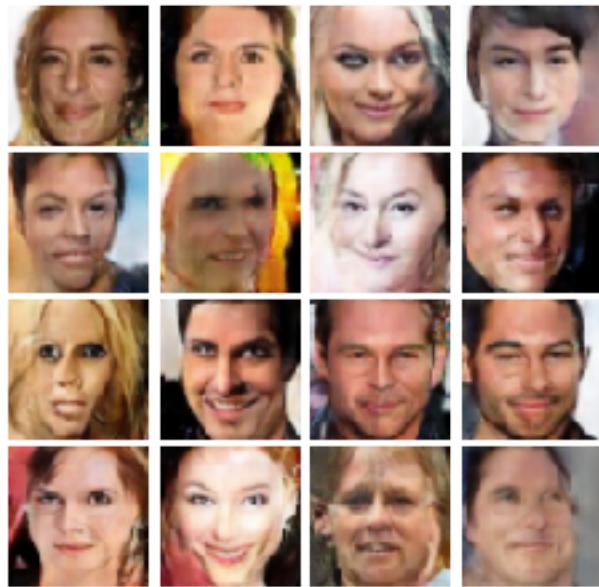
Iter: 26500, D: 0.2072, G: 2.16



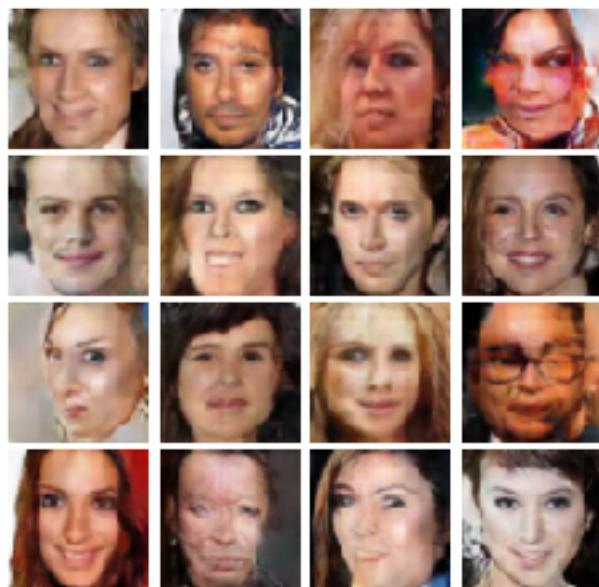
Saving the model as a checkpoint...

EPOCH: 28

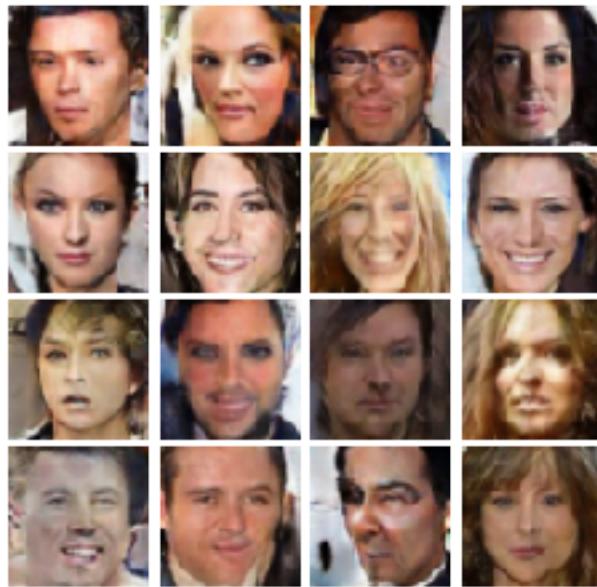
Iter: 26750, D: 0.06536, G: 3.777



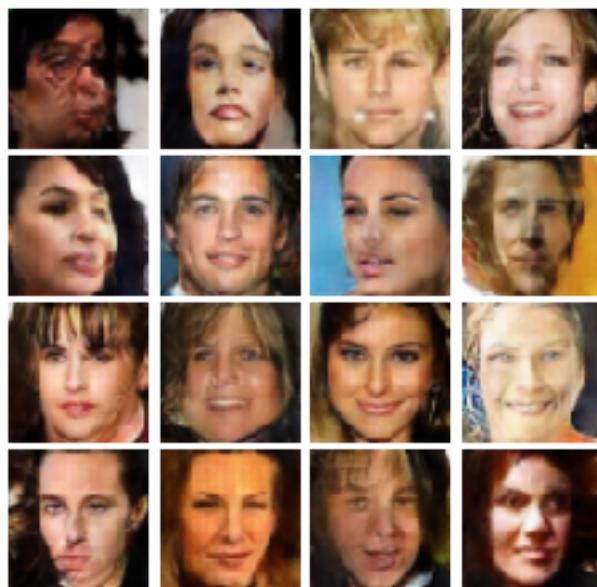
Iter: 27000, D: 0.2835, G:2.045



Iter: 27250, D: 0.1272, G:2.954



Iter: 27500, D: 0.2794, G:1.757



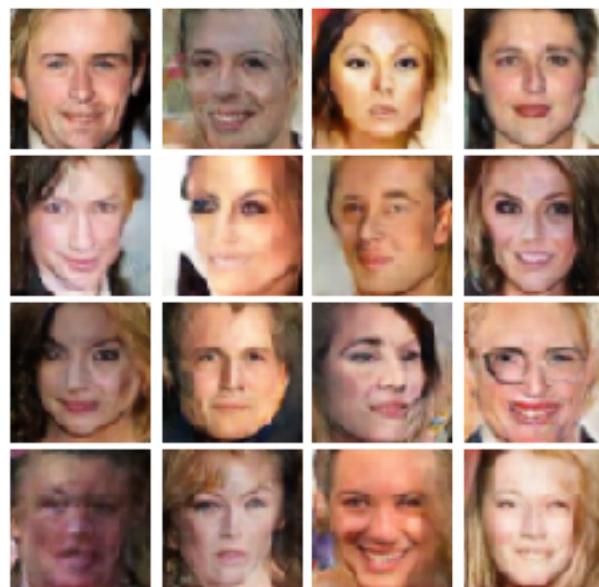
Saving the model as a checkpoint...

EPOCH: 29

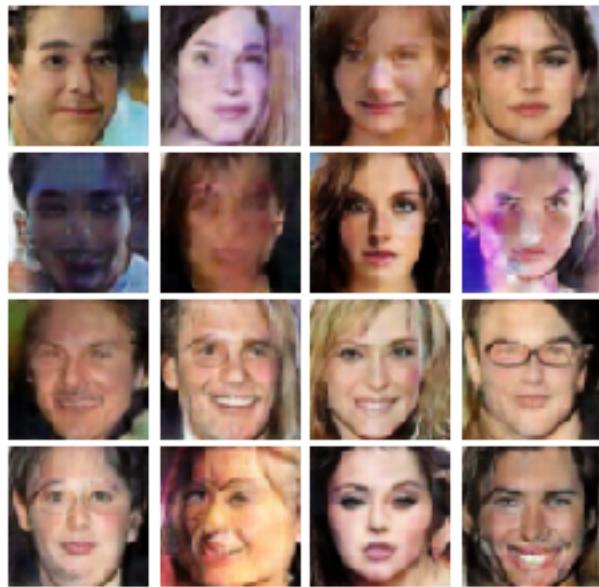
Iter: 27750, D: 0.1335, G:2.786



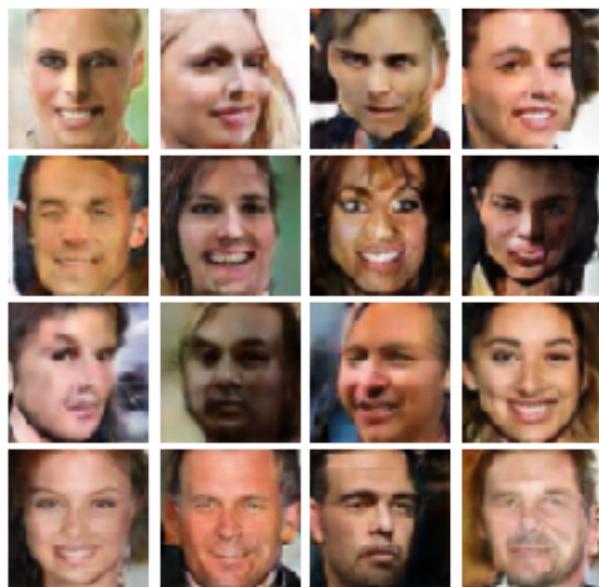
Iter: 28000, D: 0.0718, G:3.895



Iter: 28250, D: 0.2281, G:1.907



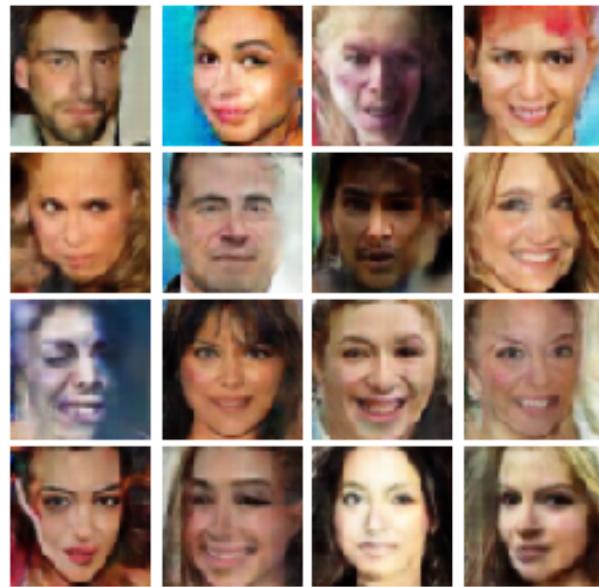
Iter: 28500, D: 0.1668, G:2.927



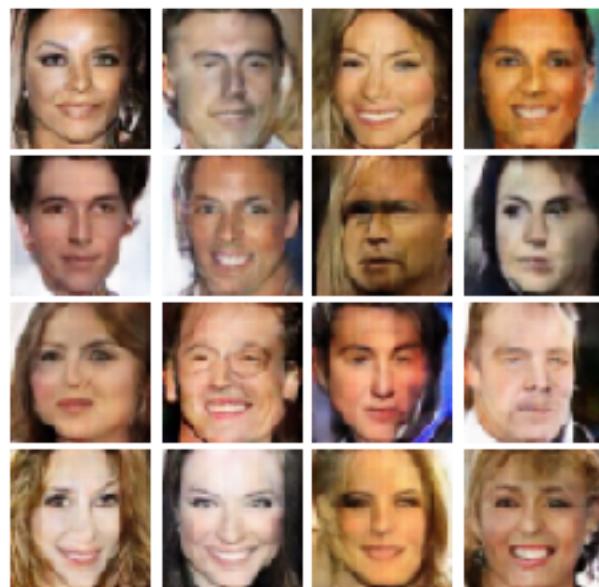
Saving the model as a checkpoint...

EPOCH: 30

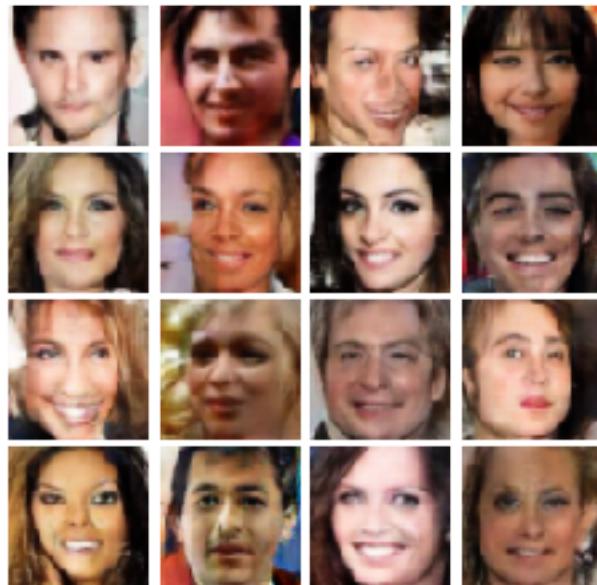
Iter: 28750, D: 0.2712, G:2.028



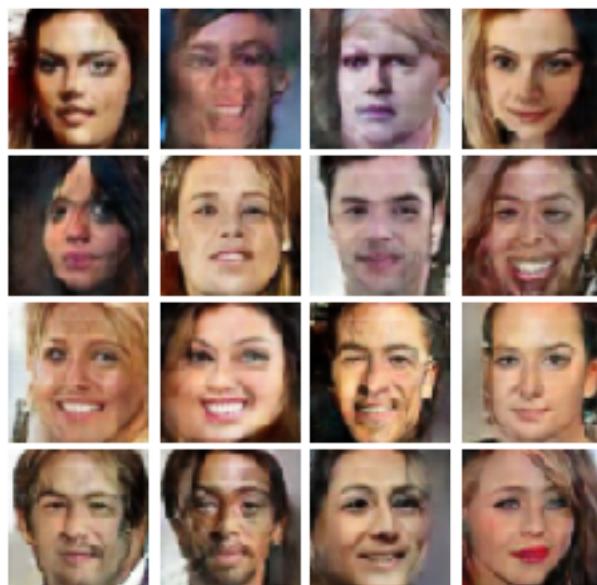
Iter: 29000, D: 0.08962, G:3.355



Iter: 29250, D: 0.5942, G:1.3



Iter: 29500, D: 0.08878, G: 3.395



Saving the model as a checkpoint...

4.0.4 Section 2.4 Train a LSGAN on CelebA [13 pts]

- Call discriminator and generator for training.
- Call optimizers for both discriminator and generator for training. (Use Adam with betas = (0.5, 0.999))
- Call train function to train.
- Train for 30 epochs.

Now, train your GAN model with LSGAN loss. **Observe the visualized result of your model, and describe what you see.**

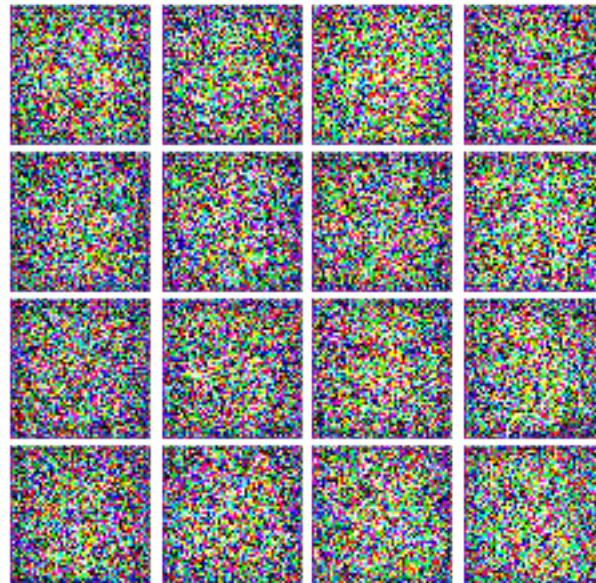
```
[ ]: # LSGAN
# Add code here:
Discriminator = Discriminator().to(device)
Generator = Generator(noise_dim=NOISE_DIM).to(device)

Discriminator_optimizer = torch.optim.Adam(Discriminator.parameters(), lr=learning_rate, betas = (0.5, 0.999))
Generator_optimizer = torch.optim.Adam(Generator.parameters(), lr=learning_rate, betas = (0.5, 0.999))

train(Discriminator, Generator, Discriminator_optimizer, Generator_optimizer, ls_discriminator_loss, ls_generator_loss, num_epochs=NUM_EPOCHS, train_loader=celeba_loader_train, device=device)
```

EPOCH: 1

Iter: 0, D: 0.6173, G:0.8



Iter: 200, D: 0.1789, G:0.6515



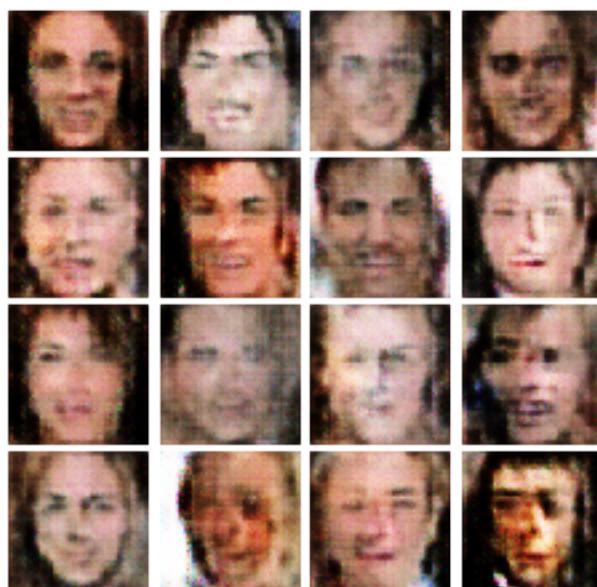
Iter: 400, D: 0.3168, G:0.3989



Iter: 600, D: 0.2022, G:0.2309



Iter: 800, D: 0.3394, G: 0.3744



Saving the model as a checkpoint...

EPOCH: 2

Iter: 1000, D: 0.2386, G: 0.332



Iter: 1200, D: 0.3897, G:0.2698



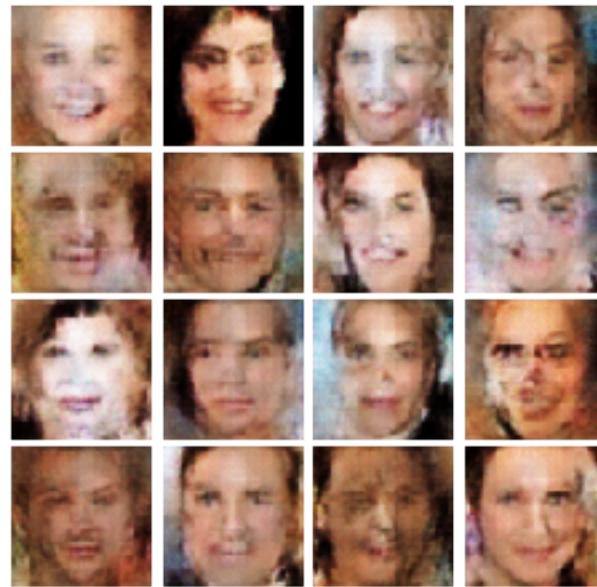
Iter: 1400, D: 0.17, G:0.3141



Iter: 1600, D: 0.2985, G: 0.2513



Iter: 1800, D: 0.2072, G: 0.1496



Saving the model as a checkpoint...

EPOCH: 3

Iter: 2000, D: 0.2835, G:0.3917



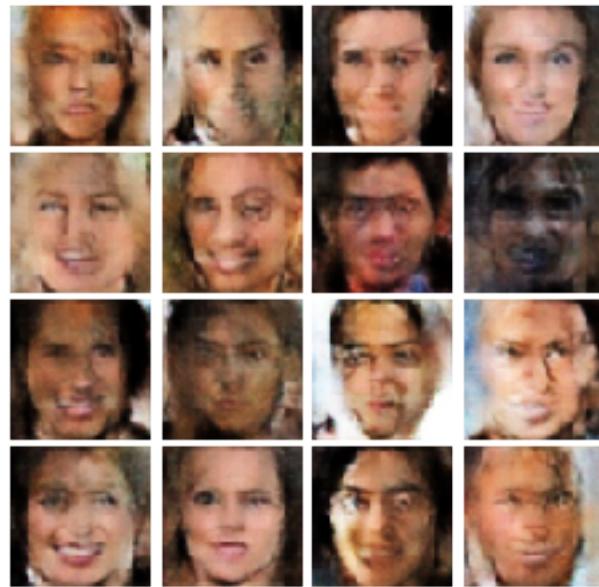
Iter: 2200, D: 0.2569, G:0.2652



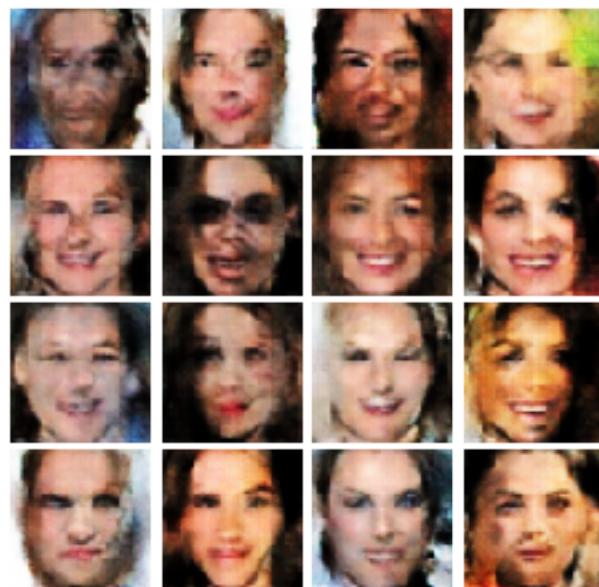
Iter: 2400, D: 0.1718, G: 0.4543



Iter: 2600, D: 0.297, G: 0.4744



Iter: 2800, D: 0.2909, G: 0.3571



Saving the model as a checkpoint...

EPOCH: 4

Iter: 3000, D: 0.6384, G: 0.6455



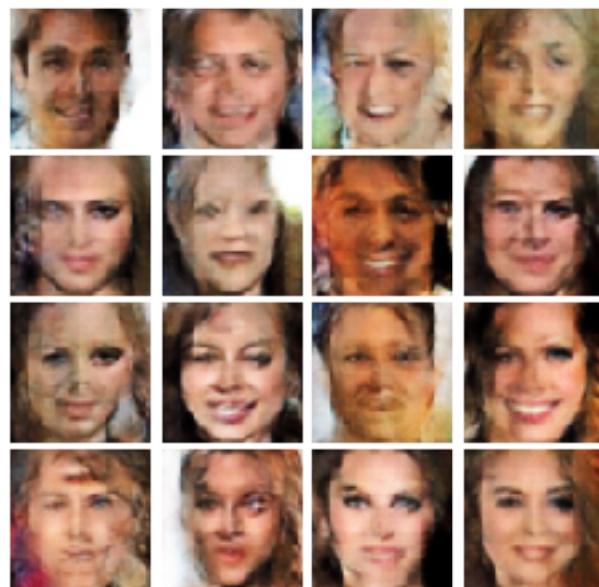
Iter: 3200, D: 0.8428, G: 0.6925



Iter: 3400, D: 0.2726, G: 0.3374



Iter: 3600, D: 0.1845, G: 0.3391



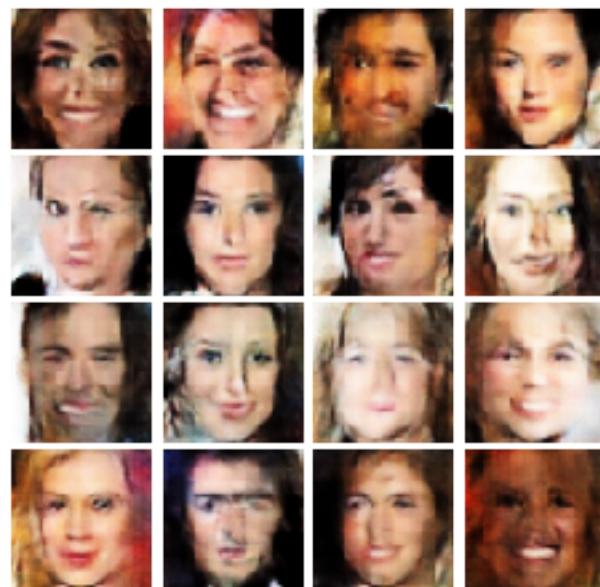
Iter: 3800, D: 0.2282, G: 0.3251



Saving the model as a checkpoint...

EPOCH: 5

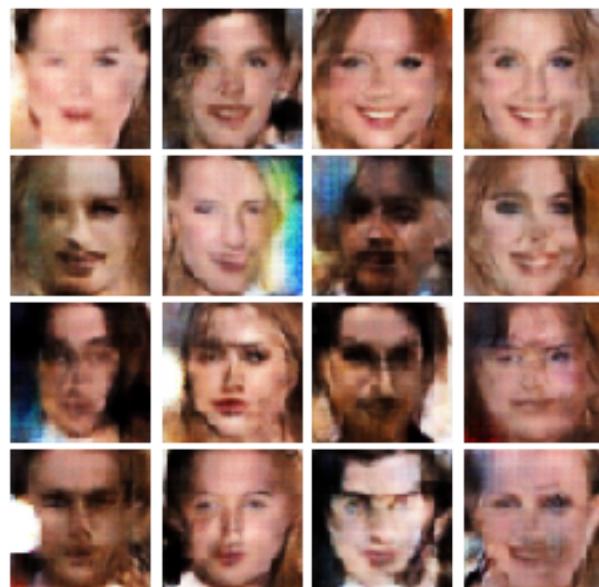
Iter: 4000, D: 0.2777, G: 0.2091



Iter: 4200, D: 0.2756, G: 0.1804



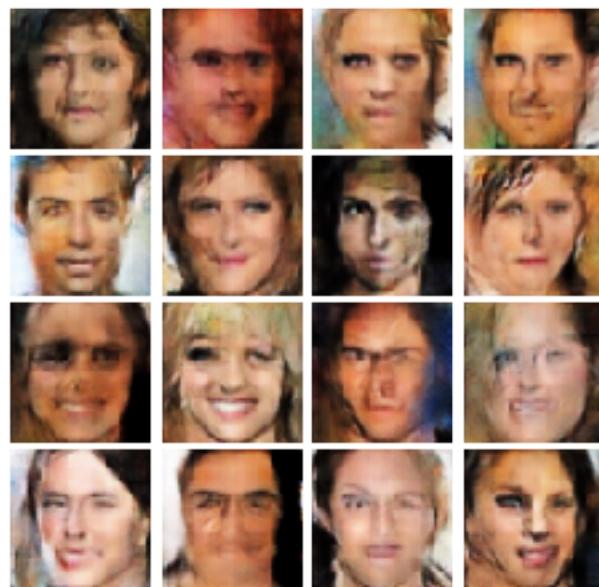
Iter: 4400, D: 0.4627, G: 0.2294



Iter: 4600, D: 0.3075, G: 0.3127



Iter: 4800, D: 0.209, G: 0.2595



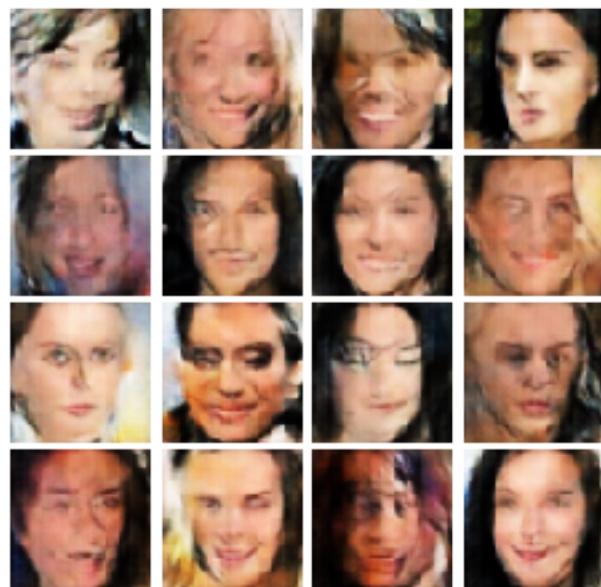
Saving the model as a checkpoint...

EPOCH: 6

Iter: 5000, D: 0.2472, G: 0.3804



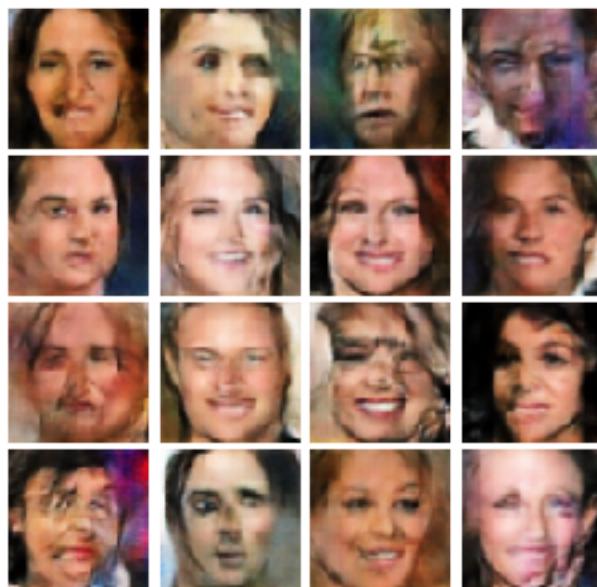
Iter: 5200, D: 0.3141, G: 0.2208



Iter: 5400, D: 0.2378, G: 0.5414



Iter: 5600, D: 0.2329, G: 0.3406



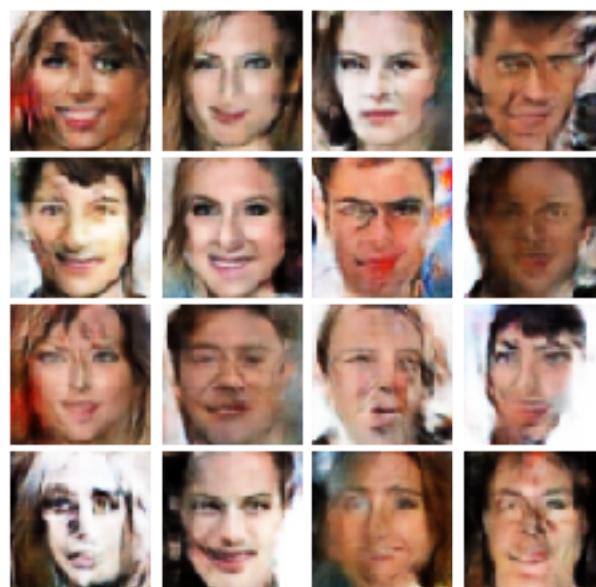
Iter: 5800, D: 0.2715, G: 0.1634



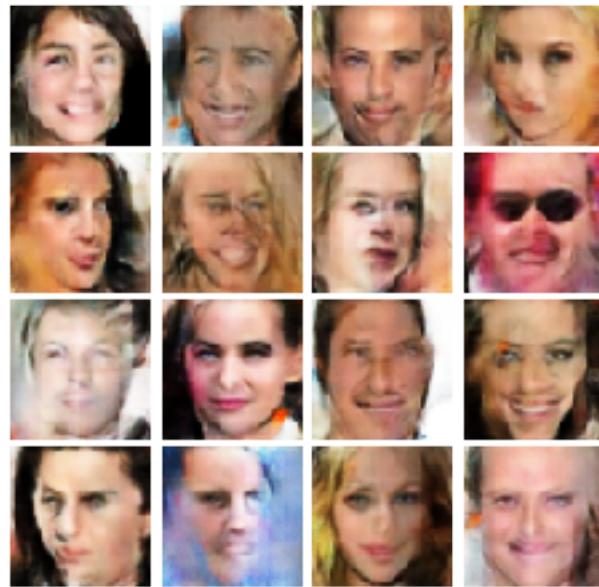
Saving the model as a checkpoint...

EPOCH: 7

Iter: 6000, D: 0.2105, G: 0.2738



Iter: 6200, D: 0.2062, G: 0.3741



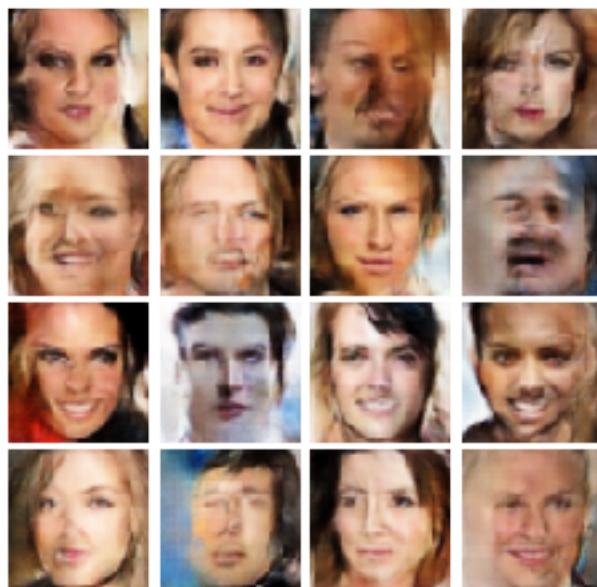
Iter: 6400, D: 0.1721, G: 0.2557



Iter: 6600, D: 0.3084, G: 0.389



Iter: 6800, D: 0.1698, G: 0.4296



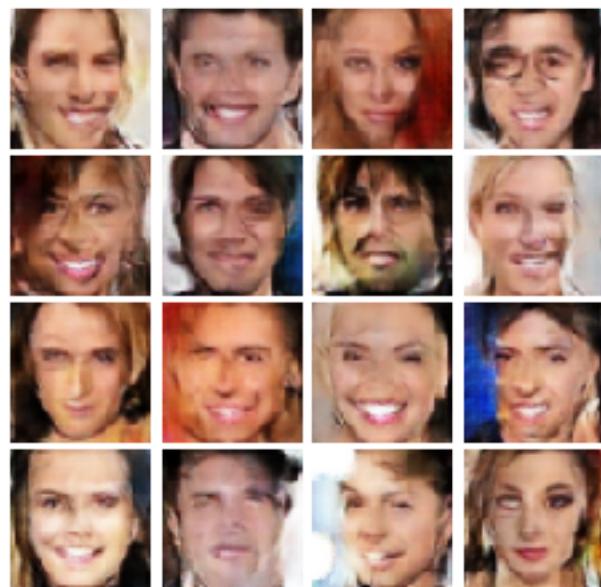
Saving the model as a checkpoint...

EPOCH: 8

Iter: 7000, D: 0.1234, G: 0.3987



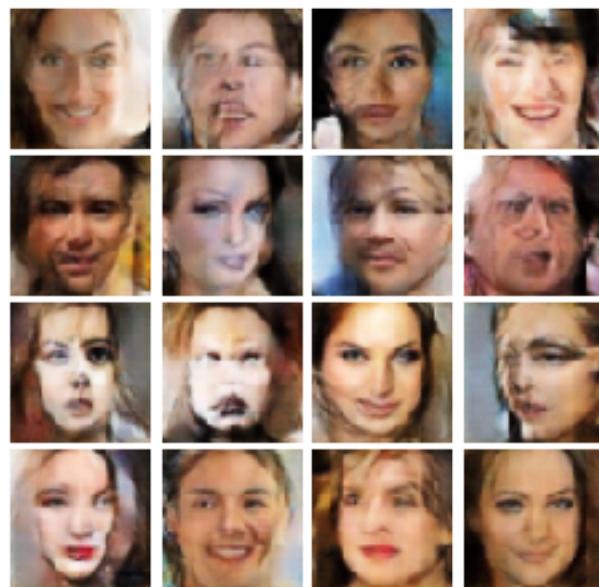
Iter: 7200, D: 0.24, G: 0.2989



Iter: 7400, D: 0.2233, G: 0.3457



Iter: 7600, D: 0.1863, G: 0.268



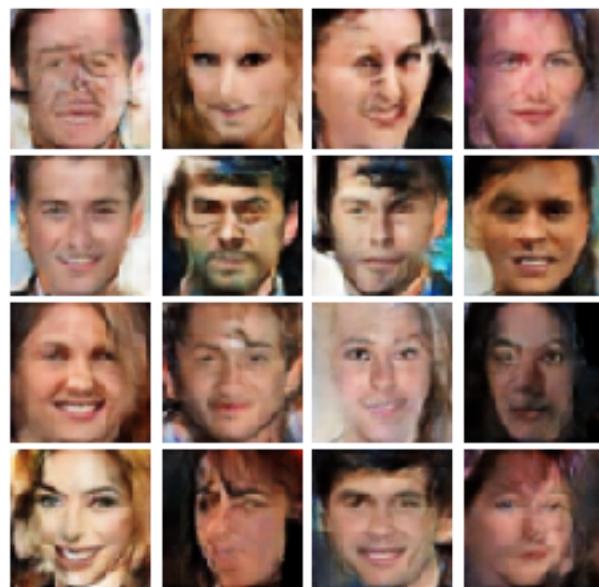
Iter: 7800, D: 0.1122, G: 0.4481



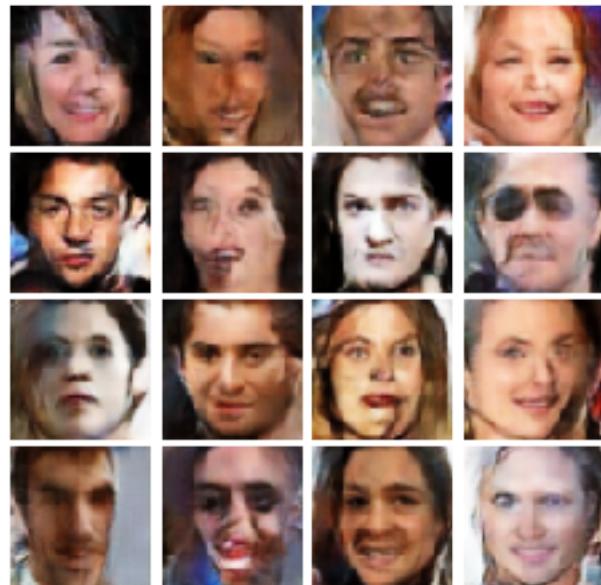
Saving the model as a checkpoint...

EPOCH: 9

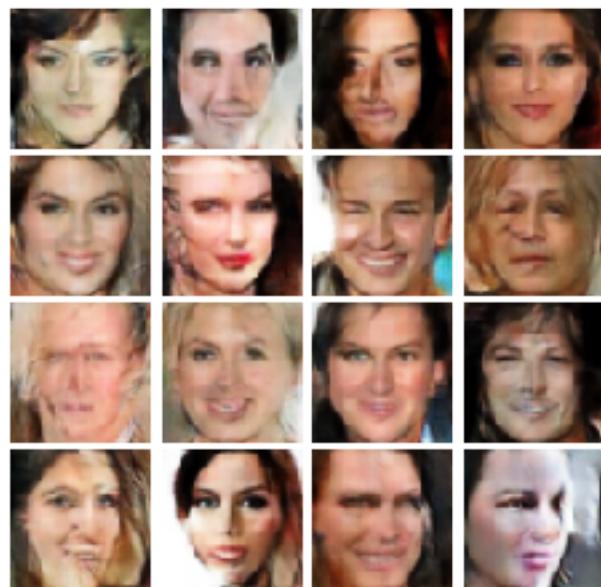
Iter: 8000, D: 0.2978, G: 0.362



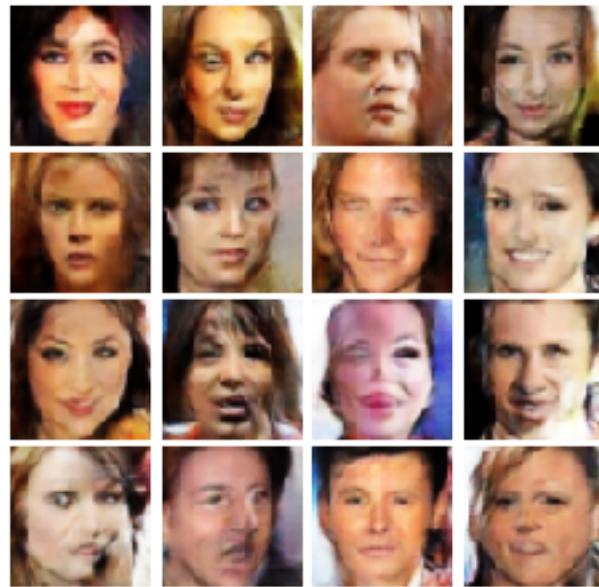
Iter: 8200, D: 0.176, G: 0.4461



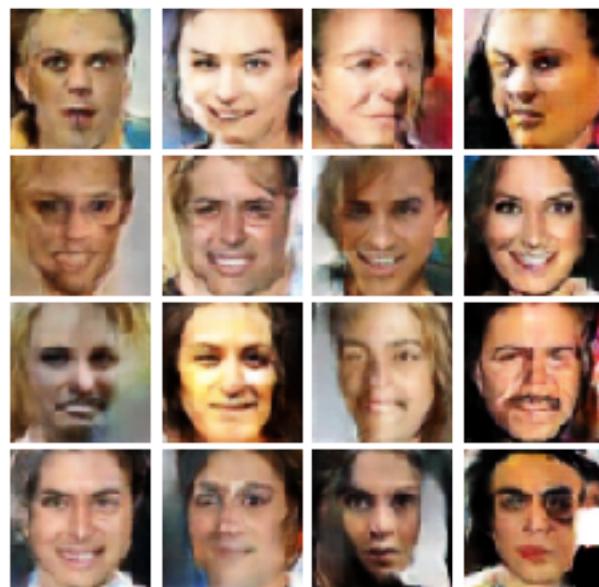
Iter: 8400, D: 0.2866, G:0.3153



Iter: 8600, D: 0.1465, G:0.481



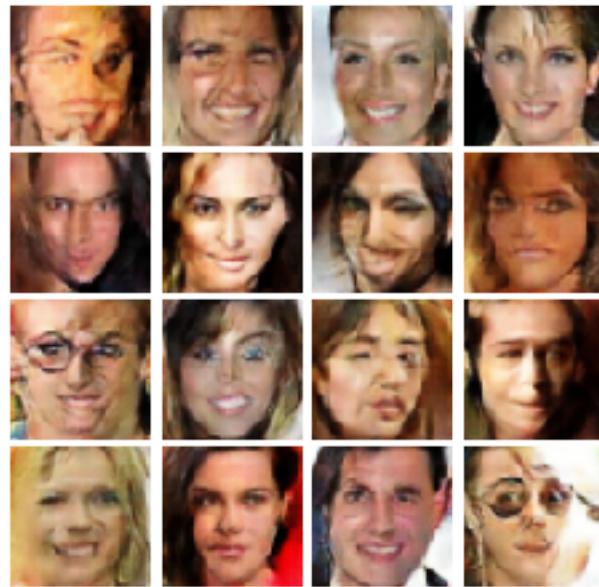
Iter: 8800, D: 0.1176, G:0.5146



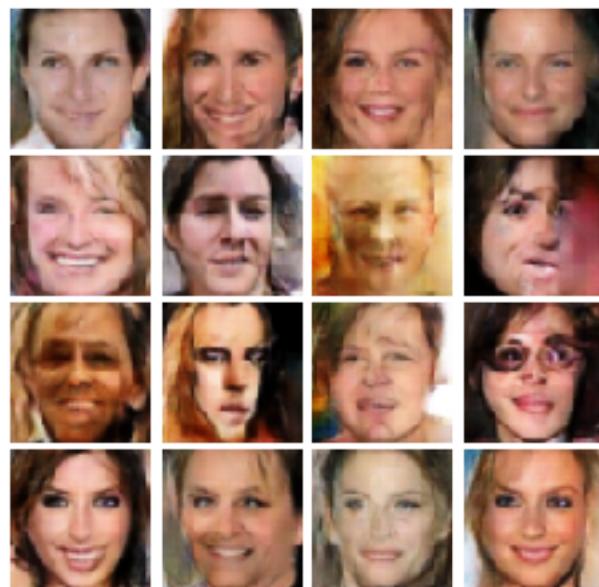
Saving the model as a checkpoint...

EPOCH: 10

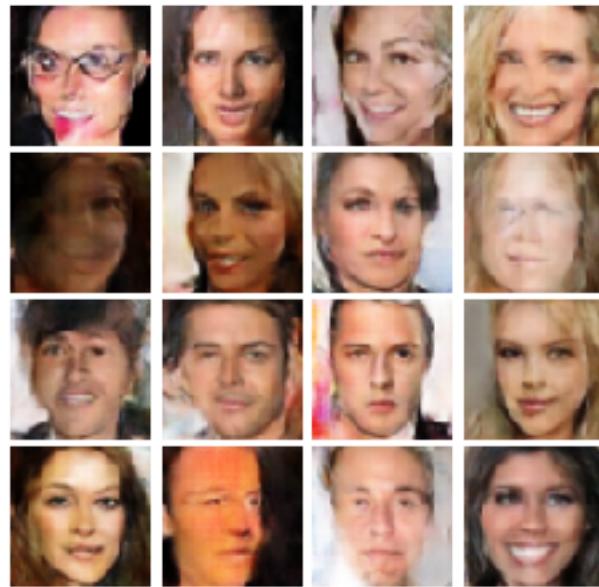
Iter: 9000, D: 0.2112, G:0.5374



Iter: 9200, D: 0.1241, G: 0.3211



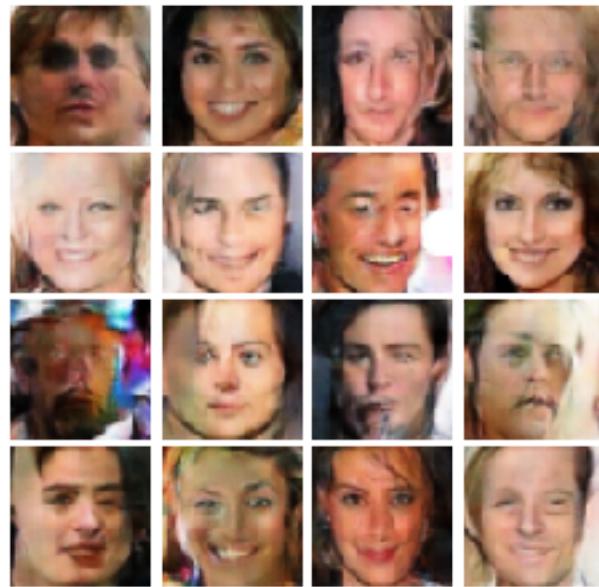
Iter: 9400, D: 0.1038, G: 0.494



Iter: 9600, D: 0.2842, G: 0.6619



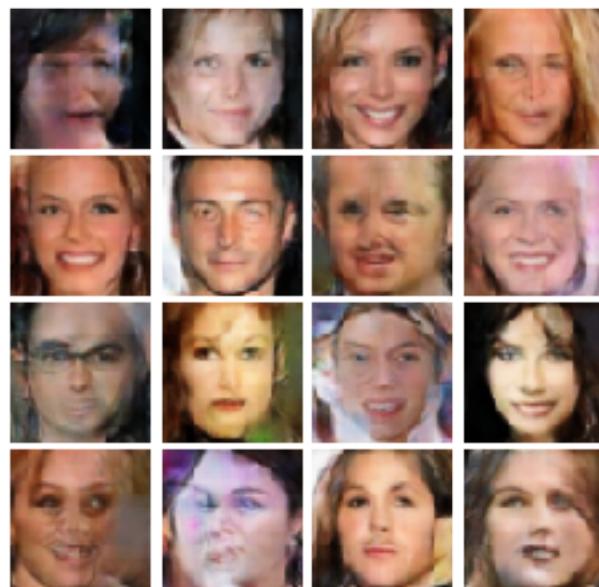
Iter: 9800, D: 0.1257, G: 0.3256



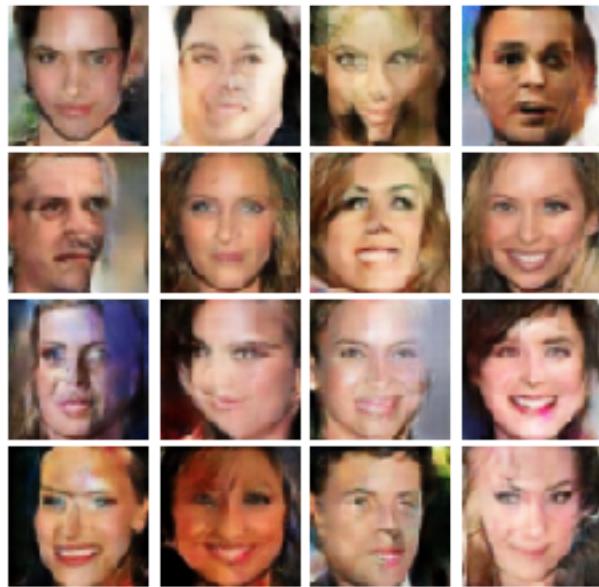
Saving the model as a checkpoint...

EPOCH: 11

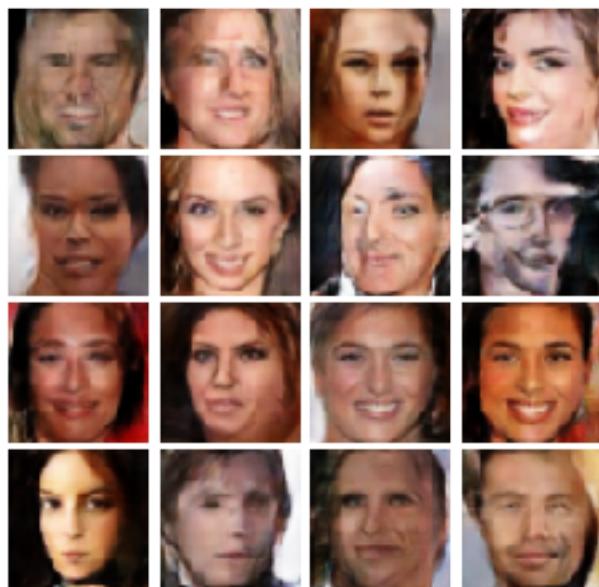
Iter: 10000, D: 0.2099, G: 0.2879



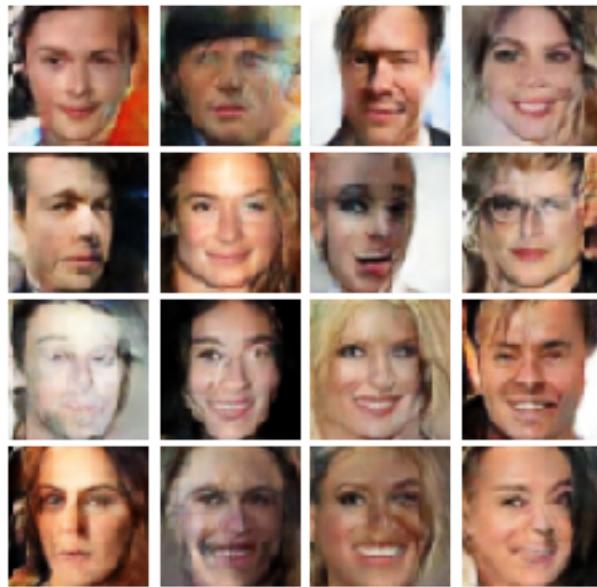
Iter: 10200, D: 0.2098, G: 0.5236



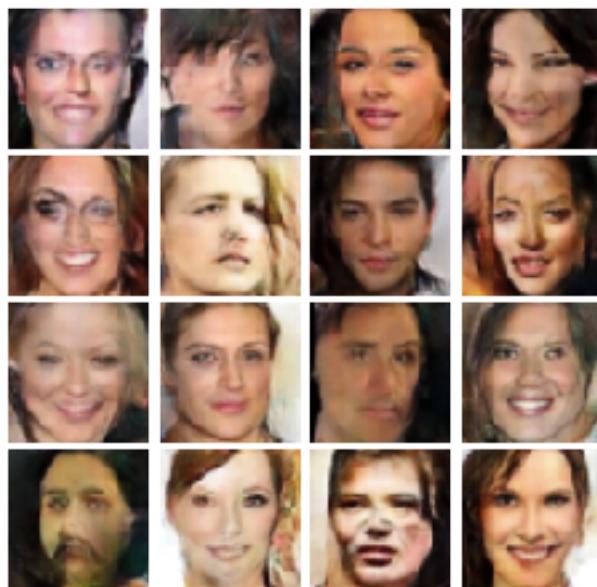
Iter: 10400, D: 0.1164, G:0.259



Iter: 10600, D: 0.1292, G:0.4572



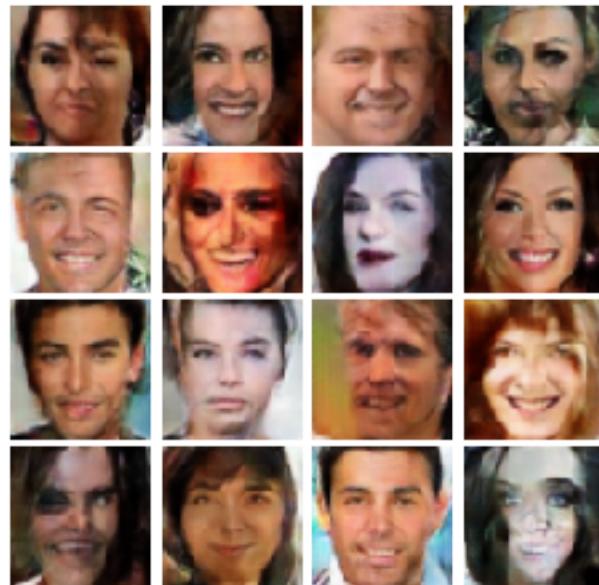
Iter: 10800, D: 0.07276, G: 0.4648



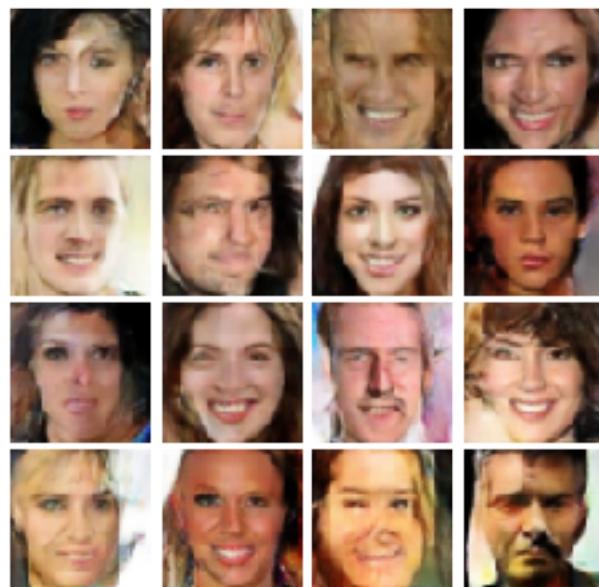
Saving the model as a checkpoint...

EPOCH: 12

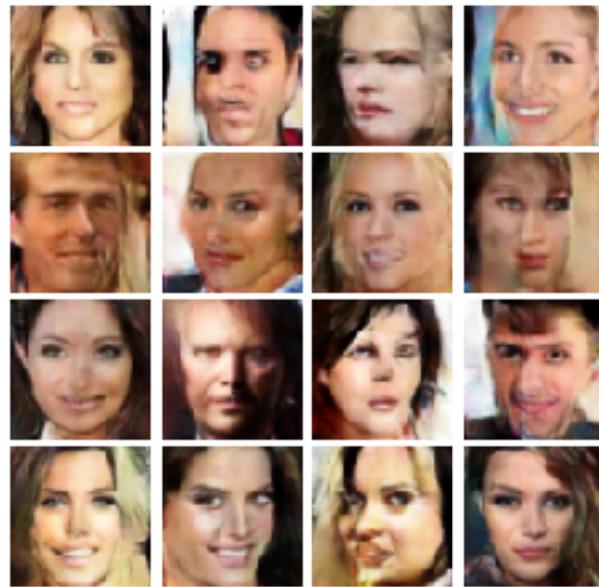
Iter: 11000, D: 0.09715, G: 0.4284



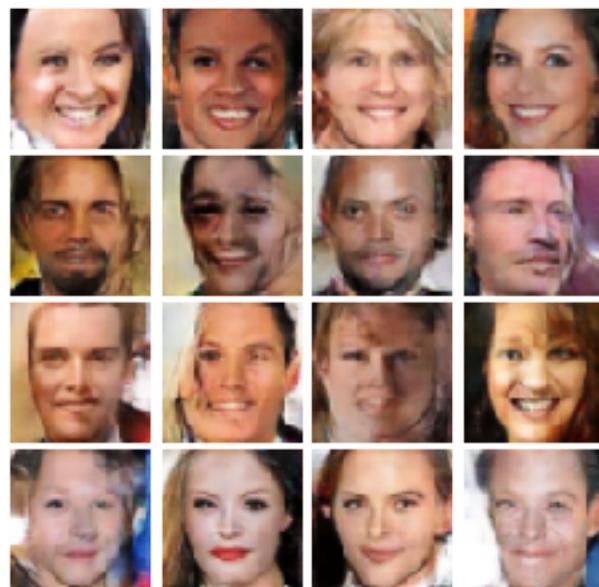
Iter: 11200, D: 0.0942, G: 0.4098



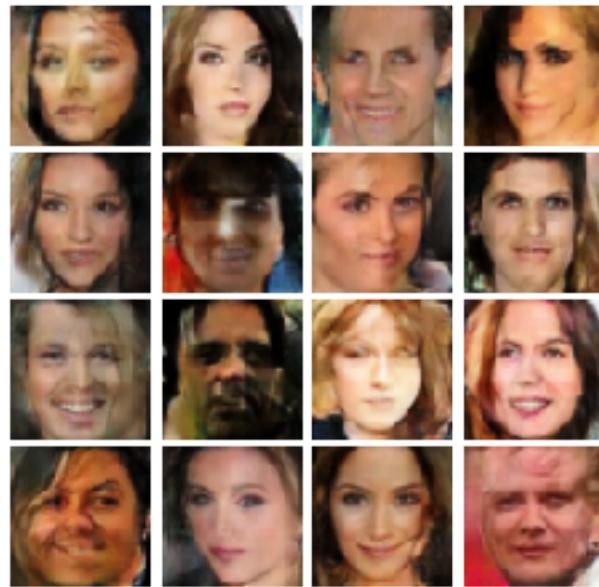
Iter: 11400, D: 0.09395, G: 0.5904



Iter: 11600, D: 0.0373, G: 0.2524



Iter: 11800, D: 0.1111, G: 0.3459



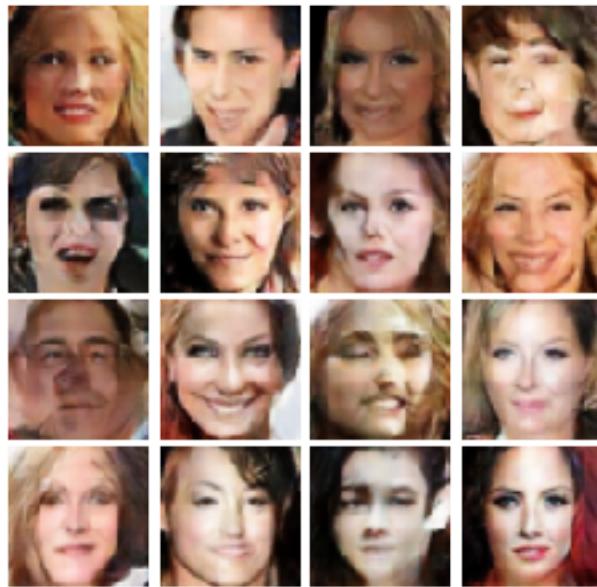
Saving the model as a checkpoint...

EPOCH: 13

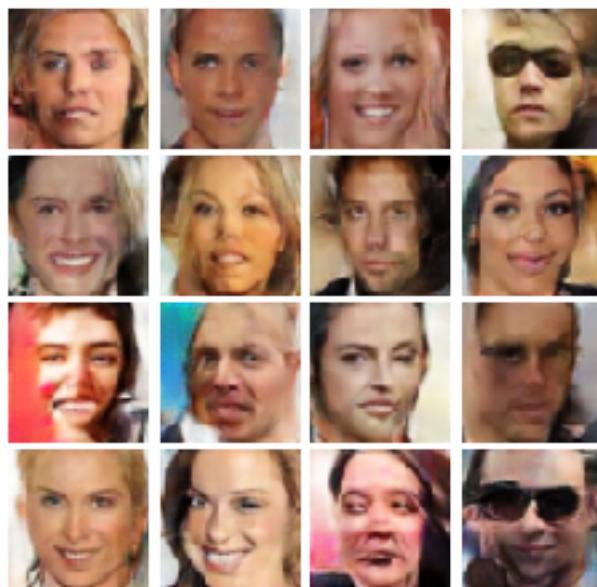
Iter: 12000, D: 0.1485, G: 0.5799



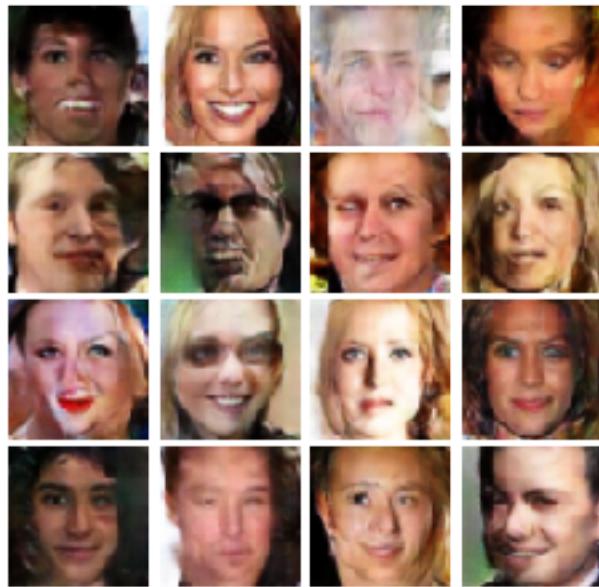
Iter: 12200, D: 0.2403, G: 0.5882



Iter: 12400, D: 0.1716, G:0.3003



Iter: 12600, D: 0.07261, G:0.4543



Iter: 12800, D: 0.1268, G:0.5717



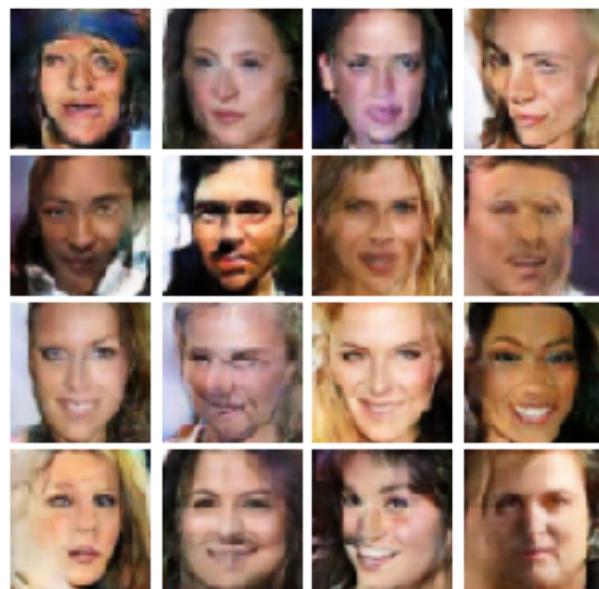
Saving the model as a checkpoint...

EPOCH: 14

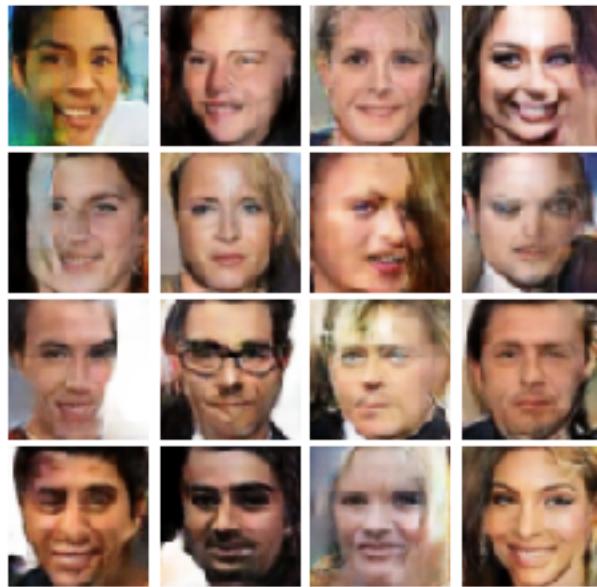
Iter: 13000, D: 0.04627, G:0.3827



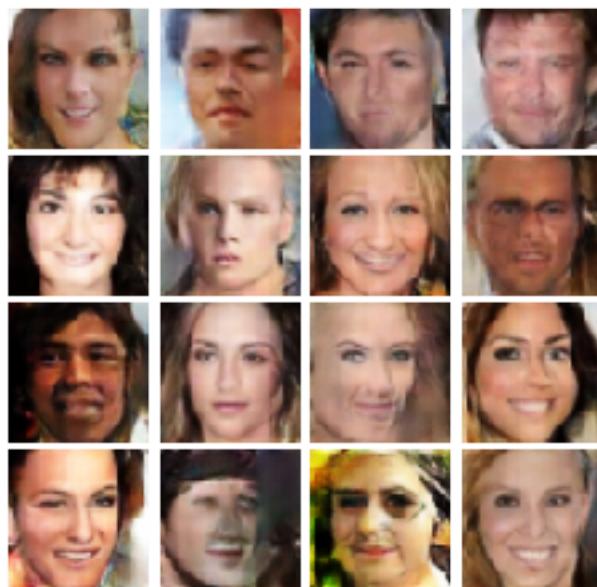
Iter: 13200, D: 0.08964, G:0.5204



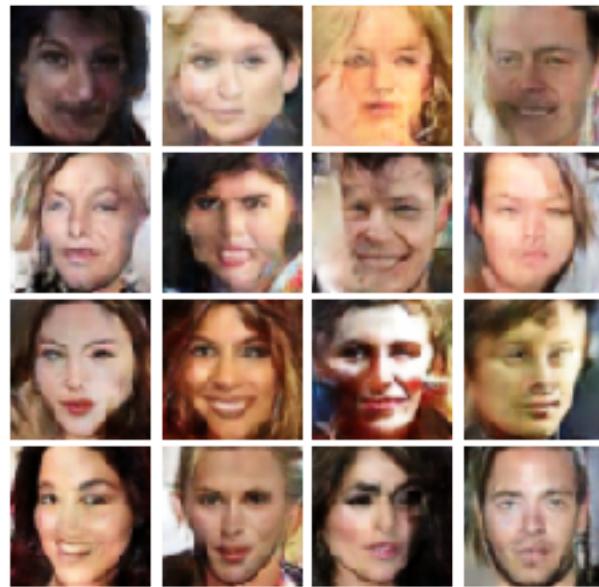
Iter: 13400, D: 0.05401, G:0.442



Iter: 13600, D: 0.1019, G: 0.4889



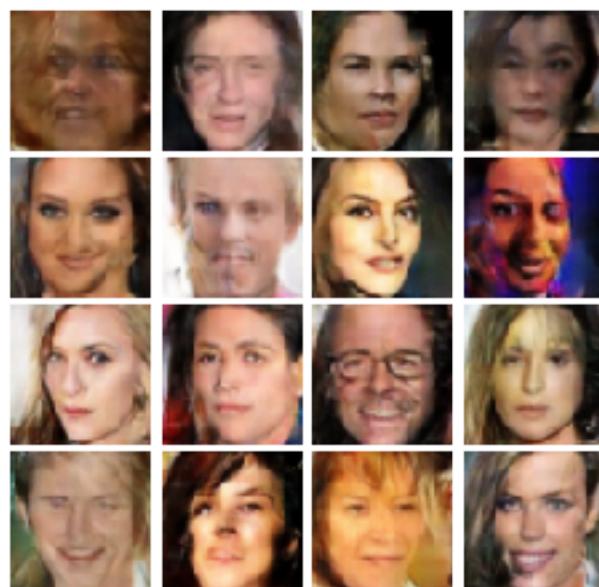
Iter: 13800, D: 0.05001, G: 0.4037



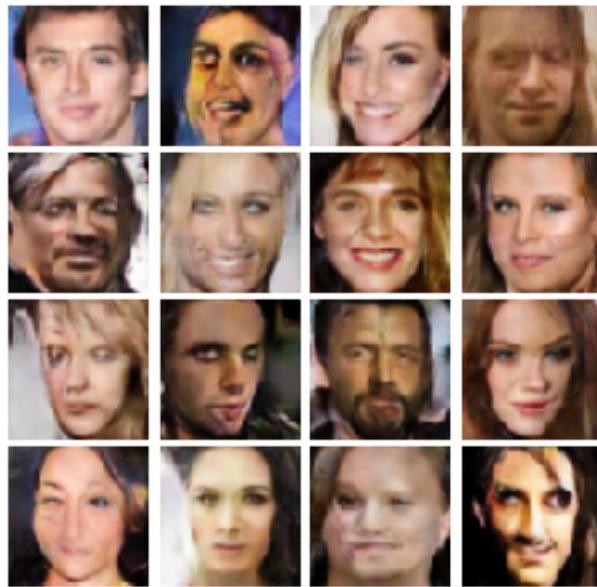
Saving the model as a checkpoint...

EPOCH: 15

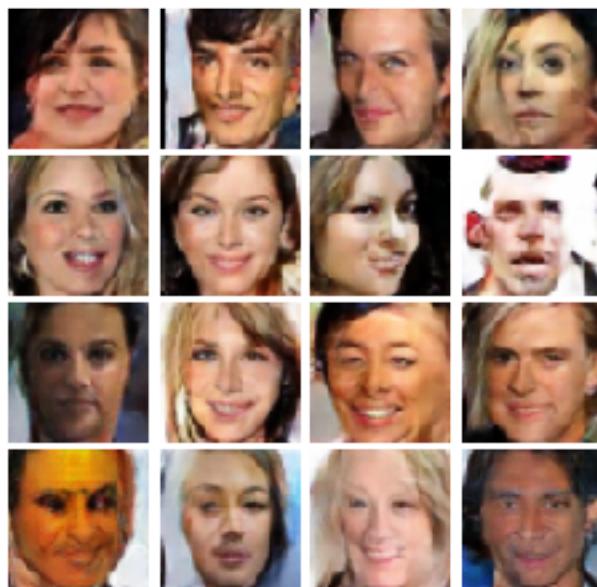
Iter: 14000, D: 0.04998, G:0.4675



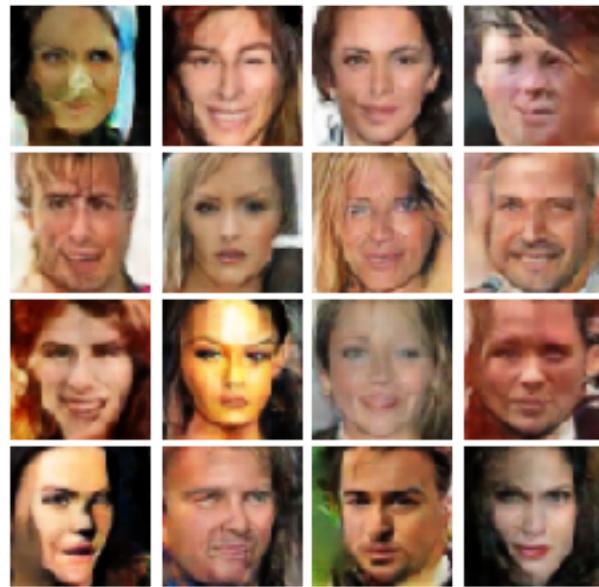
Iter: 14200, D: 0.1061, G:0.3915



Iter: 14400, D: 0.06562, G:0.5335



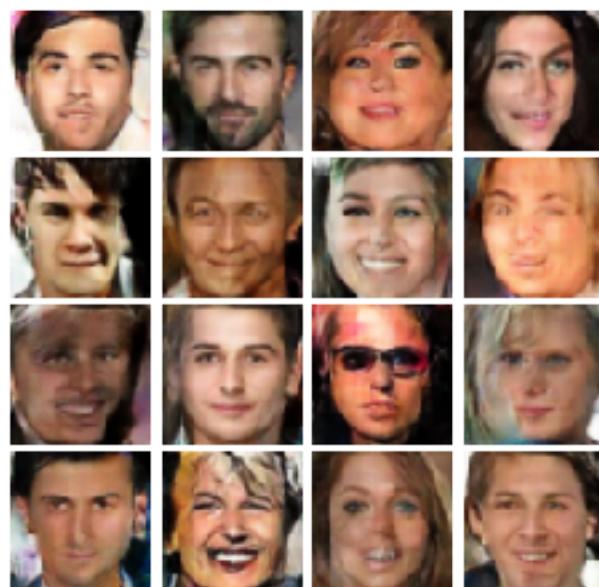
Iter: 14600, D: 0.08865, G:0.4253



Saving the model as a checkpoint...

EPOCH: 16

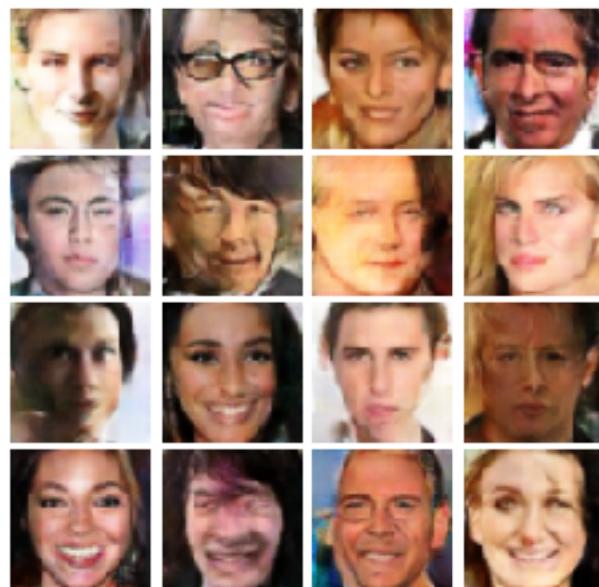
Iter: 14800, D: 0.04332, G:0.3664



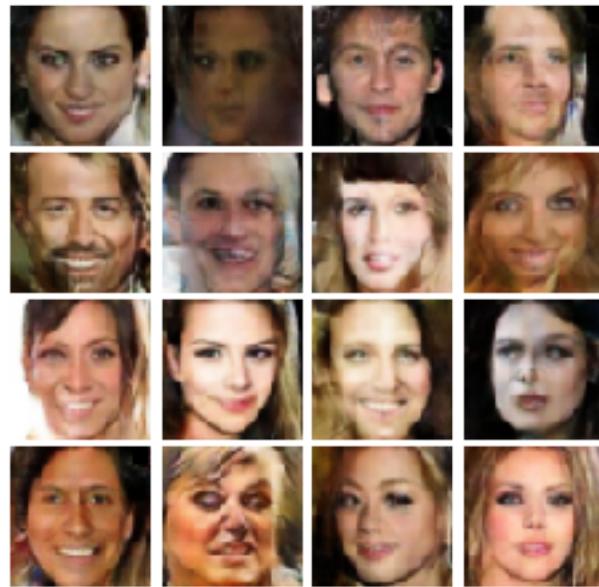
Iter: 15000, D: 0.05888, G:0.5161



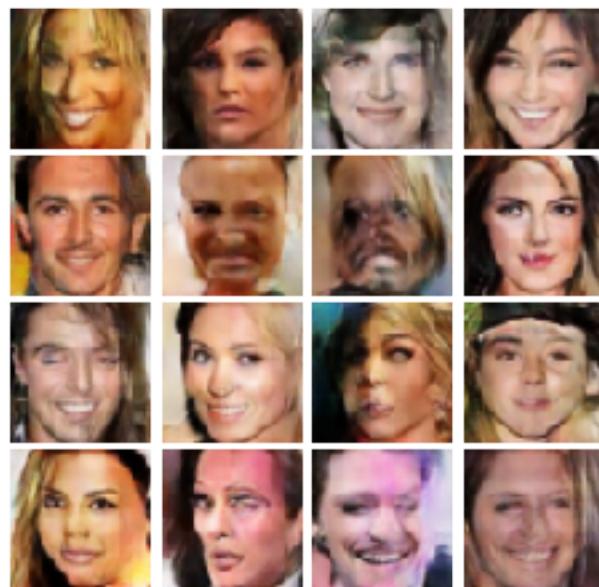
Iter: 15200, D: 0.1082, G: 0.4069



Iter: 15400, D: 0.04184, G: 0.5134



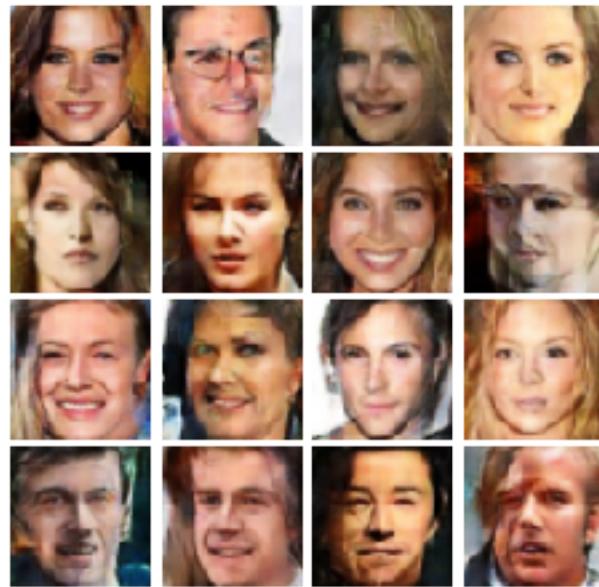
Iter: 15600, D: 0.07799, G: 0.4987



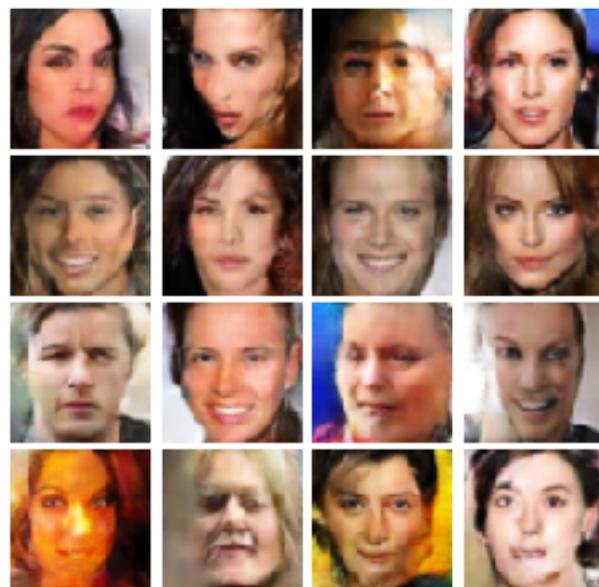
Saving the model as a checkpoint...

EPOCH: 17

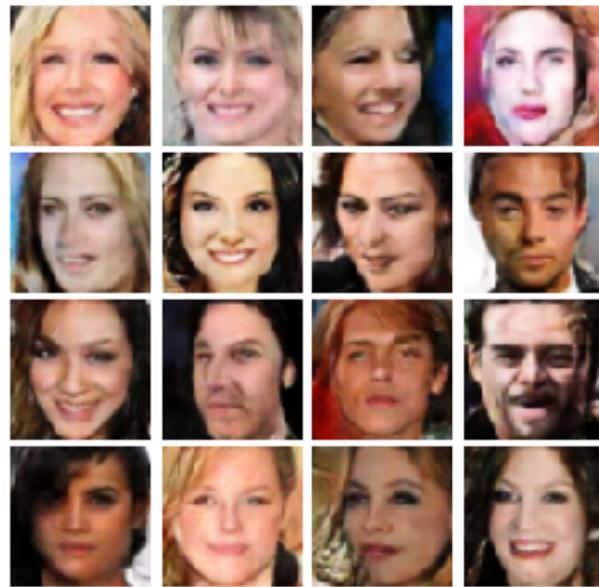
Iter: 15800, D: 0.03528, G: 0.5378



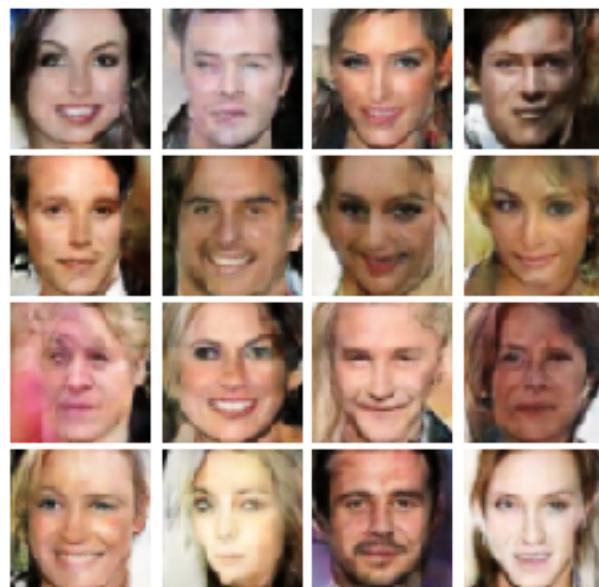
Iter: 16000, D: 0.06807, G: 0.2614



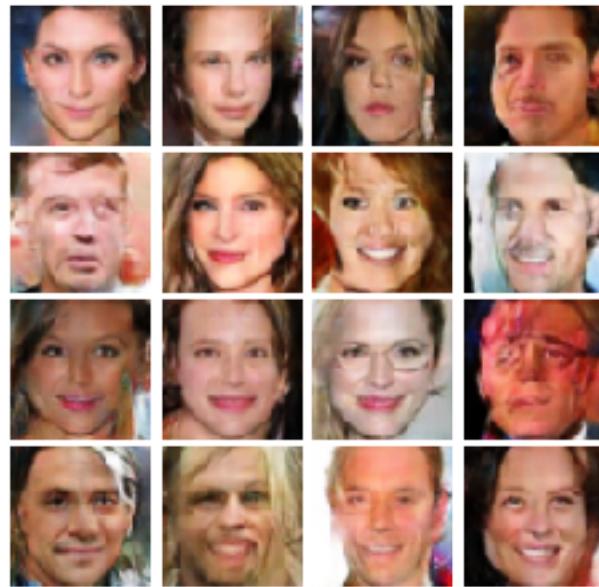
Iter: 16200, D: 0.04126, G: 0.5134



Iter: 16400, D: 0.0379, G: 0.3274



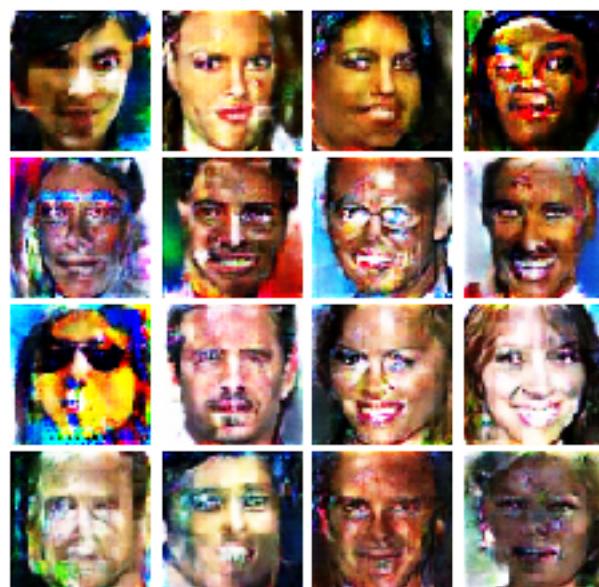
Iter: 16600, D: 0.04021, G: 0.3096



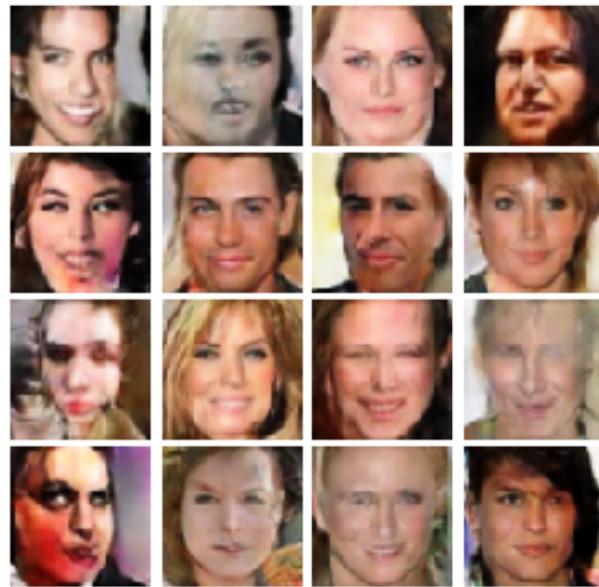
Saving the model as a checkpoint...

EPOCH: 18

Iter: 16800, D: 0.106, G: 0.6336



Iter: 17000, D: 0.03643, G: 0.4047



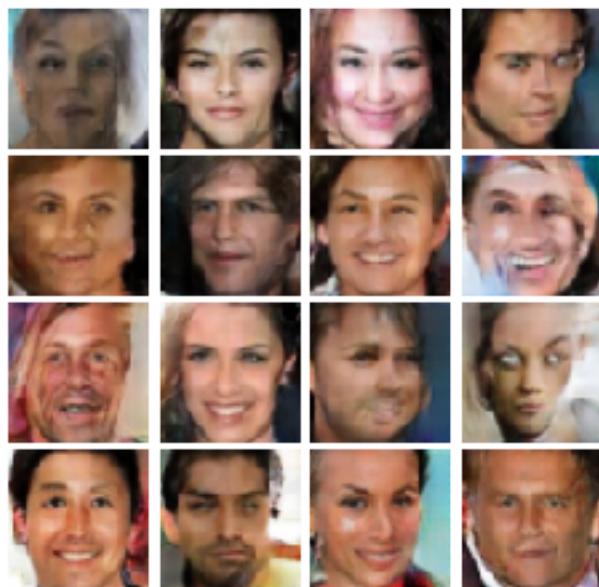
Iter: 17200, D: 0.05067, G:0.4933



Iter: 17400, D: 0.03425, G:0.4734



Iter: 17600, D: 0.09341, G: 0.1841



Saving the model as a checkpoint...

EPOCH: 19

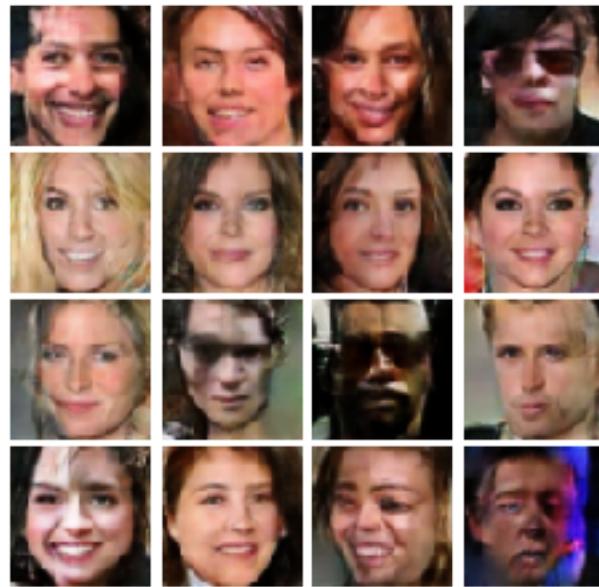
Iter: 17800, D: 0.1213, G: 0.2111



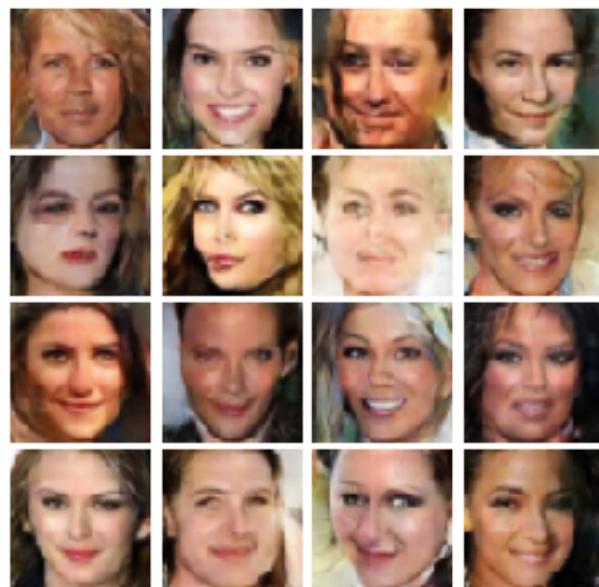
Iter: 18000, D: 0.07009, G: 0.4818



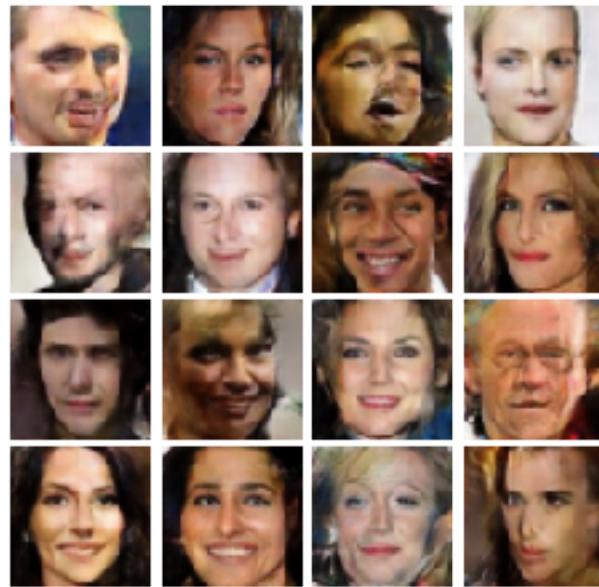
Iter: 18200, D: 0.05736, G: 0.2766



Iter: 18400, D: 0.1009, G: 0.4774



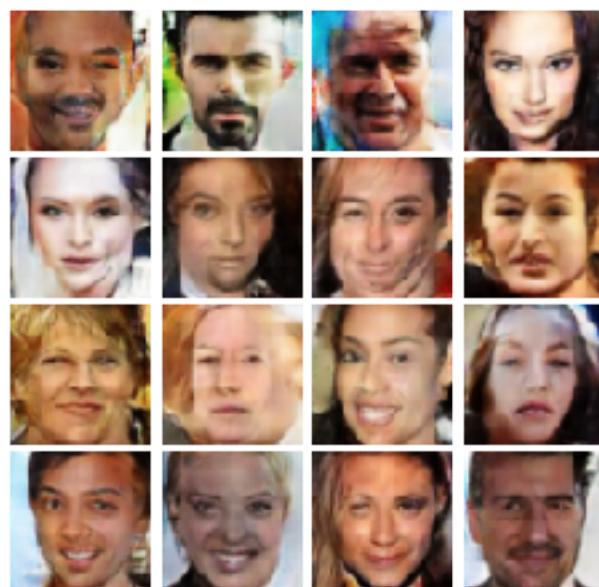
Iter: 18600, D: 0.03963, G: 0.2947



Saving the model as a checkpoint...

EPOCH: 20

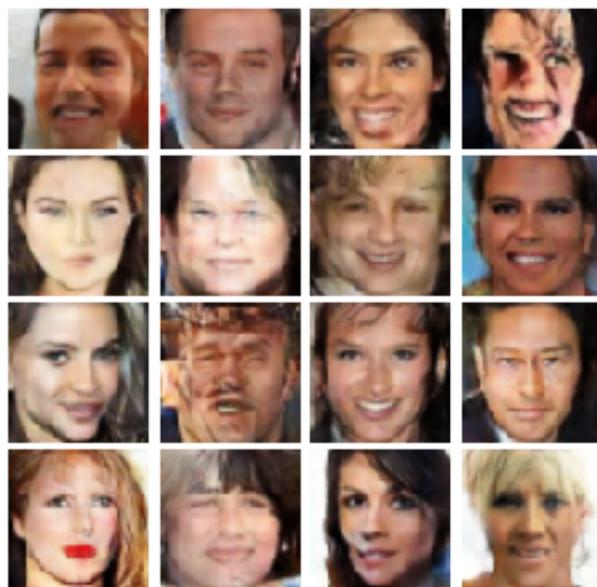
Iter: 18800, D: 0.0744, G: 0.4753



Iter: 19000, D: 0.04323, G: 0.3754



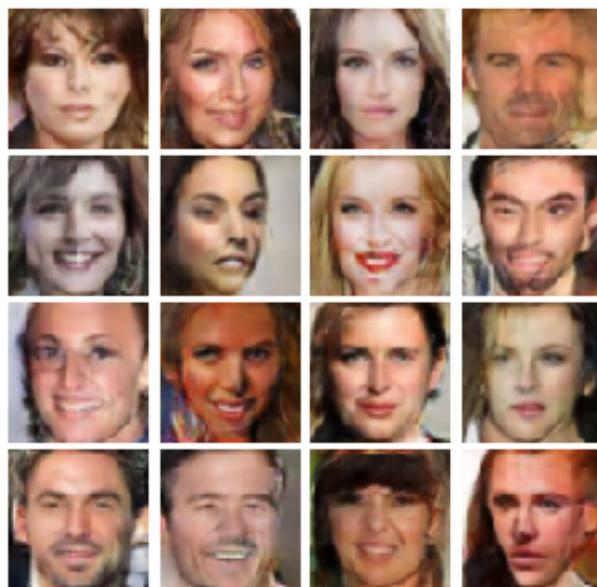
Iter: 19200, D: 0.1576, G: 0.5583



Iter: 19400, D: 0.0613, G: 0.4025



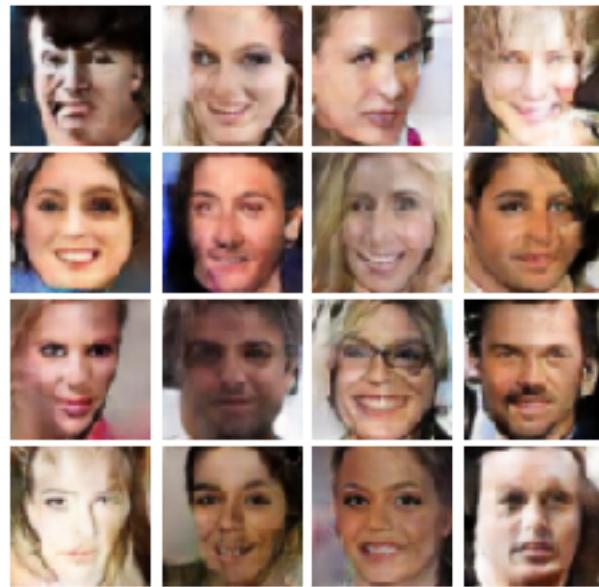
Iter: 19600, D: 0.02756, G: 0.4819



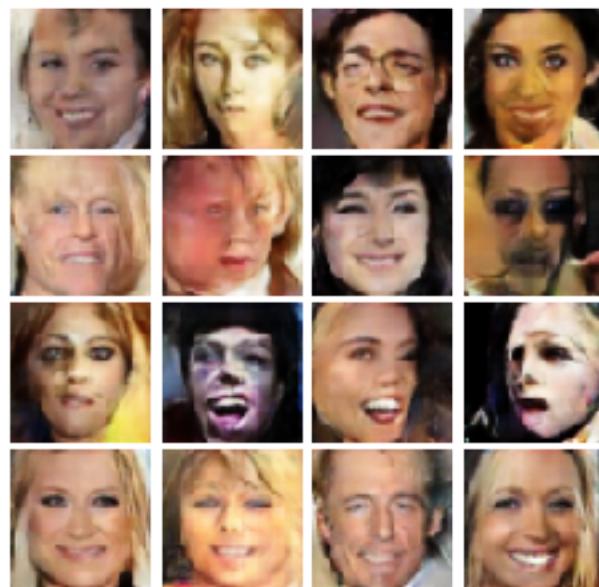
Saving the model as a checkpoint...

EPOCH: 21

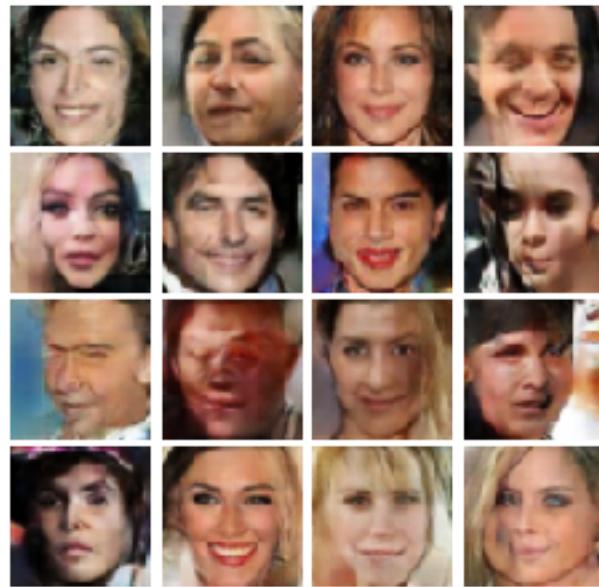
Iter: 19800, D: 0.04858, G: 0.5545



Iter: 20000, D: 0.05066, G:0.4018



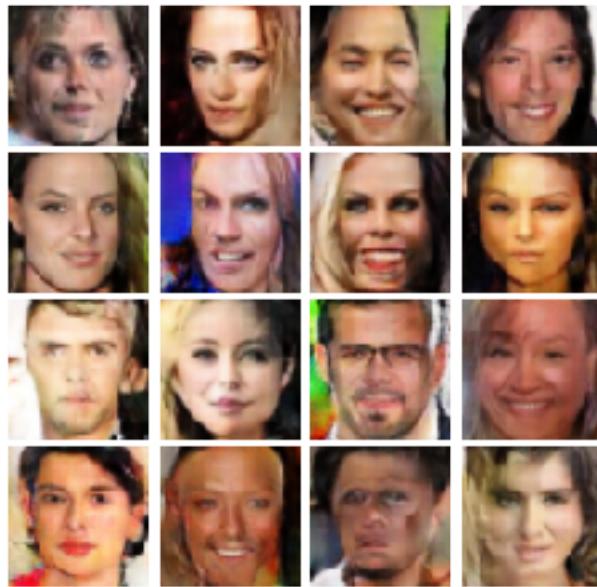
Iter: 20200, D: 0.0214, G:0.5585



Iter: 20400, D: 0.03784, G: 0.4776



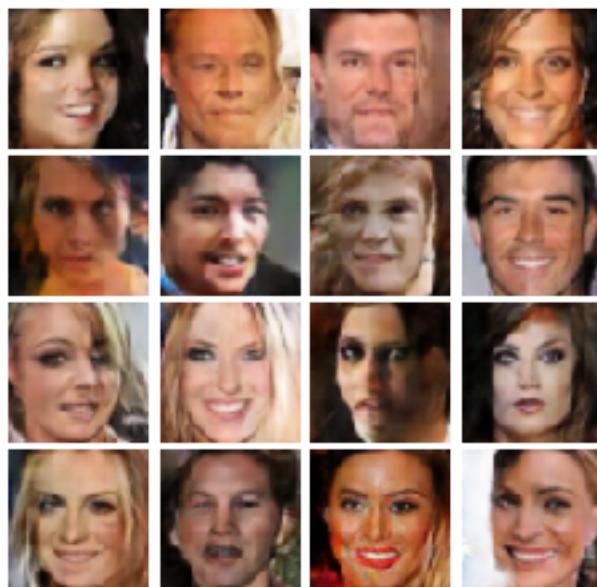
Iter: 20600, D: 0.06417, G: 0.2686



Saving the model as a checkpoint...

EPOCH: 22

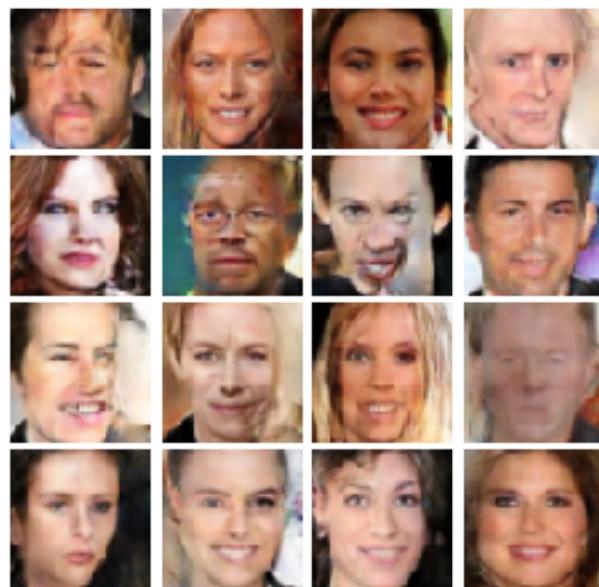
Iter: 20800, D: 0.02492, G: 0.5305



Iter: 21000, D: 0.01938, G: 0.4636



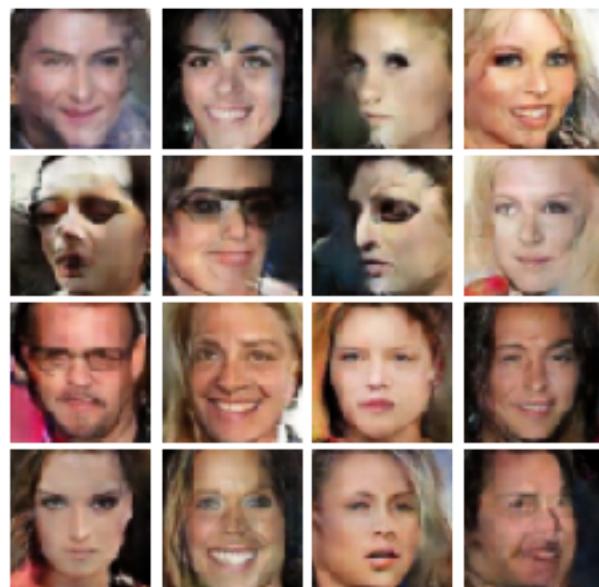
Iter: 21200, D: 0.05781, G: 0.4464



Iter: 21400, D: 0.04155, G: 0.5412



Iter: 21600, D: 0.04313, G: 0.4238



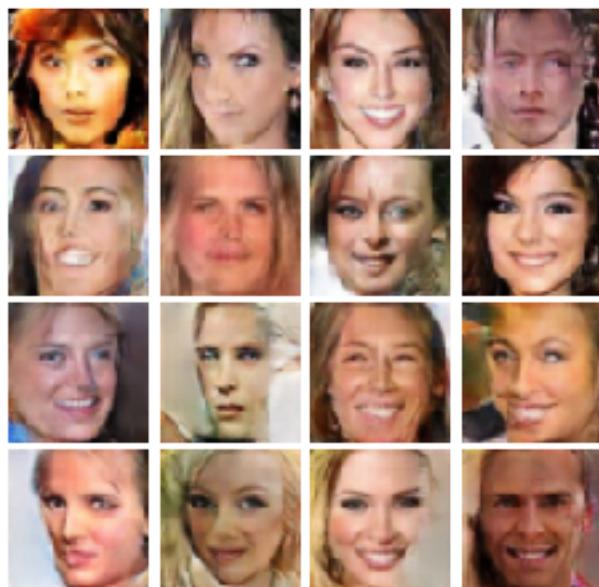
Saving the model as a checkpoint...

EPOCH: 23

Iter: 21800, D: 0.04868, G: 0.5443



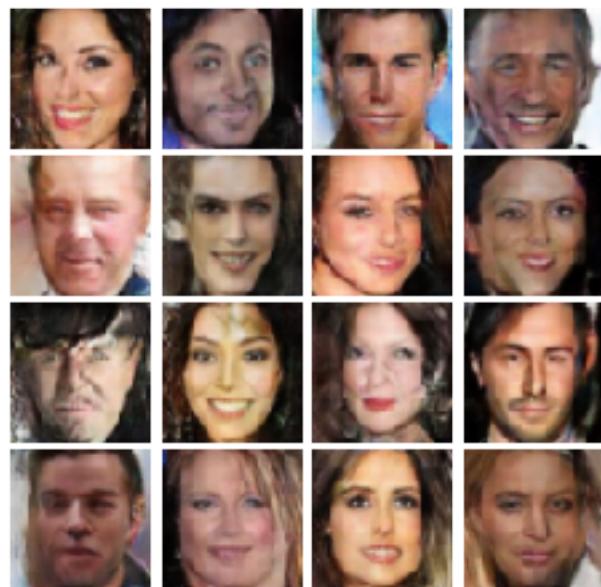
Iter: 22000, D: 0.02429, G: 0.5061



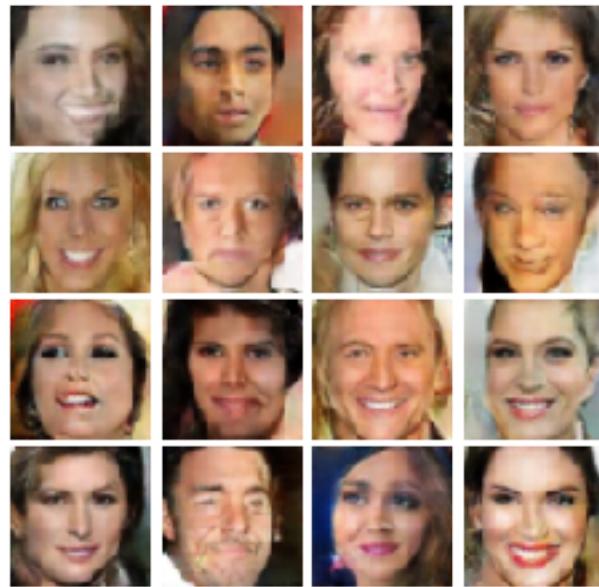
Iter: 22200, D: 0.03327, G: 0.3419



Iter: 22400, D: 0.04036, G: 0.2092



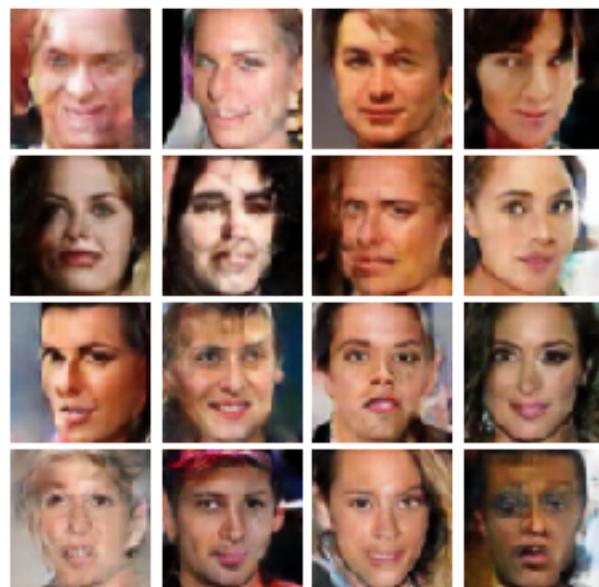
Iter: 22600, D: 0.02953, G: 0.4085



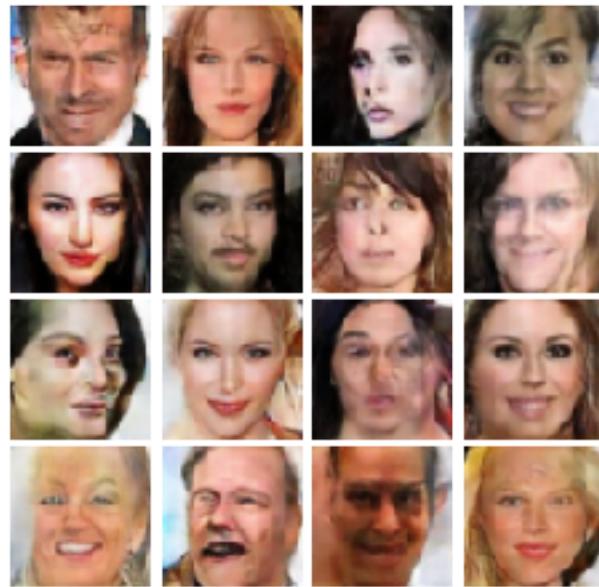
Saving the model as a checkpoint...

EPOCH: 24

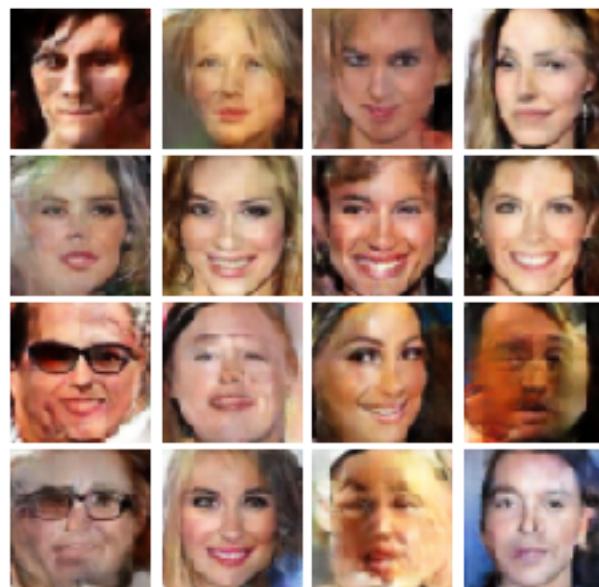
Iter: 22800, D: 0.04062, G: 0.4497



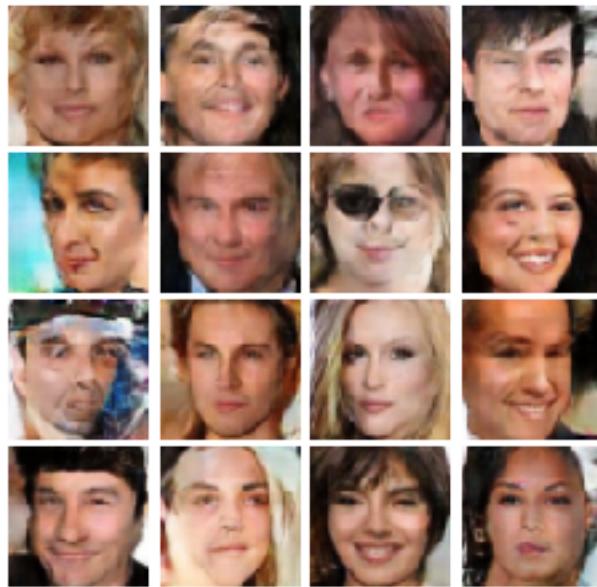
Iter: 23000, D: 0.06767, G: 0.3397



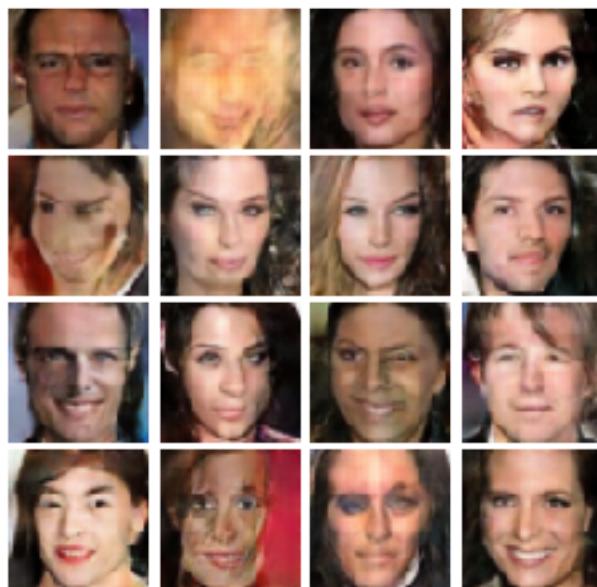
Iter: 23200, D: 0.1264, G: 0.5443



Iter: 23400, D: 0.03788, G: 0.5158



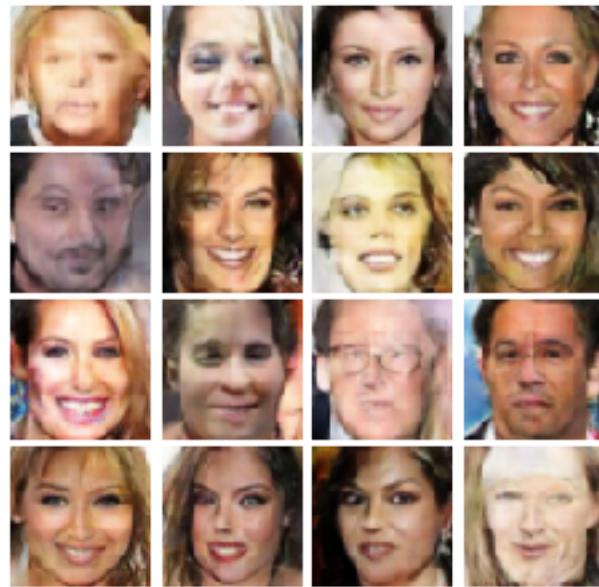
Iter: 23600, D: 0.02618, G:0.452



Saving the model as a checkpoint...

EPOCH: 25

Iter: 23800, D: 0.04964, G:0.5466



Iter: 24000, D: 0.1095, G:0.606



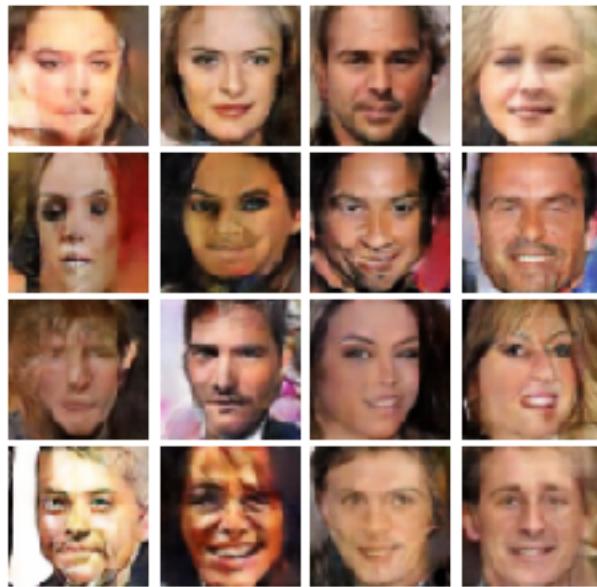
Iter: 24200, D: 0.02165, G:0.5196



Iter: 24400, D: 0.08117, G: 0.4152



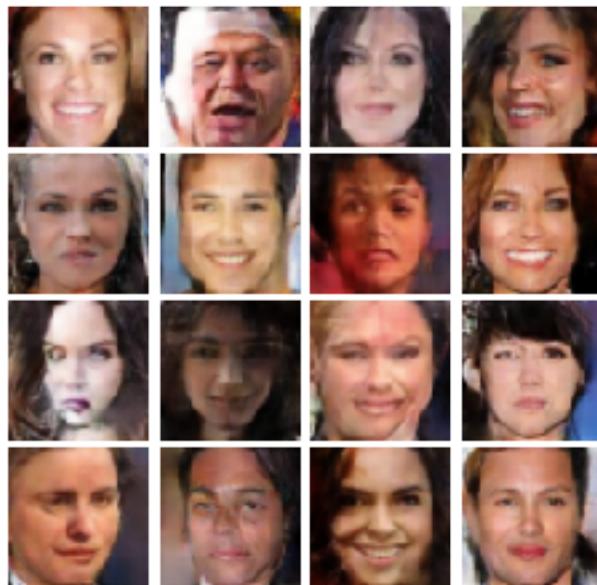
Iter: 24600, D: 0.07538, G: 0.3322



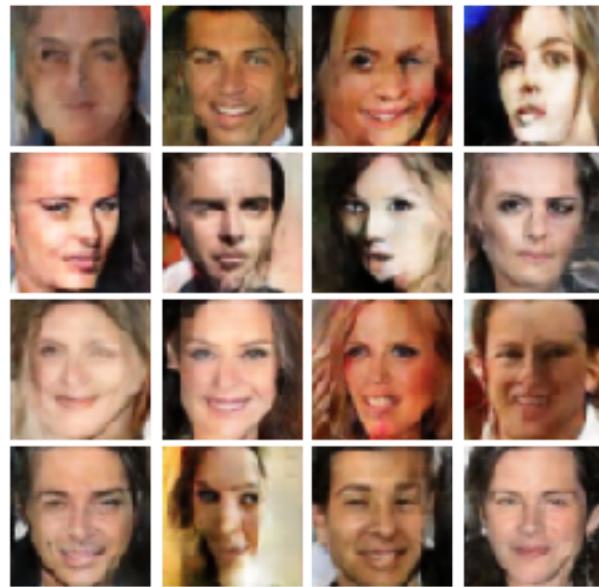
Saving the model as a checkpoint...

EPOCH: 26

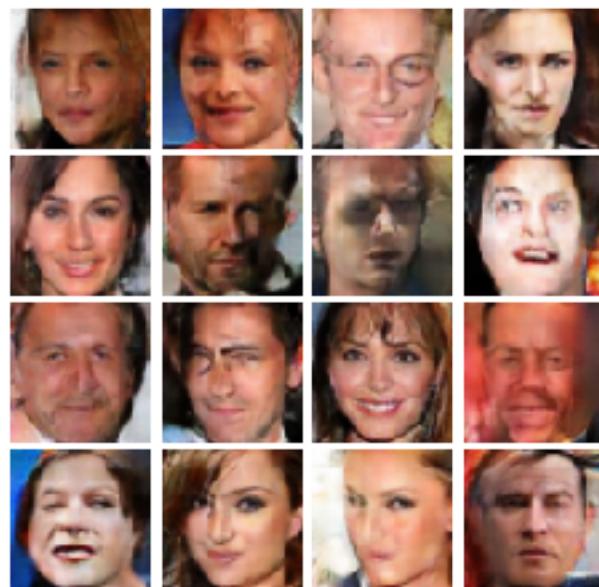
Iter: 24800, D: 0.02229, G: 0.4868



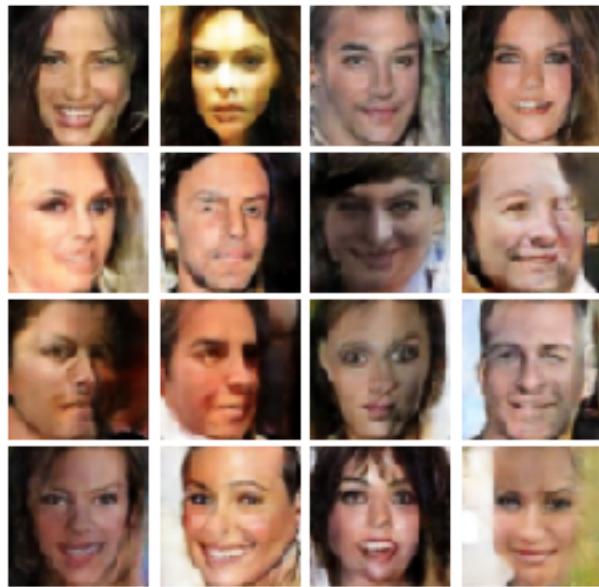
Iter: 25000, D: 0.03097, G: 0.5348



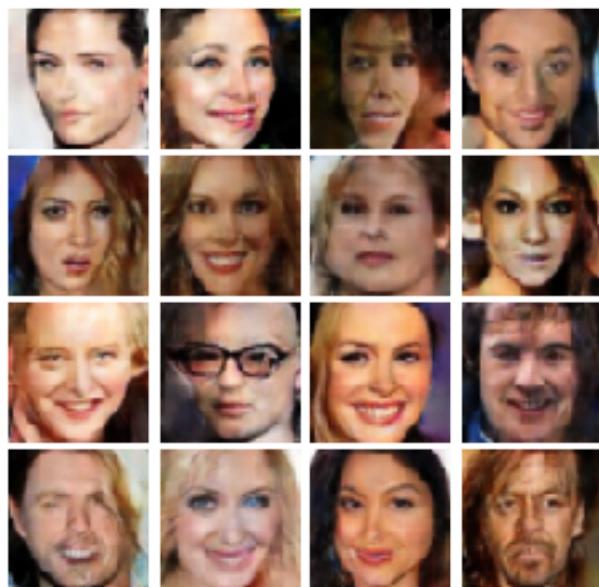
Iter: 25200, D: 0.09722, G: 0.442



Iter: 25400, D: 0.03887, G: 0.3697



Iter: 25600, D: 0.04209, G: 0.5733



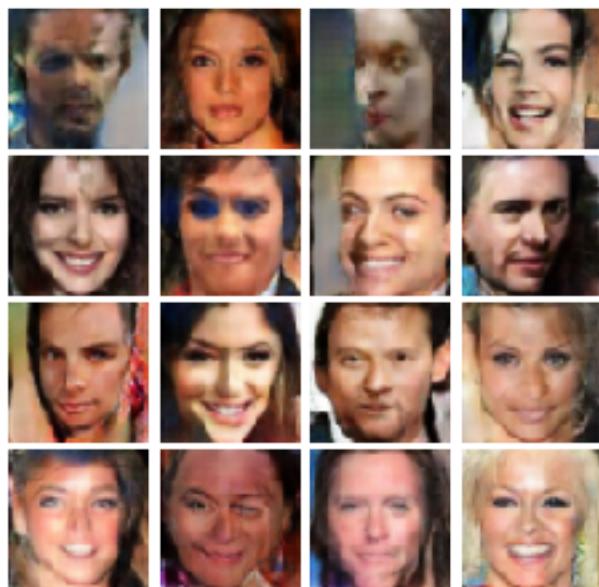
Saving the model as a checkpoint...

EPOCH: 27

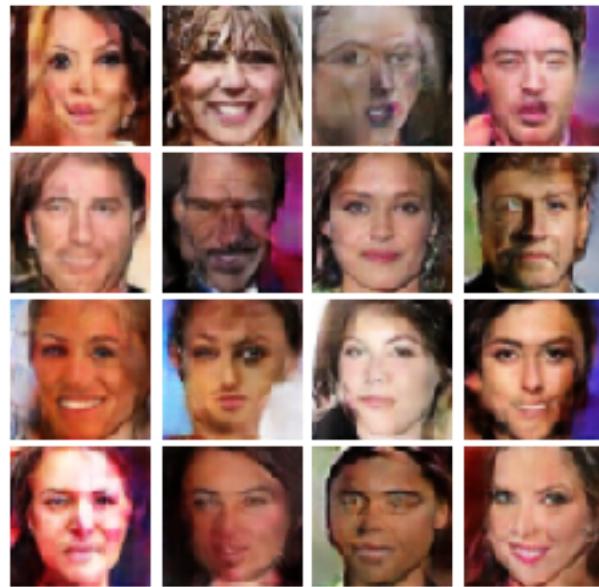
Iter: 25800, D: 0.02997, G: 0.5028



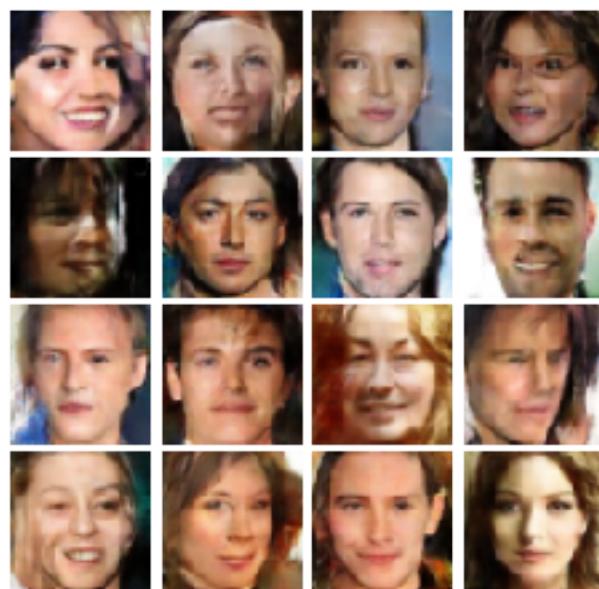
Iter: 26000, D: 0.0237, G: 0.3771



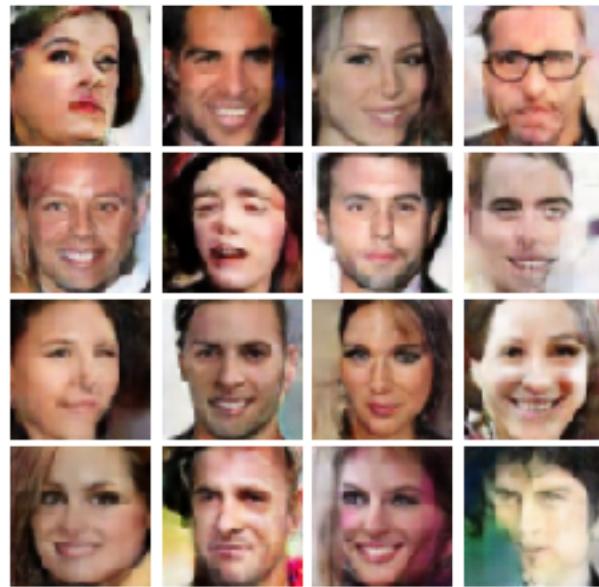
Iter: 26200, D: 0.138, G: 0.5709



Iter: 26400, D: 0.03218, G: 0.2654



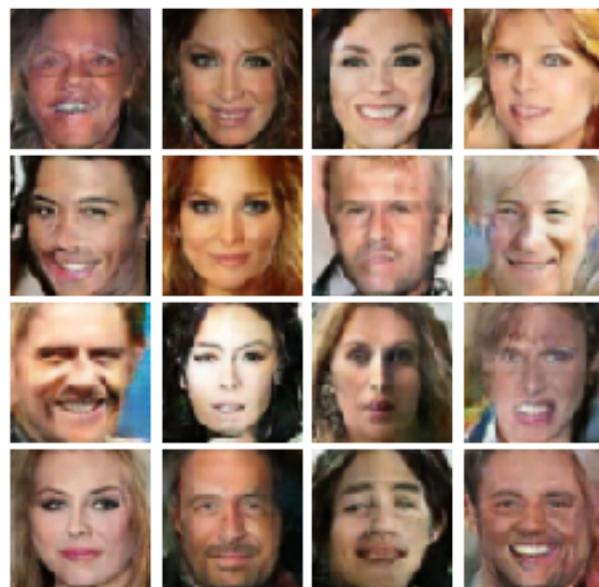
Iter: 26600, D: 0.03086, G: 0.2195



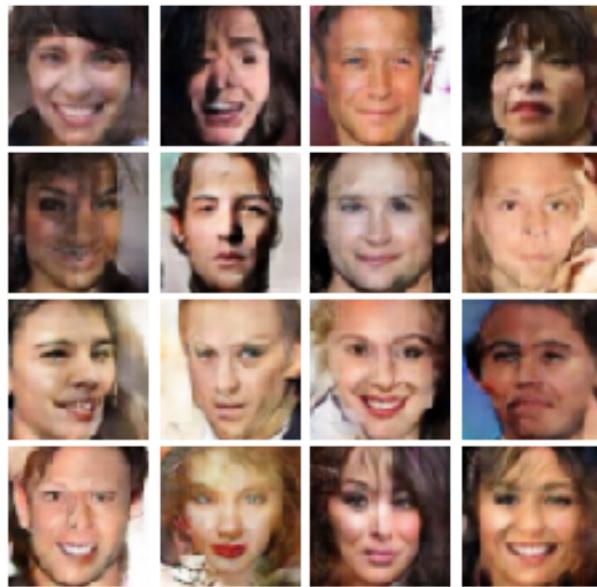
Saving the model as a checkpoint...

EPOCH: 28

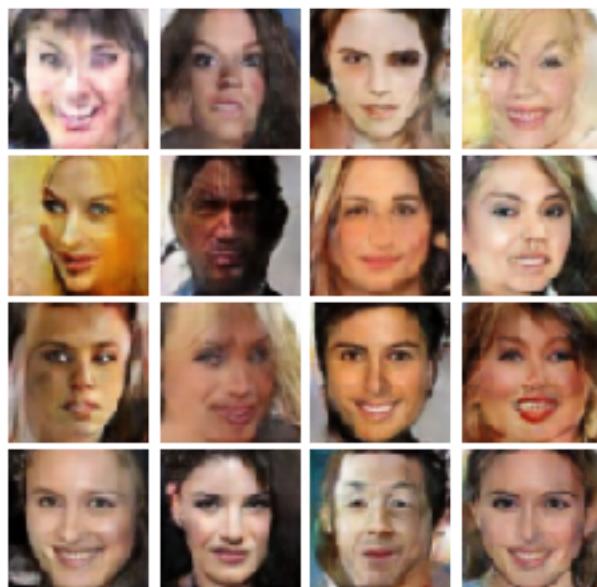
Iter: 26800, D: 0.0234, G: 0.4709



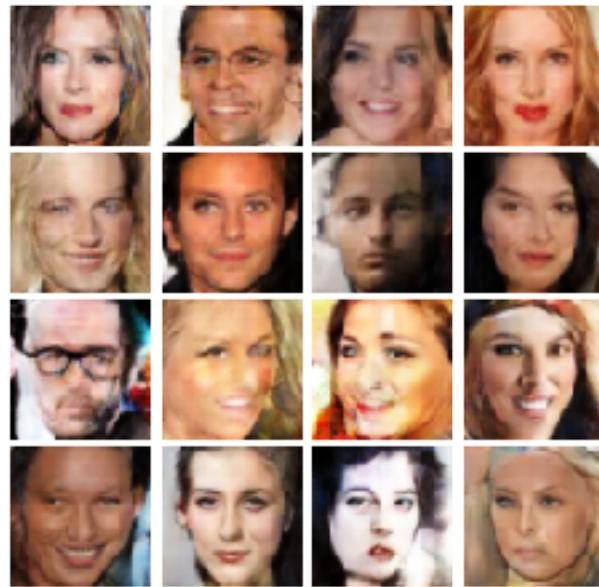
Iter: 27000, D: 0.01155, G: 0.4213



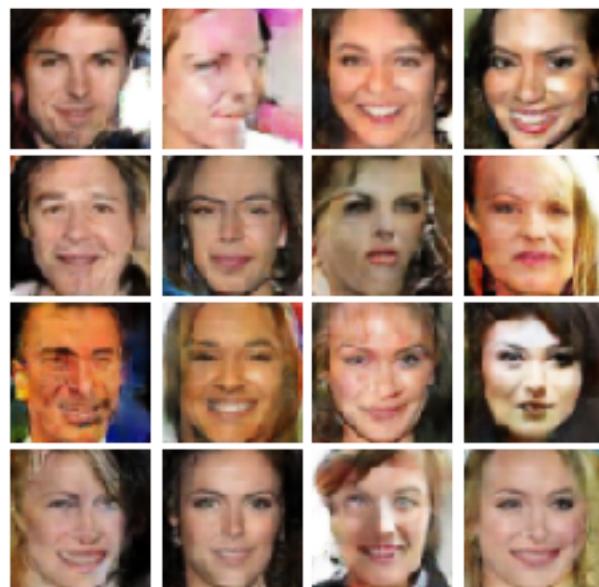
Iter: 27200, D: 0.01559, G: 0.4572



Iter: 27400, D: 0.02154, G: 0.4623



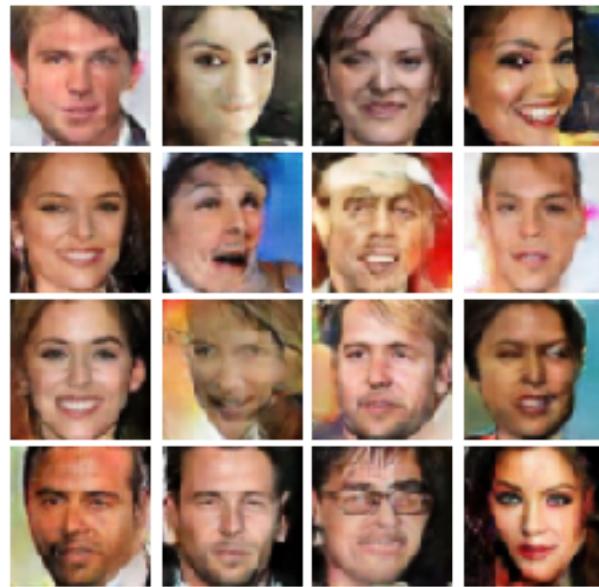
Iter: 27600, D: 0.01426, G: 0.5482



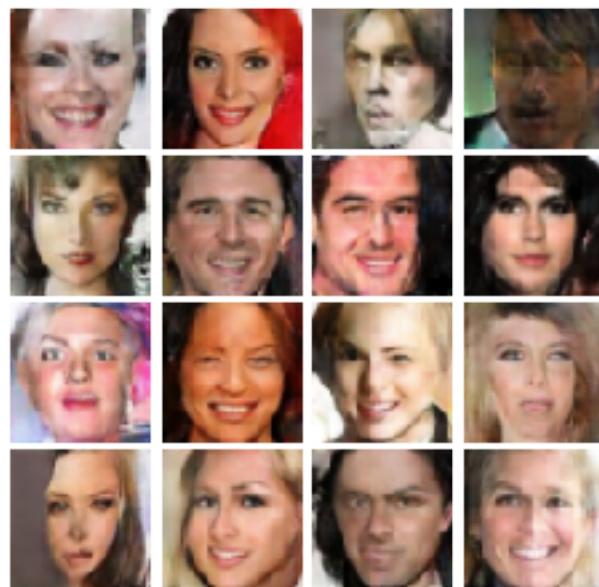
Saving the model as a checkpoint...

EPOCH: 29

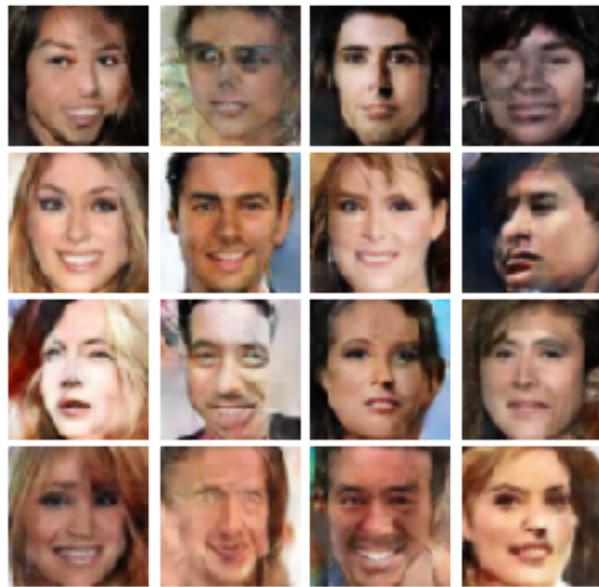
Iter: 27800, D: 0.04394, G: 0.4892



Iter: 28000, D: 0.01373, G: 0.5806



Iter: 28200, D: 0.02286, G: 0.3249



Iter: 28400, D: 0.04206, G: 0.3032



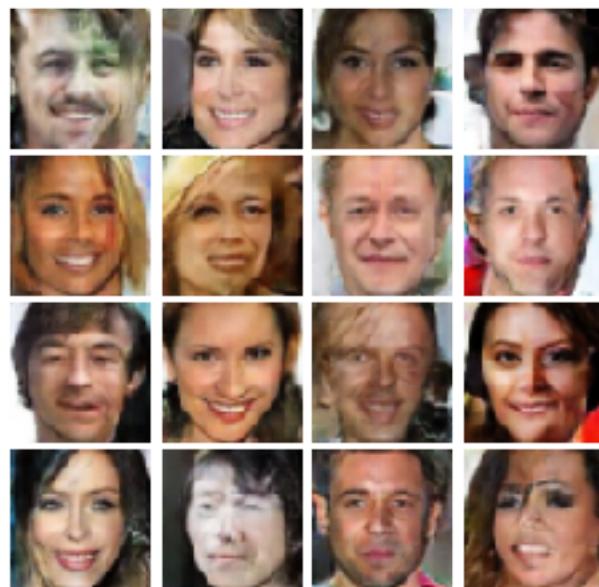
Saving the model as a checkpoint...

EPOCH: 30

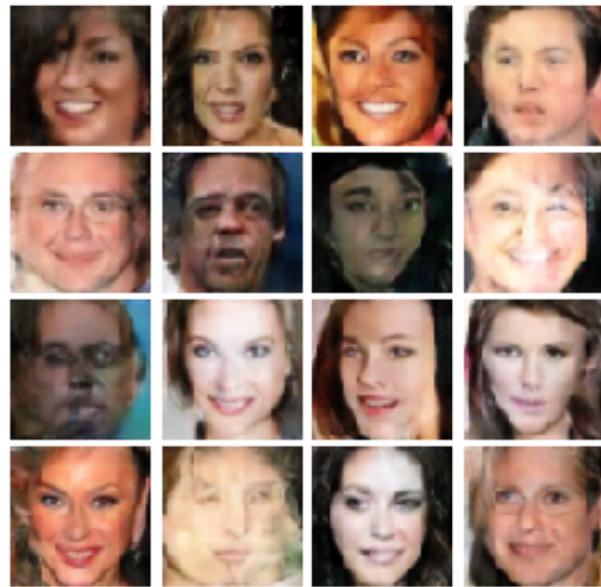
Iter: 28600, D: 0.02368, G: 0.4357



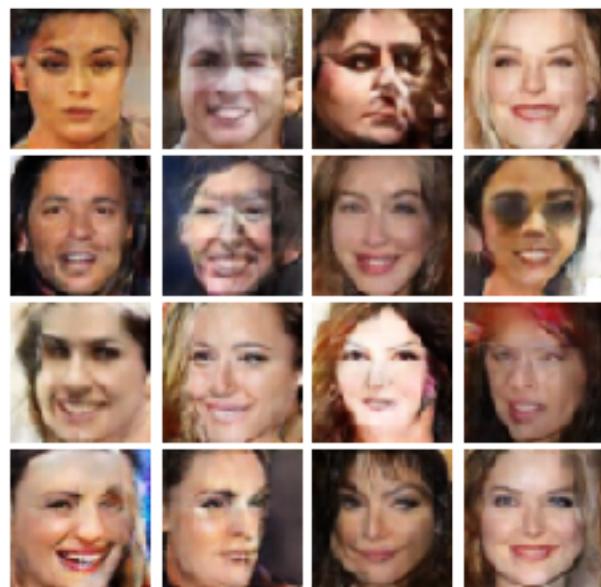
Iter: 28800, D: 0.1061, G:0.5641



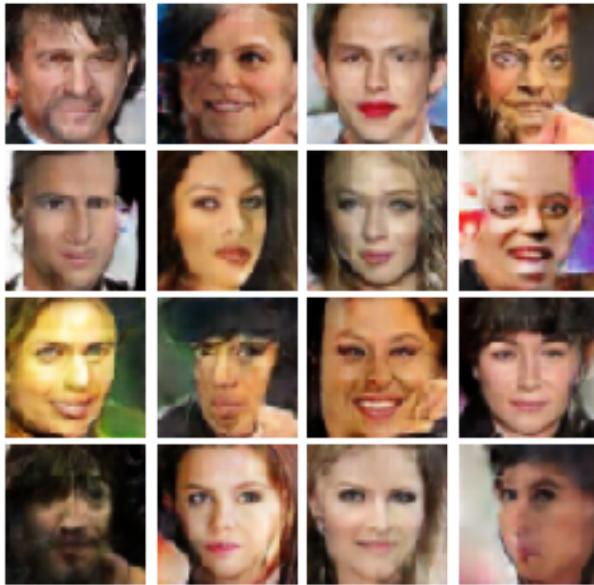
Iter: 29000, D: 0.03452, G:0.5016



Iter: 29200, D: 0.1048, G:0.6234



Iter: 29400, D: 0.0151, G:0.528



Saving the model as a checkpoint...

4.0.5 Section 2.6 Conclusion? [10 pts]

You have successfully trained your GAN models. Hurray! Now, we want you to answer the following question in few sentences. 1. Have you observed any difference between GAN and LSGAN (e.g. training, generation quality)?

The LSGAN are more stable than GAN at training. However, I didn't observe significant differences for results of GAN and LSGAN.

2. Did your GAN models generate diverse faces? Or did it always give you similar stuff all the time?

Yes, my GAN models generate diverse faces.

3. How can you quantitatively evaluate your GAN model? Is the metric(s) meaningful? Is there any drawback about this metric(s)?

Inception Score (IS) is perhaps the most widely adopted score for GAN evaluation. A higher inception score indicates better-quality generated images.

FID performs well in terms of discriminability, robustness and computational efficiency. It has been shown that FID is consistent with human judgments and is more robust to noise than IS. A lower FID score indicates more realistic images that match the statistical properties of real images.

4.0.6 Guidelines for Downloading PDF in Google Colab

- Run below cells only in Google Colab, Comment out in case of Jupyter notebook

```
[ ]: #Run below two lines (in google colab), installation steps to get .pdf of the  
→notebook
```

```
!apt-get install texlive texlive-xetex texlive-latex-extra pandoc  
!pip install pypandoc
```

```
# After installation, comment above two lines and run again to remove  
→installation comments from the notebook.
```

```
[ ]: # Find path to your notebook file in drive and enter in below line
```

```
!jupyter nbconvert --to PDF "your_notebook_path_here/DL_Assignment_5.ipynb"
```

```
#Example: "/content/drive/My Drive/DL_Fall_2020/Assignment_5/DL_Assignment_5.  
→ipynb"
```